



Red Hat

Red Hat Enterprise Linux 7

Kernel Administration Guide

Examples of Tasks for Managing the Kernel

Red Hat Enterprise Linux 7 Kernel Administration Guide

Examples of Tasks for Managing the Kernel

Jaroslav Klech
Red Hat Customer Content Services
jklech@redhat.com

Marie Dolezelova
Red Hat Customer Content Services

Mark Flitter
Red Hat Customer Content Services

Douglas Silas
Red Hat Customer Content Services

Eliska Slobodova
Red Hat Customer Content Services

Jaromir Hradilek
Red Hat Customer Content Services

Maxim Svistunov
Red Hat Customer Content Services

Robert Krátký
Red Hat Customer Content Services

Stephen Wadeley
Red Hat Customer Content Services

Florian Nadge
Red Hat Customer Content Services

Legal Notice

Copyright © 2019 Red Hat, Inc.

The text of and illustrations in this document are licensed by Red Hat under a Creative Commons Attribution–Share Alike 3.0 Unported license ("CC-BY-SA"). An explanation of CC-BY-SA is available at

<http://creativecommons.org/licenses/by-sa/3.0/>

. In accordance with CC-BY-SA, if you distribute this document or an adaptation of it, you must provide the URL for the original version.

Red Hat, as the licensor of this document, waives the right to enforce, and agrees not to assert, Section 4d of CC-BY-SA to the fullest extent permitted by applicable law.

Red Hat, Red Hat Enterprise Linux, the Shadowman logo, the Red Hat logo, JBoss, OpenShift, Fedora, the Infinity logo, and RHCE are trademarks of Red Hat, Inc., registered in the United States and other countries.

Linux® is the registered trademark of Linus Torvalds in the United States and other countries.

Java® is a registered trademark of Oracle and/or its affiliates.

XFS® is a trademark of Silicon Graphics International Corp. or its subsidiaries in the United States and/or other countries.

MySQL® is a registered trademark of MySQL AB in the United States, the European Union and other countries.

Node.js® is an official trademark of Joyent. Red Hat is not formally related to or endorsed by the official Joyent Node.js open source or commercial project.

The OpenStack® Word Mark and OpenStack logo are either registered trademarks/service marks or trademarks/service marks of the OpenStack Foundation, in the United States and other countries and are used with the OpenStack Foundation's permission. We are not affiliated with, endorsed or sponsored by the OpenStack Foundation, or the OpenStack community.

All other trademarks are the property of their respective owners.

Abstract

The Kernel Administration Guide documents tasks for maintaining the Red Hat Enterprise Linux 7 kernel. This release, includes information on using kpatch, managing kernel modules, and manually updating the kernel.

Table of Contents

PREFACE	5
CHAPTER 1. WORKING WITH KERNEL MODULES	6
1.1. WHAT IS A KERNEL MODULE?	6
1.2. LISTING CURRENTLY-LOADED MODULES	6
1.3. DISPLAYING INFORMATION ABOUT A MODULE	7
CHAPTER 2. WORKING WITH SYSCTL AND KERNEL TUNABLES	9
2.1. WHAT IS A KERNEL TUNABLE?	9
2.2. HOW TO WORK WITH KERNEL TUNABLES	9
2.2.1. Using the sysctl command	9
2.2.2. Modifying files in /etc/sysctl.d	9
2.3. WHAT TUNABLES CAN BE CONTROLLED?	10
2.3.1. Network interface tunables	10
2.3.2. Global kernel tunables	20
CHAPTER 3. LISTING OF KERNEL PARAMETERS AND VALUES	24
3.1. KERNEL COMMAND-LINE PARAMETERS	24
3.1.1. Setting kernel command-line parameters	24
3.1.2. What kernel command-line parameters can be controlled	25
3.1.2.1. Hardware specific kernel command-line parameters	25
CHAPTER 4. KERNEL FEATURES	27
4.1. CONTROL GROUPS	27
4.1.1. What is a control group?	27
4.1.2. What is a namespace?	27
4.1.3. Supported namespaces	27
4.2. KERNEL SOURCE CHECKER	28
4.2.1. Usage	28
4.3. DIRECT ACCESS FOR FILES (DAX)	28
4.4. MEMORY PROTECTION KEYS FOR USERSPACE (ALSO KNOWN AS PKU, OR PKEYS)	29
4.5. KERNEL ADDRESS SPACE LAYOUT RANDOMIZATION	29
4.6. ADVANCED ERROR REPORTING (AER)	30
CHAPTER 5. MANUALLY UPGRADING THE KERNEL	31
5.1. OVERVIEW OF KERNEL PACKAGES	31
5.2. PREPARING TO UPGRADE	32
5.3. DOWNLOADING THE UPGRADED KERNEL	34
5.4. PERFORMING THE UPGRADE	34
5.5. VERIFYING THE INITIAL RAM FILE SYSTEM IMAGE	34
Verifying the initial RAM file system image and kernel on IBM eServer System i	36
Reversing the changes made to the initial RAM file system image	37
Listing the contents of the initial RAM file system image	37
5.6. VERIFYING THE BOOT LOADER	38
CHAPTER 6. APPLYING PATCHES WITH KERNEL LIVE PATCHING	39
6.1. ACCESS TO KERNEL LIVE PATCHES	39
6.2. COMPONENTS OF KERNEL LIVE PATCHING	39
6.3. HOW KERNEL LIVE PATCHING WORKS	40
6.4. ENABLING KERNEL LIVE PATCHING	40
6.4.1. Subscribing to the live patching stream	41
Prerequisites	41
Procedure	41

Additional resources	42
6.5. UPDATING KERNEL PATCH MODULES	42
Prerequisites	42
Procedure	42
6.6. DISABLING KERNEL LIVE PATCHING	42
6.6.1. Removing the live patching package	42
Prerequisites	42
Procedure	43
Additional resources	44
6.6.2. Uninstalling the kernel patch module	44
Prerequisites	44
Procedure	44
Additional resources	45
6.6.3. Disabling the kpatch service unit file	45
Prerequisites	45
Procedure	45
Additional resources	46
6.7. KPATCH SUPPORT	46
6.8. SUPPORT FOR THIRD-PARTY LIVE PATCHING	47
6.9. LIMITATIONS OF KPATCH	47
CHAPTER 7. KERNEL CRASH DUMP GUIDE	48
7.1. INTRODUCTION TO KDUMP	48
7.1.1. About kdump and kexec	48
7.1.2. Memory requirements	48
7.2. INSTALLING AND CONFIGURING KDUMP	49
7.2.1. Installing kdump	49
7.2.2. Configuring kdump on the command line	50
7.2.2.1. Configuring the memory usage	50
7.2.2.2. Configuring the kdump type	51
7.2.2.3. Configuring the core collector	52
7.2.2.4. Configuring the default action	52
7.2.2.5. Enabling the service	52
7.2.3. Configuring kdump in the graphical user interface	53
7.2.3.1. Configuring the memory usage	53
7.2.3.2. Configuring the kdump type	54
7.2.3.3. Configuring the core collector	55
7.2.3.4. Configuring the default action	56
7.2.3.5. Enabling the service	57
7.3. TESTING THE KDUMP CONFIGURATION	58
7.3.1. Additional resources	58
7.3.1.1. Installed documentation	58
7.3.1.2. Online documentation	58
7.4. FIRMWARE ASSISTED DUMP MECHANISMS	59
7.4.1. The case for firmware assisted dump	59
7.4.2. Using fadump on IBM PowerPC hardware	59
7.4.3. Firmware assisted dump methods on IBM Z	60
7.4.4. Using sadump on Fujitsu PRIMEQUEST systems	60
7.5. ANALYZING A CORE DUMP	61
7.5.1. Installing the crash utility	61
7.5.2. Running the crash utility	61
7.5.3. Displaying the message buffer	62
7.5.4. Displaying a backtrace	63

7.5.5. Displaying a process status	64
7.5.6. Displaying virtual memory information	64
7.5.7. Displaying open files	65
7.5.8. Exiting the utility	65
7.6. FREQUENTLY ASKED QUESTIONS	66
7.7. SUPPORTED KDUMP CONFIGURATIONS AND TARGETS	68
7.7.1. Memory requirements for kdump	68
7.7.2. Minimum threshold for automatic memory reservation	68
7.7.3. Supported kdump targets	69
7.7.4. Supported kdump filtering levels	70
7.7.5. Supported default actions	70
7.7.6. Estimating kdump size	71
7.8. PORTAL LABS RELEVANT TO KDUMP	72
7.8.1. Kdump helper	72
7.8.2. Kernel oops analyzer	72
CHAPTER 8. REVISION HISTORY	73

PREFACE

The *Kernel Administration Guide* describes working with the kernel and shows several practical tasks. Beginning with information on using kernel modules, the guide then covers interaction with the **sysfs** facility, manual upgrade of the kernel and using kpatch. The guide also introduces the crash dump mechanism, which steps through the process of setting up and testing **vmcore** collection in the event of a kernel failure.

The *Kernel Administration Guide* also covers selected use cases of managing the kernel and includes reference material about command line options, kernel tunables (also known as switches), and a brief discussion of kernel features.

CHAPTER 1. WORKING WITH KERNEL MODULES

This Chapter explains:

- What is a kernel module.
- How to use the **kmod** utilities to manage modules and their dependencies.
- How to configure module parameters to control behavior of the kernel modules.
- How to load modules at boot time.



NOTE

In order to use the kernel module utilities described in this chapter, first ensure the **kmod** package is installed on your system by running, as root:

```
# yum install kmod
```

1.1. WHAT IS A KERNEL MODULE?

The Linux kernel is monolithic by design. However, it is compiled with optional or additional modules as required by each use case. This means that you can extend the kernel’s capabilities through the use of dynamically-loaded *kernel modules*. A kernel module can provide:

- a device driver which adds support for new hardware,
- support for a file system such as **GFS2** or **NFS**.

Like the kernel itself, modules can take parameters that customize their behavior, though the default parameters work well in most cases. User-space tools can list the modules currently loaded into a running kernel; query all available modules for available parameters and module-specific information; and load or unload (remove) modules dynamically into or from a running kernel. Many of these utilities, which are provided by the **kmod** package, take module dependencies into account when performing operations so that manual dependency-tracking is rarely necessary.

On modern systems, kernel modules are automatically loaded by various mechanisms when needed. However, there are occasions when it is necessary to load or unload modules manually. For example when one module is preferred over another although either is able to provide basic functionality, or when a module is misbehaving.

1.2. LISTING CURRENTLY-LOADED MODULES

You can list all kernel modules that are currently loaded into the kernel by running the **lsmod** command, for example:

```
# lsmod
Module      Size Used by
tcp_ip      12663  0
bnep        19704  2
bluetooth   372662 7 bnep
rfkill      26536  3 bluetooth
fuse         87661  3
ebtable_broute 12731  0
```

```

bridge          110196  1 ebttable_broute
stp             12976   1 bridge
llc             14552   2 stp,bridge
ebttable_filter      12827  0
ebtables        30913   3 ebttable_broute,ebttable_nat,ebttable_filter
ip6table_nat    13015   1
nf_nat_ipv6     13279   1 ip6table_nat
iptable_nat     13011   1
nf_conntrack_ipv4 14862   4
nf_defrag_ipv4   12729   1 nf_conntrack_ipv4
nf_nat_ipv4     13263   1 iptable_nat
nf_nat          21798   4 nf_nat_ipv4,nf_nat_ipv6,ip6table_nat,iptable_nat
[output truncated]

```


The **lsmod** output specifies three columns:

- **Module**
 - The name of a kernel module currently loaded in memory.
- **Size**
 - The amount of memory the kernel module uses in kilobytes.
- **Used by**
 - A decimal number representing how many dependencies there are on the **Module** field.
 - A comma separated string of dependent **Module** names. Using this list, you can first unload all the modules depending the module you want to unload.

Finally, note that **lsmod** output is less verbose and considerably easier to read than the content of the **/proc/modules** pseudo-file.

1.3. DISPLAYING INFORMATION ABOUT A MODULE

You can display detailed information about a kernel module using the **modinfo module_name** command.



NOTE

When entering the name of a kernel module as an argument to one of the **kmod** utilities, do not append a **.ko** extension to the end of the name. Kernel module names do not have extensions; their corresponding files do.

Example 1.1. Listing information about a kernel module with lsmod

To display information about the **e1000e** module, which is the Intel PRO/1000 network driver, enter the following command as **root**:

```
# modinfo e1000e
filename: /lib/modules/3.10.0-
121.el7.x86_64/kernel/drivers/net/ethernet/intel/e1000e/e1000e.ko
version: 2.3.2-k
```

license: GPL
description: Intel(R) PRO/1000 Network Driver
author: Intel Corporation,

CHAPTER 2. WORKING WITH SYSCTL AND KERNEL TUNABLES

2.1. WHAT IS A KERNEL TUNABLE?

Kernel tunables are used to customize the behavior of Red Hat Enterprise Linux at boot, or on demand while the system is running. Some hardware parameters are specified at boot time only and cannot be altered once the system is running, most however, can be altered as required and set permanent for the next boot.

2.2. HOW TO WORK WITH KERNEL TUNABLES

There are three ways to modify kernel tunables.

1. Using the **sysctl** command
2. By manually modifying configuration files in the **/etc/sysctl.d** directory
3. Through a shell, interacting with the virtual file system mounted at **/proc/sys**



NOTE

Not all boot time parameters are under control of the sysfs subsystem, some hardware specific option must be set on the kernel command line, the Kernel Parameters section of this guide addresses those options

2.2.1. Using the **sysctl** command

The **sysctl** command is used to list, read, and set kernel tunables. It can filter tunables when listing or reading and set tunables temporarily or permanently.

1. Listing variables

```
# sysctl -a
```

2. Reading variables

```
# sysctl kernel.version +
kernel.version = #1 SMP Fri Jan 19 13:19:54 UTC 2018
```

3. Writing variables temporarily

```
# sysctl <tunable class>.<tunable>=<value>
```

4. Writing variables permanently

```
# sysctl -w <tunable class>.<tunable>=<value> >> /etc/sysctl.conf
```

2.2.2. Modifying files in **/etc/sysctl.d**

To override a default at boot, you can also manually populate files in **/etc/sysctl.d**.

1. Create a new file in **/etc/sysctl.d**

```
# vim /etc/sysctl.d/99-custom.conf
```

2. Include the variables you wish to set, one per line, in the following form

```
<tunable class>.<tunable> = <value> +
<tunable class>.<tunable> = <value>
```

3. Save the file

4. Either reboot the machine to make the changes take effect

or

Execute **sysctl -p /etc/sysctl.d/99-custom.conf** to apply the changes without rebooting

2.3. WHAT TUNABLES CAN BE CONTROLLED?

Tunables are divided into groups by kernel subsystem. A Red Hat Enterprise Linux system has the following classes of tunables:

Table 2.1. Table of sysctl interfaces

Class	Subsystem
abi	Execution domains and personalities
crypto	Cryptographic interfaces
debug	Kernel debugging interfaces
dev	Device specific information
fs	Global and specific filesystem tunables
kernel	Global kernel tunables
net	Network tunables
sunrpc	Sun Remote Procedure Call (NFS)
user	User Namespace limits
vm	Tuning and management of memory, buffer, and cache

2.3.1. Network interface tunables

System administrators are able to adjust the network configuration on a running system through the networking tunables.

Networking tunables are included in the **/proc/sys/net** directory, which contains multiple subdirectories for various networking topics. To adjust the network configuration, system administrators need to modify the files within such subdirectories.

The most frequently used directories are:

1. **/proc/sys/net/core/**
2. **/proc/sys/net/ipv4/**

The **/proc/sys/net/core/** directory contains a variety of settings that control the interaction between the kernel and networking layers. By adjusting some of those tunables, you can improve performance of a system, for example by increasing the size of a receive queue, increasing the maximum connections or the memory dedicated to network interfaces. Note that the performance of a system depends on different aspects according to the individual issues.

The **/proc/sys/net/ipv4/** directory contains additional networking settings, which are useful when preventing attacks on the system or when using the system to act as a router. The directory contains both IP and TCP variables. For detailed explanation of those variables, see **/usr/share/doc/kernel-doc-<version>/Documentation/networking/ip-sysctl.txt**.

Other directories within the **/proc/sys/net/ipv4/** directory cover different aspects of the network stack:

1. **/proc/sys/net/ipv4/conf/** - allows you to configure each system interface in different ways, including the use of default settings for unconfigured devices and settings that override all special configurations
2. **/proc/sys/net/ipv4/neigh/** - contains settings for communicating with a host directly connected to the system and also contains different settings for systems more than one step away
3. **/proc/sys/net/ipv4/route/** - contains specifications that apply to routing with any interfaces on the system

This list of network tunables is relevant to IPv4 interfaces and are accessible from the **/proc/sys/net/ipv4/{all,<interface_name>}/** directory.

Description of the following parameters have been adopted from the kernel documentation sites.^[1]

log_martians

Log packets with impossible addresses to kernel log.

Type	Default
Boolean	0

Enabled if one or more of **conf/{all,interface}/log_martians** is set to TRUE

+ Further Resources

- [What is the kernel parameter net.ipv4.conf.all.log_martians for?](#)
- [Why do I see "martian source" logs in the messages file ?](#)

accept_redirects

Accept ICMP redirect messages.

Type	Default
Boolean	1

accept_redirects for the interface is enabled under the following conditions:

- Both **conf/{all,interface}/accept_redirects** are TRUE (when forwarding for the interface is enabled)
- At least one of **conf/{all,interface}/accept_redirects** is TRUE (forwarding for the interface is disabled)

For more information refer to [How to enable or disable ICMP redirects](#)

forwarding

Enable IP forwarding on an interface.

Type	Default
Boolean	0

Further Resources

- [Turning on Packet Forwarding and Nonlocal Binding](#)

mc_forwarding

Do multicast routing.

Type	Default
Boolean	0

- Read only value
- A multicast routing daemon is required.
- **conf/all/mc_forwarding** must also be set to TRUE to enable multicast routing for the interface

Further Resources

- For an explanation of the read only behavior, see [Why system reports "permission denied on key" while setting the kernel parameter "net.ipv4.conf.all.mc_forwarding"?](#)

medium_id

Arbitrary value used to differentiate the devices by the medium they are attached to.

Type	Default
Integer	0

Notes

- Two devices on the same medium can have different id values when the broadcast packets are received only on one of them.
- The default value 0 means that the device is the only interface to its medium
- value of -1 means that medium is not known.
- Currently, it is used to change the proxy_arp behavior:
- the proxy_arp feature is enabled for packets forwarded between two devices attached to different media.

Further Resources - For examples, see [Using the "medium_id" feature in Linux 2.2 and 2.4](#)

proxy_arp

Do proxy arp.

Type	Default
Boolean	0

proxy_arp for the interface is enabled if at least one of **conf/{all,interface}/proxy_arp** is set to TRUE, otherwise it is disabled

proxy_arp_pvlan

Private VLAN proxy arp.

Type	Default
Boolean	0

Allow proxy arp replies back to the same interface, to support features like [RFC 3069](#)

shared_media

Send(router) or accept(host) RFC1620 shared media redirects.

Type	Default
Boolean	1

Notes

- Overrides secure_redirects.
- **shared_media** for the interface is enabled if at least one of **conf/{all,interface}/shared_media** is set to TRUE

secure_redirects

Accept ICMP redirect messages only to gateways listed in the interface's current gateway list.

Type	Default
Boolean	1

Notes

- Even if disabled, RFC1122 redirect rules still apply.
- Overridden by shared_media.
- **secure_redirects** for the interface is enabled if at least one of **conf/{all,interface}/secure_redirects** is set to TRUE

send_redirects

Send redirects, if router.

Type	Default
Boolean	1

Notes

send_redirects for the interface is enabled if at least one of **conf/{all,interface}/send_redirects** is set to TRUE

bootp_relay

Accept packets with source address 0.b.c.d destined not to this host as local ones.

Type	Default
Boolean	0

Notes

- A BOOTP daemon must be enabled to manage these packets
- **conf/all/bootp_relay** must also be set to TRUE to enable BOOTP relay for the interface
- Not implemented, see [DHCP Relay Agent](#) in the Red Hat Enterprise Linux Networking Guide

accept_source_route

Accept packets with SRR option.

Type	Default
Boolean	1

Notes

- **conf/all/accept_source_route** must also be set to TRUE to accept packets with SRR option on the interface

accept_local

Accept packets with local source addresses.

Type	Default
Boolean	0

Notes

- In combination with suitable routing, this can be used to direct packets between two local interfaces over the wire and have them accepted properly.
- **rp_filter** must be set to a non-zero value in order for accept_local to have an effect.

route_localnet

Do not consider loopback addresses as martian source or destination while routing.

Type	Default
Boolean	0

Notes

- This enables the use of 127/8 for local routing purposes.

rp_filter

Enable source Validation

Type	Default
Integer	0

Value	Effect
0	No source validation.
1	Strict mode as defined in RFC3704 Strict Reverse Path

Value	Effect
2	Loose mode as defined in RFC3704 Loose Reverse Path

Notes

- Current recommended practice in RFC3704 is to enable strict mode to prevent IP spoofing from DDoS attacks.
- If using asymmetric routing or other complicated routing, then loose mode is recommended.
- The highest value from **conf/{all,interface}/rp_filter** is used when doing source validation on the {interface}

arp_filter

Type	Default
Boolean	0

Value	Effect
0	(default) The kernel can respond to arp requests with addresses from other interfaces. It usually makes sense, because it increases the chance of successful communication.
1	Allows you to have multiple network interfaces on the same subnet, and have the ARPs for each interface be answered based on whether or not the kernel would route a packet from the ARP'd IP out that interface (therefore you must use source based routing for this to work). In other words it allows control of cards (usually 1) that respond to an arp request.

Note

- IP addresses are owned by the complete host on Linux, not by particular interfaces. Only for more complex setups like load-balancing, does this behavior cause problems.
- **arp_filter** for the interface is enabled if at least one of **conf/{all,interface}/arp_filter** is set to TRUE

arp_announce

Define different restriction levels for announcing the local source IP address from IP packets in ARP requests sent on interface

Type	Default
Integer	0

Value	Effect
0	(default) Use any local address, configured on any interface
1	Try to avoid local addresses that are not in the target's subnet for this interface. This mode is useful when target hosts reachable via this interface require the source IP address in ARP requests to be part of their logical network configured on the receiving interface. When we generate the request we check all our subnets that include the target IP and preserve the source address if it is from such subnet. If there is no such subnet we select source address according to the rules for level 2.
2	Always use the best local address for this target. In this mode we ignore the source address in the IP packet and try to select local address that we prefer for talks with the target host. Such local address is selected by looking for primary IP addresses on all our subnets on the outgoing interface that include the target IP address. If no suitable local address is found we select the first local address we have on the outgoing interface or on all other interfaces, with the hope we receive reply for our request and even sometimes no matter the source IP address we announce.

Notes

- The highest value from **conf/{all,interface}/arp_announce** is used.
- Increasing the restriction level gives more chance for receiving answer from the resolved target while decreasing the level announces more valid sender's information.

arp_ignore

Define different modes for sending replies in response to received ARP requests that resolve local target IP addresses

Type	Default
Integer	0

Value	Effect
0	(default): reply for any local target IP address, configured on any interface
1	reply only if the target IP address is local address configured on the incoming interface
2	reply only if the target IP address is local address configured on the incoming interface and both with the sender's IP address are part from same subnet on this interface
3	do not reply for local addresses configured with scope host, only resolutions for global and link addresses are replied
4-7	reserved

Value	Effect
8	do not reply for all local addresses. The max value from conf/{all,interface}/arp_ignore is used when ARP request is received on the {interface}

Notes**arp_notify**

Define mode for notification of address and device changes.

Type	Default
Boolean	0

Value	Effect
0	do nothing
1	Generate gratuitous arp requests when device is brought up or hardware address changes.

Notes**arp_accept**

Define behavior for gratuitous ARP frames who's IP is not already present in the ARP table

Type	Default
Boolean	0

Value	Effect
0	do not create new entries in the ARP table
1	create new entries in the ARP table.

Notes

Both replies and requests type gratuitous arp trigger the ARP table to be updated, if this setting is on. If the ARP table already contains the IP address of the gratuitous arp frame, the arp table is updated regardless if this setting is on or off.

app_solicit

The maximum number of probes to send to the user space ARP daemon via netlink before dropping back to multicast probes (see mcast_solicit).

Type	Default
Integer	0

NotesSee `mcast_solicit`**disable_policy**

Disable IPSEC policy (SPD) for this interface

Type	Default
Boolean	0

needinfo

disable_xfrm

Disable IPSEC encryption on this interface, whatever the policy

Type	Default
Boolean	0

needinfo

igmpv2_unsolicited_report_interval

The interval in milliseconds in which the next unsolicited IGMPv1 or IGMPv2 report retransmit takes place.

Type	Default
Integer	10000

Notes

Milliseconds

igmpv3_unsolicited_report_interval

The interval in milliseconds in which the next unsolicited IGMPv3 report retransmit takes place.

Type	Default
Integer	1000

Notes

Milliseconds

tag

Allows you to write a number, which can be used as required.

Type	Default
Integer	0

xfrm4_gc_thresh

The threshold at which we start garbage collecting for IPv4 destination cache entries.

Type	Default
Integer	1

Notes

At twice this value the system refuses new allocations.

2.3.2. Global kernel tunables

System administrators are able to configure and monitor general settings on a running system through the global kernel tunables.

Global kernel tunables are included in the **/proc/sys/kernel/** directory either directly as named control files or grouped in further subdirectories for various configuration topics. To adjust the global kernel tunables, system administrators need to modify the control files.

Descriptions of the following parameters have been adopted from the kernel documentation sites.^[2]

dmesg_restrict

Indicates whether unprivileged users are prevented from using the **dmesg** command to view messages from the kernel's log buffer.

For further information, see [Kernel sysctl documentation](#).

core_pattern

Specifies a core dumpfile pattern name.

Max length	Default
128 characters	"core"

For further information, see [Kernel sysctl documentation](#).

hardlockup_panic

Controls the kernel panic when a hard lockup is detected.

Type	Value	Effect
Integer	0	kernel does not panic on hard lockup
Integer	1	kernel panics on hard lockup

In order to panic, the system needs to detect a hard lockup first. The detection is controlled by the [nmi_watchdog](#) parameter.

Further Resources

- [Kernel sysctl documentation](#)
- [Softlockup detector and hardlockup detector](#)

softlockup_panic

Controls the kernel panic when a soft lockup is detected.

Type	Value	Effect
Integer	0	kernel does not panic on soft lockup
Integer	1	kernel panics on soft lockup

For more information about **softlockup_panic**, see [kernel_parameters](#).

kptr_restrict

Indicates whether restrictions are placed on exposing kernel addresses via **/proc** and other interfaces.

Type	Default
Integer	0
Value	Effect
0	hashes the kernel address before printing
1	replaces printed kernel pointers with 0's under certain conditions
2	replaces printed kernel pointers with 0's unconditionally

To learn more, see [Kernel sysctl documentation](#).

nmi_watchdog

Controls the hard lockup detector on x86 systems.

Type	Default
Integer	0

Value	Effect
0	disables the hard lockup detector
1	enables the hard lockup detector

The hard lockup detector monitors each CPU for its ability to respond to interrupts.

For more details, see [Kernel sysctl documentation](#).

watchdog_thresh

Controls frequency of **hrtimer**, NMI events, soft/hard lockup thresholds.

Default threshold	Soft lockup threshold
10 seconds	$2 * \text{watchdog_thresh}$

Setting this tunable to zero disables lockup detection altogether.

For more info, consult [Kernel sysctl documentation](#).

panic, panic_on_oops, panic_on_stackoverflow, panic_on_unrecoverable_nmi, panic_on_warn, panic_on_rcu_stall, hung_task_panic

These tunables specify under what circumstances the kernel should panic.

To see more details about a group of **panic** parameters, see [Kernel sysctl documentation](#).

printk, printk_delay, printk_ratelimit, printk_ratelimit_burst, printk_devkmsg

These tunables control logging or printing of kernel error messages.

For more details about a group of **printk** parameters, see [Kernel sysctl documentation](#).

shmall, shmmmax, shm_rmid_forced

These tunables control limits for shared memory.

For more information about a group of **shm** parameters, see [Kernel sysctl documentation](#).

threads-max

Controls the maximum number of threads created by the **fork()** system call.

Min value	Max value
20	Given by FUTEX_TID_MASK (0xffffffff)

The **threads-max** value is checked against the available RAM pages. If the thread structures occupy too much of the available RAM pages, **threads-max** is reduced accordingly.

For more details, see [Kernel sysctl documentation](#).

pid_max

PID allocation wrap value.

To see more information, refer to [Kernel sysctl documentation](#).

numa_balancing

This parameter enables or disables automatic NUMA memory balancing. On NUMA machines, there is a performance penalty if remote memory is accessed by a CPU.

For more details, see [Kernel sysctl documentation](#).

`numa_balancing_scan_period_min_ms`, `numa_balancing_scan_delay_ms`,
`numa_balancing_scan_period_max_ms`, `numa_balancing_scan_size_mb`

These tunables detect if pages are properly placed or if the data should be migrated to a memory node local to where the task is running.

For more details about a group of **numa_balancing_scan** parameters, see [Kernel sysctl documentation](#).

[1] <https://www.kernel.org/doc/Documentation/>

[2] <https://www.kernel.org/doc/Documentation/>

CHAPTER 3. LISTING OF KERNEL PARAMETERS AND VALUES

3.1. KERNEL COMMAND-LINE PARAMETERS

Kernel command-line parameters, also known as kernel arguments, are used to customize the behavior of Red Hat Enterprise Linux at boot time only.

3.1.1. Setting kernel command-line parameters

This section explains how to change a kernel command-line parameter on AMD64 and Intel 64 systems and IBM Power Systems servers using the **GRUB2** boot loader, and on IBM Z using **zipl**.

Kernel command-line parameters are saved in the **boot/grub/grub.cfg** configuration file, which is generated by the **GRUB2** boot loader. Do not edit this configuration file. Changes to this file are only made by configuration scripts.

Changing kernel command-line parameters in GRUB2 for AMD64 and Intel 64 systems and IBM Power Systems Hardware.

1. Open the **/etc/default/grub** configuration file as **root** using a plain text editor such as **vim** or **Gedit**.
2. In this file, locate the line beginning with **GRUB_CMDLINE_LINUX** similar to the following:

```
GRUB_CMDLINE_LINUX="rd.lvm.lv=rhel/swap crashkernel=auto rd.lvm.lv=rhel/root rhgb quiet"
```

3. Change the value of the required kernel command-line parameter. Then, save the file and exit the editor.
4. Regenerate the **GRUB2** configuration using the edited **default** file. If your system uses BIOS firmware, execute the following command:

```
# grub2-mkconfig -o /boot/grub2/grub.cfg
```

On a system with UEFI firmware, execute the following instead:

```
# grub2-mkconfig -o /boot/efi/EFI/redhat/grub.cfg
```

After finishing the procedure above, the boot loader is reconfigured, and the kernel command-line parameter that you have specified in its configuration file is applied after the next reboot.

Changing kernel command-line parameters in zipl for IBM Z Hardware

1. Open the **/etc/zipl.conf** configuration file as **root** using a plain text editor such as **vim** or **Gedit**.
2. In this file, locate the **parameters=** section, and edit the required parameter, or add it if not present. Then, save the file and exit the editor.
3. Regenerate the **zipl** configuration:

```
# zipl
```

**NOTE**

Executing only the **zipl** command with no additional options uses default values. See the **zipl(8)** man page for information about available options.

After finishing the procedure above, the boot loader is reconfigured, and the kernel command-line parameter that you have specified in its configuration file is applied after the next reboot.

3.1.2. What kernel command-line parameters can be controlled

For complete list of kernel command-line parameters, see <https://www.kernel.org/doc/Documentation/admin-guide/kernel-parameters.txt>.

3.1.2.1. Hardware specific kernel command-line parameters

pci=option[,option...]

Specify behavior of the PCI hardware subsystem

Setting	Effect
earlydump	[X86] Dump the PCI configuration space before the kernel changes anything
off	[X86] Do not probe for the PCI bus
noaer	[PCIE] If the PCIEAER kernel parameter is enabled, this kernel boot option can be used to disable the use of PCIE advanced error reporting.
noacpi	[X86] Do not use the Advanced Configuration and Power Interface (ACPI) for Interrupt Request (IRQ) routing or for PCI scanning.
bfsort	Sort PCI devices into breadth-first order. This sorting is done to get a device order compatible with older (≤ 2.4) kernels.
nobfsort	Do not sort PCI devices into breadth-first order.

Additional PCI options are documented in the on disk documentation found in the **kernel-doc-<version>.noarch** package. Where '<version>' needs to be replaced with the corresponding kernel version.

acpi=option

Specify behavior of the Advanced Configuration and Power Interface

Setting	Effect
acpi=off	Disable ACPI

Setting	Effect
acpi=ht	<p>Use ACPI boot table parsing, but do not enable ACPI interpreter</p> <p>This disables any ACPI functionality that is not required for Hyper Threading.</p>
acpi=force	Require the ACPI subsystem to be enabled
acpi=strict	Make the ACPI layer be less tolerant of platforms that are not fully compliant with the ACPI specification.
acpi_sci=<value>	Set up ACPI SCI interrupt, where <value> is one of edge,level,high,low.
acpi=noirq	Do not use ACPI for IRQ routing
acpi=nocmcff	Disable firmware first (FF) mode for corrected errors. This disables parsing the HEST CMC error source to check if firmware has set the FF flag. This can result in duplicate corrected error reports.

CHAPTER 4. KERNEL FEATURES

This chapter explains the purpose and use of kernel features that enable many user space tools and includes resources for further investigation of those tools.

4.1. CONTROL GROUPS

4.1.1. What is a control group?



NOTE

Control Group Namespaces are a Technology Preview in Red Hat Enterprise Linux 7.5

Linux Control Groups (cgroups) enable limits on the use of system hardware, ensuring that an individual process running inside a **cgroup** only utilizes as much as has been allowed in the **cgroups** configuration.

Control Groups restrict the volume of usage on a resource that has been enabled by a **namespace**. For example, the network namespace allows a process to access a particular network card, the cgroup ensures that the process does not exceed 50% usage of that card, ensuring bandwidth is available for other processes.

Control Group Namespaces provide a virtualized view of individual cgroups through the /proc/self/ns/cgroup interface.

The purpose is to prevent leakage of privileged data from the global namespaces to the cgroup and to enable other features, such as container migration.

Because it is now much easier to associate a container with a single cgroup, containers have a much more coherent cgroup view, it also enables tasks inside the container to have a virtualized view of the cgroup it belongs to.

4.1.2. What is a namespace?

Namespaces are a kernel feature that allow a virtual view of isolated system resources. By isolating a process from system resources, you can specify and control what a process is able to interact with. Namespaces are an essential part of Control Groups.

4.1.3. Supported namespaces

The following namespaces are supported from Red Hat Enterprise Linux 7.5 and later

- Mount
 - The mount namespace isolates file system mount points, enabling each process to have a distinct filesystem space within which to operate.
- UTS
 - Hostname and NIS domain name
- IPC
 - System V IPC, POSIX message queues

- PID
 - Process IDs
- Network
 - Network devices, stacks, ports, etc.
- User
 - User and group IDs
- Control Groups
 - Isolates cgroups



NOTE

Usage of Control Groups is documented in the [Resource Management Guide](#)

4.2. KERNEL SOURCE CHECKER

The Linux Kernel Module Source Checker (ksc) is a tool to check for non whitelist symbols in a given kernel module. Red Hat Partners can also use the tool to request review of a symbol for whitelist inclusion, by filing a bug in Red Hat bugzilla database.

4.2.1. Usage

The tool accepts the path to a module with the "-k" option

```
# ksc -k e1000e.ko
Checking against architecture x86_64
Total symbol usage: 165 Total Non white list symbol usage: 74

# ksc -k /path/to/module
```

Output is saved in **\$HOME/ksc-result.txt**. If review of the symbols for whitelist addition is requested, then the usage description for each non-whitelisted symbol must be added to the **ksc-result.txt** file. The request bug can then be filed by running **ksc** with the "-p" option.



NOTE

KSC currently does not support **xz** compression. The **ksc** tool is unable to process the **xz** compression method and reports the following error:

Invalid architecture, (Only kernel object files are supported)

Until this limitation is resolved, system administrators need to manually uncompress any third party modules using **xz** compression, before running the **ksc** tool.

4.3. DIRECT ACCESS FOR FILES (DAX)

Direct Access for files, known as 'file system dax', or 'fs dax', enables applications to read and write data on a dax-capable storage device without using the page cache to buffer access to the device.

This functionality is available when using the 'ext4' or 'xfs' file system, and is enabled either by mounting the file system with **-o dax** or by adding **dax** to the options section for the mount entry in **/etc/fstab**.

Further information, including code examples can be found in the **kernel-doc** package and is stored at **/usr/share/doc/kernel-doc-<version>/Documentation/filesystems/dax.txt** where '<version>' is the corresponding kernel version number.

4.4. MEMORY PROTECTION KEYS FOR USERSPACE (ALSO KNOWN AS PKU, OR PKEYS)

Memory Protection Keys provide a mechanism for enforcing page-based protections, but without requiring modification of the page tables when an application changes protection domains. It works by dedicating 4 previously ignored bits in each page table entry to a "protection key", giving 16 possible keys.

Memory Protection Keys are hardware feature of some Intel CPU chipsets. To determine if your processor supports this feature, check for the presence of **pku** in **/proc/cpuinfo**

```
$ grep pku /proc/cpuinfo
```

To support this feature, the CPUs provide a new user-accessible register (PKRU) with two separate bits (Access Disable and Write Disable) for each key. Two new instructions (RDPKRU and WRPKRU) exist for reading and writing to the new register.

Further documentation, including programming examples can be found in **/usr/share/doc/kernel-doc-*/Documentation/x86/protection-keys.txt** which is provided by the **kernel-doc** package.

4.5. KERNEL ADDRESS SPACE LAYOUT RANDOMIZATION

Kernel Address Space Layout Randomization (KASLR) consists of two parts which work together to enhance the security of the Linux kernel:

- kernel text KASLR
- memory management KASLR

The physical address and virtual address of kernel text itself are randomized to a different position separately. The physical address of the kernel can be anywhere under 64TB, while the virtual address of the kernel is restricted between [0xffffffff80000000, 0xffffffffc0000000], the 1GB space.

Memory management KASLR has three sections whose starting address is randomized in a specific area. KASLR can thus prevent inserting and redirecting the execution of the kernel to a malicious code if this code relies on knowing where symbols of interest are located in the kernel address space.

Memory management KASLR sections are:

- direct mapping section
- vmalloc section
- vmemmap section

KASLR code is now compiled into the Linux kernel, and it is enabled by default. To disable it explicitly, add the **nokaslr** kernel option to the kernel command line.

4.6. ADVANCED ERROR REPORTING (AER)

Define what AER is and what it is used for.

Describe the right way to enable AER, mention parameters "pcie_ports=native" and "acpi=nocmcff". Explain under which circumstances they should be used and what are possible "side effects".

CHAPTER 5. MANUALLY UPGRADING THE KERNEL

The Red Hat Enterprise Linux kernel is custom-built by the Red Hat Enterprise Linux kernel team to ensure its integrity and compatibility with supported hardware. Before Red Hat releases a kernel, it must first pass a rigorous set of quality assurance tests.

Red Hat Enterprise Linux kernels are packaged in the RPM format so that they are easy to upgrade and verify using the **Yum** or **PackageKit** package managers. **PackageKit** automatically queries the Red Hat Content Delivery Network servers and informs you of packages with available updates, including kernel packages.

This chapter is therefore **only** useful for users who need to manually update a kernel package using the **rpm** command instead of **yum**.



WARNING

Whenever possible, use either the **Yum** or **PackageKit** package manager to install a new kernel because they always **install** a new kernel instead of replacing the current one, which could potentially leave your system unable to boot.



WARNING

Custom kernels are **not** supported by Red Hat. However, guidance can be obtained from the knowledgebase document <https://access.redhat.com/solutions/25039>.

For more information on installing kernel packages with **yum**, see
https://access.redhat.com/documentation/en-US/Red_Hat_Enterprise_Linux/7/html/System_Administrators_Guide/ch-yum.html#sec-Updating_Packages.

For information on Red Hat Content Delivery Network, see
https://access.redhat.com/documentation/en-US/Red_Hat_Enterprise_Linux/7/html/System_Administrators_Guide/chap-Subscription_and_Support-Registering_a_System_and_Managing_Subscriptions.html.

5.1. OVERVIEW OF KERNEL PACKAGES

Red Hat Enterprise Linux contains the following kernel packages:

- **kernel** – Contains the kernel for single-core, multi-core, and multi-processor systems.
- **kernel-debug** – Contains a kernel with numerous debugging options enabled for kernel diagnosis, at the expense of reduced performance.

- **kernel-devel** – Contains the kernel headers and makefiles sufficient to build modules against the **kernel** package.
- **kernel-debug-devel** – Contains the development version of the kernel with numerous debugging options enabled for kernel diagnosis, at the expense of reduced performance.
- **kernel-doc** – Documentation files from the kernel source. Various portions of the Linux kernel and the device drivers shipped with it are documented in these files. Installation of this package provides a reference to the options that can be passed to Linux kernel modules at load time. By default, these files are placed in the **/usr/share/doc/kernel-doc-kernel_version/** directory.
- **kernel-headers** – Includes the C header files that specify the interface between the Linux kernel and user-space libraries and programs. The header files define structures and constants that are needed for building most standard programs.
- **linux-firmware** – Contains all of the firmware files that are required by various devices to operate.
- **perf** – This package contains the **perf** tool, which enables performance monitoring of the Linux kernel.
- **kernel-abi-whitelists** – Contains information pertaining to the Red Hat Enterprise Linux kernel ABI, including a lists of kernel symbols that are needed by external Linux kernel modules and a **yum** plug-in to aid enforcement.
- **kernel-tools** – Contains tools for manipulating the Linux kernel and supporting documentation.

5.2. PREPARING TO UPGRADE

Before upgrading the kernel, it is recommended that you take some precautionary steps.

First, ensure that working boot media exists for the system in case a problem occurs. If the boot loader is not configured properly to boot the new kernel, you can use this media to boot into Red Hat Enterprise Linux

USB media often comes in the form of flash devices sometimes called *pen drives*, *thumb disks*, or *keys*, or as an externally-connected hard disk device. Almost all media of this type is formatted as a **VFAT** file system. You can create bootable USB media on media formatted as **ext2**, **ext3**, **ext4**, or **VFAT**.

You can transfer a distribution image file or a minimal boot media image file to USB media. Make sure that sufficient free space is available on the device. Around 4 GB is required for a distribution DVD image, around 700 MB for a distribution CD image, or around 10 MB for a minimal boot media image.

You must have a copy of the **boot.iso** file from a Red Hat Enterprise Linux installation DVD, or installation CD-ROM #1, and you need a USB storage device formatted with the **VFAT** file system and around 16 MB of free space.

For more information on using USB storage devices, review
[https://access.redhat.com/solutions/624423 How to format a USB key](https://access.redhat.com/solutions/624423) and
[https://access.redhat.com/solutions/39373 How to manually mount a USB flash drive in a non-graphical environment](https://access.redhat.com/solutions/39373).

The following procedure does not affect existing files on the USB storage device unless they have the same path names as the files that you copy onto it. To create USB boot media, perform the following commands as the **root** user:

1. Install the **syslinux** package if it is not installed on your system. To do so, as root, run the **yum install syslinux** command.
2. Install the **SYSLINUX** bootloader on the USB storage device:

```
# syslinux /dev/sdX1
```

...where *sdX* is the device name.

3. Create mount points for **boot.iso** and the USB storage device:

```
# mkdir /mnt/isoboot /mnt/diskboot
```

4. Mount **boot.iso**:

```
# mount -o loop boot.iso /mnt/isoboot
```

5. Mount the USB storage device:

```
# mount /dev/sdX1 /mnt/diskboot
```

6. Copy the **ISOLINUX** files from the **boot.iso** to the USB storage device:

```
# cp /mnt/isoboot/isolinux/* /mnt/diskboot
```

7. Use the **isolinux.cfg** file from **boot.iso** as the **syslinux.cfg** file for the USB device:

```
# grep -v local /mnt/isoboot/isolinux/isolinux.cfg > /mnt/diskboot/syslinux.cfg
```

8. Unmount **boot.iso** and the USB storage device:

```
# umount /mnt/isoboot /mnt/diskboot
```

9. Reboot the machine with the boot media and verify that you are able to boot with it before continuing.

Alternatively, on systems with a floppy drive, you can create a boot diskette by installing the **mkbootdisk** package and running the **mkbootdisk** command as **root**. See **man mkbootdisk** man page after installing the package for usage information.

To determine which kernel packages are installed, execute the command **yum list installed "kernel-*"** at a shell prompt. The output comprises some or all of the following packages, depending on the system's architecture, and the version numbers might differ:

```
# yum list installed "kernel-*"
kernel.x86_64      3.10.0-54.0.1.el7      @rhel7/7.0
kernel-devel.x86_64 3.10.0-54.0.1.el7      @rhel7
kernel-headers.x86_64 3.10.0-54.0.1.el7    @rhel7/7.0
```

From the output, determine which packages need to be downloaded for the kernel upgrade. For a single processor system, the only required package is the **kernel** package. See [Section 5.1, “Overview of kernel packages”](#) for descriptions of the different packages.

5.3. DOWNLOADING THE UPGRADED KERNEL

There are several ways to determine if an updated kernel is available for the system.

- Security Errata – See <https://access.redhat.com/site/security/updates/active/> for information on security errata, including kernel upgrades that fix security issues.
- The Red Hat Content Delivery Network – For a system subscribed to the Red Hat Content Delivery Network, the **yum** package manager can download the latest kernel and upgrade the kernel on the system. The **Dracut** utility creates an initial RAM file system image if needed, and configure the boot loader to boot the new kernel. For more information on installing packages from the Red Hat Content Delivery Network, see https://access.redhat.com/documentation/en-US/Red_Hat_Enterprise_Linux/7/html/System_Administrators_Guide/ch-yum.html. For more information on subscribing a system to the Red Hat Content Delivery Network, see https://access.redhat.com/documentation/en-US/Red_Hat_Enterprise_Linux/7/html/System_Administrators_Guide/chap-Subscription_and_Support-Registering_a_System_and_Managing_Subscriptions.html.

If **yum** was used to download and install the updated kernel from the Red Hat Network, follow the instructions in [Section 5.5, “Verifying the initial RAM file system image”](#) and [Section 5.6, “Verifying the boot loader”](#) only, **do not** change the kernel to boot by default. Red Hat Network automatically changes the default kernel to the latest version. To install the kernel manually, continue to [Section 5.4, “Performing the upgrade”](#).

5.4. PERFORMING THE UPGRADE

After retrieving all of the necessary packages, it is time to upgrade the existing kernel.



IMPORTANT

It is strongly recommended that you keep the old kernel in case there are problems with the new kernel.

At a shell prompt, change to the directory that contains the kernel RPM packages. Use **-i** argument with the **rpm** command to keep the old kernel. Do **not** use the **-U** option, since it overwrites the currently installed kernel, which creates boot loader problems. For example:

```
# rpm -vh kernel-kernel_version.arch.rpm
```

The next step is to verify that the initial RAM file system image has been created. See [Section 5.5, “Verifying the initial RAM file system image”](#) for details.

5.5. VERIFYING THE INITIAL RAM FILE SYSTEM IMAGE

The job of the initial RAM file system image is to preload the block device modules, such as for IDE, SCSI or RAID, so that the root file system, on which those modules normally reside, can then be accessed and mounted. On Red Hat Enterprise Linux 7 systems, whenever a new kernel is installed using either the **Yum**, **PackageKit**, or **RPM** package manager, the **Dracut** utility is always called by the installation scripts to create an *initramfs* (initial RAM file system image).

If you make changes to the kernel attributes by modifying the **/etc/sysctl.conf** file or another **sysctl** configuration file, and if the changed settings are used early in the boot process, then rebuilding the Initial RAM File System Image by running the **dracut -f** command might be necessary. An example is if

you have made changes related to networking and are booting from network-attached storage.

On all architectures other than IBM eServer System i (see [the section called “Verifying the initial RAM file system image and kernel on IBM eServer System i”](#)), you can create an **initramfs** by running the **dracut** command. However, you usually do not need to create an **initramfs** manually: this step is automatically performed if the kernel and its associated packages are installed or upgraded from RPM packages distributed by Red Hat.

You can verify that an **initramfs** corresponding to your current kernel version exists and is specified correctly in the **grub.cfg** configuration file by following this procedure:

Verifying the initial RAM file system image

- As **root**, list the contents in the **/boot** directory and find the kernel (**vmlinuz-kernel_version**) and **initramfs-kernel_version** with the latest (most recent) version number:

Example 5.1. Ensuring that the kernel and initramfs versions match

```
# ls /boot
config-3.10.0-67.el7.x86_64
config-3.10.0-78.el7.x86_64
efi
grub
grub2
initramfs-0-rescue-07f43f20a54c4ce8ada8b70d33fd001c.img
initramfs-3.10.0-67.el7.x86_64.img
initramfs-3.10.0-67.el7.x86_64kdump.img
initramfs-3.10.0-78.el7.x86_64.img
initramfs-3.10.0-78.el7.x86_64kdump.img
initrd-plymouth.img
symvers-3.10.0-67.el7.x86_64.gz
symvers-3.10.0-78.el7.x86_64.gz
System.map-3.10.0-67.el7.x86_64
System.map-3.10.0-78.el7.x86_64
vmlinuz-0-rescue-07f43f20a54c4ce8ada8b70d33fd001c
vmlinuz-3.10.0-67.el7.x86_64
vmlinuz-3.10.0-78.el7.x86_64
```

[Example 5.1, “Ensuring that the kernel and initramfs versions match”](#) shows that:

- we have three kernels installed (or, more correctly, three kernel files are present in the **/boot** directory),
- the latest kernel is **vmlinuz-3.10.0-78.el7.x86_64**, and
- an **initramfs** file matching our kernel version, **initramfs-3.10.0-78.el7.x86_64kdump.img**, also exists.



IMPORTANT

In the **/boot** directory you might find several **initramfs-kernel_versionkdump.img** files. These are special files created by the **Kdump** mechanism for kernel debugging purposes, are not used to boot the system, and can safely be ignored. For more information on **kdump**, see the [Red Hat Enterprise Linux 7 Kernel Crash Dump Guide](#).

- If your **initramfs-kernel_version** file does not match the version of the latest kernel in the **/boot** directory, or, in certain other situations, you might need to generate an **initramfs** file with the **Dracut** utility. Simply invoking **dracut** as **root** without options causes it to generate an **initramfs** file in **/boot** for the latest kernel present in that directory:

```
# dracut
```

You must use the **-f, --force** option if you want **dracut** to overwrite an existing **initramfs** (for example, if your **initramfs** has become corrupt). Otherwise **dracut** refuses to overwrite the existing **initramfs** file:

```
# dracut
```

Does not override existing initramfs (/boot/initramfs-3.10.0-78.el7.x86_64.img) without -force

You can create an initramfs in the current directory by calling **dracut** **initramfs_name kernel_version**:

```
# dracut "initramfs-$(uname -r).img" $(uname -r)
```

If you need to specify specific kernel modules to be preloaded, add the names of those modules (minus any file name suffixes such as **.ko**) inside the parentheses of the **add_dracutmodules+="module more_modules"** directive of the **/etc/dracut.conf** configuration file. You can list the file contents of an **initramfs** image file created by dracut by using the **lsinitrd initramfs_file** command:

```
# lsinitrd /boot/initramfs-3.10.0-78.el7.x86_64.img
Image: /boot/initramfs-3.10.0-78.el7.x86_64.img: 11M
=====
```

```
dracut-033-68.el7
=====
```

```
drwxr-xr-x 12 root  root  0 Feb 5 06:35 .
drwxr-xr-x  2 root  root  0 Feb 5 06:35 proc
lrwxrwxrwx  1 root  root  24 Feb 5 06:35 init -> /usr/lib/systemd/systemd
drwxr-xr-x 10 root  root  0 Feb 5 06:35 etc
drwxr-xr-x  2 root  root  0 Feb 5 06:35 usr/lib/modprobe.d
[output truncated]
```

See **man dracut** and **man dracut.conf** for more information on options and usage.

- Examine the **/boot/grub2/grub.cfg** configuration file to ensure that an **initramfs-kernel_version.img** file exists for the kernel version you are booting. For example:

```
# grep initramfs /boot/grub2/grub.cfg
initrd16 /initramfs-3.10.0-123.el7.x86_64.img
initrd16 /initramfs-0-rescue-6d547dbfd01c46f6a4c1baa8c4743f57.img
```

See [Section 5.6, “Verifying the boot loader”](#) for more information.

Verifying the initial RAM file system image and kernel on IBM eServer System i

On IBM eServer System i machines, the initial RAM file system and kernel files are combined into a single

file, which is created with the **addRamDisk** command. This step is performed automatically if the kernel and its associated packages are installed or upgraded from the RPM packages distributed by Red Hat thus, it does not need to be executed manually. To verify that it was created, run the following command as **root** to make sure the **/boot/vmlinird-kernel_version** file already exists:

```
# ls -l /boot/
```

The *kernel_version* needs to match the version of the kernel just installed.

Reversing the changes made to the initial RAM file system image

In some cases, for example, if you misconfigure the system and it no longer boots, you need to reverse the changes made to the Initial RAM File System Image by following this procedure:

Reversing Changes Made to the Initial RAM File System Image

1. Reboot the system choosing the rescue kernel in the GRUB menu.
2. Change the incorrect setting that caused the **initramfs** to malfunction.
3. Recreate the **initramfs** with the correct settings by running the following command as root:

```
# dracut --kver kernel_version --force
```

The above procedure might be useful if, for example, you incorrectly set the **vm.nr_hugepages** in the **sysctl.conf** file. Because the **sysctl.conf** file is included in **initramfs**, the new **vm.nr_hugepages** setting gets applied in **initramfs** and causes rebuilding of the **initramfs**. However, because the setting is incorrect, the new **initramfs** is broken and the newly built kernel does not boot, which necessitates correcting the setting using the above procedure.

Listing the contents of the initial RAM file system image

To list the files that are included in the **initramfs**, run the following command as root:

```
# lsinitrd
```

To only list files in the **/etc** directory, use the following command:

```
# lsinitrd | grep etc/
```

To output the contents of a specific file stored in the **initramfs** for the current kernel, use the **-f** option:

```
# lsinitrd -f filename
```

For example, to output the contents of **sysctl.conf**, use the following command:

```
# lsinitrd -f /etc/sysctl.conf
```

To specify a kernel version, use the **--kver** option:

```
# lsinitrd --kver kernel_version -f /etc/sysctl.conf
```

For example, to list the information about kernel version 3.10.0-327.10.1.el7.x86_64, use the following command:

```
# lsinitrd --kver 3.10.0-327.10.1.el7.x86_64 -f /etc/sysctl.conf
```

5.6. VERIFYING THE BOOT LOADER

You can install a kernel either with the **yum** command or with the **rpm** command.

When you install a kernel using **rpm**, the kernel package creates an entry in the boot loader configuration file for that new kernel.

Note that both commands configure the new kernel to boot as the default kernel only when you include the following setting in the **/etc/sysconfig/kernel** configuration file:

```
DEFAULTKERNEL=kernel  
UPDATEDefault=yes
```

The **DEFAULTKERNEL** option specifies the default kernel package type. The **UPDATEDefault** option specifies whether the new kernel package makes the new kernels the default.

CHAPTER 6. APPLYING PATCHES WITH KERNEL LIVE PATCHING

The Red Hat Enterprise Linux kernel live patching solution allows you to patch a running kernel without rebooting or restarting any processes. The solution enables system administrators to apply critical security patches to the kernel immediately, without having to wait for long-running tasks to complete, for users to log off, or for scheduled downtime. It gives more control over uptime without sacrificing security or stability.

Note that not every critical or important CVE will be resolved using the kernel live patching solution. Our goal is to reduce the required reboots for security-related patches, not to eliminate them entirely.



WARNING

Some incompatibilities exist between kernel live patching and other kernel subcomponents. Read the [Section 6.9, “Limitations of kpatch”](#) section carefully before using kernel live patching.

6.1. ACCESS TO KERNEL LIVE PATCHES

Kernel live patching capability is implemented as a kernel module (**.ko** file) that is delivered as an RPM package.

All customers have access to kernel live patches, which are delivered through the usual channels. However, customers who do not subscribe to an extended support offering will lose access to new patches for the current minor release once the next minor release becomes available. For example, customers with standard subscriptions will only be able to live patch the RHEL 8.2 kernels until RHEL 8.3 is released.

6.2. COMPONENTS OF KERNEL LIVE PATCHING

The components of kernel live patching are as follows:

Kernel patch module

- The delivery mechanism for kernel live patches.
- This is a kernel module which is built specifically for the kernel being patched.
- The patch module contains the code of the desired fixes for the kernel.
- The patch modules register with the **livepatch** kernel subsystem and provide information about original functions to be replaced, with corresponding pointers to the replacement functions. Kernel patch modules are delivered in **kpatch-patch-<kernel-version>** RPMs.

The **kpatch** utility

A command-line tool which allows you to manage patch modules.

The **kpatch** service


A **systemd** service required by **multiuser.target**. This target loads the kernel patch module at boot time.

6.3. HOW KERNEL LIVE PATCHING WORKS

The **kpatch** kernel patching solution uses the **livepatch** kernel subsystem to redirect old functions to new ones. When a live kernel patch is applied to a system, the following things happen:

1. The kernel patch module is copied to the **/var/lib/kpatch/** directory and registered for re-application to the kernel via **systemd** on next boot.
2. The kpatch module is loaded into the running kernel and the patched functions are registered to the **ftrace** mechanism with a pointer to the location in memory of the new code.
3. When the kernel accesses the patched function, it is redirected by the **ftrace** mechanism which bypasses the original functions and redirects the kernel to patched version of the function.

Figure 6.1. How kernel live patching works



RHEL_424549_0119

6.4. ENABLING KERNEL LIVE PATCHING

A kernel patch module is delivered in an RPM package, specific to the version of the kernel being patched. Each RPM package will be cumulatively updated over time.

The following subsections describe how to ensure you receive all future cumulative live patching updates for a given kernel.



WARNING

Red Hat strongly discourages from using any patch modules which were not provided for your specific usage by Red Hat itself.

6.4.1. Subscribing to the live patching stream

This procedure describes installing a particular live patching package. By doing so you subscribe to the live patching stream for a given kernel and ensure that you receive all future cumulative live patching updates for that kernel.

Prerequisites

- Root permissions
- You are familiar with [basic package administration](#) in Red Hat Enterprise Linux 7.

Procedure

1. Optionally, check your kernel version:

```
$ uname -r
3.10.0-1062.el7.x86_64
```

2. Search for a live patching package that corresponds to the version of your kernel:

```
# yum search $(uname -r)
```

The live patching package contains a patch module, if the package's version is 1-1 or higher. In that case the kernel will be automatically patched during the installation of the live patching package.

3. Install the live patching package:

```
# yum install "kpatch-patch = $(uname -r)"
```

The example command above installs and applies the latest cumulative live patches for that specific kernel only. Later, if you install a new kernel and reboot into it, you will have to re-enable the live patching for the new kernel by running **yum install kpatch-patch-<new_kernel_version>**.

The kernel patch module is also installed into the **/var/lib/kpatch/** directory to be loaded by the **systemd** system and service manager during the future reboots.



NOTE

If there are not yet any live patches available for the given kernel, an empty live patching package will be installed. An empty live patching package will have a Version-Release of 0-0, for example **kpatch-patch-3_10_0-1062-0-0.el7.x86_64.rpm**. The installation of the empty RPM subscribes the system to all future live patches for the given kernel.

4. Optionally, verify the kernel is patched:

```
$ kpatch list
Loaded patch modules:
kpatch_3_10_0_1062_1_1 [enabled]
```

Installed patch modules:
 kpatch_3_10_0_1062_1_1 (3.10.0-1062.el7.x86_64)
 ...

The output shows that the kernel patch module has been loaded into the kernel, which is now patched with the latest fixes from the **kpatch-patch-3_10_0-1062-1-1.el7.x86_64.rpm** package.

Additional resources

- For more information about the **kpatch** command-line utility, refer to the **[kpatch\(1\)](#)** manual page.

6.5. UPDATING KERNEL PATCH MODULES

Since kernel patch modules are delivered and applied through RPM packages, updating a cumulative kernel patch module is like updating any other RPM package.

Prerequisites

- Root permissions
- You are familiar with [Section 6.4.1, “Subscribing to the live patching stream”](#) and are subscribed to the live patching stream.
- You obtained an updated live patching package.
- You are familiar with [updating packages](#) in Red Hat Enterprise Linux 7.

Procedure

- Run **yum update** to update to a new cumulative version.
 The command above automatically installs and applies any future released cumulative live patches to the running kernel.



NOTE

When the system reboots into the same kernel, the kernel is automatically live patched again by the **kpatch.service** systemd unit file.

6.6. DISABLING KERNEL LIVE PATCHING

In case system administrators encountered some unanticipated negative effects connected with the Red Hat Enterprise Linux kernel live patching solution and they no longer wish to immediately apply critical security patches to their running kernel, they can disable the mechanism in three distinct ways:

6.6.1. Removing the live patching package

The following procedure describes how to disable the Red Hat Enterprise Linux kernel live patching solution by removing the live patching package that you installed in your system.

Prerequisites

- Root permissions

- You are familiar with [basic package administration](#) in Red Hat Enterprise Linux 7.
- The live patching package is installed.

Procedure

1. Select the live patching package:

```
$ yum list installed | grep kpatch-patch
kpatch-patch-3_10_0-1062.x86_64      1-1.el7      @@commandline
...
...
```

The example output above lists live patching packages that you installed.

2. Remove the live patching package:

```
# yum remove kpatch-patch-3_10_0-1062.x86_64
Updating Subscription Management repositories.
Dependencies resolved.
=====
Package           Arch    Version     Repository   Size
=====
Removing:
kpatch-patch-3_10_0-1062  x86_64  1-1.el7      @@commandline  14 k
Transaction Summary
=====
Remove 1 Package

Freed space: 14 k
Is this ok [y/N]: y
Running transaction check
Transaction check succeeded.
Running transaction test
Transaction test succeeded.
Running transaction
Preparing : 1/1
Erasing   : kpatch-patch-3_10_0-1062-1-1.el7.x86_64      1/1
Running scriptlet: kpatch-patch-3_10_0-1062-1-1.el7.x86_64      1/1
uninstalling kpatch-3_10_0-1062-1-1.ko (3.10.0-1062.el7.x86_64)

Verifying   : kpatch-patch-3_10_0-1062-1-1.el7.x86_64      1/1
Installed products updated.

Removed:
kpatch-patch-3_10_0-1062-1-1.el7.x86_64

Complete!
```

When a live patching package is removed, the kernel remains patched until the next reboot, but the kernel patch module is removed from disk. After the next reboot, the corresponding kernel will no longer be patched.

3. Reboot your system.
4. Verify that the live patching package has been removed:

```
$ yum list installed | grep kpatch-patch  
<NO_RESULT>
```

5. Verify that the kernel live patching solution is disabled:

```
$ kpatch list  
Loaded patch modules:  
<NO_RESULT>
```

The example output above shows that the kernel is not patched and the kernel live patching solution is not active.

Additional resources

- For more information about the **kpatch** command-line utility, refer to the **[kpatch\(1\)](#)** manual page.

6.6.2. Uninstalling the kernel patch module

The following procedure describes how to prevent the Red Hat Enterprise Linux kernel live patching solution from applying a kernel patch module on subsequent boots.

Prerequisites

- Root permissions
- A live patching package is installed.
- A kernel patch module is installed and loaded.

Procedure

1. Select a kernel patch module:

```
$ kpatch list  
Loaded patch modules:  
kpatch_3_10_0_1062_1_1 [enabled]  
  
Installed patch modules:  
kpatch_3_10_0_1062_1_1 (3.10.0-1062.el7.x86_64)  
...
```

2. Uninstall the selected kernel patch module:

```
# kpatch uninstall kpatch_3_10_0_1062_1_1  
uninstalling kpatch_3_10_0_1062_1_1 (3.10.0-1062.el7.x86_64)
```

- Note that the uninstalled kernel patch module is still loaded:

```
$ kpatch list  
Loaded patch modules:
```

```
kpatch_3_10_0_1062_1_1 [enabled]
```

```
Installed patch modules:  
<NO_RESULT>
```

When the selected module is uninstalled, kernel remains patched until the next reboot, but the kernel patch module is removed from disk.

3. Reboot your system.
4. Verify that the kernel patch module has been uninstalled:

```
$ kpatch list  
Loaded patch modules:  
<NO_RESULT>
```

The example output above shows no loaded or installed kernel patch modules, therefore the kernel is not patched and the kernel live patching solution is not active.

Additional resources

- For more information about the **kpatch** command-line utility, refer to the **[kpatch\(1\)](#)** manual page.

6.6.3. Disabling the kpatch service unit file

The following procedure describes how to prevent the Red Hat Enterprise Linux kernel live patching solution from applying all kernel patch modules globally on subsequent boots.

Prerequisites

- Root permissions
- A live patching package is installed.
- A kernel patch module is installed and loaded.

Procedure

1. Verify **kpatch.service** is enabled and running:

```
# systemctl status kpatch.service
● kpatch.service - "Apply kpatch kernel patches"
  Loaded: loaded (/usr/lib/systemd/system/kpatch.service; enabled; vendor preset: disabled)
  Active: active (exited) since Sat 2019-07-27 13:46:07 CEST; 34min ago
    Process: 831 ExecStart=/usr/sbin/kpatch load --all (code=exited, status=0/SUCCESS)
   Main PID: 831 (code=exited, status=0/SUCCESS)

Jul 27 13:46:07 redhat systemd[1]: Starting "Apply kpatch kernel patches"...
Jul 27 13:46:07 redhat kpatch[831]: loading patch module: /var/lib/kpatch/3.10.0-
1062.el7.x86_64/kpatch-3_10_0-1062-1-1.ko
Jul 27 13:46:07 redhat systemd[1]: Started "Apply kpatch kernel patches".
```

2. Disable **kpatch.service**:

```
# systemctl disable kpatch.service
Removed /etc/systemd/system/multi-user.target.wants/kpatch.service.
```

- Note that the applied kernel patch module is still loaded:

```
$ kpatch list
Loaded patch modules:
kpatch_3_10_0_1062_1_1 [enabled]

Installed patch modules:
kpatch_3_10_0_1062_1_1 (3.10.0-1062.el7.x86_64)
```

3. Reboot your system.
4. Optionally, verify the status of **kpatch.service**:

```
# systemctl status kpatch.service
● kpatch.service - "Apply kpatch kernel patches"
  Loaded: loaded (/usr/lib/systemd/system/kpatch.service; disabled; vendor preset: disabled)
  Active: inactive (dead)
```

The example output testifies that **kpatch.service** has been disabled and is not running. Thereby the kernel live patching solution is not active.

5. Verify that the kernel patch module has been unloaded:

```
# kpatch list
Loaded patch modules:
<NO_RESULT>

Installed patch modules:
kpatch_3_10_0_1062_1_1 (3.10.0-1062.el7.x86_64)
```

The example output above shows that a kernel patch module is still installed but the kernel is not patched.

Additional resources

- For more information about the **kpatch** command-line utility, refer to the **[kpatch\(1\)](#)** manual page.
- For more information about the **systemd** system and service manager, unit configuration files, their locations, as well as a complete list of **systemd** unit types, see the relevant sections in [System Administrator's Guide](#).

6.7. KPATCH SUPPORT

- Live kernel patches are available for all RHEL subscriptions and on AMD64 and Intel 64 and IBM Power Systems architectures.
- Not all CVE remediation patches that are kernel related may be able to be live patched
- Red Hat will generate, maintain and support live kernel patches for kernels that are up to 1 year old.

- Live kernel patches will be produced and supported for RHEL Releases throughout their EUS coverage period.
- Live patches for RHEL kernels will be produced and supported for RHEL for SAP Solutions Releases throughout their Update Services (E4S) coverage period.
- Live kernel patches will be for the following RHEL minor releases with current or planned EUS streams: RHEL 7.6, 7.7 RHEL 8.1, 8.2 ,8.4.

6.8. SUPPORT FOR THIRD-PARTY LIVE PATCHING

The **kpatch** utility is the only kernel live patching utility supported by Red Hat with the RPM modules supplied through your Red Hat support contract. Red Hat will not support any live patches which were not provided by Red Hat itself.

If you require support for an issue that arises with a third-party live patch, Red Hat recommends that you open a case with the live patching vendor at the outset of any investigation in which a root cause determination is necessary. This allows the source code to be supplied if the vendor allows, and for their support organization to provide assistance in root cause determination prior to escalating the investigation to Red Hat Support.

For any system running with third-party live patches, Red Hat reserves the right to ask for reproduction with Red Hat shipped and supported software. In the event that this is not possible, we require a similar system and workload be deployed on your test environment without live patches applied, to confirm if the same behavior is observed.

For more information about third-party software support policies, see [How does Red Hat Global Support Services handle third-party software, drivers, and/or uncertified hardware/hypervisors or guest operating systems?](#)

6.9. LIMITATIONS OF KPATCH

- **kpatch** is not a general-purpose kernel upgrade mechanism. It is used for applying simple security and bug fix updates when rebooting the system is not immediately possible.
- Do not use the **SystemTap** or **kprobe** tools during or after loading a patch. The patch could fail to take effect until after such probes have been removed.

CHAPTER 7. KERNEL CRASH DUMP GUIDE

7.1. INTRODUCTION TO KDUMP

7.1.1. About kdump and kexec

Kdump is a kernel crash dumping mechanism that allows you to save the contents of the system's memory for later analysis. It relies on **kexec**, which can be used to boot a Linux kernel from the context of another kernel, bypass BIOS, and preserve the contents of the first kernel's memory that would otherwise be lost.

In case of a system crash, kdump uses kexec to boot into a second kernel (a *capture kernel*). This second kernel resides in a reserved part of the system memory that is inaccessible to the first kernel. The second kernel then captures the contents of the crashed kernel's memory (a *crash dump*) and saves it.



IMPORTANT

A kernel crash dump can be the only information available in the event of a failure, the importance of having this data in a business critical environment cannot be underestimated. Red Hat advise that System Administrators regularly update and test **kexec-tools** in your normal kernel update cycle. This is especially important when new kernel features are implemented.

7.1.2. Memory requirements

In order for kdump to be able to capture a kernel crash dump and save it for further analysis, a part of the system memory has to be permanently reserved for the capture kernel. When reserved, this part of the system memory is not available to main kernel.

The memory requirements vary based on certain system parameters. One of the major factors is the system's hardware architecture. To find out the exact name of the machine architecture (such as **x86_64**) and print it to standard output, type the following command at a shell prompt:

```
uname -m
```

Another factor which influences the amount of memory to be reserved is the total amount of installed system memory. For example, on the **x86_64** architecture, the amount of reserved memory is 160 MB + 2 bits for every 4 KB of RAM. On a system with 1 TB of total physical memory installed, this means 224 MB (160 MB + 64 MB). For a complete list of memory requirements for kdump based on the system architecture and the amount of physical memory, see [Section 7.7.1, "Memory requirements for kdump"](#).

On many systems, kdump can estimate the amount of required memory and reserve it automatically. This behavior is enabled by default, but only works on systems that have more than a certain amount of total available memory, which varies based on the system architecture. See [Section 7.7.2, "Minimum threshold for automatic memory reservation"](#) for a list of minimum requirements for automatic memory reservation based on the system architecture.

If the system has less than the minimum amount of memory required for the automatic allocation to work or if your use case requires a different value, you can configure the amount of reserved memory manually. For information on how to do so on the command line, see [Section 7.2.2.1, "Configuring the memory usage"](#). For information on how to configure the amount of reserved memory in the graphical user interface, see [Section 7.2.3.1, "Configuring the memory usage"](#).



IMPORTANT

It is highly recommended to test the configuration after setting up the kdump service, even when using the automatic memory reservation. For instructions on how to test your configuration, see [Section 7.3, “Testing the kdump configuration”](#).

7.2. INSTALLING AND CONFIGURING KDUMP

7.2.1. Installing kdump

In many cases, the **kdump** service is installed and activated by default on new Red Hat Enterprise Linux 7 installations. The **Anaconda** installer provides a screen for kdump configuration when performing an interactive installation using the graphical or text interface. The installer screen is titled **Kdump** and is available from the main **Installation Summary** screen, and only allows limited configuration – you can only select whether kdump is enabled and how much memory is reserved. Information about memory requirements for kdump is available in [Section 7.7.1, “Memory requirements for kdump”](#). The Kdump configuration screen in the installer is documented in the [*Red Hat Enterprise Linux 7 Installation Guide*](#).




NOTE

In previous releases of Red Hat Enterprise Linux, kdump configuration was available in the **Firstboot** utility which was automatically executed **after** the installation finished and the system rebooted for the first time. Starting with Red Hat Enterprise Linux 7.1, kdump configuration has been moved into the installer.

Some installation options, such as custom Kickstart installations, do not have to install or enable kdump by default. If this is the case on your system, and you want to install kdump additionally, execute the following command as **root** at a shell prompt:

```
# yum install kexec-tools
```

The command above secures installation of kdump and all other necessary packages, assuming your system has an active subscription or a custom repository containing the **kexec-tools** package for your system’s architecture.



NOTE

If you do not know whether kdump is installed on your system, you can check using **rpm**:

```
$ rpm -q kexec-tools
```

Additionally, a graphical configuration tool is available, but not installed by default if you use the command described above. To install this utility, which is described in [Section 7.2.3, “Configuring kdump in the graphical user interface”](#), use the following command as **root**:

```
# yum install system-config-kdump
```

For more information on how to install new packages in Red Hat Enterprise Linux 7 using the **Yum** package manager, see the [*Red Hat Enterprise Linux 7 System Administrator’s Guide*](#).



IMPORTANT

Starting with Red Hat Enterprise Linux 7.4 the **Intel IOMMU** driver is supported with **kdump**. When running kernels from version 7.3 or earlier, it is advised that **Intel IOMMU** support is disabled.

7.2.2. Configuring kdump on the command line

7.2.2.1. Configuring the memory usage

Memory reserved for the kdump kernel is always reserved during system boot, which means that the amount of memory is specified in the system's boot loader configuration.

To specify the memory reserved for kdump kernel, set the **crashkernel=** option to the required value. For example, to reserve 128 MB of memory, use the following:

```
crashkernel=128M
```

For information about how to change the **crashkernel=** option on AMD64 and Intel 64 systems and IBM Power Systems servers using the **GRUB2** boot loader, and on IBM Z using **zipl**, see [Section 3.1.1, "Setting kernel command-line parameters"](#).

The **crashkernel=** option can be defined in multiple ways. The **auto** value enables automatic configuration of reserved memory based on the total amount of memory in the system, following the guidelines described in [Section 7.7.1, "Memory requirements for kdump"](#). Larger memory systems, up to the established limits of the operating system are calculated according to architecture with the **crashkernel=auto** option.

Replace the **auto** value with a specific amount of memory to change this behavior.

The **crashkernel=** option can be particularly useful with smaller memory systems. For example, to reserve 128 MB of memory, use the following:

```
crashkernel=128M
```

You can also set the amount of reserved memory to be variable, depending on the total amount of installed memory. The syntax for variable memory reservation is **crashkernel=<range1>:<size1>,<range2>:<size2>**. For example:

```
crashkernel=512M-2G:64M,2G-:128M
```

The above example reserves 64 MB of memory if the total amount of system memory is 512 MB or higher and lower than 2 GB. If the total amount of memory is more than 2 GB, 128 MB is reserved for kdump instead.

Some systems require to reserve memory with a certain fixed offset. If the offset is set, the reserved memory begins there. To offset the reserved memory, use the following syntax:

```
crashkernel=128M@16M
```

The example above means that kdump reserves 128 MB of memory starting at 16 MB (physical address 0x01000000). If the offset parameter is set to 0 or omitted entirely, kdump offsets the reserved memory automatically. This syntax can also be used when setting a variable memory reservation as

described above; in this case, the offset is always specified last (for example, **crashkernel=512M-2G:64M,2G-:128M@16M**).

7.2.2.2. Configuring the kdump type

When a kernel crash is captured, the core dump can be either stored as a file in a local file system, written directly to a device, or sent over a network using the **NFS** (Network File System) or **SSH** (Secure Shell) protocol. Only one of these options can be set at the moment, and the default option is to store the **vmcore** file in the **/var/crash/** directory of the local file system. To change this, as **root**, open the **/etc/kdump.conf** configuration file in a text editor and edit the options as described below.

To change the local directory in which the core dump is to be saved, remove the hash sign ("#") from the beginning of the **#path /var/crash** line, and replace the value with a desired directory path.

```
path /usr/local/cores
```



IMPORTANT

In Red Hat Enterprise Linux 7, the directory defined as the kdump target using the **path** directive must exist when the **kdump** systemd service is started – otherwise the service fails. This behavior is different from earlier releases of Red Hat Enterprise Linux, where the directory was being created automatically if it did not exist when starting the service.

Optionally, if you wish to write the file to a different partition, follow the same procedure with the one of the lines beginning with **#ext4**. Here, you can use either a device name (the **#ext4 /dev/vg/lv_kdump** line), a file system label (the **#ext4 LABEL=/boot** line) or a UUID (the **#ext4 UUID=03138356-5e61-4ab3-b58e-27507ac41937** line). Change the file system type as well as the device name, label or UUID to the desired values. For example:

```
ext4 UUID=03138356-5e61-4ab3-b58e-27507ac41937
```



IMPORTANT

Specifying storage devices using a **LABEL=** or **UUID=** is recommended. Disk device names such as **/dev/sda3** are not guaranteed to be consistent across reboot. See the [Red Hat Enterprise Linux 7 Storage Administration Guide](#) for information about persistent disk device naming.



IMPORTANT

When dumping to DASD on s390x hardware, it is essential that the dump devices are correctly specified in **/etc/dasd.conf** before proceeding.

To write the dump directly to a device, remove the hash sign ("#") from the beginning of the **#raw /dev/vg/lv_kdump** line, and replace the value with a desired device name. For example:

```
raw /dev/sdb1
```

To store the dump to a remote machine using the **NFS** protocol, remove the hash sign ("#") from the beginning of the **#nfs my.server.com:/export/tmp** line, and replace the value with a valid hostname and directory path. For example:

```
nfs penguin.example.com:/export/cores
```

To store the dump to a remote machine using the **SSH** protocol, remove the hash sign ("#") from the beginning of the **#ssh user@my.server.com** line, and replace the value with a valid username and hostname. To include your SSH key in the configuration as well, remove the hash sign from the beginning of the **#sshkey /root/.ssh/kdump_id_rsa** line and change the value to the location of a key valid on the server you are trying to dump to. For example:

```
ssh john@penguin.example.com
sshkey /root/.ssh/mykey
```

For information on how to configure an SSH server and set up a key-based authentication, see the [Red Hat Enterprise Linux 7 System Administrator's Guide](#).

For a complete list of currently supported and unsupported targets sorted by type, see [Table 7.3, "Supported kdump Targets"](#).

7.2.2.3. Configuring the core collector

To reduce the size of the **vmcore** dump file, **kdump** allows you to specify an external application (a **core collector**) to compress the data, and optionally leave out all irrelevant information. Currently, the only fully supported core collector is **makedumpfile**.

To enable the core collector, as **root**, open the **/etc/kdump.conf** configuration file in a text editor, remove the hash sign ("#") from the beginning of the **#core_collector makedumpfile -l --message-level 1 -d 31** line, and edit the command line options as described below.

To enable the dump file compression, add the **-l** parameter. For example:

```
core_collector makedumpfile -l
```

To remove certain pages from the dump, add the **-d value** parameter, where *value* is a sum of values of pages you want to omit as described in [Table 7.4, "Supported Filtering Levels"](#). For example, to remove both zero and free pages, use the following:

```
core_collector makedumpfile -d 17 -c
```

See the **[makedumpfile\(8\)](#)** man page for a complete list of available options.

7.2.2.4. Configuring the default action

By default, when **kdump** fails to create a core dump at the target location specified in [Section 7.2.2.2, "Configuring the kdump type"](#), **kdump** reboots the system without saving the **vmcore**. To change this behavior, as **root**, open the **/etc/kdump.conf** configuration file in a text editor, remove the hash sign ("#") from the beginning of the **#default shell** line, and replace the value with a desired action as described in [Table 7.5, "Supported Default Actions"](#).

For example:

```
default reboot
```

7.2.2.5. Enabling the service

To start the **kdump** daemon at boot time, type the following at a shell prompt as **root**:

```
systemctl enable kdump.service
```

This enables the service for **multi-user.target**. Similarly, typing **systemctl disable kdump** disables **kdump**. To start the service in the current session, use the following command as **root**:

```
systemctl start kdump.service
```



IMPORTANT

In Red Hat Enterprise Linux 7, the directory defined as the kdump target must exist when the **kdump** systemd service is started – otherwise the service fails. This behavior is different from earlier releases of Red Hat Enterprise Linux, where the directory was being created automatically if it did not exist when starting the service.

For more information on systemd and configuring services in general, see the *Red Hat Enterprise Linux 7 System Administrator's Guide*.

7.2.3. Configuring kdump in the graphical user interface

To start the Kernel Dump Configuration utility, select **Activities → Other → Kernel crash dumps** from the panel, or type **system-config-kdump** at a shell prompt. As a result a window appears as shown in Figure 7.1, "Basic Settings".

The utility allows you to configure **kdump** as well as to enable or disable starting the service at boot time. When you are done, click **Apply** to save the changes. Unless you are already authenticated, enter the superuser password. The utility presents you with a reminder that you must reboot the system in order to apply any changes you have made to the configuration.



IMPORTANT

On IBM Z or PowerPC systems with **SELinux** running in Enforcing mode, the **kdumpgui_run_bootloader** Boolean must be enabled before launching the Kernel Dump Configuration utility. This Boolean allows **system-config-kdump** to run the boot loader in the **bootloader_t** SELinux domain. To permanently enable the Boolean, run the following command as root:

```
# setsebool -P kdumpgui_run_bootloader 1
```



IMPORTANT


When dumping to DASD on s390x hardware, it is essential that the dump devices are correctly specified in **/etc/dasd.conf** before proceeding.

7.2.3.1. Configuring the memory usage

The **Basic Settings** tab enables you to configure the amount of memory that is reserved for the **kdump** kernel. To do so, select the **Manual settings** radio button, and click the up and down arrow buttons next to the **New kdump Memory** field to increase or decrease the amount of memory to be reserved. Notice


that the **Usable Memory** field changes accordingly showing you the remaining memory that is available to the system. See [Section 7.1.2, “Memory requirements”](#) for more information on kdump’s memory requirements.

Figure 7.1. Basic Settings



7.2.3.2. Configuring the kdump type

The **Target Settings** tab allows you to specify the target location for the **vmcore** dump. The dump can be either stored as a file in a local file system, written directly to a device, or sent over a network using the **NFS** (Network File System) or **SSH** (Secure Shell) protocol.

Figure 7.2. Target Settings

To save the dump to the local file system, select the **Local filesystem** radio button. Optionally, you can customize the settings by choosing a different partition from the **Partition** drop-down list and a target directory using the **Path** field.



IMPORTANT

In Red Hat Enterprise Linux 7, the directory defined as the kdump target must exist when the **kdump** systemd service is started – otherwise the service fails. This behavior is different from earlier releases of Red Hat Enterprise Linux, where the directory was being created automatically if it did not exist when starting the service.

To write the dump directly to a device, select the **Raw device** radio button, and choose the desired target device from the drop-down list next to it.


To send the dump to a remote machine over a network connection, select the **Network** radio button. To use the **NFS** protocol, select the **NFS** radio button, and fill the **Server name** and **Path to directory** fields. To use the **SSH** protocol, select the **SSH** radio button, and fill the **Server name**, **Path to directory**, and **User name** fields with the remote server address, target directory, and a valid user name respectively.

For information on how to configure an SSH server and set up a key-based authentication, see the [Red Hat Enterprise Linux 7 System Administrator's Guide](#). For a complete list of currently supported targets, see [Table 7.3, "Supported kdump Targets"](#).

7.2.3.3. Configuring the core collector

The **Filtering Settings** tab enables you to select the filtering level for the **vmcore** dump.

Figure 7.3. Filtering Settings




To exclude the **zero page**, **cache page**, **cache private**, **user data**, or **free page** from the dump, select the checkbox next to the appropriate label.

7.2.3.4. Configuring the default action

To choose which action should be performed when **kdump** fails to create a core dump, select an appropriate option from the **Action if dumping fails** drop-down list. Available options are:

- *Dump to rootfs and reboot* attempts to save the core locally and then reboots the system
- *Reboot* the default action which reboots the system
- *Start a Shell* to present a user with an interactive shell prompt
- *Halt* to halt the system
- *Poweroff* to power the system off

Figure 7.4. Filtering Settings



To customize the options that are passed to the **makedumpfile** core collector, edit the **Core collector** text field; see [Section 7.2.2.3, “Configuring the core collector”](#) for more information.

7.2.3.5. Enabling the service

To start the **kdump** service at boot time, click the **Enable** button on the toolbar and then click the **Apply** button. This enables and activates the service for **multi-user.target**. Click the **Disable** button and confirm by clicking the **Apply** button to disable the service immediately.



IMPORTANT

In Red Hat Enterprise Linux 7, the directory defined as the kdump target must exist when the **kdump** systemd service is started – otherwise the service fails. This behavior is different from earlier releases of Red Hat Enterprise Linux, where the directory was being created automatically if it did not exist when starting the service.

For more information on systemd targets and configuring services in general, see the [Red Hat Enterprise Linux 7 System Administrator’s Guide](#).

7.3. TESTING THE KDUMP CONFIGURATION



WARNING

The commands below cause the kernel to crash. Use caution when following these steps, and by no means use them on a production system.


To test the configuration, reboot the system with **kdump** enabled, and make sure that the service is running:

```
~]# systemctl is-active kdump  
active
```

Then type the following commands at a shell prompt:

```
echo 1 > /proc/sys/kernel/sysrq  
echo c > /proc/sysrq-trigger
```

This forces the Linux kernel to crash, and the **address-YYYY-MM-DD-HH:MM:SS/vmcore** file is copied to the location you have selected in the configuration (that is, to **/var/crash/** by default).



NOTE

In addition to confirming the validity of the configuration, this action can also be used to record how long it takes to a crash dump to complete if it is performed under a representative test load.

7.3.1. Additional resources

7.3.1.1. Installed documentation

- **kdump.conf(5)** – a manual page for the **/etc/kdump.conf** configuration file containing the full documentation of available options.
- **zipl.conf(5)** – a manual page for the **/etc/zipl.conf** configuration file.
- **zipl(8)** – a manual page for the **zipl** boot loader utility for IBM Z.
- **makedumpfile(8)** – a manual page for the **makedumpfile** core collector.
- **kexec(8)** – a manual page for **kexec**.
- **crash(8)** – a manual page for the **crash** utility.
- **/usr/share/doc/kexec-tools-version/kexec-kdump-howto.txt** – an overview of the **kdump** and **kexec** installation and usage.

7.3.1.2. Online documentation

<https://access.redhat.com/site/solutions/6038>

The Red Hat Knowledgebase article about the **kexec** and **kdump** configuration.

<https://access.redhat.com/site/solutions/223773>

The Red Hat Knowledgebase article about supported **kdump** targets.

<http://people.redhat.com/anderson/>

The **crash** utility homepage.

<https://www.gnu.org/software/grub/>

The **GRUB2** boot loader homepage and documentation.

7.4. FIRMWARE ASSISTED DUMP MECHANISMS

7.4.1. The case for firmware assisted dump

The **kexec** and **kdump** mechanisms are a reliable and proven method of capturing a core dump on AMD64 and Intel 64 systems. However, some hardware with a longer history, particularly mini and mainframe systems, allows us to leverage the onboard firmware to isolate regions of memory and prevent any accidental overwriting of data that is important to the crash analysis.


This chapter covers some of the available firmware assisted dump methods and how they integrate with Red Hat Enterprise Linux.

7.4.2. Using fadump on IBM PowerPC hardware

Firmware-assisted dump (**fadump**) is a reliable alternative to **kexec-kdump** available on IBM PowerPC LPARS. It captures vmcore from a fully-reset system with PCI and I/O devices reinitialized. While this mechanism uses the firmware to preserve the memory in case of a crash, it reuses the **kdump** userspace scripts to save the vmcore"

To achieve this, **fadump** registers the regions of memory that must be preserved in the event of a crash with the system firmware. These regions consist of all the system memory contents, except the boot memory, system registers and hardware Page Table Entries (PTEs).

For further details about the **fadump** mechanism, including PowerPC-specific methods of resetting hardware, review `/usr/share/doc/kexec-tools-X.y.z/fadump-howto.txt` where "X.y.z" correspond to the version number of **kexec-tools** installed on your system.



NOTE

The area of memory not preserved and known as **boot memory** is the amount of RAM required to successfully boot the kernel after a crash event. By default, the boot memory size is 256MB or 5% of total system RAM, whichever is larger.

Unlike a **kexec**-initiated event, the **fadump** process uses the production kernel to recover a crash dump. When booting after a crash, PowerPC hardware makes the device node `/proc/device-tree/rtas(ibm, kernel-dump)` available to **procfs**, which the fadump-aware **kdump** scripts check for to save the vmcore. After this has completed, the system is rebooted cleanly.

Enabling fadump

1. Install and configure **kdump** as described in [Section 7.2, “Installing and configuring kdump”](#).

2. Add **fadump=on** to the **GRUB_CMDLINE_LINUX** line in **/etc/default/grub**:

```
GRUB_CMDLINE_LINUX="rd.lvm.lv=rhel/swap crashkernel=auto rd.lvm.lv=rhel/root rhgb
quiet fadump=on"
```

3. (optional) If you want to specify reserved boot memory instead of accepting the defaults, configure **crashkernel=xxM** to **GRUB_CMDLINE_LINUX** in **/etc/default/grub**, where *xx* is the amount of the memory required in megabytes:

```
GRUB_CMDLINE_LINUX="rd.lvm.lv=rhel/swap crashkernel=xxM rd.lvm.lv=rhel/root rhgb
quiet fadump=on"
```



IMPORTANT

As with all boot configuration options, it is strongly recommended that you test the configuration before it is needed. If you observe Out of Memory (OOM) errors when booting from the crash kernel, increase the value specified in **crashkernel=** until the crash kernel can boot cleanly. Some trial and error may be required in this case.

7.4.3. Firmware assisted dump methods on IBM Z

There are two firmware assisted dump mechanisms on IBM Z. They are **Stand-alone Dump** and **VMDUMP**.

The **kdump** infrastructure is supported and utilized on these systems and configuration from Red Hat Enterprise Linux is described in [Section 7.2, "Installing and configuring kdump"](#). However, there are potentially some advantages to using either of the firmware assisted methods IBM Z hardware provides.

The Stand-alone Dump (SADMP) mechanism is initiated and controlled from the system console, and must be stored on an IPL bootable device.

Similar to SADMP is VMDUMP. This tool is also initiated from the system console, but has a mechanism to retrieve the resulting dump from hardware and copy it to a system for analysis.

One advantage of these methods (and similarly to other hardware based dump mechanisms), is the ability to capture the state of a machine in the Early Boot phase (before the kdump service is started)

Although VMDUMP contains a mechanism to receive the dump file into a Red Hat Enterprise Linux system, the configuration and control of both SADMP and VMDUMP are managed from the IBM Z Hardware console.

IBM discuss SADMP in detail, at

http://www.ibm.com/support/knowledgecenter/SSLTBW_2.1.0/com.ibm.zos.v2r1.ieav100/standa.htm and VMDUMP at

http://www.ibm.com/support/knowledgecenter/en/linuxonibm/com.ibm.linux.z.lgdt/lgdt_t_vmdump.html

IBM also have a documentation set for using the dump tools on Red Hat Enterprise Linux 7 at http://www.ibm.com/support/knowledgecenter/linuxonibm/com.ibm.linux.z.lgdt/lgdt_t_usingdump-tools.html

7.4.4. Using sadump on Fujitsu PRIMEQUEST systems

The Fujitsu **sadump** mechanism is designed to provide a fallback dump capture in the event **kdump** is unable to complete successfully.

The **sadump** process is invoked manually from the system ManageMent Board (MMB) interface.

With this system, configure kdump as normal for an X86_64 server and then perform the following additional steps to enable **sadump**.

Add or edit the following lines in **/etc/sysctl.conf** to ensure that **kdump** starts as expected for **sadump**.


```
kernel.panic=0
kernel.unknown_nmi_panic=1
```

In addition to the above, you must also add some options to **/etc/kdump.conf** to ensure that **kdump** behaves correctly for sadump.

In particular, ensure that after **kdump**, the system does not reboot. If the system reboots after **kdump** has failed to save core, then you have no opportunity to invoke **sadump**.

Set the **default** action in **/etc/kdump.conf** to be either *halt* or *shell* to achieve this.

```
default shell
```



IMPORTANT

For details on configuring your hardware for **sadump**, see the FUJITSU Server PRIMEQUEST 2000 Series Installation Manual.

7.5. ANALYZING A CORE DUMP

To determine the cause of the system crash, you can use the **crash** utility, which provides an interactive prompt very similar to the GNU Debugger (GDB). This utility allows you to interactively analyze a running Linux system as well as a core dump created by **netdump**, **diskdump**, **xendump**, or **kdump**.

7.5.1. Installing the crash utility

To install the **crash** analyzing tool, execute the following command from a shell prompt as **root**:

```
yum install crash
```

In addition to **crash**, it is also necessary to install the **kernel-debuginfo** package that corresponds to your running kernel, which provides the data necessary for dump analysis. To install **kernel-debuginfo** we use the **debuginfo-install** command as **root**:

```
debuginfo-install kernel
```

For more information on how to install new packages in Red Hat Enterprise Linux using the **Yum** package manager, see the [Red Hat Enterprise Linux 7 System Administrator's Guide](#).

7.5.2. Running the crash utility

To start the utility, type the command in the following form at a shell prompt:

```
crash /usr/lib/debug/lib/modules/<kernel>/vmlinux \ /var/crash/<timestamp>/vmcore
```

Use the same <kernel> version that was captured by **kdump**. To find out which kernel you are currently running, use the **uname -r** command.

Example 7.1. Running the crash utility

```
~]# crash /usr/lib/debug/lib/modules/2.6.32-69.el6.i686/vmlinux \
/var/crash/127.0.0.1-2010-08-25-08:45:02/vmcore

crash 5.0.0-23.el6
Copyright (C) 2002-2010 Red Hat, Inc.
Copyright (C) 2004, 2005, 2006 IBM Corporation
Copyright (C) 1999-2006 Hewlett-Packard Co
Copyright (C) 2005, 2006 Fujitsu Limited
Copyright (C) 2006, 2007 VA Linux Systems Japan K.K.
Copyright (C) 2005 NEC Corporation
Copyright (C) 1999, 2002, 2007 Silicon Graphics, Inc.
Copyright (C) 1999, 2000, 2001, 2002 Mission Critical Linux, Inc.
This program is free software, covered by the GNU General Public License,
and you are welcome to change it and/or distribute copies of it under
certain conditions. Enter "help copying" to see the conditions.
This program has absolutely no warranty. Enter "help warranty" for details.

GNU gdb (GDB) 7.0
Copyright (C) 2009 Free Software Foundation, Inc.
License GPLv3+: GNU GPL version 3 or later <http://gnu.org/licenses/gpl.html>
This is free software: you are free to change and redistribute it.
There is NO WARRANTY, to the extent permitted by law. Type "show copying"
and "show warranty" for details.
This GDB was configured as "i686-pc-linux-gnu"...

KERNEL: /usr/lib/debug/lib/modules/2.6.32-69.el6.i686/vmlinux
DUMPFILE: /var/crash/127.0.0.1-2010-08-25-08:45:02/vmcore [PARTIAL DUMP]
CPUS: 4
DATE: Wed Aug 25 08:44:47 2010
UPTIME: 00:09:02
LOAD AVERAGE: 0.00, 0.01, 0.00
TASKS: 140
NODENAME: hp-dl320g5-02.lab.bos.redhat.com
RELEASE: 2.6.32-69.el6.i686
VERSION: #1 SMP Tue Aug 24 10:31:45 EDT 2010
MACHINE: i686 (2394 Mhz)
MEMORY: 8 GB
PANIC: "Oops: 0002 [#1] SMP " (check log for details)
PID: 5591
COMMAND: "bash"
TASK: f196d560 [THREAD_INFO: ef4da000]
CPU: 2
STATE: TASK_RUNNING (PANIC)

crash>
```

7.5.3. Displaying the message buffer

To display the kernel message buffer, type the **log** command at the interactive prompt.


Example 7.2. Displaying the kernel message buffer

```

crash> log
... several lines omitted ...
EIP: 0060:[<c068124f>] EFLAGS: 00010096 CPU: 2
EIP is at sysrq_handle_crash+0xf/0x20
EAX: 00000063 EBX: 00000063 ECX: c09e1c8c EDX: 00000000
ESI: c0a09ca0 EDI: 00000286 EBP: 00000000 ESP: ef4dbf24
DS: 007b ES: 007b FS: 00d8 GS: 00e0 SS: 0068
Process bash (pid: 5591, ti=ef4da000 task=f196d560 task.ti=ef4da000)
Stack:
c068146b c0960891 c0968653 00000003 00000000 00000002 efade5c0 c06814d0
<0> ffffffb c068150f b7776000 f2600c40 c0569ec4 ef4dbf9c 00000002 b7776000
<0> efade5c0 00000002 b7776000 c0569e60 c051de50 ef4dbf9c f196d560 ef4dbfb4
Call Trace:
[<c068146b>] ? __handle_sysrq+0xfb/0x160
[<c06814d0>] ? write_sysrq_trigger+0x0/0x50
[<c068150f>] ? write_sysrq_trigger+0x3f/0x50
[<c0569ec4>] ? proc_reg_write+0x64/0xa0
[<c0569e60>] ? proc_reg_write+0x0/0xa0
[<c051de50>] ? vfs_write+0xa0/0x190
[<c051e8d1>] ? sys_write+0x41/0x70
[<c0409adc>] ? syscall_call+0x7/0xb
Code: a0 c0 01 0f b6 41 03 19 d2 f7 d2 83 e2 03 83 e0 cf c1 e2 04 09 d0 88 41 03 f3 c3 90 c7 05
c8 1b 9e c0 01 00 00 00 0f ae f8 89 f6 <c6> 05 00 00 00 00 01 c3 89 f6 8d bc 27 00 00 00 00 8d 50
d0 83
EIP: [<c068124f>] sysrq_handle_crash+0xf/0x20 SS:ESP 0068:ef4dbf24
CR2: 0000000000000000

```

Type **help log** for more information on the command usage.

**NOTE**

The kernel message buffer includes the most essential information about the system crash and, as such, it is always dumped first into the **vmcore-dmesg.txt** file. This is useful when an attempt to get the full **vmcore** file failed, for example because of lack of space on the target location. By default, **vmcore-dmesg.txt** is located in the **/var/crash/** directory.

7.5.4. Displaying a backtrace

To display the kernel stack trace, type the **bt** command at the interactive prompt. You can use **bt <pid>** to display the backtrace of a single process.

Example 7.3. Displaying the kernel stack trace

```

crash> bt
PID: 5591 TASK: f196d560 CPU: 2 COMMAND: "bash"
#0 [ef4dbdcc] crash_kexec at c0494922
#1 [ef4dbe20] oops_end at c080e402
#2 [ef4dbe34] no_context at c043089d
#3 [ef4dbe58] bad_area at c0430b26
#4 [ef4dbe6c] do_page_fault at c080fb9b

```

```
#5 [ef4dbee4] error_code (via page_fault) at c080d809
  EAX: 00000063 EBX: 00000063 ECX: c09e1c8c EDX: 00000000 EBP: 00000000
  DS: 007b   ESI: c0a09ca0 ES: 007b   EDI: 00000286 GS: 00e0
  CS: 0060   EIP: c068124f ERR: ffffffff EFLAGS: 00010096
#6 [ef4dbf18] sysrq_handle_crash at c068124f
#7 [ef4dbf24] __handle_sysrq at c0681469
#8 [ef4dbf48] write_sysrq_trigger at c068150a
#9 [ef4dbf54] proc_reg_write at c0569ec2
#10 [ef4dbf74] vfs_write at c051de4e
#11 [ef4dbf94] sys_write at c051e8cc
#12 [ef4dbfb0] system_call at c0409ad5
  EAX: fffffda EBX: 00000001 ECX: b7776000 EDX: 00000002
  DS: 007b   ESI: 00000002 ES: 007b   EDI: b7776000
  SS: 007b   ESP: bfcb2088 EBP: bfcb20b4 GS: 0033
  CS: 0073   EIP: 00edc416 ERR: 00000004 EFLAGS: 00000246
```

Type **help bt** for more information on the command usage.

7.5.5. Displaying a process status

To display status of processes in the system, type the **ps** command at the interactive prompt. You can use **ps <pid>** to display the status of a single process.

Example 7.4. Displaying the status of processes in the system

```
crash> ps
 PID  PPID CPU TASK ST %MEM  VSZ  RSS COMM
 > 0  0 0 c09dc560 RU 0.0 0 0 [swapper]
 > 0  0 1 f7072030 RU 0.0 0 0 [swapper]
 0  0 2 f70a3a90 RU 0.0 0 0 [swapper]
 > 0  0 3 f70ac560 RU 0.0 0 0 [swapper]
 1  0 1 f705ba90 IN 0.0 2828 1424 init
 ... several lines omitted ...
 5566  1 1 f2592560 IN 0.0 12876 784 auditd
 5567  1 2 ef427560 IN 0.0 12876 784 auditd
 5587 5132 0 f196d030 IN 0.0 11064 3184 sshd
 > 5591 5587 2 f196d560 RU 0.0 5084 1648 bash
```

Type **help ps** for more information on the command usage.

7.5.6. Displaying virtual memory information

To display basic virtual memory information, type the **vm** command at the interactive prompt. You can use **vm <pid>** to display information on a single process.

Example 7.5. Displaying virtual memory information of the current context

```
crash> vm
PID: 5591 TASK: f196d560 CPU: 2 COMMAND: "bash"
 MM  PGD  RSS  TOTAL_V
f19b5900 ef9c6000 1648k 5084k
```

VMA	START	END	FLAGS	FILE
f1bb0310	242000	260000	8000875	/lib/ld-2.12.so
f26af0b8	260000	261000	8100871	/lib/ld-2.12.so
efbc275c	261000	262000	8100873	/lib/ld-2.12.so
efbc2a18	268000	3ed000	8000075	/lib/libc-2.12.so
efbc23d8	3ed000	3ee000	8000070	/lib/libc-2.12.so
efbc2888	3ee000	3f0000	8100071	/lib/libc-2.12.so
efbc2cd4	3f0000	3f1000	8100073	/lib/libc-2.12.so
efbc243c	3f1000	3f4000	100073	
efbc28ec	3f6000	3f9000	8000075	/lib/libdl-2.12.so
efbc2568	3f9000	3fa000	8100071	/lib/libdl-2.12.so
efbc2f2c	3fa000	3fb000	8100073	/lib/libdl-2.12.so
f26af888	7e6000	7fc000	8000075	/lib/libtinfo.so.5.7
f26aff2c	7fc000	7ff000	8100073	/lib/libtinfo.so.5.7
efbc211c	d83000	d8f000	8000075	/lib/libnss_files-2.12.so
efbc2504	d8f000	d90000	8100071	/lib/libnss_files-2.12.so
efbc2950	d90000	d91000	8100073	/lib/libnss_files-2.12.so
f26afe00	edc000	edd000	4040075	
f1bb0a18	8047000	8118000	8001875	/bin/bash
f1bb01e4	8118000	811d000	8101873	/bin/bash
f1bb0c70	811d000	8122000	100073	
f26afae0	9fd9000	9ffa000	100073	
... several lines omitted ...				

Type **help vm** for more information on the command usage.

7.5.7. Displaying open files

To display information about open files, type the **files** command at the interactive prompt. You can use **files <pid>** to display files opened by only one selected process.

Example 7.6. Displaying information about open files of the current context

```
crash> files
PID: 5591 TASK: f196d560 CPU: 2 COMMAND: "bash"
ROOT: / CWD: /root
FD FILE DENTRY INODE TYPE PATH
0 f734f640 eedc2c6c eecd6048 CHR /pts/0
1 efade5c0 eee14090 f00431d4 REG /proc/sysrq-trigger
2 f734f640 eedc2c6c eecd6048 CHR /pts/0
10 f734f640 eedc2c6c eecd6048 CHR /pts/0
255 f734f640 eedc2c6c eecd6048 CHR /pts/0
```

Type **help files** for more information on the command usage.

7.5.8. Exiting the utility

To exit the interactive prompt and terminate **crash**, type **exit** or **q**.

Example 7.7. Exiting the crash utility

```
crash> exit  
~]#
```

7.6. FREQUENTLY ASKED QUESTIONS

What considerations need to be made for using Kdump in a clustered environment?

[How do I configure kdump for use with the RHEL 6, 7 High Availability Add-On?](#) shows the options available to system administrators using the High Availability Add-On.

Kdump fails during early boot, How do I capture the boot log?

If there is a problem booting the second kernel, it is necessary to review the early boot logs, these can be obtained by enabling a serial console to the affected machine.

[How do I setup serial console in RHEL7?](#) shows the configuration needed to enable access to the early boot messages.

How do I increase the messaging from makedumpfile for debugging?

In the event that **makedumpfile** fails, then it is necessary to increase the log level to understand what is going wrong. This is different from setting the dump level and is achieved by editing **/etc/kdump.conf** and increasing the **message_level** option to **makedumpfile** on the **core_collector** line entry.

By default **makedumpfile** is set to level 7, which includes the progress indicator, common message, and error message output. Set this level to 31 to get further debugging information.

Ensure that your **core_collector** config line looks similar to this when set:

```
core_collector makedumpfile -l --message-level 1 -d 31
```

How do I debug Dracut?

Sometimes **dracut** can fail to build an initramfs. If this happens, increase the log level in **dracut** to isolate the issue.

Edit **/etc/kdump.conf** and change the **dracut_args** line to include the option **-L 5** in addition to any other dracut arguments you require.

If you have no other options configured in **dracut_args**, the result looks similar to this:

```
dracut_args -L 5
```

What methods of dumping are available for virtual machines?

In most cases, the **kdump** mechanism is sufficient for obtaining a memory dump from a machine after a crash or panic. This can be set up in the same manner as installations to bare metal.

However, in some cases, it is necessary to work directly with the hypervisor to obtain a crash dump. There are two mechanisms available with **libvirt** to achieve this; **pvpanic** and **virsh dump**. Both of these methods are described in the [Virtualization Deployment and Administration Guide](#).

The **pvpanic** mechanism can be found at [Virtualization Deployment and Administration Guide - Setting a Panic Device](#).

The **virsh dump** command is discussed in [Virtualization Deployment and Administration Guide - Creating a Dump File of a Domain's Core](#).

How do I upload a large dump file to Red Hat Support Services?

In some cases, it might be necessary to send a kernel crash dump file to Red Hat Global Support Services for analysis. However, the dump file can be very large, even after being filtered. Since files larger than 250 MB cannot be uploaded directly through the Red Hat Customer Portal when opening a new support case, an FTP server is provided by Red Hat for uploading large files.

The FTP server's address is **dropbox.redhat.com** and the files are to be uploaded in the **/incoming/** directory. Your FTP client needs to be set into passive mode; if your firewall does not allow this mode, use the **origin-dropbox.redhat.com** server using active mode.

Make sure that the uploaded files are compressed using a program such as **gzip** and properly and descriptively named. Using your support case number in the file name is recommended. After successfully uploading all necessary files, provide the engineer in charge of your support case with the exact file name and its SHA1 or MD5 checksum.

For more specific instructions and additional information, see [How to provide files to Red Hat Support](#).

How much time is needed for a crash dump to complete?

It is often necessary, for the purposes of disaster recovery planning, to know how long a dump takes to complete. However, the length of time it takes is highly dependent on the amount of memory being copied to disk and the speed of the interfaces between RAM and storage.

For any test of timings, the system must be operating under a representative load, otherwise the page exclusion choices can present a false view of kdump behavior with a fully loaded production system. This discrepancy is present especially when working with very large quantities of RAM.

Also consider storage interfaces in your planning when assessing time to dump. Because of network constraints, a connection dumping over **ssh** for example, can take longer to complete than a locally attached SATA disk.

How is Kdump configured during installation?

You can configure **kdump** during installation with a limited set of options in kickstart or the interactive GUI.

The **kdump** configuration using the **anaconda** installation GUI is documented in the [KDUMP section](#) of the Installation Guide.

The **kickstart** syntax is:

```
%addon com_redhat_kdump [--disable,enable] [--reserve-mb=[auto,value]]
%end
```

With this add-on to Kickstart, you can disable or enable kdump functionality, optionally defining the reserved memory size, either by explicitly invoking the default option of auto (which is also the case if the entire switch is omitted) or specifying a numeric value in megabytes.

To learn how Kickstart can be used to automate system deployments, read [Kickstart Installations](#) in the Installation Guide.

For further details about Kickstart add-on syntax, review the [Kickstart Syntax Reference](#) in the Installation Guide.

7.7. SUPPORTED KDUMP CONFIGURATIONS AND TARGETS

7.7.1. Memory requirements for kdump

In order for kdump to be able to capture a kernel crash dump and save it for further analysis, a part of the system memory has to be permanently reserved for the capture kernel. The table below contains a list of minimum memory requirements for kdump based on the system's architecture and total available physical memory.

For information on how to change memory settings on the command line, see [Section 7.2.2.1, “Configuring the memory usage”](#). For instructions on how to set up the amount of reserved memory in the graphical user interface, see [Section 7.2.3.1, “Configuring the memory usage”](#).

Table 7.1. Minimum Amount of Reserved Memory Required for kdump

Architecture	Available Memory	Minimum Reserved Memory
AMD64 and Intel 64 (x86_64)	2 GB and more	160 MB + 2 bits for every 4 KB of RAM. For a system with 1 TB of memory, 224 MB is the minimum (160 + 64 MB).
IBM POWER (ppc64)	2 GB to 4 GB	384 MB of RAM.
	4 GB to 16 GB	512 MB of RAM.
	16 GB to 64 GB	1 GB of RAM.
	64 GB to 128 GB	2 GB of RAM.
	128 GB and more	4 GB of RAM.
IBM Z (s390x)	2 GB and more	160 MB + 2 bits for every 4 KB of RAM. For a system with 1 TB of memory, 224 MB is the minimum (160 + 64 MB).

For more information about various Red Hat Enterprise Linux technology capabilities and limits, see <https://access.redhat.com/articles/rhel-limits>.

7.7.2. Minimum threshold for automatic memory reservation

On some systems, it is possible to allocate memory for kdump automatically, either by using the **crashkernel=auto** parameter in the bootloader's configuration file, or by enabling this option in the graphical configuration utility. For this automatic reservation to work, however, a certain amount of total memory needs to be available in the system. This amount differs based on the system's architecture.

The table below lists the thresholds for automatic memory allocation. If the system has less memory than specified in the table, memory needs to be reserved manually.

For information on how to change these settings on the command line, see [Section 7.2.2.1, "Configuring the memory usage"](#). For instructions on how to change the amount of reserved memory in the graphical user interface, see [Section 7.2.3.1, "Configuring the memory usage"](#).

Table 7.2. Minimum Amount of Memory Required for Automatic Memory Reservation

Architecture	Required Memory
AMD64 and Intel 64 (x86_64)	2 GB
IBM POWER (ppc64)	2 GB
IBM Z (s390x)	4 GB

7.7.3. Supported kdump targets

When a kernel crash is captured, the core dump can be either written directly to a device, stored as a file on a local file system, or sent over a network. The table below contains a complete list of dump targets that are currently supported or explicitly unsupported by kdump.

For information on how to configure the target type on the command line, see [Section 7.2.2.2, "Configuring the kdump type"](#). For information on how to do so in the graphical user interface, see [Section 7.2.3.2, "Configuring the kdump type"](#).

Table 7.3. Supported kdump Targets

Type	Supported Targets	Unsupported Targets
Raw device	All locally attached raw disks and partitions.	
Local file system	ext2 , ext3 , ext4 , and xfs file systems on directly attached disk drives, hardware RAID logical drives, LVM devices, and mdraid arrays.	Any local file system not explicitly listed as supported in this table, including the auto type (automatic file system detection).
Remote directory	Remote directories accessed using the NFS or SSH protocol over IPv4 .	Remote directories on the rootfs file system accessed using the NFS protocol.
Remote directories accessed using the FCoE (<i>Fibre Channel over Ethernet</i>) protocol.	qla2xxx , lpfc and bfa hardware FCoE targets. bnx2fc and ixgbe software FCoE targets.	
Remote directories accessed using the iSCSI protocol over both hardware and software initiators.	Remote directories accessed using the iSCSI protocol on be2iscsi hardware.	Multipath-based storages.

Type	Supported Targets	Unsupported Targets
		Remote directories accessed over IPv6 .
		Remote directories accessed using the SMB or CIFS protocol.
		Remote directories accessed using wireless network interfaces.



NOTE

When dumping to a software **FCoE** target, you can encounter Out of Memory (OOM) issue. In such cases, increase the default **crashkernel=auto** parameter value. For more information on how to set up this kernel boot parameter, see [Section 7.2.2.1, “Configuring the memory usage”](#).

7.7.4. Supported kdump filtering levels

To reduce the size of the dump file, kdump uses the **makedumpfile** core collector to compress the data and optionally leave out irrelevant information. The table below contains a complete list of filtering levels that are currently supported by the **makedumpfile** utility.

For instructions on how to configure the core collector on the command line, see [Section 7.2.2.3, “Configuring the core collector”](#). For information on how to do so in the graphical user interface, see [Section 7.2.3.3, “Configuring the core collector”](#).

Table 7.4. Supported Filtering Levels

Option	Description
1	Zero pages
2	Cache pages
4	Cache private
8	User pages
16	Free pages



NOTE

The **makedumpfile** command supports removal of transparent huge pages and hugetlfs pages on Red Hat Enterprise Linux 7.3 and later. Consider both these types of hugepages User Pages and remove them using the **-8** level.

7.7.5. Supported default actions

By default, when kdump fails to create a core dump, the operating system reboots. You can, however, configure kdump to perform a different operation in case it fails to save the core dump to the primary target. The table below lists all default actions that are currently supported by kdump.

For detailed information on how to set up the default action on the command line, see [Section 7.2.2.4, “Configuring the default action”](#). For information on how to do so in the graphical user interface, see [Section 7.2.3.4, “Configuring the default action”](#).

Table 7.5. Supported Default Actions

Option	Description
dump_to_rootfs	Attempt to save the core dump to the root file system. This option is especially useful in combination with a network target: if the network target is unreachable, this option configures kdump to save the core dump locally. The system is rebooted afterwards.
reboot	Reboot the system, losing the core dump in the process.
halt	Halt the system, losing the core dump in the process.
poweroff	Power off the system, losing the core dump in the process.
shell	Run a shell session from within the initramfs, allowing the user to record the core dump manually.

7.7.6. Estimating kdump size

When planning and building your **kdump** environment it is necessary to know how much space is required for the dump file before one is produced. The **makedumpfile** command can help with this.

The **--mem-usage** option provides a useful report about excludable pages, that can be used to determine which dump level you want to assign. Run this command when the system is under representative load, otherwise **makedumpfile** returns a smaller value than is expected in your production environment.

```
[root@hostname ~]# makedumpfile --mem-usage /proc/kcore
```

TYPE	PAGES	EXCLUDABLE	DESCRIPTION
ZERO	501635	yes	Pages filled with zero
CACHE	51657	yes	Cache pages
CACHE_PRIVATE	5442	yes	Cache pages + private
USER	16301	yes	User process pages
FREE	77738211	yes	Free pages
KERN_DATA	1333192	no	Dumpable kernel data



IMPORTANT

The **makedumpfile** command reports in **pages**. This means that you must calculate the size of memory in use against the kernel page size, which in the Red Hat Enterprise Linux kernel, is 4 kilobytes for AMD64 and Intel 64 architectures, and 64 kilobytes for IBM POWER architecture.

7.8. PORTAL LABS RELEVANT TO KDUMP

The [Portal Labs](#) are small web applications that can help system administrators perform several system tasks. There are currently two labs focused on Kdump. The Kdump Helper and the Kernel Oops Analyzer.

7.8.1. Kdump helper

The [Kdump Helper](#) is a series of questions and actions that assist in preparing the configuration files for **kdump**.

The Lab's workflow includes steps for both clustered and standalone environments.

7.8.2. Kernel oops analyzer

The [Kernel Oops Analyzer](#) is a tool to process Oops messages and search for known solutions without having to unwind the crash dump stack.

The Kernel Oops Analyzer uses information from **makedumpfile** to compare the oops message from a crashed machine with known issues in the knowledge base. This can enable System Administrators to rule out known issues quickly after an unexpected outage, and before opening a support ticket for a further analysis.

CHAPTER 8. REVISION HISTORY

0.1-6

Tue Aug 6 2019, Jaroslav Klech (jklech@redhat.com)

- Document version for 7.7 GA publication.

0.1-5

Fri Oct 19 2018, Jaroslav Klech (jklech@redhat.com)

- Document version for 7.6 GA publication.

0.1-4

Mon Mar 26 2018, Marie Doleželová (mdolezel@redhat.com)

- Document version for 7.5 GA publication.

0.1-3

Mon Jan 5 2018, Mark Flitter (mflitter@redhat.com)

- Document version for 7.5 Beta publication.

0.1-2

Mon Jul 31 2017, Mark Flitter (mflitter@redhat.com)

- Document version for 7.4 GA publication.

0.1-0

Thu Apr 20 2017, Mark Flitter (mflitter@redhat.com)

- Initial build for review