

Red Hat

Red Hat Enterprise Linux 7

High Availability Add-On Reference

Reference guide for configuration and management of the High Availability Add-On

Red Hat Enterprise Linux 7 High Availability Add-On Reference

Reference guide for configuration and management of the High Availability Add-On

Steven Levine
Red Hat Customer Content Services
slevine@redhat.com

Legal Notice

Copyright © 2018 Red Hat, Inc.

This document is licensed by Red Hat under the [Creative Commons Attribution-ShareAlike 3.0 Unported License](#). If you distribute this document, or a modified version of it, you must provide attribution to Red Hat, Inc. and provide a link to the original. If the document is modified, all Red Hat trademarks must be removed.

Red Hat, as the licensor of this document, waives the right to enforce, and agrees not to assert, Section 4d of CC-BY-SA to the fullest extent permitted by applicable law.

Red Hat, Red Hat Enterprise Linux, the Shadowman logo, the Red Hat logo, JBoss, OpenShift, Fedora, the Infinity logo, and RHCE are trademarks of Red Hat, Inc., registered in the United States and other countries.

Linux® is the registered trademark of Linus Torvalds in the United States and other countries.

Java® is a registered trademark of Oracle and/or its affiliates.

XFS® is a trademark of Silicon Graphics International Corp. or its subsidiaries in the United States and/or other countries.

MySQL® is a registered trademark of MySQL AB in the United States, the European Union and other countries.

Node.js® is an official trademark of Joyent. Red Hat is not formally related to or endorsed by the official Joyent Node.js open source or commercial project.

The OpenStack® Word Mark and OpenStack logo are either registered trademarks/service marks or trademarks/service marks of the OpenStack Foundation, in the United States and other countries and are used with the OpenStack Foundation's permission. We are not affiliated with, endorsed or sponsored by the OpenStack Foundation, or the OpenStack community.

All other trademarks are the property of their respective owners.

Abstract

Red Hat High Availability Add-On Reference provides reference information about installing, configuring, and managing the Red Hat High Availability Add-On for Red Hat Enterprise Linux 7.

Table of Contents

CHAPTER 1. RED HAT HIGH AVAILABILITY ADD-ON CONFIGURATION AND MANAGEMENT REFERENCE	
OVERVIEW	5
1.1. NEW AND CHANGED FEATURES	5
1.2. INSTALLING PACEMAKER CONFIGURATION TOOLS	8
1.3. CONFIGURING THE IPTABLES FIREWALL TO ALLOW CLUSTER COMPONENTS	8
1.4. THE CLUSTER AND PACEMAKER CONFIGURATION FILES	9
1.5. CLUSTER CONFIGURATION CONSIDERATIONS	10
1.6. UPDATING A RED HAT ENTERPRISE LINUX HIGH AVAILABILITY CLUSTER	10
1.7. ISSUES WITH LIVE MIGRATION OF VMS IN A RHEL CLUSTER	11
CHAPTER 2. THE PCSD WEB UI	12
2.1. PCSD WEB UI SETUP	12
2.2. CREATING A CLUSTER WITH THE PCSD WEB UI	13
2.3. CONFIGURING CLUSTER COMPONENTS	15
2.4. CONFIGURING A HIGH AVAILABILITY PCSD WEB UI	17
CHAPTER 3. THE PCS COMMAND LINE INTERFACE	18
3.1. THE PCS COMMANDS	18
3.2. PCS USAGE HELP DISPLAY	18
3.3. VIEWING THE RAW CLUSTER CONFIGURATION	19
3.4. SAVING A CONFIGURATION CHANGE TO A FILE	19
3.5. DISPLAYING STATUS	19
3.6. DISPLAYING THE FULL CLUSTER CONFIGURATION	20
3.7. DISPLAYING THE CURRENT PCS VERSION	20
3.8. BACKING UP AND RESTORING A CLUSTER CONFIGURATION	20
CHAPTER 4. CLUSTER CREATION AND ADMINISTRATION	21
4.1. CLUSTER CREATION	21
4.2. CONFIGURING TIMEOUT VALUES FOR A CLUSTER	22
4.3. CONFIGURING REDUNDANT RING PROTOCOL (RRP)	23
4.4. MANAGING CLUSTER NODES	23
4.5. SETTING USER PERMISSIONS	26
4.6. REMOVING THE CLUSTER CONFIGURATION	28
4.7. DISPLAYING CLUSTER STATUS	28
4.8. CLUSTER MAINTENANCE	29
CHAPTER 5. FENCING: CONFIGURING STONITH	30
5.1. AVAILABLE STONITH (FENCING) AGENTS	30
5.2. GENERAL PROPERTIES OF FENCING DEVICES	30
5.3. DISPLAYING DEVICE-SPECIFIC FENCING OPTIONS	31
5.4. CREATING A FENCING DEVICE	32
5.5. DISPLAYING FENCING DEVICES	32
5.6. MODIFYING AND DELETING FENCING DEVICES	32
5.7. MANAGING NODES WITH FENCE DEVICES	33
5.8. ADDITIONAL FENCING CONFIGURATION OPTIONS	33
5.9. CONFIGURING FENCING LEVELS	38
5.10. CONFIGURING FENCING FOR REDUNDANT POWER SUPPLIES	39
5.11. CONFIGURING ACPI FOR USE WITH INTEGRATED FENCE DEVICES	40
5.12. TESTING A FENCE DEVICE	43
CHAPTER 6. CONFIGURING CLUSTER RESOURCES	46
6.1. RESOURCE CREATION	46
6.2. RESOURCE PROPERTIES	47

6.3. RESOURCE-SPECIFIC PARAMETERS	47
6.4. RESOURCE META OPTIONS	48
6.5. RESOURCE GROUPS	51
6.6. RESOURCE OPERATIONS	52
6.7. DISPLAYING CONFIGURED RESOURCES	55
6.8. MODIFYING RESOURCE PARAMETERS	56
6.9. MULTIPLE MONITORING OPERATIONS	56
6.10. ENABLING AND DISABLING CLUSTER RESOURCES	57
6.11. CLUSTER RESOURCES CLEANUP	57
CHAPTER 7. RESOURCE CONSTRAINTS	58
7.1. LOCATION CONSTRAINTS	58
7.2. ORDER CONSTRAINTS	63
7.3. COLOCATION OF RESOURCES	65
7.4. DISPLAYING CONSTRAINTS	68
CHAPTER 8. MANAGING CLUSTER RESOURCES	69
8.1. MANUALLY MOVING RESOURCES AROUND THE CLUSTER	69
8.2. MOVING RESOURCES DUE TO FAILURE	71
8.3. MOVING RESOURCES DUE TO CONNECTIVITY CHANGES	71
8.4. ENABLING, DISABLING, AND BANNING CLUSTER RESOURCES	72
8.5. DISABLING A MONITOR OPERATION	73
8.6. MANAGED RESOURCES	74
CHAPTER 9. ADVANCED CONFIGURATION	75
9.1. RESOURCE CLONES	75
9.2. MULTISTATE RESOURCES: RESOURCES THAT HAVE MULTIPLE MODES	78
9.3. CONFIGURING A VIRTUAL DOMAIN AS A RESOURCE	79
9.4. THE PACEMAKER_REMOTE SERVICE	81
9.5. PACEMAKER SUPPORT FOR DOCKER CONTAINERS (TECHNOLOGY PREVIEW)	89
9.6. UTILIZATION AND PLACEMENT STRATEGY	96
9.7. CONFIGURING STARTUP ORDER FOR RESOURCE DEPENDENCIES NOT MANAGED BY PACEMAKER (RED HAT ENTERPRISE LINUX 7.4 AND LATER)	100
9.8. QUERYING A PACEMAKER CLUSTER WITH SNMP (RED HAT ENTERPRISE LINUX 7.5 AND LATER)	100
9.9. CONFIGURING RESOURCES TO REMAIN STOPPED ON CLEAN NODE SHUTDOWN (RED HAT ENTERPRISE LINUX 7.8 AND LATER)	103
CHAPTER 10. CLUSTER QUORUM	107
10.1. CONFIGURING QUORUM OPTIONS	107
10.2. QUORUM ADMINISTRATION COMMANDS (RED HAT ENTERPRISE LINUX 7.3 AND LATER)	107
10.3. MODIFYING QUORUM OPTIONS (RED HAT ENTERPRISE LINUX 7.3 AND LATER)	108
10.4. THE QUORUM UNBLOCK COMMAND	109
10.5. QUORUM DEVICES	109
CHAPTER 11. PACEMAKER RULES	117
11.1. NODE ATTRIBUTE EXPRESSIONS	117
11.2. TIME/DATE BASED EXPRESSIONS	118
11.3. DATE SPECIFICATIONS	119
11.4. DURATIONS	120
11.5. CONFIGURING RULES WITH PCS	120
CHAPTER 12. PACEMAKER CLUSTER PROPERTIES	121
12.1. SUMMARY OF CLUSTER PROPERTIES AND OPTIONS	121
12.2. SETTING AND REMOVING CLUSTER PROPERTIES	123
12.3. QUERYING CLUSTER PROPERTY SETTINGS	124

CHAPTER 13. TRIGGERING SCRIPTS FOR CLUSTER EVENTS	125
13.1. PACEMAKER ALERT AGENTS (RED HAT ENTERPRISE LINUX 7.3 AND LATER)	125
13.2. EVENT NOTIFICATION WITH MONITORING RESOURCES	132
CHAPTER 14. CONFIGURING MULTI-SITE CLUSTERS WITH PACEMAKER	135
APPENDIX A. OCF RETURN CODES	139
APPENDIX B. CLUSTER CREATION IN RED HAT ENTERPRISE LINUX 6 AND RED HAT ENTERPRISE LINUX 7	142
B.1. CLUSTER CREATION WITH RGMANAGER AND WITH PACEMAKER	142
B.2. PACEMAKER INSTALLATION IN RED HAT ENTERPRISE LINUX 6 AND RED HAT ENTERPRISE LINUX 7	146
APPENDIX C. REVISION HISTORY	148
INDEX	149

CHAPTER 1. RED HAT HIGH AVAILABILITY ADD-ON CONFIGURATION AND MANAGEMENT REFERENCE OVERVIEW

This document provides descriptions of the options and features that the Red Hat High Availability Add-On using Pacemaker supports. For a step by step basic configuration example, see *Red Hat High Availability Add-On Administration*.

You can configure a Red Hat High Availability Add-On cluster with the **pcs** configuration interface or with the **pcsd** GUI interface.

1.1. NEW AND CHANGED FEATURES

This section lists features of the Red Hat High Availability Add-On that are new since the initial release of Red Hat Enterprise Linux 7.

1.1.1. New and Changed Features for Red Hat Enterprise Linux 7.1

Red Hat Enterprise Linux 7.1 includes the following documentation and feature updates and changes.

- The **pcs resource cleanup** command can now reset the resource status and **failcount** for all resources, as documented in [Section 6.11, “Cluster Resources Cleanup”](#).
- You can specify a **lifetime** parameter for the **pcs resource move** command, as documented in [Section 8.1, “Manually Moving Resources Around the Cluster”](#).
- As of Red Hat Enterprise Linux 7.1, you can use the **pcs acl** command to set permissions for local users to allow read-only or read-write access to the cluster configuration by using access control lists (ACLs). For information on ACLs, see [Section 4.5, “Setting User Permissions”](#).
- [Section 7.2.3, “Ordered Resource Sets”](#) and [Section 7.3, “Colocation of Resources”](#) have been extensively updated and clarified.
- [Section 6.1, “Resource Creation”](#) documents the **disabled** parameter of the **pcs resource create** command, to indicate that the resource being created is not started automatically.
- [Section 10.1, “Configuring Quorum Options”](#) documents the new **cluster quorum unlock** feature, which prevents the cluster from waiting for all nodes when establishing quorum.
- [Section 6.1, “Resource Creation”](#) documents the **before** and **after** parameters of the **pcs resource create** command, which can be used to configure resource group ordering.
- As of the Red Hat Enterprise Linux 7.1 release, you can backup the cluster configuration in a tarball and restore the cluster configuration files on all nodes from backup with the **backup** and **restore** options of the **pcs config** command. For information on this feature, see [Section 3.8, “Backing Up and Restoring a Cluster Configuration”](#).
- Small clarifications have been made throughout this document.

1.1.2. New and Changed Features for Red Hat Enterprise Linux 7.2

Red Hat Enterprise Linux 7.2 includes the following documentation and feature updates and changes.

- You can now use the **pcs resource relocate run** command to move a resource to its preferred

node, as determined by current cluster status, constraints, location of resources and other settings. For information on this command, see [Section 8.1.2, “Moving a Resource to its Preferred Node”](#).

- [Section 13.2, “Event Notification with Monitoring Resources”](#) has been modified and expanded to better document how to configure the **ClusterMon** resource to execute an external program to determine what to do with cluster notifications.
- When configuring fencing for redundant power supplies, you now are only required to define each device once and to specify that both devices are required to fence the node. For information on configuring fencing for redundant power supplies, see [Section 5.10, “Configuring Fencing for Redundant Power Supplies”](#).
- This document now provides a procedure for adding a node to an existing cluster in [Section 4.4.3, “Adding Cluster Nodes”](#).
- The new **resource-discovery** location constraint option allows you to indicate whether Pacemaker should perform resource discovery on a node for a specified resource, as documented in [Table 7.1, “Simple Location Constraint Options”](#).
- Small clarifications and corrections have been made throughout this document.

1.1.3. New and Changed Features for Red Hat Enterprise Linux 7.3

Red Hat Enterprise Linux 7.3 includes the following documentation and feature updates and changes.

- [Section 9.4, “The pacemaker_remote Service”](#), has been wholly rewritten for this version of the document.
- You can configure Pacemaker alerts by means of alert agents, which are external programs that the cluster calls in the same manner as the cluster calls resource agents to handle resource configuration and operation. Pacemaker alert agents are described in [Section 13.1, “Pacemaker Alert Agents \(Red Hat Enterprise Linux 7.3 and later\)”](#).
- New quorum administration commands are supported with this release which allow you to display the quorum status and to change the **expected_votes** parameter. These commands are described in [Section 10.2, “Quorum Administration Commands \(Red Hat Enterprise Linux 7.3 and Later\)”](#).
- You can now modify general quorum options for your cluster with the **pcs quorum update** command, as described in [Section 10.3, “Modifying Quorum Options \(Red Hat Enterprise Linux 7.3 and later\)”](#).
- You can configure a separate quorum device which acts as a third-party arbitration device for the cluster. The primary use of this feature is to allow a cluster to sustain more node failures than standard quorum rules allow. This feature is provided for technical preview only. For information on quorum devices, see [Section 10.5, “Quorum Devices”](#).
- Red Hat Enterprise Linux release 7.3 provides the ability to configure high availability clusters that span multiple sites through the use of a Booth cluster ticket manager. This feature is provided for technical preview only. For information on the Booth cluster ticket manager, see [Chapter 14, Configuring Multi-Site Clusters with Pacemaker](#).
- When configuring a KVM guest node running a the **pacemaker_remote** service, you can include guest nodes in groups, which allows you to group a storage device, file system, and VM. For information on configuring KVM guest nodes, see [Section 9.4.5, “Configuration Overview: KVM Guest Node”](#).

Additionally, small clarifications and corrections have been made throughout this document.

1.1.4. New and Changed Features for Red Hat Enterprise Linux 7.4

Red Hat Enterprise Linux 7.4 includes the following documentation and feature updates and changes.

- Red Hat Enterprise Linux release 7.4 provides full support for the ability to configure high availability clusters that span multiple sites through the use of a Booth cluster ticket manager. For information on the Booth cluster ticket manager, see [Chapter 14, Configuring Multi-Site Clusters with Pacemaker](#).
- Red Hat Enterprise Linux 7.4 provides full support for the ability to configure a separate quorum device which acts as a third-party arbitration device for the cluster. The primary use of this feature is to allow a cluster to sustain more node failures than standard quorum rules allow. For information on quorum devices, see [Section 10.5, “Quorum Devices”](#).
- You can now specify nodes in fencing topology by a regular expression applied on a node name and by a node attribute and its value. For information on configuring fencing levels, see [Section 5.9, “Configuring Fencing Levels”](#).
- Red Hat Enterprise Linux 7.4 supports the **NodeUtilization** resource agent, which can detect the system parameters of available CPU, host memory availability, and hypervisor memory availability and add these parameters into the CIB. For information on this resource agent, see [Section 9.6.5, “The NodeUtilization Resource Agent \(Red Hat Enterprise Linux 7.4 and later\)”](#).
- For Red Hat Enterprise Linux 7.4, the **cluster node add-guest** and the **cluster node remove-guest** commands replace the **cluster remote-node add** and **cluster remote-node remove** commands. The **pcs cluster node add-guest** command sets up the **authkey** for guest nodes and the **pcs cluster node add-remote** command sets up the **authkey** for remote nodes. For updated guest and remote node configuration procedures, see [Section 9.3, “Configuring a Virtual Domain as a Resource”](#).
- Red Hat Enterprise Linux 7.4 supports the **systemd resource-agents-deps** target. This allows you to configure the appropriate startup order for a cluster that includes resources with dependencies that are not themselves managed by the cluster, as described in [Section 9.7, “Configuring Startup Order for Resource Dependencies not Managed by Pacemaker \(Red Hat Enterprise Linux 7.4 and later\)”](#).
- The format for the command to create a resource as a master/slave clone has changed for this release. For information on creating a master/slave clone, see [Section 9.2, “Multistate Resources: Resources That Have Multiple Modes”](#).

1.1.5. New and Changed Features for Red Hat Enterprise Linux 7.5

Red Hat Enterprise Linux 7.5 includes the following documentation and feature updates and changes.

- As of Red Hat Enterprise Linux 7.5, you can use the **pcs_snmp_agent** daemon to query a Pacemaker cluster for data by means of SNMP. For information on querying a cluster with SNMP, see [Section 9.8, “Querying a Pacemaker Cluster with SNMP \(Red Hat Enterprise Linux 7.5 and later\)”](#).

1.1.6. New and Changed Features for Red Hat Enterprise Linux 7.8

Red Hat Enterprise Linux 7.8 includes the following documentation and feature updates and changes.

- As of Red Hat Enterprise Linux 7.8, you can configure Pacemaker so that when a node shuts

down cleanly, the resources attached to the node will be locked to the node and unable to start elsewhere until they start again when the node that has shut down rejoins the cluster. This allows you to power down nodes during maintenance windows when service outages are acceptable without causing that node's resources to fail over to other nodes in the cluster. For information on configuring resources to remain stopped on clean node shutdown, see [Section 9.9, “Configuring Resources to Remain Stopped on Clean Node Shutdown \(Red Hat Enterprise Linux 7.8 and later\)”](#).

1.2. INSTALLING PACEMAKER CONFIGURATION TOOLS

You can use the following **yum install** command to install the Red Hat High Availability Add-On software packages along with all available fence agents from the High Availability channel.

```
# yum install pcs pacemaker fence-agents-all
```

Alternately, you can install the Red Hat High Availability Add-On software packages along with only the fence agent that you require with the following command.

```
# yum install pcs pacemaker fence-agents-model
```

The following command displays a listing of the available fence agents.

```
# rpm -q -a | grep fence
fence-agents-rhevm-4.0.2-3.el7.x86_64
fence-agents-ilo-mp-4.0.2-3.el7.x86_64
fence-agents-ipmilan-4.0.2-3.el7.x86_64
...
...
```

The **lvm2-cluster** and **gfs2-utils** packages are part of ResilientStorage channel. You can install them, as needed, with the following command.

```
# yum install lvm2-cluster gfs2-utils
```



WARNING

After you install the Red Hat High Availability Add-On packages, you should ensure that your software update preferences are set so that nothing is installed automatically. Installation on a running cluster can cause unexpected behaviors.

1.3. CONFIGURING THE IPTABLES FIREWALL TO ALLOW CLUSTER COMPONENTS



NOTE

The ideal firewall configuration for cluster components depends on the local environment, where you may need to take into account such considerations as whether the nodes have multiple network interfaces or whether off-host firewalling is present. The example here, which opens the ports that are generally required by a Pacemaker cluster, should be modified to suit local conditions.

[Table 1.1, “Ports to Enable for High Availability Add-On”](#) shows the ports to enable for the Red Hat High Availability Add-On and provides an explanation for what the port is used for. You can enable all of these ports by means of the **firewalld** daemon by executing the following commands.

```
# firewall-cmd --permanent --add-service=high-availability
# firewall-cmd --add-service=high-availability
```

Table 1.1. Ports to Enable for High Availability Add-On

Port	When Required
TCP 2224	<p>Required on all nodes (needed by the pcsd Web UI and required for node-to-node communication)</p> <p>It is crucial to open port 2224 in such a way that pcs from any node can talk to all nodes in the cluster, including itself. When using the Booth cluster ticket manager or a quorum device you must open port 2224 on all related hosts, such as Booth arbiters or the quorum device host.</p>
TCP 3121	<p>Required on all nodes if the cluster has any Pacemaker Remote nodes</p> <p>Pacemaker's crmd daemon on the full cluster nodes will contact the pacemaker_remoted daemon on Pacemaker Remote nodes at port 3121. If a separate interface is used for cluster communication, the port only needs to be open on that interface. At a minimum, the port should open on Pacemaker Remote nodes to full cluster nodes. Because users may convert a host between a full node and a remote node, or run a remote node inside a container using the host's network, it can be useful to open the port to all nodes. It is not necessary to open the port to any hosts other than nodes.</p>
TCP 5403	Required on the quorum device host when using a quorum device with corosync-qnetd . The default value can be changed with the -p option of the corosync-qnetd command.
UDP 5404	Required on corosync nodes if corosync is configured for multicast UDP
UDP 5405	Required on all corosync nodes (needed by corosync)
TCP 21064	Required on all nodes if the cluster contains any resources requiring DLM (such as clvm or GFS2)
TCP 9929, UDP 9929	Required to be open on all cluster nodes and booth arbitrator nodes to connections from any of those same nodes when the Booth ticket manager is used to establish a multi-site cluster.

1.4. THE CLUSTER AND PACEMAKER CONFIGURATION FILES

The configuration files for the Red Hat High Availability add-on are **corosync.conf** and **cib.xml**.

The **corosync.conf** file provides the cluster parameters used by **corosync**, the cluster manager that Pacemaker is built on. In general, you should not edit the **corosync.conf** directly but, instead, use the **pcs** or **pcsd** interface. However, there may be a situation where you do need to edit this file directly. For information on editing the **corosync.conf** file, see [Editing the corosync.conf file in Red Hat Enterprise Linux 7](#).

The **cib.xml** file is an XML file that represents both the cluster's configuration and current state of all resources in the cluster. This file is used by Pacemaker's Cluster Information Base (CIB). The contents of the CIB are automatically kept in sync across the entire cluster. Do not edit the **cib.xml** file directly; use the **pcs** or **pcsd** interface instead.

1.5. CLUSTER CONFIGURATION CONSIDERATIONS

When configuring a Red Hat High Availability Add-On cluster, you must take the following considerations into account:

- Red Hat does not support cluster deployments greater than 32 nodes for RHEL 7.7 (and later). It is possible, however, to scale beyond that limit with remote nodes running the **pacemaker_remote** service. For information on the **pacemaker_remote** service, see [Section 9.4, “The pacemaker_remote Service”](#).
- The use of Dynamic Host Configuration Protocol (DHCP) for obtaining an IP address on a network interface that is utilized by the **corosync** daemons is not supported. The DHCP client can periodically remove and re-add an IP address to its assigned interface during address renewal. This will result in **corosync** detecting a connection failure, which will result in fencing activity from any other nodes in the cluster using **corosync** for heartbeat connectivity.

1.6. UPDATING A RED HAT ENTERPRISE LINUX HIGH AVAILABILITY CLUSTER

Updating packages that make up the RHEL High Availability and Resilient Storage Add-Ons, either individually or as a whole, can be done in one of two general ways:

- Rolling Updates*: Remove one node at a time from service, update its software, then integrate it back into the cluster. This allows the cluster to continue providing service and managing resources while each node is updated.
- Entire Cluster Update*: Stop the entire cluster, apply updates to all nodes, then start the cluster back up.



WARNING

It is critical that when performing software update procedures for Red Hat Enterprise Linux High Availability and Resilient Storage clusters, you ensure that any node that will undergo updates is not an active member of the cluster before those updates are initiated.

For a full description of each of these methods and the procedures to follow for the updates, see [Recommended Practices for Applying Software Updates to a RHEL High Availability or Resilient Storage Cluster](#).

1.7. ISSUES WITH LIVE MIGRATION OF VMS IN A RHEL CLUSTER

Information on support policies for RHEL high availability clusters with virtualized cluster members can be found in [Support Policies for RHEL High Availability Clusters - General Conditions with Virtualized Cluster Members](#). As noted, Red Hat does not support live migration of active cluster nodes across hypervisors or hosts. If you need to perform a live migration, you will first need to stop the cluster services on the VM to remove the node from the cluster, and then start the cluster back up after performing the migration.

The following steps outline the procedure for removing a VM from a cluster, migrating the VM, and restoring the VM to the cluster.



NOTE

Before performing this procedure, consider the effect on cluster quorum of removing a cluster node. For example, if you have a three node cluster and you remove one node, your cluster can withstand only one more node failure. If one node of a three node cluster is already down, removing a second node will lose quorum.

1. If any preparations need to be made before stopping or moving the resources or software running on the VM to migrate, perform those steps.
2. Move any managed resources off the VM. If there are specific requirements or preferences for where resources should be relocated, then consider creating new location constraints to place the resources on the correct node.
3. Place the VM in standby mode to ensure it is not considered in service, and to cause any remaining resources to be relocated elsewhere or stopped.

```
# pcs cluster standby VM
```

4. Run the following command on the VM to stop the cluster software on the VM.

```
# pcs cluster stop
```

5. Perform the live migration of the VM.
6. Start cluster services on the VM.

```
# pcs cluster start
```

7. Take the VM out of standby mode.
8. If you created any temporary location constraints before putting the VM in standby mode, adjust or remove those constraints to allow resources to go back to their normally preferred locations.

```
# pcs cluster unstandby VM
```

CHAPTER 2. THE PCSD WEB UI

This chapter provides an overview of configuring a Red Hat High Availability cluster with the **pcsd** Web UI.

2.1. PCSD WEB UI SETUP

To set up your system to use the **pcsd** Web UI to configure a cluster, use the following procedure.

1. Install the Pacemaker configuration tools, as described in [Section 1.2, “Installing Pacemaker configuration tools”](#).
2. On each node that will be part of the cluster, use the **passwd** command to set the password for user **hacluster**, using the same password on each node.
3. Start and enable the **pcsd** daemon on each node:

```
# systemctl start pcsd.service  
# systemctl enable pcsd.service
```

4. On one node of the cluster, authenticate the nodes that will constitute the cluster with the following command. After executing this command, you will be prompted for a **Username** and a **Password**. Specify **hacluster** as the **Username**.

```
# pcs cluster auth node1 node2 ... nodeN
```

5. On any system, open a browser to the following URL, specifying one of the nodes you have authorized (note that this uses the **https** protocol). This brings up the **pcsd** Web UI login screen.

```
https://nodename:2224
```

6. Log in as user **hacluster**. This brings up the **Manage Clusters** page as shown in [Figure 2.1, “Manage Clusters page”](#).




Figure 2.1. Manage Clusters page

2.2. CREATING A CLUSTER WITH THE PCSD WEB UI

From the **Manage Clusters** page, you can create a new cluster, add an existing cluster to the Web UI, or remove a cluster from the Web UI.

- To create a cluster, click on **Create New** and enter the name of the cluster to create and the nodes that constitute the cluster. You can also configure advanced cluster options from this screen, including the transport mechanism for cluster communication, as described in [Section 2.2.1, “Advanced Cluster Configuration Options”](#). After entering the cluster information, click **Create Cluster**.
- To add an existing cluster to the Web UI, click on **Add Existing** and enter the host name or IP address of a node in the cluster that you would like to manage with the Web UI.

Once you have created or added a cluster, the cluster name is displayed on the **Manage Cluster** page. Selecting the cluster displays information about the cluster.



NOTE

When using the **pcsd** Web UI to configure a cluster, you can move your mouse over the text describing many of the options to see longer descriptions of those options as a **tooltip** display.

2.2.1. Advanced Cluster Configuration Options

When creating a cluster, you can click on **Advanced Options** to configure additional cluster options, as shown in [Figure 2.2, “Create Clusters page”](#). For information about the options displayed, move your mouse over the text for that option.

Note that you can configure a cluster with Redundant Ring Protocol by specifying the interfaces for each node. The Redundant Ring Protocol settings display will change if you select **UDP** rather than the default value of **UDPU** as the transport mechanism for the cluster.

Create Cluster

Enter the hostnames of the nodes you would like to use to create a cluster:

Cluster Name:

Node 1:

Node 2:

Node 3:

[More nodes...](#)

▼ Advanced Options:

Transport: **UDPU (default)**

Wait for All:

Auto Tie Breaker:

Last Man Standing:

Last Man Standing Window: ms

Use IPv6:

Token Timeout: ms

Token Timeout Coefficient: ms

Join Timeout: ms

Consensus Timeout: ms

Missed Messages Count:

Failures Count:

Redundant Ring Protocol settings for UDPU transport:

Node 1 (Ring 1):

Node 2 (Ring 1):

Node 3 (Ring 1):

Create Cluster **Cancel**

Figure 2.2. Create Clusters page

2.2.2. Setting Cluster Management Permissions

There are two sets of cluster permissions that you can grant to users:

- Permissions for managing the cluster with the Web UI, which also grants permissions to run **pcs** commands that connect to nodes over a network. This section describes how to configure those permissions with the Web UI.
- Permissions for local users to allow read-only or read-write access to the cluster configuration, using ACLs. Configuring ACLs with the Web UI is described in [Section 2.3.4, “Configuring ACLs”](#).

For further information on user permissions, see [Section 4.5, “Setting User Permissions”](#).

You can grant permission for specific users other than user **hacluster** to manage the cluster through the Web UI and to run **pcs** commands that connect to nodes over a network by adding them to the group **haclient**. You can then configure the permissions set for an individual member of the group **haclient** by clicking the **Permissions** tab on the **Manage Clusters** page and setting the permissions on the resulting screen. From this screen, you can also set permissions for groups.

You can grant the following permissions:

- Read permissions, to view the cluster settings
- Write permissions, to modify cluster settings (except for permissions and ACLs)
- Grant permissions, to modify cluster permissions and ACLs
- Full permissions, for unrestricted access to a cluster, including adding and removing nodes, with access to keys and certificates

2.3. CONFIGURING CLUSTER COMPONENTS

To configure the components and attributes of a cluster, click on the name of the cluster displayed on the **Manage Clusters** screen. This brings up the **Nodes** page, as described in [Section 2.3.1, “Cluster Nodes”](#). This page displays a menu along the top of the page, as shown in [Figure 2.3, “Cluster Components Menu”](#), with the following entries:

- **Nodes**, as described in [Section 2.3.1, “Cluster Nodes”](#)
- **Resources**, as described in [Section 2.3.2, “Cluster Resources”](#)
- **Fence Devices**, as described in [Section 2.3.3, “Fence Devices”](#)
- **ACLs**, as described in [Section 2.3.4, “Configuring ACLs”](#)
- **Cluster Properties**, as described in [Section 2.3.5, “Cluster Properties”](#)



Figure 2.3. Cluster Components Menu

2.3.1. Cluster Nodes

Selecting the **Nodes** option from the menu along the top of the cluster management page displays the

currently configured nodes and the status of the currently selected node, including which resources are running on the node and the resource location preferences. This is the default page that displays when you select a cluster from the **Manage Clusters** screen.

You can add or remove nodes from this page, and you can start, stop, restart, or put a node in standby mode. For information on standby mode, see [Section 4.4.5, "Standby Mode"](#).

You can also configure fence devices directly from this page, as described in [Section 2.3.3, "Fence Devices"](#). by selecting **Configure Fencing**.

2.3.2. Cluster Resources

Selecting the **Resources** option from the menu along the top of the cluster management page displays the currently configured resources for the cluster, organized according to resource groups. Selecting a group or a resource displays the attributes of that group or resource.

From this screen, you can add or remove resources, you can edit the configuration of existing resources, and you can create a resource group.

To add a new resource to the cluster, click **Add**. This brings up the **Add Resource** screen. When you select a resource type from the dropdown **Type** menu, the arguments you must specify for that resource appear in the menu. You can click **Optional Arguments** to display additional arguments you can specify for the resource you are defining. After entering the parameters for the resource you are creating, click **Create Resource**.

When configuring the arguments for a resource, a brief description of the argument appears in the menu. If you move the cursor to the field, a longer help description of that argument is displayed.

You can define a resource as a cloned resource, or as a master/slave resource. For information on these resource types, see [Chapter 9, Advanced Configuration](#).

Once you have created at least one resource, you can create a resource group. For information on resource groups, see [Section 6.5, "Resource Groups"](#).

To create a resource group, select a resource that will be part of the group from the **Resources** screen, then click **Create Group**. This displays the **Create Group** screen. Enter a group name and click **Create Group**. This returns you to the **Resources** screen, which now displays the group name for the resource. After you have created a resource group, you can indicate that group name as a resource parameter when you create or modify additional resources.

2.3.3. Fence Devices

Selecting the **Fence Devices** option from the menu along the top of the cluster management page displays **Fence Devices** screen, showing the currently configured fence devices.

To add a new fence device to the cluster, click **Add**. This brings up the **Add Fence Device** screen. When you select a fence device type from the drop-down **Type** menu, the arguments you must specify for that fence device appear in the menu. You can click on **Optional Arguments** to display additional arguments you can specify for the fence device you are defining. After entering the parameters for the new fence device, click **Create Fence Instance**.

For information on configuring fence devices with Pacemaker, see [Chapter 5, Fencing: Configuring STONITH](#).

2.3.4. Configuring ACLs

Selecting the **ACLs** option from the menu along the top of the cluster management page displays a screen from which you can set permissions for local users, allowing read-only or read-write access to the cluster configuration by using access control lists (ACLs).

To assign ACL permissions, you create a role and define the access permissions for that role. Each role can have an unlimited number of permissions (read/write/deny) applied to either an XPath query or the ID of a specific element. After defining the role, you can assign it to an existing user or group.

2.3.5. Cluster Properties

Selecting the **Cluster Properties** option from the menu along the top of the cluster management page displays the cluster properties and allows you to modify these properties from their default values. For information on the Pacemaker cluster properties, see [Chapter 12, Pacemaker Cluster Properties](#).

2.4. CONFIGURING A HIGH AVAILABILITY PCSD WEB UI

When you use the **pcsd** Web UI, you connect to one of the nodes of the cluster to display the cluster management pages. If the node to which you are connecting goes down or becomes unavailable, you can reconnect to the cluster by opening your browser to a URL that specifies a different node of the cluster. It is possible, however, to configure the **pcsd** Web UI itself for high availability, in which case you can continue to manage the cluster without entering a new URL.

To configure the **pcsd** Web UI for high availability, perform the following steps.

1. Ensure that **PCSD_SSL_CERT_SYNC_ENABLED** is set to **true** in the **/etc/sysconfig/pcsd** configuration file, which is the default value in RHEL 7. Enabling certificate syncing causes **pcsd** to sync the **pcsd** certificates for the cluster setup and node add commands.
2. Create an **IPaddr2** cluster resource, which is a floating IP address that you will use to connect to the **pcsd** Web UI. The IP address must not be one already associated with a physical node. If the **IPaddr2** resource's NIC device is not specified, the floating IP must reside on the same network as one of the node's statically assigned IP addresses, otherwise the NIC device to assign the floating IP address cannot be properly detected.
3. Create custom SSL certificates for use with **pcsd** and ensure that they are valid for the addresses of the nodes used to connect to the **pcsd** Web UI.
 1. To create custom SSL certificates, you can use either wildcard certificates or you can use the Subject Alternative Name certificate extension. For information on the Red Hat Certificate System, see the [Red Hat Certificate System Administration Guide](#) .
 2. Install the custom certificates for **pcsd** with the **pcs pcسد certkey** command.
 3. Sync the **pcsd** certificates to all nodes in the cluster with the **pcs pcسد sync-certificates** command.
4. Connect to the **pcsd** Web UI using the floating IP address you configured as a cluster resource.



NOTE

Even when you configure the **pcsd** Web UI for high availability, you will be asked to log in again when the node to which you are connecting goes down.

CHAPTER 3. THE PCS COMMAND LINE INTERFACE

The **pcs** command line interface controls and configures **corosync** and Pacemaker by providing an interface to the **corosync.conf** and **cib.xml** files.

The general format of the **pcs** command is as follows.

```
pcs [-f file] [-h] [commands]...
```

3.1. THE PCS COMMANDS

The **pcs** commands are as follows.

- **cluster**

Configure cluster options and nodes. For information on the **pcs cluster** command, see [Chapter 4, Cluster Creation and Administration](#).

- **resource**

Create and manage cluster resources. For information on the **pcs cluster** command, see [Chapter 6, Configuring Cluster Resources](#), [Chapter 8, Managing Cluster Resources](#), and [Chapter 9, Advanced Configuration](#).

- **stonith**

Configure fence devices for use with Pacemaker. For information on the **pcs stonith** command, see [Chapter 5, Fencing: Configuring STONITH](#).

- **constraint**

Manage resource constraints. For information on the **pcs constraint** command, see [Chapter 7, Resource Constraints](#).

- **property**

Set Pacemaker properties. For information on setting properties with the **pcs property** command, see [Chapter 12, Pacemaker Cluster Properties](#).

- **status**

View current cluster and resource status. For information on the **pcs status** command, see [Section 3.5, “Displaying Status”](#).

- **config**

Display complete cluster configuration in user readable form. For information on the **pcs config** command, see [Section 3.6, “Displaying the Full Cluster Configuration”](#).

3.2. PCS USAGE HELP DISPLAY

You can use the **-h** option of **pcs** to display the parameters of a **pcs** command and a description of those parameters. For example, the following command displays the parameters of the **pcs resource** command. Only a portion of the output is shown.

```
# pcs resource -h
Usage: pcs resource [commands]...
Manage pacemaker resources
Commands:
  show [resource id] [--all]
    Show all currently configured resources or if a resource is specified
    show the options for the configured resource. If --all is specified
    resource options will be displayed

  start <resource id>
    Start resource specified by resource_id
  ...

```

3.3. VIEWING THE RAW CLUSTER CONFIGURATION

Although you should not edit the cluster configuration file directly, you can view the raw cluster configuration with the **pcs cluster cib** command.

You can save the raw cluster configuration to a specified file with the **pcs cluster cib filename** command as described in [Section 3.4, “Saving a Configuration Change to a File”](#).

3.4. SAVING A CONFIGURATION CHANGE TO A FILE

When using the **pcs** command, you can use the **-f** option to save a configuration change to a file without affecting the active CIB.

If you have previously configured a cluster and there is already an active CIB, you use the following command to save the raw xml file.

```
pcs cluster cib filename
```

For example, the following command saves the raw xml from the CIB into a file name **testfile**.

```
pcs cluster cib testfile
```

The following command creates a resource in the file **testfile1** but does not add that resource to the currently running cluster configuration.

```
# pcs -f testfile1 resource create VirtualIP ocf:heartbeat:IPaddr2 ip=192.168.0.120 cidr_netmask=24
  op monitor interval=30s
```

You can push the current content of **testfile** to the CIB with the following command.

```
pcs cluster cib-push filename
```

3.5. DISPLAYING STATUS

You can display the status of the cluster and the cluster resources with the following command.

```
pcs status commands
```

If you do not specify a *commands* parameter, this command displays all information about the cluster and the resources. You display the status of only particular cluster components by specifying **resources**, **groups**, **cluster**, **nodes**, or **pcsd**.

3.6. DISPLAYING THE FULL CLUSTER CONFIGURATION

Use the following command to display the full current cluster configuration.

```
pcs config
```

3.7. DISPLAYING THE CURRENT PCS VERSION

The following command displays the current version of **pcs** that is running.

```
pcs --version
```

3.8. BACKING UP AND RESTORING A CLUSTER CONFIGURATION

As of the Red Hat Enterprise Linux 7.1 release, you can back up the cluster configuration in a tarball with the following command. If you do not specify a file name, the standard output will be used.

```
pcs config backup filename
```

Use the following command to restore the cluster configuration files on all nodes from the backup. If you do not specify a file name, the standard input will be used. Specifying the **--local** option restores only the files on the current node.

```
pcs config restore [--local] [filename]
```

CHAPTER 4. CLUSTER CREATION AND ADMINISTRATION

This chapter describes how to perform basic cluster administration with Pacemaker, including creating the cluster, managing the cluster components, and displaying cluster status.

4.1. CLUSTER CREATION

To create a running cluster, perform the following steps:

1. Start the **pcsd** on each node in the cluster.
2. Authenticate the nodes that will constitute the cluster.
3. Configure and sync the cluster nodes.
4. Start cluster services on the cluster nodes.

The following sections described the commands that you use to perform these steps.

4.1.1. Starting the **pcsd** daemon

The following commands start the **pcsd** service and enable **pcsd** at system start. These commands should be run on each node in the cluster.

```
# systemctl start pcsd.service
# systemctl enable pcsd.service
```

4.1.2. Authenticating the Cluster Nodes

The following command authenticates **pcs** to the **pcs** daemon on the nodes in the cluster.

- The user name for the **pcs** administrator must be **hacluster** on every node. It is recommended that the password for user **hacluster** be the same on each node.
- If you do not specify **username** or **password**, the system will prompt you for those parameters for each node when you execute the command.
- If you do not specify any nodes, this command will authenticate **pcs** on the nodes that are specified with a **pcs cluster setup** command, if you have previously executed that command.

```
pcs cluster auth [node] [...] [-u username] [-p password]
```

For example, the following command authenticates user **hacluster** on **z1.example.com** for both of the nodes in the cluster that consist of **z1.example.com** and **z2.example.com**. This command prompts for the password for user **hacluster** on the cluster nodes.

```
root@z1 ~]# pcs cluster auth z1.example.com z2.example.com
Username: hacluster
Password:
z1.example.com: Authorized
z2.example.com: Authorized
```

Authorization tokens are stored in the file **~/.pcs/tokens** (or **/var/lib/pcsd/tokens**).

4.1.3. Configuring and Starting the Cluster Nodes

The following command configures the cluster configuration file and syncs the configuration to the specified nodes.

- If you specify the **--start** option, the command will also start the cluster services on the specified nodes. If necessary, you can also start the cluster services with a separate **pcs cluster start** command.

When you create a cluster with the **pcs cluster setup --start** command or when you start cluster services with the **pcs cluster start** command, there may be a slight delay before the cluster is up and running. Before performing any subsequent actions on the cluster and its configuration, it is recommended that you use the **pcs cluster status** command to be sure that the cluster is up and running.

- If you specify the **--local** option, the command will perform changes on the local node only.

```
pcs cluster setup [--start] [--local] --name cluster_name node1 [node2] [...]
```

The following command starts cluster services on the specified node or nodes.

- If you specify the **--all** option, the command starts cluster services on all nodes.
- If you do not specify any nodes, cluster services are started on the local node only.

```
pcs cluster start [--all] [node] [...]
```

4.2. CONFIGURING TIMEOUT VALUES FOR A CLUSTER

When you create a cluster with the **pcs cluster setup** command, timeout values for the cluster are set to default values that should be suitable for most cluster configurations. If your system requires different timeout values, however, you can modify these values with the **pcs cluster setup** options summarized in [Table 4.1, “Timeout Options”](#)

Table 4.1. Timeout Options

Option	Description
--token <i>timeout</i>	Sets time in milliseconds until a token loss is declared after not receiving a token (default 1000 ms)
--join <i>timeout</i>	sets time in milliseconds to wait for join messages (default 50 ms)
--consensus <i>timeout</i>	sets time in milliseconds to wait for consensus to be achieved before starting a new round of membership configuration (default 1200 ms)
--miss_count_const <i>count</i>	sets the maximum number of times on receipt of a token a message is checked for retransmission before a retransmission occurs (default 5 messages)

Option	Description
--fail_recv_const <i>failures</i>	specifies how many rotations of the token without receiving any messages when messages should be received may occur before a new configuration is formed (default 2500 failures)

For example, the following command creates the cluster **new_cluster** and sets the token timeout value to 10000 milliseconds (10 seconds) and the join timeout value to 100 milliseconds.

```
# pcs cluster setup --name new_cluster nodeA nodeB --token 10000 --join 100
```

4.3. CONFIGURING REDUNDANT RING PROTOCOL (RRP)



NOTE

Red Hat supports the configuration of Redundant Ring Protocol (RRP) in clusters subject to the conditions described in the "Redundant Ring Protocol (RRP)" section of [Support Policies for RHEL High Availability Clusters – Cluster Interconnect Network Interfaces](#).

When you create a cluster with the **pcs cluster setup** command, you can configure a cluster with Redundant Ring Protocol by specifying both interfaces for each node. When using the default **udpu** transport, when you specify the cluster nodes you specify the ring 0 address followed by a **,**, then the ring 1 address.

For example, the following command configures a cluster named **my_rrp_clusterM** with two nodes, node A and node B. Node A has two interfaces, **nodeA-0** and **nodeA-1**. Node B has two interfaces, **nodeB-0** and **nodeB-1**. To configure these nodes as a cluster using RRP, execute the following command.

```
# pcs cluster setup --name my_rrp_cluster nodeA-0,nodeA-1 nodeB-0,nodeB-1
```

For information on configuring RRP in a cluster that uses **udp** transport, see the help screen for the **pcs cluster setup** command.

4.4. MANAGING CLUSTER NODES

The following sections describe the commands you use to manage cluster nodes, including commands to start and stop cluster services and to add and remove cluster nodes.

4.4.1. Stopping Cluster Services

The following command stops cluster services on the specified node or nodes. As with the **pcs cluster start**, the **--all** option stops cluster services on all nodes and if you do not specify any nodes, cluster services are stopped on the local node only.

```
pcs cluster stop [--all] [node] [...]
```

You can force a stop of cluster services on the local node with the following command, which performs a **kill -9** command.

```
pcs cluster kill
```

4.4.2. Enabling and Disabling Cluster Services

Use the following command to configure the cluster services to run on startup on the specified node or nodes.

- If you specify the **--all** option, the command enables cluster services on all nodes.
- If you do not specify any nodes, cluster services are enabled on the local node only.

```
pcs cluster enable [--all] [node] [...]
```

Use the following command to configure the cluster services not to run on startup on the specified node or nodes.

- If you specify the **--all** option, the command disables cluster services on all nodes.
- If you do not specify any nodes, cluster services are disabled on the local node only.

```
pcs cluster disable [--all] [node] [...]
```

4.4.3. Adding Cluster Nodes



NOTE

It is highly recommended that you add nodes to existing clusters only during a production maintenance window. This allows you to perform appropriate resource and deployment testing for the new node and its fencing configuration.

Use the following procedure to add a new node to an existing cluster. In this example, the existing cluster nodes are **clusternode-01.example.com**, **clusternode-02.example.com**, and **clusternode-03.example.com**. The new node is **newnode.example.com**.

On the new node to add to the cluster, perform the following tasks.

1. Install the cluster packages. If the cluster uses SBD, the Booth ticket manager, or a quorum device, you must manually install the respective packages (**sbd**, **booth-site**, **corosync-qdevice**) on the new node as well.

```
[root@newnode ~]# yum install -y pcs fence-agents-all
```

2. If you are running the **firewalld** daemon, execute the following commands to enable the ports that are required by the Red Hat High Availability Add-On.

```
# firewall-cmd --permanent --add-service=high-availability
# firewall-cmd --add-service=high-availability
```

3. Set a password for the user ID **hacluster**. It is recommended that you use the same password for each node in the cluster.

```
[root@newnode ~]# passwd hacluster
Changing password for user hacluster.
New password:
Retype new password:
passwd: all authentication tokens updated successfully.
```

4. Execute the following commands to start the **pcsd** service and to enable **pcsd** at system start.

```
# systemctl start pcsd.service
# systemctl enable pcsd.service
```

On a node in the existing cluster, perform the following tasks.

1. Authenticate user **hacluster** on the new cluster node.

```
[root@clusternode-01 ~]# pcs cluster auth newnode.example.com
Username: hacluster
Password:
newnode.example.com: Authorized
```

2. Add the new node to the existing cluster. This command also syncs the cluster configuration file **corosync.conf** to all nodes in the cluster, including the new node you are adding.

```
[root@clusternode-01 ~]# pcs cluster node add newnode.example.com
```

On the new node to add to the cluster, perform the following tasks.

1. Start and enable cluster services on the new node.

```
[root@newnode ~]# pcs cluster start
Starting Cluster...
[root@newnode ~]# pcs cluster enable
```

2. Ensure that you configure and test a fencing device for the new cluster node. For information on configuring fencing devices, see [Chapter 5, Fencing: Configuring STONITH](#).

4.4.4. Removing Cluster Nodes

The following command shuts down the specified node and removes it from the cluster configuration file, **corosync.conf**, on all of the other nodes in the cluster. For information on removing all information about the cluster from the cluster nodes entirely, thereby destroying the cluster permanently, see [Section 4.6, “Removing the Cluster Configuration”](#).

```
pcs cluster node remove node
```

4.4.5. Standby Mode

The following command puts the specified node into standby mode. The specified node is no longer able to host resources. Any resources currently active on the node will be moved to another node. If you specify the **--all**, this command puts all nodes into standby mode.

You can use this command when updating a resource's packages. You can also use this command when testing a configuration, to simulate recovery without actually shutting down a node.

```
pcs cluster standby node | --all
```

The following command removes the specified node from standby mode. After running this command, the specified node is then able to host resources. If you specify the **--all**, this command removes all nodes from standby mode.

```
pcs cluster unstandby node | --all
```

Note that when you execute the **pcs cluster standby** command, this prevents resources from running on the indicated node. When you execute the **pcs cluster unstandby** command, this allows resources to run on the indicated node. This does not necessarily move the resources back to the indicated node; where the resources can run at that point depends on how you have configured your resources initially. For information on resource constraints, see [Chapter 7, Resource Constraints](#).

4.5. SETTING USER PERMISSIONS

You can grant permission for specific users other than user **hacluster** to manage the cluster. There are two sets of permissions that you can grant to individual users:

- Permissions that allow individual users to manage the cluster through the Web UI and to run **pcs** commands that connect to nodes over a network, as described in [Section 4.5.1, "Setting Permissions for Node Access Over a Network"](#). Commands that connect to nodes over a network include commands to set up a cluster, or to add or remove nodes from a cluster.
- Permissions for local users to allow read-only or read-write access to the cluster configuration, as described in [Section 4.5.2, "Setting Local Permissions Using ACLs"](#). Commands that do not require connecting over a network include commands that edit the cluster configuration, such as those that create resources and configure constraints.

In situations where both sets of permissions have been assigned, the permissions for commands that connect over a network are applied first, and then permissions for editing the cluster configuration on the local node are applied. Most **pcs** commands do not require network access and in those cases the network permissions will not apply.

4.5.1. Setting Permissions for Node Access Over a Network

To grant permission for specific users to manage the cluster through the Web UI and to run **pcs** commands that connect to nodes over a network, add those users to the group **haclient**. You can then use the Web UI to grant permissions for those users, as described in [Section 2.2.2, "Setting Cluster Management Permissions"](#).

4.5.2. Setting Local Permissions Using ACLs

As of Red Hat Enterprise Linux 7.1, you can use the **pcs acl** command to set permissions for local users to allow read-only or read-write access to the cluster configuration by using access control lists (ACLs). You can also configure ACLs using the **pcsd** Web UI, as described in [Section 2.3.4, "Configuring ACLs"](#). By default, the root user and any user who is a member of the group **haclient** has full local read/write access to the cluster configuration.

Setting permissions for local users is a two step process:

1. Execute the **pcs acl role create...** command to create a *role* which defines the permissions for that role.
2. Assign the role you created to a user with the **pcs acl user create** command.

The following example procedure provides read-only access for a cluster configuration to a local user named **rouser**.

1. This procedure requires that the user **rouser** exists on the local system and that the user **rouser** is a member of the group **haclient**.

```
# adduser rouser
# usermod -a -G haclient rouser
```

2. Enable Pacemaker ACLs with the **enable-acl** cluster property.

```
# pcs property set enable-acl=true --force
```

3. Create a role named **read-only** with read-only permissions for the cib.

```
# pcs acl role create read-only description="Read access to cluster" read xpath /cib
```

4. Create the user **rouser** in the pcs ACL system and assign that user the **read-only** role.

```
# pcs acl user create rouser read-only
```

5. View the current ACLs.

```
# pcs acl
User: rouser
Roles: read-only
Role: read-only
Description: Read access to cluster
Permission: read xpath /cib (read-only-read)
```

The following example procedure provides write access for a cluster configuration to a local user named **wuser**.

1. This procedure requires that the user **wuser** exists on the local system and that the user **wuser** is a member of the group **haclient**.

```
# adduser wuser
# usermod -a -G haclient wuser
```

2. Enable Pacemaker ACLs with the **enable-acl** cluster property.

```
# pcs property set enable-acl=true --force
```

3. Create a role named **write-access** with write permissions for the cib.

```
# pcs acl role create write-access description="Full access" write xpath /cib
```

4. Create the user **wuser** in the pcs ACL system and assign that user the **write-access** role.

```
# pcs acl user create wuser write-access
```

5. View the current ACLs.

```
# pcs acl
User: rouser
Roles: read-only
User: wuser
Roles: write-access
Role: read-only
Description: Read access to cluster
Permission: read xpath /cib (read-only-read)
Role: write-access
Description: Full Access
Permission: write xpath /cib (write-access-write)
```

For further information about cluster ACLs, see the help screen for the **pcs acl** command.

4.6. REMOVING THE CLUSTER CONFIGURATION

To remove all cluster configuration files and stop all cluster services, thus permanently destroying a cluster, use the following command.



WARNING

This command permanently removes any cluster configuration that has been created. It is recommended that you run **pcs cluster stop** before destroying the cluster.

```
pcs cluster destroy
```

4.7. DISPLAYING CLUSTER STATUS

The following command displays the current status of the cluster and the cluster resources.

```
pcs status
```

You can display a subset of information about the current status of the cluster with the following commands.

The following command displays the status of the cluster, but not the cluster resources.

```
pcs cluster status
```

The following command displays the status of the cluster resources.

```
pcs status resources
```

4.8. CLUSTER MAINTENANCE

In order to perform maintenance on the nodes of your cluster, you may need to stop or move the resources and services running on that cluster. Or you may need to stop the cluster software while leaving the services untouched. Pacemaker provides a variety of methods for performing system maintenance.

- If you need to stop a node in a cluster while continuing to provide the services running on that cluster on another node, you can put the cluster node in standby mode. A node that is in standby mode is no longer able to host resources. Any resource currently active on the node will be moved to another node, or stopped if no other node is eligible to run the resource.

For information on standby mode, see [Section 4.4.5, “Standby Mode”](#).

- If you need to move an individual resource off the node on which it is currently running without stopping that resource, you can use the **pcs resource move** command to move the resource to a different node. For information on the **pcs resource move** command, see [Section 8.1, “Manually Moving Resources Around the Cluster”](#).

When you execute the **pcs resource move** command, this adds a constraint to the resource to prevent it from running on the node on which it is currently running. When you are ready to move the resource back, you can execute the **pcs resource clear** or the **pcs constraint delete** command to remove the constraint. This does not necessarily move the resources back to the original node, however, since where the resources can run at that point depends on how you have configured your resources initially. You can relocate a resource to a specified node with the **pcs resource relocate run** command, as described in [Section 8.1.1, “Moving a Resource from its Current Node”](#).

- If you need to stop a running resource entirely and prevent the cluster from starting it again, you can use the **pcs resource disable** command. For information on the **pcs resource disable** command, see [Section 8.4, “Enabling, Disabling, and Banning Cluster Resources”](#).
- If you want to prevent Pacemaker from taking any action for a resource (for example, if you want to disable recovery actions while performing maintenance on the resource, or if you need to reload the `/etc/sysconfig/pacemaker` settings), use the **pcs resource unmanage** command, as described in [Section 8.6, “Managed Resources”](#). Pacemaker Remote connection resources should never be unmanaged.
- If you need to put the cluster in a state where no services will be started or stopped, you can set the **maintenance-mode** cluster property. Putting the cluster into maintenance mode automatically unmanages all resources. For information on setting cluster properties, see [Table 12.1, “Cluster Properties”](#).
- If you need to perform maintenance on a Pacemaker remote node, you can remove that node from the cluster by disabling the remote node resource, as described in [Section 9.4.8, “System Upgrades and pacemaker_remote”](#).

CHAPTER 5. FENCING: CONFIGURING STONITH

STONITH is an acronym for "Shoot The Other Node In The Head" and it protects your data from being corrupted by rogue nodes or concurrent access.

Just because a node is unresponsive, this does not mean it is not accessing your data. The only way to be 100% sure that your data is safe, is to fence the node using STONITH so we can be certain that the node is truly offline, before allowing the data to be accessed from another node.

STONITH also has a role to play in the event that a clustered service cannot be stopped. In this case, the cluster uses STONITH to force the whole node offline, thereby making it safe to start the service elsewhere.

For more complete general information on fencing and its importance in a Red Hat High Availability cluster, see [Fencing in a Red Hat High Availability Cluster](#).

5.1. AVAILABLE STONITH (FENCING) AGENTS

Use the following command to view of list of all available STONITH agents. You specify a filter, then this command displays only the STONITH agents that match the filter.

```
pcs stonith list [filter]
```

5.2. GENERAL PROPERTIES OF FENCING DEVICES

Any cluster node can fence any other cluster node with any fence device, regardless of whether the fence resource is started or stopped. Whether the resource is started controls only the recurring monitor for the device, not whether it can be used, with the following exceptions:

- You can disable a fencing device by running the **pcs stonith disable stonith_id** command. This will prevent any node from using that device
- To prevent a specific node from using a fencing device, you can configure location constraints for the fencing resource with the **pcs constraint location ... avoids** command.
- Configuring **stonith-enabled=false** will disable fencing altogether. Note, however, that Red Hat does not support clusters when fencing is disabled, as it is not suitable for a production environment.

[Table 5.1, “General Properties of Fencing Devices”](#) describes the general properties you can set for fencing devices. Refer to [Section 5.3, “Displaying Device-Specific Fencing Options”](#) for information on fencing properties you can set for specific fencing devices.



NOTE

For information on more advanced fencing configuration properties, see [Section 5.8, “Additional Fencing Configuration Options”](#)

Table 5.1. General Properties of Fencing Devices

Field	Type	Default	Description
pcmk_host_map	string		A mapping of host names to port numbers for devices that do not support host names. For example: node1:1;node2:2,3 tells the cluster to use port 1 for node1 and ports 2 and 3 for node2
pcmk_host_list	string		A list of machines controlled by this device (Optional unless pcmk_host_check=static-list).
pcmk_host_check	string	dynamic-list	How to determine which machines are controlled by the device. Allowed values: dynamic-list (query the device), static-list (check the pcmk_host_list attribute), none (assume every device can fence every machine)

5.3. DISPLAYING DEVICE-SPECIFIC FENCING OPTIONS

Use the following command to view the options for the specified STONITH agent.

```
pcs stonith describe stonith_agent
```

For example, the following command displays the options for the fence agent for APC over telnet/SSH.

```
# pcs stonith describe fence_apc
Stonith options for: fence_apc
ipaddr (required): IP Address or Hostname
login (required): Login Name
passwd: Login password or passphrase
passwd_script: Script to retrieve password
cmd_prompt: Force command prompt
secure: SSH connection
port (required): Physical plug number or name of virtual machine
identity_file: Identity file for ssh
switch: Physical switch number on device
inet4_only: Forces agent to use IPv4 addresses only
inet6_only: Forces agent to use IPv6 addresses only
ipport: TCP port to use for connection with device
action (required): Fencing Action
verbose: Verbose mode
debug: Write debug information to given file
version: Display version information and exit
help: Display help and exit
separator: Separator for CSV created by operation list
power_timeout: Test X seconds for status change after ON/OFF
shell_timeout: Wait X seconds for cmd prompt after issuing command
login_timeout: Wait X seconds for cmd prompt after login
```

power_wait: Wait X seconds after issuing ON/OFF
delay: Wait X seconds before fencing is started
retry_on: Count of attempts to retry power on



WARNING

For fence agents that provide a **method** option, a value of **cycle** is unsupported and should not be specified, as it may cause data corruption.

5.4. CREATING A FENCING DEVICE

The following command creates a stonith device.

```
pcs stonith create stonith_id stonith_device_type [stonith_device_options]  

# pcs stonith create MyStonith fence_virt pcmk_host_list=f1 op monitor interval=30s
```

Some fence devices can fence only a single node, while other devices can fence multiple nodes. The parameters you specify when you create a fencing device depend on what your fencing device supports and requires.

- Some fence devices can automatically determine what nodes they can fence.
- You can use the **pcmk_host_list** parameter when creating a fencing device to specify all of the machines that are controlled by that fencing device.
- Some fence devices require a mapping of host names to the specifications that the fence device understands. You can map host names with the **pcmk_host_map** parameter when creating a fencing device.

For information on the **pcmk_host_list** and **pcmk_host_map** parameters, see [Table 5.1, “General Properties of Fencing Devices”](#).

After configuring a fence device, it is imperative that you test the device to ensure that it is working correctly. For information on testing fence devices, see [Section 5.12, “Testing a Fence Device”](#).

5.5. DISPLAYING FENCING DEVICES

The following command shows all currently configured fencing devices. If a *stonith_id* is specified, the command shows the options for that configured stonith device only. If the **--full** option is specified, all configured stonith options are displayed.

```
pcs stonith show [stonith_id] [--full]
```

5.6. MODIFYING AND DELETING FENCING DEVICES

Use the following command to modify or add options to a currently configured fencing device.

`pcs stonith update stonith_id [stonith_device_options]`

Use the following command to remove a fencing device from the current configuration.

`pcs stonith delete stonith_id`

5.7. MANAGING NODES WITH FENCE DEVICES

You can fence a node manually with the following command. If you specify **--off** this will use the **off** API call to stonith which will turn the node off instead of rebooting it.

`pcs stonith fence node [--off]`

In a situation where no stonith device is able to fence a node even if it is no longer active, the cluster may not be able to recover the resources on the node. If this occurs, after manually ensuring that the node is powered down you can enter the following command to confirm to the cluster that the node is powered down and free its resources for recovery.



WARNING

If the node you specify is not actually off, but running the cluster software or services normally controlled by the cluster, data corruption/cluster failure will occur.

`pcs stonith confirm node`

5.8. ADDITIONAL FENCING CONFIGURATION OPTIONS

[Table 5.2, “Advanced Properties of Fencing Devices”](#) summarizes additional properties you can set for fencing devices. Note that these properties are for advanced use only.

Table 5.2. Advanced Properties of Fencing Devices

Field	Type	Default	Description
pcmk_host_argument	string	port	An alternate parameter to supply instead of port. Some devices do not support the standard port parameter or may provide additional ones. Use this to specify an alternate, device-specific, parameter that should indicate the machine to be fenced. A value of none can be used to tell the cluster not to supply any additional parameters.

Field	Type	Default	Description
pcmk_reboot_action	string	reboot	An alternate command to run instead of reboot . Some devices do not support the standard commands or may provide additional ones. Use this to specify an alternate, device-specific, command that implements the reboot action.
pcmk_reboot_timeout	time	60s	Specify an alternate timeout to use for reboot actions instead of stonith-timeout . Some devices need much more/less time to complete than normal. Use this to specify an alternate, device-specific, timeout for reboot actions.
pcmk_reboot_retries	integer	2	The maximum number of times to retry the reboot command within the timeout period. Some devices do not support multiple connections. Operations may fail if the device is busy with another task so Pacemaker will automatically retry the operation, if there is time remaining. Use this option to alter the number of times Pacemaker retries reboot actions before giving up.
pcmk_off_action	string	off	An alternate command to run instead of off . Some devices do not support the standard commands or may provide additional ones. Use this to specify an alternate, device-specific, command that implements the off action.
pcmk_off_timeout	time	60s	Specify an alternate timeout to use for off actions instead of stonith-timeout . Some devices need much more or much less time to complete than normal. Use this to specify an alternate, device-specific, timeout for off actions.
pcmk_off_retries	integer	2	The maximum number of times to retry the off command within the timeout period. Some devices do not support multiple connections. Operations may fail if the device is busy with another task so Pacemaker will automatically retry the operation, if there is time remaining. Use this option to alter the number of times Pacemaker retries off actions before giving up.

Field	Type	Default	Description
pcmk_list_action	string	list	An alternate command to run instead of list . Some devices do not support the standard commands or may provide additional ones. Use this to specify an alternate, device-specific, command that implements the list action.
pcmk_list_timeout	time	60s	Specify an alternate timeout to use for list actions instead of stonith-timeout . Some devices need much more or much less time to complete than normal. Use this to specify an alternate, device-specific, timeout for list actions.
pcmk_list_retries	integer	2	The maximum number of times to retry the list command within the timeout period. Some devices do not support multiple connections. Operations may fail if the device is busy with another task so Pacemaker will automatically retry the operation, if there is time remaining. Use this option to alter the number of times Pacemaker retries list actions before giving up.
pcmk_monitor_action	string	monitor	An alternate command to run instead of monitor . Some devices do not support the standard commands or may provide additional ones. Use this to specify an alternate, device-specific, command that implements the monitor action.
pcmk_monitor_timeout	time	60s	Specify an alternate timeout to use for monitor actions instead of stonith-timeout . Some devices need much more or much less time to complete than normal. Use this to specify an alternate, device-specific, timeout for monitor actions.
pcmk_monitor_retries	integer	2	The maximum number of times to retry the monitor command within the timeout period. Some devices do not support multiple connections. Operations may fail if the device is busy with another task so Pacemaker will automatically retry the operation, if there is time remaining. Use this option to alter the number of times Pacemaker retries monitor actions before giving up.

Field	Type	Default	Description
pcmk_status_action	string	status	An alternate command to run instead of status . Some devices do not support the standard commands or may provide additional ones. Use this to specify an alternate, device-specific, command that implements the status action.
pcmk_status_timeout	time	60s	Specify an alternate timeout to use for status actions instead of stonith-timeout . Some devices need much more or much less time to complete than normal. Use this to specify an alternate, device-specific, timeout for status actions.
pcmk_status_retries	integer	2	The maximum number of times to retry the status command within the timeout period. Some devices do not support multiple connections. Operations may fail if the device is busy with another task so Pacemaker will automatically retry the operation, if there is time remaining. Use this option to alter the number of times Pacemaker retries status actions before giving up.
pcmk_delay_base	time	0s	<p>Enable a base delay for stonith actions and specify a base delay value. In a cluster with an even number of nodes, configuring a delay can help avoid nodes fencing each other at the same time in an even split. A random delay can be useful when the same fence device is used for all nodes, and differing static delays can be useful on each fencing device when a separate device is used for each node. The overall delay is derived from a random delay value adding this static delay so that the sum is kept below the maximum delay. If you set pcmk_delay_base but do not set pcmk_delay_max, there is no random component to the delay and it will be the value of pcmk_delay_base.</p> <p>Some individual fence agents implement a "delay" parameter, which is independent of delays configured with a pcmk_delay_* property. If both of these delays are configured, they are added together and thus would generally not be used in conjunction.</p>

Field	Type	Default	Description
pcmk_delay_max	time	0s	<p>Enable a random delay for stonith actions and specify the maximum of random delay. In a cluster with an even number of nodes, configuring a delay can help avoid nodes fencing each other at the same time in an even split. A random delay can be useful when the same fence device is used for all nodes, and differing static delays can be useful on each fencing device when a separate device is used for each node. The overall delay is derived from this random delay value adding a static delay so that the sum is kept below the maximum delay. If you set pcmk_delay_max but do not set pcmk_delay_base there is no static component to the delay.</p> <p>Some individual fence agents implement a "delay" parameter, which is independent of delays configured with a pcmk_delay_* property. If both of these delays are configured, they are added together and thus would generally not be used in conjunction.</p>
pcmk_action_limit	integer	1	The maximum number of actions that can be performed in parallel on this device. The cluster property concurrent-fencing=true needs to be configured first. A value of -1 is unlimited.
pcmk_on_action	string	on	For advanced use only: An alternate command to run instead of on . Some devices do not support the standard commands or may provide additional ones. Use this to specify an alternate, device-specific, command that implements the on action.
pcmk_on_timeout	time	60s	For advanced use only: Specify an alternate timeout to use for on actions instead of stonith-timeout . Some devices need much more or much less time to complete than normal. Use this to specify an alternate, device-specific, timeout for on actions.
pcmk_on_retries	integer	2	For advanced use only: The maximum number of times to retry the on command within the timeout period. Some devices do not support multiple connections. Operations may fail if the device is busy with another task so Pacemaker will automatically retry the operation, if there is time remaining. Use this option to alter the number of times Pacemaker retries on actions before giving up.

Field	Type	Default	Description

You can determine how a cluster node should react if notified of its own fencing by setting the **fence-reaction** cluster property, as described in [Table 12.1, “Cluster Properties”](#). A cluster node may receive notification of its own fencing if fencing is misconfigured, or if fabric fencing is in use that does not cut cluster communication. Although the default value for this property is **stop**, which attempts to immediately stop Pacemaker and keep it stopped, the safest choice for this value is **panic**, which attempts to immediately reboot the local node. If you prefer the stop behavior, as is most likely to be the case in conjunction with fabric fencing, it is recommended that you set this explicitly.

5.9. CONFIGURING FENCING LEVELS

Pacemaker supports fencing nodes with multiple devices through a feature called fencing topologies. To implement topologies, create the individual devices as you normally would and then define one or more fencing levels in the fencing topology section in the configuration.

- Each level is attempted in ascending numeric order, starting at 1.
- If a device fails, processing terminates for the current level. No further devices in that level are exercised and the next level is attempted instead.
- If all devices are successfully fenced, then that level has succeeded and no other levels are tried.
- The operation is finished when a level has passed (success), or all levels have been attempted (failed).

Use the following command to add a fencing level to a node. The devices are given as a comma-separated list of stonith ids, which are attempted for the node at that level.

```
pcs stonith level add level node devices
```

The following command lists all of the fencing levels that are currently configured.

```
pcs stonith level
```

In the following example, there are two fence devices configured for node **rh7-2**: an ilo fence device called **my_il0** and an apc fence device called **my_apc**. These commands sets up fence levels so that if the device **my_il0** fails and is unable to fence the node, then Pacemaker will attempt to use the device **my_apc**. This example also shows the output of the **pcs stonith level** command after the levels are configured.

```
# pcs stonith level add 1 rh7-2 my_il0
```

```
# pcs stonith level add 2 rh7-2 my_apc
# pcs stonith level
Node: rh7-2
  Level 1 - my_il0
  Level 2 - my_apc
```

The following command removes the fence level for the specified node and devices. If no nodes or devices are specified then the fence level you specify is removed from all nodes.

```
pcs stonith level remove level [node_id] [stonith_id] ... [stonith_id]
```

The following command clears the fence levels on the specified node or stonith id. If you do not specify a node or stonith id, all fence levels are cleared.

```
pcs stonith level clear [node|stonith_id(s)]
```

If you specify more than one stonith id, they must be separated by a comma and no spaces, as in the following example.

```
# pcs stonith level clear dev_a,dev_b
```

The following command verifies that all fence devices and nodes specified in fence levels exist.

```
pcs stonith level verify
```

As of Red Hat Enterprise Linux 7.4, you can specify nodes in fencing topology by a regular expression applied on a node name and by a node attribute and its value. For example, the following commands configure nodes **node1**, **node2**, and ` **node3** to use fence devices **apc1** and ` **apc2**, and nodes ` **node4**, **node5**, and ` **node6** to use fence devices **apc3** and ` **apc4**.

```
pcs stonith level add 1 "regexp%node[1-3]" apc1,apc2
pcs stonith level add 1 "regexp%node[4-6]" apc3,apc4
```

The following commands yield the same results by using node attribute matching.

```
pcs node attribute node1 rack=1
pcs node attribute node2 rack=1
pcs node attribute node3 rack=1
pcs node attribute node4 rack=2
pcs node attribute node5 rack=2
pcs node attribute node6 rack=2
pcs stonith level add 1 attrib%rack=1 apc1,apc2
pcs stonith level add 1 attrib%rack=2 apc3,apc4
```

5.10. CONFIGURING FENCING FOR REDUNDANT POWER SUPPLIES

When configuring fencing for redundant power supplies, the cluster must ensure that when attempting to reboot a host, both power supplies are turned off before either power supply is turned back on.

If the node never completely loses power, the node may not release its resources. This opens up the possibility of nodes accessing these resources simultaneously and corrupting them.

Prior to Red Hat Enterprise Linux 7.2, you needed to explicitly configure different versions of the

devices which used either the 'on' or 'off' actions. Since Red Hat Enterprise Linux 7.2, it is now only required to define each device once and to specify that both are required to fence the node, as in the following example.

```
# pcs stonith create apc1 fence_apc_snmp ipaddr=apc1.example.com login=user
passwd='7a4D#1j!pz864' pcmk_host_map="node1.example.com:1;node2.example.com:2"

# pcs stonith create apc2 fence_apc_snmp ipaddr=apc2.example.com login=user
passwd='7a4D#1j!pz864' pcmk_host_map="node1.example.com:1;node2.example.com:2"

# pcs stonith level add 1 node1.example.com apc1,apc2
# pcs stonith level add 1 node2.example.com apc1,apc2
```

5.11. CONFIGURING ACPI FOR USE WITH INTEGRATED FENCE DEVICES

If your cluster uses integrated fence devices, you must configure ACPI (Advanced Configuration and Power Interface) to ensure immediate and complete fencing.

If a cluster node is configured to be fenced by an integrated fence device, disable ACPI Soft-Off for that node. Disabling ACPI Soft-Off allows an integrated fence device to turn off a node immediately and completely rather than attempting a clean shutdown (for example, **shutdown -h now**). Otherwise, if ACPI Soft-Off is enabled, an integrated fence device can take four or more seconds to turn off a node (see the note that follows). In addition, if ACPI Soft-Off is enabled and a node panics or freezes during shutdown, an integrated fence device may not be able to turn off the node. Under those circumstances, fencing is delayed or unsuccessful. Consequently, when a node is fenced with an integrated fence device and ACPI Soft-Off is enabled, a cluster recovers slowly or requires administrative intervention to recover.



NOTE

The amount of time required to fence a node depends on the integrated fence device used. Some integrated fence devices perform the equivalent of pressing and holding the power button; therefore, the fence device turns off the node in four to five seconds. Other integrated fence devices perform the equivalent of pressing the power button momentarily, relying on the operating system to turn off the node; therefore, the fence device turns off the node in a time span much longer than four to five seconds.

- The preferred way to disable ACPI Soft-Off is to change the BIOS setting to "instant-off" or an equivalent setting that turns off the node without delay, as described in [Section 5.11.1, "Disabling ACPI Soft-Off with the BIOS"](#).

Disabling ACPI Soft-Off with the BIOS may not be possible with some systems. If disabling ACPI Soft-Off with the BIOS is not satisfactory for your cluster, you can disable ACPI Soft-Off with one of the following alternate methods:

- Setting **HandlePowerKey=ignore** in the **/etc/systemd/logind.conf** file and verifying that the node turns off immediately when fenced, as described in [Section 5.11.2, "Disabling ACPI Soft-Off in the logind.conf file"](#). This is the first alternate method of disabling ACPI Soft-Off.
- Appending **acpi=off** to the kernel boot command line, as described in [Section 5.11.3, "Disabling ACPI Completely in the GRUB 2 File"](#). This is the second alternate method of disabling ACPI Soft-Off, if the preferred or the first alternate method is not available.

**IMPORTANT**

This method completely disables ACPI; some computers do not boot correctly if ACPI is completely disabled. Use this method *only* if the other methods are not effective for your cluster.

5.11.1. Disabling ACPI Soft-Off with the BIOS

You can disable ACPI Soft-Off by configuring the BIOS of each cluster node with the following procedure.

**NOTE**

The procedure for disabling ACPI Soft-Off with the BIOS may differ among server systems. You should verify this procedure with your hardware documentation.

1. Reboot the node and start the **BIOS CMOS Setup Utility** program.
2. Navigate to the **Power** menu (or equivalent power management menu).
3. At the **Power** menu, set the **Soft-Off by PWR-BTTN** function (or equivalent) to **Instant-Off** (or the equivalent setting that turns off the node by means of the power button without delay). [Example 5.1, “BIOS CMOS Setup Utility: Soft-Off by PWR-BTTN set to Instant-Off”](#) shows a **Power** menu with **ACPI Function** set to **Enabled** and **Soft-Off by PWR-BTTN** set to **Instant-Off**.

**NOTE**

The equivalents to **ACPI Function**, **Soft-Off by PWR-BTTN**, and **Instant-Off** may vary among computers. However, the objective of this procedure is to configure the BIOS so that the computer is turned off by means of the power button without delay.

4. Exit the **BIOS CMOS Setup Utility** program, saving the BIOS configuration.
5. Verify that the node turns off immediately when fenced. For information on testing a fence device, see [Section 5.12, “Testing a Fence Device”](#).

Example 5.1. BIOS CMOS Setup Utility: Soft-Off by PWR-BTTN set to Instant-Off

		Item Help	
ACPI Function	[Enabled]		
ACPI Suspend Type	[S1(POS)]		
x Run VGABIOS if S3 Resume	Auto	Menu Level *	
Suspend Mode	[Disabled]		
HDD Power Down	[Disabled]		
Soft-Off by PWR-BTTN	[Instant-Off]		
CPU THRM-Throttling	[50.0%]		
Wake-Up by PCI card	[Enabled]		
Power On by Ring	[Enabled]		
Wake Up On LAN	[Enabled]		
x USB KB Wake-Up From S3	Disabled		
Resume by Alarm	[Disabled]		

x Date(of Month) Alarm	0			
x Time(hh:mm:ss) Alarm	0 : 0 :			
POWER ON Function	[BUTTON ONLY			
x KB Power ON Password	Enter			
x Hot Key Power ON	Ctrl-F1			
+-----+-----+				

This example shows **ACPI Function** set to **Enabled**, and **Soft-Off by PWR-BTTN** set to **Instant-Off**.

5.11.2. Disabling ACPI Soft-Off in the logind.conf file

To disable power-key handing in the **/etc/systemd/logind.conf** file, use the following procedure.

1. Define the following configuration in the **/etc/systemd/logind.conf** file:

```
HandlePowerKey=ignore
```

2. Reload the **systemd** configuration:

```
# systemctl daemon-reload
```

3. Verify that the node turns off immediately when fenced. For information on testing a fence device, see [Section 5.12, “Testing a Fence Device”](#).

5.11.3. Disabling ACPI Completely in the GRUB 2 File

You can disable ACPI Soft-Off by appending **acpi=off** to the GRUB menu entry for a kernel.



IMPORTANT

This method completely disables ACPI; some computers do not boot correctly if ACPI is completely disabled. Use this method *only* if the other methods are not effective for your cluster.

Use the following procedure to disable ACPI in the GRUB 2 file:

1. Use the **--args** option in combination with the **--update-kernel** option of the **grubby** tool to change the **grub.cfg** file of each cluster node as follows:

```
# grubby --args=acpi=off --update-kernel=ALL
```

For general information on GRUB 2, see the [Working with GRUB 2](#) chapter in the [System Administrator’s Guide](#).

2. Reboot the node.
3. Verify that the node turns off immediately when fenced. For information on testing a fence device, see [Section 5.12, “Testing a Fence Device”](#).

5.12. TESTING A FENCE DEVICE

Fencing is a fundamental part of the Red Hat Cluster infrastructure and it is therefore important to validate or test that fencing is working properly.

Use the following procedure to test a fence device.

1. Use ssh, telnet, HTTP, or whatever remote protocol is used to connect to the device to manually log in and test the fence device or see what output is given. For example, if you will be configuring fencing for an IPMI-enabled device, then try to log in remotely with **ipmitool**. Take note of the options used when logging in manually because those options might be needed when using the fencing agent.

If you are unable to log in to the fence device, verify that the device is pingable, there is nothing such as a firewall configuration that is preventing access to the fence device, remote access is enabled on the fencing agent, and the credentials are correct.

2. Run the fence agent manually, using the fence agent script. This does not require that the cluster services are running, so you can perform this step before the device is configured in the cluster. This can ensure that the fence device is responding properly before proceeding.



NOTE

The examples in this section use the **fence_ilo** fence agent script for an iLO device. The actual fence agent you will use and the command that calls that agent will depend on your server hardware. You should consult the man page for the fence agent you are using to determine which options to specify. You will usually need to know the login and password for the fence device and other information related to the fence device.

The following example shows the format you would use to run the **fence_ilo** fence agent script with **-o status** parameter to check the status of the fence device interface on another node without actually fencing it. This allows you to test the device and get it working before attempting to reboot the node. When running this command, you specify the name and password of an iLO user that has power on and off permissions for the iLO device.

```
# fence_ilo -a ipaddress -l username -p password -o status
```

The following example shows the format you would use to run the **fence_ilo** fence agent script with the **-o reboot** parameter. Running this command on one node reboots another node on which you have configured the fence agent.

```
# fence_ilo -a ipaddress -l username -p password -o reboot
```

If the fence agent failed to properly do a status, off, on, or reboot action, you should check the hardware, the configuration of the fence device, and the syntax of your commands. In addition, you can run the fence agent script with the debug output enabled. The debug output is useful for some fencing agents to see where in the sequence of events the fencing agent script is failing when logging into the fence device.

```
# fence_ilo -a ipaddress -l username -p password -o status -D /tmp/$(hostname)-fence_agent.debug
```

When diagnosing a failure that has occurred, you should ensure that the options you specified when manually logging in to the fence device are identical to what you passed on to the fence agent with the fence agent script.

For fence agents that support an encrypted connection, you may see an error due to certificate validation failing, requiring that you trust the host or that you use the fence agent's **ssl-insecure** parameter. Similarly, if SSL/TLS is disabled on the target device, you may need to account for this when setting the SSL parameters for the fence agent.



NOTE

If the fence agent that is being tested is a **fence_drac**, **fence_ilc**, or some other fencing agent for a systems management device that continues to fail, then fall back to trying **fence_ipmilan**. Most systems management cards support IPMI remote login and the only supported fencing agent is **fence_ipmilan**.

- Once the fence device has been configured in the cluster with the same options that worked manually and the cluster has been started, test fencing with the **pcs stonith fence** command from any node (or even multiple times from different nodes), as in the following example. The **pcs stonith fence** command reads the cluster configuration from the CIB and calls the fence agent as configured to execute the fence action. This verifies that the cluster configuration is correct.

```
# pcs stonith fence node_name
```

If the **pcs stonith fence** command works properly, that means the fencing configuration for the cluster should work when a fence event occurs. If the command fails, it means that cluster management cannot invoke the fence device through the configuration it has retrieved. Check for the following issues and update your cluster configuration as needed.

- Check your fence configuration. For example, if you have used a host map you should ensure that the system can find the node using the host name you have provided.
- Check whether the password and user name for the device include any special characters that could be misinterpreted by the bash shell. Making sure that you enter passwords and user names surrounded by quotation marks could address this issue.
- Check whether you can connect to the device using the exact IP address or host name you specified in the **pcs stonith** command. For example, if you give the host name in the stonith command but test by using the IP address, that is not a valid test.
- If the protocol that your fence device uses is accessible to you, use that protocol to try to connect to the device. For example many agents use ssh or telnet. You should try to connect to the device with the credentials you provided when configuring the device, to see if you get a valid prompt and can log in to the device.

If you determine that all your parameters are appropriate but you still have trouble connecting to your fence device, you can check the logging on the fence device itself, if the device provides that, which will show if the user has connected and what command the user issued. You can also search through the **/var/log/messages** file for instances of stonith and error, which could give some idea of what is transpiring, but some agents can provide additional information.

- Once the fence device tests are working and the cluster is up and running, test an actual failure. To do this, take an action in the cluster that should initiate a token loss.

- Take down a network. How you take a network depends on your specific configuration. In many cases, you can physically pull the network or power cables out of the host.



NOTE

Disabling the network interface on the local host rather than physically disconnecting the network or power cables is not recommended as a test of fencing because it does not accurately simulate a typical real-world failure.

- Block corosync traffic both inbound and outbound using the local firewall.

The following example blocks corosync, assuming the default corosync port is used, **firewalld** is used as the local firewall, and the network interface used by corosync is in the default firewall zone:

```
# firewall-cmd --direct --add-rule ipv4 filter OUTPUT 2 -p udp --dport=5405 -j DROP  
# firewall-cmd --add-rich-rule='rule family="ipv4" port port="5405" protocol="udp" drop'
```

- Simulate a crash and panic your machine with **sysrq-trigger**. Note, however, that triggering a kernel panic can cause data loss; it is recommended that you disable your cluster resources first.

```
# echo c > /proc/sysrq-trigger
```

CHAPTER 6. CONFIGURING CLUSTER RESOURCES

This chapter provides information on configuring resources in a cluster.

6.1. RESOURCE CREATION

Use the following command to create a cluster resource.

```
pcs resource create resource_id [standard:[provider:]] type [resource_options] [op operation_action operation_options [operation_action operation options]... ] [meta meta_options...] [clone clone_options] | master [master_options] | --group group_name [-before resource_id] | --after resource_id] | [bundle bundle_id] [--disabled] [--wait[=n]]
```

When you specify the **--group** option, the resource is added to the resource group named. If the group does not exist, this creates the group and adds this resource to the group. For information on resource groups, see [Section 6.5, “Resource Groups”](#).

The **--before** and **--after** options specify the position of the added resource relative to a resource that already exists in a resource group.

Specifying the **--disabled** option indicates that the resource is not started automatically.

The following command creates a resource with the name **VirtualIP** of standard **ocf**, provider **heartbeat**, and type **IPAddr2**. The floating address of this resource is 192.168.0.120, the system will check whether the resource is running every 30 seconds.

```
# pcs resource create VirtualIP ocf:heartbeat:IPAddr2 ip=192.168.0.120 cidr_netmask=24 op monitor interval=30s
```

Alternately, you can omit the *standard* and *provider* fields and use the following command. This will default to a standard of **ocf** and a provider of **heartbeat**.

```
# pcs resource create VirtualIP IPAddr2 ip=192.168.0.120 cidr_netmask=24 op monitor interval=30s
```

Use the following command to delete a configured resource.

```
pcs resource delete resource_id
```

For example, the following command deletes an existing resource with a resource ID of **VirtualIP**

```
# pcs resource delete VirtualIP
```

- For information on the *resource_id*, *standard*, *provider*, and *type* fields of the **pcs resource create** command, see [Section 6.2, “Resource Properties”](#).
- For information on defining resource parameters for individual resources, see [Section 6.3, “Resource-Specific Parameters”](#).
- For information on defining resource meta options, which are used by the cluster to decide how a resource should behave, see [Section 6.4, “Resource Meta Options”](#).
- For information on defining the operations to perform on a resource, see [Section 6.6, “Resource Operations”](#).

- Specifying the **clone** option creates a clone resource. Specifying the **master** option creates a master/slave resource. For information on resource clones and resources with multiple modes, see [Chapter 9, Advanced Configuration](#).

6.2. RESOURCE PROPERTIES

The properties that you define for a resource tell the cluster which script to use for the resource, where to find that script and what standards it conforms to. [Table 6.1, “Resource Properties”](#) describes these properties.

Table 6.1. Resource Properties

Field	Description
resource_id	Your name for the resource
standard	The standard the script conforms to. Allowed values: ocf , service , upstart , systemd , lsb , stonith
type	The name of the Resource Agent you wish to use, for example IPAddr or Filesystem
provider	The OCF spec allows multiple vendors to supply the same resource agent. Most of the agents shipped by Red Hat use heartbeat as the provider.

[Table 6.2, “Commands to Display Resource Properties”](#) summarizes the commands that display the available resource properties.

Table 6.2. Commands to Display Resource Properties

pcs Display Command	Output
pcs resource list	Displays a list of all available resources.
pcs resource standards	Displays a list of available resources agent standards.
pcs resource providers	Displays a list of available resources agent providers.
pcs resource list <i>string</i>	Displays a list of available resources filtered by the specified string. You can use this command to display resources filtered by the name of a standard, a provider, or a type.

6.3. RESOURCE-SPECIFIC PARAMETERS

For any individual resource, you can use the following command to display the parameters you can set for that resource.

```
# pcs resource describe standard:provider:type|type
```

For example, the following command displays the parameters you can set for a resource of type **LVM**.

```
# pcs resource describe LVM
Resource options for: LVM
  volgrpname (required): The name of volume group.
  exclusive: If set, the volume group will be activated exclusively.
  partial_activation: If set, the volume group will be activated even
    only partial of the physical volumes available. It helps to set to
    true, when you are using mirroring logical volumes.
```

6.4. RESOURCE META OPTIONS

In addition to the resource-specific parameters, you can configure additional resource options for any resource. These options are used by the cluster to decide how your resource should behave. [Table 6.3, "Resource Meta Options"](#) describes these options.

Table 6.3. Resource Meta Options

Field	Default	Description
priority	0	If not all resources can be active, the cluster will stop lower priority resources in order to keep higher priority ones active.
target-role	Started	What state should the cluster attempt to keep this resource in? Allowed values: * <i>Stopped</i> - Force the resource to be stopped * <i>Started</i> - Allow the resource to be started (In the case of multistate resources, they will not promoted to master) * <i>Master</i> - Allow the resource to be started and, if appropriate, promoted
is-managed	true	Is the cluster allowed to start and stop the resource? Allowed values: true, false
resource-stickiness	0	Value to indicate how much the resource prefers to stay where it is.

Field	Default	Description
requires	Calculated	<p>Indicates under what conditions the resource can be started.</p> <p>Defaults to fencing except under the conditions noted below. Possible values:</p> <ul style="list-style-type: none"> * nothing – The cluster can always start the resource. * quorum – The cluster can only start this resource if a majority of the configured nodes are active. This is the default value if stonith-enabled is false or the resource's standard is stonith. * fencing – The cluster can only start this resource if a majority of the configured nodes are active <i>and</i> any failed or unknown nodes have been powered off. * unfencing – The cluster can only start this resource if a majority of the configured nodes are active <i>and</i> any failed or unknown nodes have been powered off <i>and</i> only on nodes that have been <i>unfenced</i>. This is the default value if the provides=unfencing stonith meta option has been set for a fencing device.
migration-threshold	INFINITY	<p>How many failures may occur for this resource on a node, before this node is marked ineligible to host this resource. A value of 0 indicates that this feature is disabled (the node will never be marked ineligible); by contrast, the cluster treats INFINITY (the default) as a very large but finite number. This option has an effect only if the failed operation has on-fail=restart (the default), and additionally for failed start operations if the cluster property start-failure-is-fatal is false. For information on configuring the migration-threshold option, see Section 8.2, “Moving Resources Due to Failure”. For information on the start-failure-is-fatal option, see Table 12.1, “Cluster Properties”.</p>
failure-timeout	0 (disabled)	<p>Used in conjunction with the migration-threshold option, indicates how many seconds to wait before acting as if the failure had not occurred, and potentially allowing the resource back to the node on which it failed. As with any time-based actions, this is not guaranteed to be checked more frequently than the value of the cluster-recheck-interval cluster parameter. For information on configuring the failure-timeout option, see Section 8.2, “Moving Resources Due to Failure”.</p>

Field	Default	Description
multiple-active	stop_start	What should the cluster do if it ever finds the resource active on more than one node. Allowed values: * block - mark the resource as unmanaged * stop_only - stop all active instances and leave them that way * stop_start - stop all active instances and start the resource in one location only

To change the default value of a resource option, use the following command.

```
pcs resource defaults options
```

For example, the following command resets the default value of **resource-stickiness** to 100.

```
# pcs resource defaults resource-stickiness=100
```

Omitting the *options* parameter from the **pcs resource defaults** displays a list of currently configured default values for resource options. The following example shows the output of this command after you have reset the default value of **resource-stickiness** to 100.

```
# pcs resource defaults
resource-stickiness:100
```

Whether you have reset the default value of a resource meta option or not, you can set a resource option for a particular resource to a value other than the default when you create the resource. The following shows the format of the **pcs resource create** command you use when specifying a value for a resource meta option.

```
pcs resource create resource_id standard:provider:type|type [resource options] [meta meta_options...]
```

For example, the following command creates a resource with a **resource-stickiness** value of 50.

```
# pcs resource create VirtualIP ocf:heartbeat:IPaddr2 ip=192.168.0.120 cidr_netmask=24 meta
resource-stickiness=50
```

You can also set the value of a resource meta option for an existing resource, group, cloned resource, or master resource with the following command.

```
pcs resource meta resource_id | group_id | clone_id | master_id meta_options
```

In the following example, there is an existing resource named **dummy_resource**. This command sets the **failure-timeout** meta option to 20 seconds, so that the resource can attempt to restart on the same node in 20 seconds.

```
# pcs resource meta dummy_resource failure-timeout=20s
```

After executing this command, you can display the values for the resource to verify that **failure-timeout=20s** is set.

```
# pcs resource show dummy_resource
Resource: dummy_resource (class=ocf provider=heartbeat type=Dummy)
Meta Attrs: failure-timeout=20s
Operations: start interval=0s timeout=20 (dummy_resource-start-timeout-20)
              stop interval=0s timeout=20 (dummy_resource-stop-timeout-20)
              monitor interval=10 timeout=20 (dummy_resource-monitor-interval-10)
```

For information on resource clone meta options, see [Section 9.1, “Resource Clones”](#). For information on resource master meta options, see [Section 9.2, “Multistate Resources: Resources That Have Multiple Modes”](#).

6.5. RESOURCE GROUPS

One of the most common elements of a cluster is a set of resources that need to be located together, start sequentially, and stop in the reverse order. To simplify this configuration, Pacemaker supports the concept of groups.

You create a resource group with the following command, specifying the resources to include in the group. If the group does not exist, this command creates the group. If the group exists, this command adds additional resources to the group. The resources will start in the order you specify them with this command, and will stop in the reverse order of their starting order.

```
pcs resource group add group_name resource_id [resource_id] ... [resource_id]
[--before resource_id] [--after resource_id]
```

You can use the **--before** and **--after** options of this command to specify the position of the added resources relative to a resource that already exists in the group.

You can also add a new resource to an existing group when you create the resource, using the following command. The resource you create is added to the group named *group_name*.

```
pcs resource create resource_id standard:provider:type/type [resource_options] [op operation_action operation_options] --group group_name
```

You remove a resource from a group with the following command. If there are no resources in the group, this command removes the group itself.

```
pcs resource group remove group_name resource_id...
```

The following command lists all currently configured resource groups.

```
pcs resource group list
```

The following example creates a resource group named **shortcut** that contains the existing resources **IPaddr** and **Email**.

```
# pcs resource group add shortcut IPaddr Email
```

There is no limit to the number of resources a group can contain. The fundamental properties of a group are as follows.

- Resources are started in the order in which you specify them (in this example, **IPaddr** first, then **Email**).
- Resources are stopped in the reverse order in which you specify them. (**Email** first, then **IPaddr**).

If a resource in the group cannot run anywhere, then no resource specified after that resource is allowed to run.

- If **IPaddr** cannot run anywhere, neither can **Email**.
- If **Email** cannot run anywhere, however, this does not affect **IPaddr** in any way.

Obviously as the group grows bigger, the reduced configuration effort of creating resource groups can become significant.

6.5.1. Group Options

A resource group inherits the following options from the resources that it contains: **priority**, **target-role**, **is-managed**. For information on resource options, see [Table 6.3, “Resource Meta Options”](#).

6.5.2. Group Stickiness

Stickiness, the measure of how much a resource wants to stay where it is, is additive in groups. Every active resource of the group will contribute its stickiness value to the group’s total. So if the default **resource-stickiness** is 100, and a group has seven members, five of which are active, then the group as a whole will prefer its current location with a score of 500.

6.6. RESOURCE OPERATIONS

To ensure that resources remain healthy, you can add a monitoring operation to a resource’s definition. If you do not specify a monitoring operation for a resource, by default the **pcs** command will create a monitoring operation, with an interval that is determined by the resource agent. If the resource agent does not provide a default monitoring interval, the **pcs** command will create a monitoring operation with an interval of 60 seconds.

[Table 6.4, “Properties of an Operation”](#) summarizes the properties of a resource monitoring operation.

Table 6.4. Properties of an Operation

Field	Description
id	Unique name for the action. The system assigns this when you configure an operation.
name	The action to perform. Common values: monitor , start , stop

Field	Description
interval	If set to a nonzero value, a recurring operation is created that repeats at this frequency, in seconds. A nonzero value makes sense only when the action name is set to monitor . A recurring monitor action will be executed immediately after a resource start completes, and subsequent monitor actions are scheduled starting at the time the previous monitor action completed. For example, if a monitor action with interval=20s is executed at 01:00:00, the next monitor action does not occur at 01:00:20, but at 20 seconds after the first monitor action completes. If set to zero, which is the default value, this parameter allows you to provide values to be used for operations created by the cluster. For example, if the interval is set to zero, the name of the operation is set to start , and the timeout value is set to 40, then Pacemaker will use a timeout of 40 seconds when starting this resource. A monitor operation with a zero interval allows you to set the timeout/on-fail/enabled values for the probes that Pacemaker does at startup to get the current status of all resources when the defaults are not desirable.
timeout	If the operation does not complete in the amount of time set by this parameter, abort the operation and consider it failed. The default value is the value of timeout if set with the pcs resource op defaults command, or 20 seconds if it is not set. If you find that your system includes a resource that requires more time than the system allows to perform an operation (such as start , stop , or monitor), investigate the cause and if the lengthy execution time is expected you can increase this value. The timeout value is not a delay of any kind, nor does the cluster wait the entire timeout period if the operation returns before the timeout period has completed.
on-fail	The action to take if this action ever fails. Allowed values: <ul style="list-style-type: none"> * ignore - Pretend the resource did not fail * block - Do not perform any further operations on the resource * stop - Stop the resource and do not start it elsewhere * restart - Stop the resource and start it again (possibly on a different node) * fence - STONITH the node on which the resource failed * standby - Move <i>all</i> resources away from the node on which the resource failed The default for the stop operation is fence when STONITH is enabled and block otherwise. All other operations default to restart .
enabled	If false , the operation is treated as if it does not exist. Allowed values true , false

6.6.1. Configuring Resource Operations

You can configure monitoring operations when you create a resource, using the following command.

```
pcs resource create resource_id standard;provider:type/type [resource_options] [op operation_action
operation_options [operation_type operation_options]...]
```

For example, the following command creates an **IPaddr2** resource with a monitoring operation. The new resource is called **VirtuallP** with an IP address of 192.168.0.99 and a netmask of 24 on **eth2**. A monitoring operation will be performed every 30 seconds.

```
# pcs resource create VirtualIP ocf:heartbeat:IPaddr2 ip=192.168.0.99 cidr_netmask=24 nic=eth2 op monitor interval=30s
```

Alternately, you can add a monitoring operation to an existing resource with the following command.

```
pcs resource op add resource_id operation_action [operation_properties]
```

Use the following command to delete a configured resource operation.

```
pcs resource op remove resource_id operation_name operation_properties
```



NOTE

You must specify the exact operation properties to properly remove an existing operation.

To change the values of a monitoring option, you can update the resource. For example, you can create a **VirtualIP** with the following command.

```
# pcs resource create VirtualIP ocf:heartbeat:IPaddr2 ip=192.168.0.99 cidr_netmask=24 nic=eth2
```

By default, this command creates these operations.

```
Operations: start interval=0s timeout=20s (VirtualIP-start-timeout-20s)
           stop interval=0s timeout=20s (VirtualIP-stop-timeout-20s)
           monitor interval=10s timeout=20s (VirtualIP-monitor-interval-10s)
```

To change the stop timeout operation, execute the following command.

```
# pcs resource update VirtualIP op stop interval=0s timeout=40s
# pcs resource show VirtualIP
Resource: VirtualIP (class=ocf provider=heartbeat type=IPaddr2)
Attributes: ip=192.168.0.99 cidr_netmask=24 nic=eth2
Operations: start interval=0s timeout=20s (VirtualIP-start-timeout-20s)
           monitor interval=10s timeout=20s (VirtualIP-monitor-interval-10s)
           stop interval=0s timeout=40s (VirtualIP-name-stop-interval-0s-timeout-40s)
```



NOTE

When you update a resource's operation with the **pcs resource update** command, any options you do not specifically call out are reset to their default values.

6.6.2. Configuring Global Resource Operation Defaults

You can use the following command to set global default values for monitoring operations.

```
pcs resource op defaults [options]
```

For example, the following command sets a global default of a **timeout** value of 240 seconds for all monitoring operations.

```
# pcs resource op defaults timeout=240s
```

To display the currently configured default values for monitoring operations, do not specify any options when you execute the **pcs resource op defaults** command.

For example, following command displays the default monitoring operation values for a cluster which has been configured with a **timeout** value of 240 seconds.

```
# pcs resource op defaults
timeout: 240s
```

Note that a cluster resource will use the global default only when the option is not specified in the cluster resource definition. By default, resource agents define the **timeout** option for all operations. For the global operation timeout value to be honored, you must create the cluster resource without the **timeout** option explicitly or you must remove the **timeout** option by updating the cluster resource, as in the following command.

```
# pcs resource update VirtualIP op monitor interval=10s
```

For example, after setting a global default of a **timeout** value of 240 seconds for all monitoring operations and updating the cluster resource **VirtualIP** to remove the timeout value for the **monitor** operation, the resource **VirtualIP** will then have timeout values for **start**, **stop**, and **monitor** operations of 20s, 40s and 240s, respectively. The global default value for timeout operations is applied here only on the **monitor** operation, where the default **timeout** option was removed by the previous command.

```
# pcs resource show VirtualIP
Resource: VirtualIP (class=ocf provider=heartbeat type=IPaddr2)
Attributes: ip=192.168.0.99 cidr_netmask=24 nic=eth2
Operations: start interval=0s timeout=20s (VirtualIP-start-timeout-20s)
            monitor interval=10s (VirtualIP-monitor-interval-10s)
            stop interval=0s timeout=40s (VirtualIP-name-stop-interval-0s-timeout-40s)
```

6.7. DISPLAYING CONFIGURED RESOURCES

To display a list of all configured resources, use the following command.

```
pcs resource show
```

For example, if your system is configured with a resource named **VirtualIP** and a resource named **WebSite**, the **pcs resource show** command yields the following output.

```
# pcs resource show
VirtualIP (ocf::heartbeat:IPaddr2): Started
WebSite (ocf::heartbeat:apache): Started
```

To display a list of all configured resources and the parameters configured for those resources, use the **-full** option of the **pcs resource show** command, as in the following example.

```
# pcs resource show --full
Resource: VirtualIP (type=IPaddr2 class=ocf provider=heartbeat)
Attributes: ip=192.168.0.120 cidr_netmask=24
Operations: monitor interval=30s
```

```
Resource: WebSite (type=apache class=ocf provider=heartbeat)
Attributes: statusurl=http://localhost/server-status configfile=/etc/httpd/conf/httpd.conf
Operations: monitor interval=1min
```

To display the configured parameters for a resource, use the following command.

```
pcs resource show resource_id
```

For example, the following command displays the currently configured parameters for resource **VirtualIP**.

```
# pcs resource show VirtualIP
Resource: VirtualIP (type=IPaddr2 class=ocf provider=heartbeat)
Attributes: ip=192.168.0.120 cidr_netmask=24
Operations: monitor interval=30s
```

6.8. MODIFYING RESOURCE PARAMETERS

To modify the parameters of a configured resource, use the following command.

```
pcs resource update resource_id [resource_options]
```

The following sequence of commands show the initial values of the configured parameters for resource **VirtualIP**, the command to change the value of the **ip** parameter, and the values following the update command.

```
# pcs resource show VirtualIP
Resource: VirtualIP (type=IPaddr2 class=ocf provider=heartbeat)
Attributes: ip=192.168.0.120 cidr_netmask=24
Operations: monitor interval=30s
# pcs resource update VirtualIP ip=192.169.0.120
# pcs resource show VirtualIP
Resource: VirtualIP (type=IPaddr2 class=ocf provider=heartbeat)
Attributes: ip=192.169.0.120 cidr_netmask=24
Operations: monitor interval=30s
```

6.9. MULTIPLE MONITORING OPERATIONS

You can configure a single resource with as many monitor operations as a resource agent supports. In this way you can do a superficial health check every minute and progressively more intense ones at higher intervals.



NOTE

When configuring multiple monitor operations, you must ensure that no two operations are performed at the same interval.

To configure additional monitoring operations for a resource that supports more in-depth checks at different levels, you add an **OCF_CHECK_LEVEL=n** option.

For example, if you configure the following **IPaddr2** resource, by default this creates a monitoring operation with an interval of 10 seconds and a timeout value of 20 seconds.

```
# pcs resource create VirtualIP ocf:heartbeat:IPaddr2 ip=192.168.0.99 cidr_netmask=24 nic=eth2
```

If the Virtual IP supports a different check with a depth of 10, the following command causes Pacemaker to perform the more advanced monitoring check every 60 seconds in addition to the normal Virtual IP check every 10 seconds. (As noted, you should not configure the additional monitoring operation with a 10-second interval as well.)

```
# pcs resource op add VirtualIP monitor interval=60s OCF_CHECK_LEVEL=10
```

6.10. ENABLING AND DISABLING CLUSTER RESOURCES

The following command enables the resource specified by **resource_id**.

```
pcs resource enable resource_id
```

The following command disables the resource specified by **resource_id**.

```
pcs resource disable resource_id
```

6.11. CLUSTER RESOURCES CLEANUP

If a resource has failed, a failure message appears when you display the cluster status. If you resolve that resource, you can clear that failure status with the **pcs resource cleanup** command. This command resets the resource status and **failcount**, telling the cluster to forget the operation history of a resource and re-detect its current state.

The following command cleans up the resource specified by *resource_id*.

```
pcs resource cleanup resource_id
```

If you do not specify a *resource_id*, this command resets the resource status and **failcount** for all resources.

As of Red Hat Enterprise Linux 7.5, the **pcs resource cleanup** command probes only the resources that display as a failed action. To probe all resources on all nodes you can enter the following command:

```
pcs resource refresh
```

By default, the **pcs resource refresh** command probes only the nodes where a resource's state is known. To probe all resources even if the state is not known, enter the following command:

```
pcs resource refresh --full
```

CHAPTER 7. RESOURCE CONSTRAINTS

You can determine the behavior of a resource in a cluster by configuring constraints for that resource. You can configure the following categories of constraints:

- **location** constraints – A location constraint determines which nodes a resource can run on. Location constraints are described in [Section 7.1, “Location Constraints”](#).
- **order** constraints – An order constraint determines the order in which the resources run. Order constraints are described in [Section 7.2, “Order Constraints”](#).
- **colocation** constraints – A colocation constraint determines where resources will be placed relative to other resources. Colocation constraints are described in [Section 7.3, “Colocation of Resources”](#).

As a shorthand for configuring a set of constraints that will locate a set of resources together and ensure that the resources start sequentially and stop in reverse order, Pacemaker supports the concept of resource groups. For information on resource groups, see [Section 6.5, “Resource Groups”](#).

7.1. LOCATION CONSTRAINTS

Location constraints determine which nodes a resource can run on. You can configure location constraints to determine whether a resource will prefer or avoid a specified node.

In addition to location constraints, the node on which a resource runs is influenced by the **resource-stickiness** value for that resource, which determines to what degree a resource prefers to remain on the node where it is currently running. For information on setting the **resource-stickiness** value, see [Section 7.1.5, “Configuring a Resource to Prefer its Current Node”](#).

7.1.1. Basic Location Constraints

You can configure a basic location constraint to specify whether a resource prefers or avoid a node, with an optional **score** value to indicate the relative degree of preference for the constraint.

The following command creates a location constraint for a resource to prefer the specified node or nodes. Note that it is possible to create constraints on a particular resource for more than one node with a single command.

```
pcs constraint location rsc prefers node[=score] [node[=score]] ...
```

The following command creates a location constraint for a resource to avoid the specified node or nodes.

```
pcs constraint location rsc avoids node[=score] [node[=score]] ...
```

[Table 7.1, “Simple Location Constraint Options”](#) summarizes the meanings of the options for configuring location constraints in their simplest form.

Table 7.1. Simple Location Constraint Options

Field	Description
rsc	A resource name

Field	Description
node	A node's name
score	<p>Positive integer value to indicate the preference for whether a resource should prefer or avoid a node. INFINITY is the default score value for a resource location constraint.</p> <p>A value of INFINITY for score in a pcs constraint location rsc prefers command indicates that the resource will prefer that node if the node is available, but does not prevent the resource from running on another node if the specified node is unavailable.</p> <p>A value of INFINITY for score in a pcs constraint location rsc avoids command indicates that the resource will never run on that node, even if no other node is available. This is the equivalent of setting a pcs constraint location add command with a score of INFINITY.</p>

The following command creates a location constraint to specify that the resource **Webserver** prefers node **node1**.

```
# pcs constraint location Webserver prefers node1
```

As of Red Hat Enterprise Linux 7.4, **pcs** supports regular expressions in location constraints on the command line. These constraints apply to multiple resources based on the regular expression matching resource name. This allows you to configure multiple location constraints with a single command line.

The following command creates a location constraint to specify that resources **dummy0** to **dummy9** prefer **node1**.

```
# pcs constraint location 'regexp%dummy[0-9]' prefers node1
```

Since Pacemaker uses POSIX extended regular expressions as documented at http://pubs.opengroup.org/onlinepubs/9699919799/basedefs/V1_chap09.html#tag_09_04, you can specify the same constraint with the following command.

```
# pcs constraint location 'regexp%dummy[[digit]]' prefers node1
```

7.1.2. Advanced Location Constraints

When configuring a location constraint on a node, you can use the **resource-discovery** option of the **pcs constraint location** command to indicate a preference for whether Pacemaker should perform resource discovery on this node for the specified resource. Limiting resource discovery to a subset of nodes the resource is physically capable of running on can significantly boost performance when a large set of nodes is present. When **pacemaker_remote** is in use to expand the node count into the hundreds of nodes range, this option should be considered.

The following command shows the format for specifying the **resource-discovery** option of the **pcs constraint location** command. Note that *id* is the constraint id. The meanings of *rsc*, *node*, and *score* are summarized in [Table 7.1, “Simple Location Constraint Options”](#). In this command, a positive value for *score* corresponds to a basic location ` constraint that configures a resource to prefer a node, while a negative value for *score* corresponds to a basic location ` constraint that configures a resource to avoid a node. As with basic location constraints, you can use regular expressions for resources with these constraints as well.

`pcs constraint location add id rsc node score [resource-discovery=option]`

Table 7.2, “Resource Discovery Values” summarizes the meanings of the values you can specify for the **resource-discovery** option.

Table 7.2. Resource Discovery Values

Value	Description
always	Always perform resource discovery for the specified resource on this node. This is the default resource-discovery value for a resource location constraint.
never	Never perform resource discovery for the specified resource on this node.
exclusive	Perform resource discovery for the specified resource only on this node (and other nodes similarly marked as exclusive). Multiple location constraints using exclusive discovery for the same resource across different nodes creates a subset of nodes resource-discovery is exclusive to. If a resource is marked for exclusive discovery on one or more nodes, that resource is only allowed to be placed within that subset of nodes.

Note that setting the **resource-discovery** option to **never** or **exclusive** allows the possibility for the resource to be active in those locations without the cluster’s knowledge. This can lead to the resource being active in more than one location if the service is started outside the cluster’s control (for example, by **systemd** or by an administrator). This can also occur if the **resource-discovery** property is changed while part of the cluster is down or suffering split-brain, or if the **resource-discovery** property is changed for a resource and node while the resource is active on that node. For this reason, using this option is appropriate only when you have more than eight nodes and there is a way to guarantee that the resource can run only in a particular location (for example, when the required software is not installed anywhere else).

7.1.3. Using Rules to Determine Resource Location

For more complicated location constraints, you can use Pacemaker rules to determine a resource’s location. For general information about Pacemaker rules and the properties you can set, see [Chapter 11, Pacemaker Rules](#).

Use the following command to configure a Pacemaker constraint that uses rules. If **score** is omitted, it defaults to INFINITY. If **resource-discovery** is omitted, it defaults to **always**. For information on the **resource-discovery** option, see [Section 7.1.2, “Advanced Location Constraints”](#). As with basic location constraints, you can use regular expressions for resources with these constraints as well.

When using rules to configure location constraints, the value of **score** can be positive or negative, with a positive value indicating “prefers” and a negative value indicating “avoids”.

`pcs constraint location rsc rule [resource-discovery=option] [role=master|slave] [score=score | score-attribute=attribute] expression`

The *expression* option can be one of the following where *duration_options* and *date_spec_options* are: hours, monthdays, weekdays, yeardays, months, weeks, years, weekyears, moon as described in [Table 11.5, “Properties of a Date Specification”](#).

- **defined|not_defined attribute**

- ***attribute* lt|gt|lte|gte|eq|ne [string|integer|version] *value***
- **date gt|lt *date***
- **date in-range *date* to *date***
- **date in-range *date* to duration *duration_options* ...**
- **date-spec *date_spec_options***
- ***expression* and|or *expression***
- **(*expression*)**

The following location constraint configures an expression that is true if now is any time in the year 2018.

```
# pcs constraint location Webserver rule score=INFINITY date-spec years=2018
```

The following command configures an expression that is true from 9 am to 5 pm, Monday through Friday. Note that the hours value of 16 matches up to 16:59:59, as the numeric value (hour) still matches.

```
# pcs constraint location Webserver rule score=INFINITY date-spec hours="9-16" weekdays="1-5"
```

The following command configures an expression that is true when there is a full moon on Friday the thirteenth.

```
# pcs constraint location Webserver rule date-spec weekdays=5 monthdays=13 moon=4
```

7.1.4. Location Constraint Strategy

Using any of the location constraints described in [Section 7.1.1, “Basic Location Constraints”](#), [Section 7.1.2, “Advanced Location Constraints”](#), and [Section 7.1.3, “Using Rules to Determine Resource Location”](#) you can configure a general strategy for specifying which nodes a resources can run on:

- Opt-In Clusters – Configure a cluster in which, by default, no resource can run anywhere and then selectively enable allowed nodes for specific resources. The procedure for configuring an opt-in cluster is described in [Section 7.1.4.1, “Configuring an “Opt-In” Cluster”](#).
- Opt-Out Clusters – Configure a cluster in which, by default, all resources can run anywhere and then create location constraints for resources that are not allowed to run on specific nodes. The procedure for configuring an opt-out cluster is described in [Section 7.1.4.2, “Configuring an “Opt-Out” Cluster”](#). This is the default Pacemaker strategy.

Whether you should choose to configure your cluster as an opt-in or opt-out cluster depends both on your personal preference and the make-up of your cluster. If most of your resources can run on most of the nodes, then an opt-out arrangement is likely to result in a simpler configuration. On the other hand, if most resources can only run on a small subset of nodes an opt-in configuration might be simpler.

7.1.4.1. Configuring an “Opt-In” Cluster

To create an opt-in cluster, set the **symmetric-cluster** cluster property to **false** to prevent resources from running anywhere by default.

```
# pcs property set symmetric-cluster=false
```

Enable nodes for individual resources. The following commands configure location constraints so that the resource **Webserver** prefers node **example-1**, the resource **Database** prefers node **example-2**, and both resources can fail over to node **example-3** if their preferred node fails. When configuring location constraints for an opt-in cluster, setting a score of zero allows a resource to run on a node without indicating any preference to prefer or avoid the node.

```
# pcs constraint location Webserver prefers example-1=200
# pcs constraint location Webserver prefers example-3=0
# pcs constraint location Database prefers example-2=200
# pcs constraint location Database prefers example-3=0
```

7.1.4.2. Configuring an "Opt-Out" Cluster

To create an opt-out cluster, set the **symmetric-cluster** cluster property to **true** to allow resources to run everywhere by default.

```
# pcs property set symmetric-cluster=true
```

The following commands will then yield a configuration that is equivalent to the example in [Section 7.1.4.1, "Configuring an "Opt-In" Cluster"](#). Both resources can fail over to node **example-3** if their preferred node fails, since every node has an implicit score of 0.

```
# pcs constraint location Webserver prefers example-1=200
# pcs constraint location Webserver avoids example-2=INFINITY
# pcs constraint location Database avoids example-1=INFINITY
# pcs constraint location Database prefers example-2=200
```

Note that it is not necessary to specify a score of INFINITY in these commands, since that is the default value for the score.

7.1.5. Configuring a Resource to Prefer its Current Node

Resources have a **resource-stickiness** value that you can set as a meta attribute when you create the resource, as described in [Section 6.4, "Resource Meta Options"](#). The **resource-stickiness** value determines how much a resource wants to remain on the node where it is currently running. Pacemaker considers the **resource-stickiness** value in conjunction with other settings (for example, the score values of location constraints) to determine whether to move a resource to another node or to leave it in place.

By default, a resource is created with a **resource-stickiness** value of 0. Pacemaker's default behavior when **resource-stickiness** is set to 0 and there are no location constraints is to move resources so that they are evenly distributed among the cluster nodes. This may result in healthy resources moving more often than you desire. To prevent this behavior, you can set the default **resource-stickiness** value to 1. This default will apply to all resources in the cluster. This small value can be easily overridden by other constraints that you create, but it is enough to prevent Pacemaker from needlessly moving healthy resources around the cluster.

The following command sets the default resource-stickiness value to 1.

```
# pcs resource defaults resource-stickiness=1
```

If the **resource-stickiness** value is set, then no resources will move to a newly-added node. If resource balancing is desired at that point, you can temporarily set the **resource-stickiness** value back to 0.

Note that if a location constraint score is higher than the resource-stickiness value, the cluster may still move a healthy resource to the node where the location constraint points.

For further information about how Pacemaker determines where to place a resource, see [Section 9.6, "Utilization and Placement Strategy"](#).

7.2. ORDER CONSTRAINTS

Order constraints determine the order in which the resources run.

Use the following command to configure an order constraint.

```
pcs constraint order [action] resource_id then [action] resource_id [options]
```

[Table 7.3, "Properties of an Order Constraint"](#). summarizes the properties and options for configuring order constraints.

Table 7.3. Properties of an Order Constraint

Field	Description
resource_id	The name of a resource on which an action is performed.
action	<p>The action to perform on a resource. Possible values of the <i>action</i> property are as follows:</p> <ul style="list-style-type: none"> * start - Start the resource. * stop - Stop the resource. * promote - Promote the resource from a slave resource to a master resource. * demote - Demote the resource from a master resource to a slave resource. <p>If no action is specified, the default action is start. For information on master and slave resources, see Section 9.2, "Multistate Resources: Resources That Have Multiple Modes".</p>
kind option	<p>How to enforce the constraint. The possible values of the kind option are as follows:</p> <ul style="list-style-type: none"> * Optional - Only applies if both resources are executing the specified action. For information on optional ordering, see Section 7.2.2, "Advisory Ordering". * Mandatory - Always (default value). If the first resource you specified is stopping or cannot be started, the second resource you specified must be stopped. For information on mandatory ordering, see Section 7.2.1, "Mandatory Ordering". * Serialize - Ensure that no two stop/start actions occur concurrently for a set of resources.
symmetrical option	If true, which is the default, stop the resources in the reverse order. Default value: true

7.2.1. Mandatory Ordering

A mandatory constraints indicates that the second resource you specify cannot run without the first resource you specify being active. This is the default value of the **kind** option. Leaving the default value ensures that the second resource you specify will react when the first resource you specify changes state.

- If the first resource you specified was running and is stopped, the second resource you specified will also be stopped (if it is running).
- If the first resource you specified resource was not running and cannot be started, the resource you specified will be stopped (if it is running).
- If the first resource you specified is (re)started while the second resource you specified is running, the second resource you specified will be stopped and restarted.

Note, however, that the cluster reacts to each state change. If the first resource is restarted and is in a started state again before the second resource initiated a stop operation, the second resource will not need to be restarted.

7.2.2. Advisory Ordering

When the **kind=Optional** option is specified for an order constraint, the constraint is considered optional and only applies if both resources are executing the specified actions. Any change in state by the first resource you specify will have no effect on the second resource you specify.

The following command configures an advisory ordering constraint for the resources named **VirtualIP** and **dummy_resource**.

```
# pcs constraint order VirtualIP then dummy_resource kind=Optional
```

7.2.3. Ordered Resource Sets

A common situation is for an administrator to create a chain of ordered resources, where, for example, resource A starts before resource B which starts before resource C. If your configuration requires that you create a set of resources that is colocated and started in order, you can configure a resource group that contains those resources, as described in [Section 6.5, “Resource Groups”](#). There are some situations, however, where configuring the resources that need to start in a specified order as a resource group is not appropriate:

- You may need to configure resources to start in order and the resources are not necessarily colocated.
- You may have a resource C that must start after either resource A or B has started but there is no relationship between A and B.
- You may have resources C and D that must start after both resources A and B have started, but there is no relationship between A and B or between C and D.

In these situations, you can create an order constraint on a set or sets of resources with the **pcs constraint order set** command.

You can set the following options for a set of resources with the **pcs constraint order set** command.

- **sequential**, which can be set to **true** or **false** to indicate whether the set of resources must be ordered relative to each other.

Setting **sequential** to **false** allows a set to be ordered relative to other sets in the ordering

constraint, without its members being ordered relative to each other. Therefore, this option makes sense only if multiple sets are listed in the constraint; otherwise, the constraint has no effect.

- **require-all**, which can be set to **true** or **false** to indicate whether all of the resources in the set must be active before continuing. Setting **require-all** to **false** means that only one resource in the set needs to be started before continuing on to the next set. Setting **require-all** to **false** has no effect unless used in conjunction with unordered sets, which are sets for which **sequential** is set to **false**.
- **action**, which can be set to **start**, **promote**, **demote** or **stop**, as described in [Table 7.3, "Properties of an Order Constraint"](#).

You can set the following constraint options for a set of resources following the **setoptions** parameter of the **pcs constraint order set** command.

- **id**, to provide a name for the constraint you are defining.
- **score**, to indicate the degree of preference for this constraint. For information on this option, see [Table 7.4, "Properties of a Colocation Constraint"](#).

```
pcs constraint order set resource1 resource2 [resourceN]... [options] [set resourceX resourceY ... [options]] [setoptions [constraint_options]]
```

If you have three resources named **D1**, **D2**, and **D3**, the following command configures them as an ordered resource set.

```
# pcs constraint order set D1 D2 D3
```

7.2.4. Removing Resources From Ordering Constraints

Use the following command to remove resources from any ordering constraint.

```
pcs constraint order remove resource1 [resourceN]...
```

7.3. COLOCATION OF RESOURCES

A colocation constraint determines that the location of one resource depends on the location of another resource.

There is an important side effect of creating a colocation constraint between two resources: it affects the order in which resources are assigned to a node. This is because you cannot place resource A relative to resource B unless you know where resource B is. So when you are creating colocation constraints, it is important to consider whether you should colocate resource A with resource B or resource B with resource A.

Another thing to keep in mind when creating colocation constraints is that, assuming resource A is colocated with resource B, the cluster will also take into account resource A's preferences when deciding which node to choose for resource B.

The following command creates a colocation constraint.

```
pcs constraint colocation add [master|slave] source_resource with [master|slave] target_resource [score] [options]
```

For information on master and slave resources, see [Section 9.2, “Multistate Resources: Resources That Have Multiple Modes”](#).

[Table 7.4, “Properties of a Colocation Constraint”](#) summarizes the properties and options for configuring colocation constraints.

Table 7.4. Properties of a Colocation Constraint

Field	Description
source_resource	The colocation source. If the constraint cannot be satisfied, the cluster may decide not to allow the resource to run at all.
target_resource	The colocation target. The cluster will decide where to put this resource first and then decide where to put the source resource.
score	Positive values indicate the resource should run on the same node. Negative values indicate the resources should not run on the same node. A value of +INFINITY , the default value, indicates that the <i>source_resource</i> must run on the same node as the <i>target_resource</i> . A value of -INFINITY indicates that the <i>source_resource</i> must not run on the same node as the <i>target_resource</i> .

7.3.1. Mandatory Placement

Mandatory placement occurs any time the constraint's score is **+INFINITY** or **-INFINITY**. In such cases, if the constraint cannot be satisfied, then the *source_resource* is not permitted to run. For **score=INFINITY**, this includes cases where the *target_resource* is not active.

If you need **myresource1** to always run on the same machine as **myresource2**, you would add the following constraint:

```
# pcs constraint colocation add myresource1 with myresource2 score=INFINITY
```

Because **INFINITY** was used, if **myresource2** cannot run on any of the cluster nodes (for whatever reason) then **myresource1** will not be allowed to run.

Alternatively, you may want to configure the opposite, a cluster in which **myresource1** cannot run on the same machine as **myresource2**. In this case use **score=-INFINITY**

```
# pcs constraint colocation add myresource1 with myresource2 score=-INFINITY
```

Again, by specifying **-INFINITY**, the constraint is binding. So if the only place left to run is where **myresource2** already is, then **myresource1** may not run anywhere.

7.3.2. Advisory Placement

If mandatory placement is about "must" and "must not", then advisory placement is the "I would prefer if" alternative. For constraints with scores greater than **-INFINITY** and less than **INFINITY**, the cluster will try to accommodate your wishes but may ignore them if the alternative is to stop some of the cluster resources. Advisory colocation constraints can combine with other elements of the configuration to behave as if they were mandatory.

7.3.3. Colocating Sets of Resources

If your configuration requires that you create a set of resources that is colocated and started in order, you can configure a resource group that contains those resources, as described in [Section 6.5, "Resource Groups"](#). There are some situations, however, where configuring the resources that need to be colocated as a resource group is not appropriate:

- You may need to colocate a set of resources but the resources do not necessarily need to start in order.
- You may have a resource C that must be colated with either resource A or B has started but there is no relationship between A and B.
- You may have resources C and D that must be colocated with both resources A and B, but there is no relationship between A and B or between C and D.

In these situations, you can create a colocation constraint on a set or sets of resources with the **pcs constraint colocation set** command.

You can set the following options for a set of resources with the **pcs constraint colocation set** command.

- **sequential**, which can be set to **true** or **false** to indicate whether the members of the set must be colocated with each other.
Setting **sequential** to **false** allows the members of this set to be colocated with another set listed later in the constraint, regardless of which members of this set are active. Therefore, this option makes sense only if another set is listed after this one in the constraint; otherwise, the constraint has no effect.
- **role**, which can be set to **Stopped**, **Started**, **Master**, or **Slave**. For information on multistate resources, see [Section 9.2, "Multistate Resources: Resources That Have Multiple Modes"](#).

You can set the following constraint options for a set of resources following the **setoptions** parameter of the **pcs constraint colocation set** command.

- **kind**, to indicate how to enforce the constraint. For information on this option, see [Table 7.3, "Properties of an Order Constraint"](#).
- **symmetrical**, to indicate the order in which to stop the resources. If true, which is the default, stop the resources in the reverse order. Default value: **true**
- **id**, to provide a name for the constraint you are defining.

When listing members of a set, each member is colocated with the one before it. For example, "set A B" means "B is colocated with A". However, when listing multiple sets, each set is colocated with the one after it. For example, "set C D sequential=false set A B" means "set C D (where C and D have no relation between each other) is colocated with set A B (where B is colocated with A)".

The following command creates a colocation constraint on a set or sets of resources.

```
pcs constraint colocation set resource1 resource2 [resourceN]... [options] [set resourceX resourceY
... [options]] [setoptions [constraint_options]]
```

7.3.4. Removing Colocation Constraints

Use the following command to remove colocation constraints with *source_resource*.

```
pcs constraint colocation remove source_resource target_resource
```

7.4. DISPLAYING CONSTRAINTS

There are several commands you can use to display constraints that have been configured.

The following command lists all current location, order, and colocation constraints.

```
pcs constraint list|show
```

The following command lists all current location constraints.

- If **resources** is specified, location constraints are displayed per resource. This is the default behavior.
- If **nodes** is specified, location constraints are displayed per node.
- If specific resources or nodes are specified, then only information about those resources or nodes is displayed.

```
pcs constraint location [show resources|nodes [specific nodes|resources]] [--full]
```

The following command lists all current ordering constraints. If the **--full** option is specified, show the internal constraint IDs.

```
pcs constraint order show [--full]
```

The following command lists all current colocation constraints. If the **--full** option is specified, show the internal constraint IDs.

```
pcs constraint colocation show [--full]
```

The following command lists the constraints that reference specific resources.

```
pcs constraint ref resource ...
```

CHAPTER 8. MANAGING CLUSTER RESOURCES

This chapter describes various commands you can use to manage cluster resources. It provides information on the following procedures.

- [Section 8.1, “Manually Moving Resources Around the Cluster”](#)
- [Section 8.2, “Moving Resources Due to Failure”](#)
- [Section 8.4, “Enabling, Disabling, and Banning Cluster Resources”](#)
- [Section 8.5, “Disabling a Monitor Operation”](#)

8.1. MANUALLY MOVING RESOURCES AROUND THE CLUSTER

You can override the cluster and force resources to move from their current location. There are two occasions when you would want to do this:

- When a node is under maintenance, and you need to move all resources running on that node to a different node
- When individually specified resources needs to be moved

To move all resources running on a node to a different node, you put the node in standby mode. For information on putting a cluster node in standby node, see [Section 4.4.5, “Standby Mode”](#).

You can move individually specified resources in either of the following ways.

- You can use the **pcs resource move** command to move a resource off a node on which it is currently running, as described in [Section 8.1.1, “Moving a Resource from its Current Node”](#).
- You can use the **pcs resource relocate run** command to move a resource to its preferred node, as determined by current cluster status, constraints, location of resources and other settings. For information on this command, see [Section 8.1.2, “Moving a Resource to its Preferred Node”](#).

8.1.1. Moving a Resource from its Current Node

To move a resource off the node on which it is currently running, use the following command, specifying the *resource_id* of the resource as defined. Specify the **destination_node** if you want to indicate on which node to run the resource that you are moving.

```
pcs resource move resource_id [destination_node] [--master] [lifetime=lifetime]
```



NOTE

When you execute the **pcs resource move** command, this adds a constraint to the resource to prevent it from running on the node on which it is currently running. You can execute the **pcs resource clear** or the **pcs constraint delete** command to remove the constraint. This does not necessarily move the resources back to the original node; where the resources can run at that point depends on how you have configured your resources initially.

If you specify the **--master** parameter of the **pcs resource move** command, the scope of the constraint is limited to the master role and you must specify *master_id* rather than *resource_id*.

You can optionally configure a **lifetime** parameter for the **pcs resource move** command to indicate a period of time the constraint should remain. You specify the units of a **lifetime** parameter according to the format defined in ISO 8601, which requires that you specify the unit as a capital letter such as Y (for years), M (for months), W (for weeks), D (for days), H (for hours), M (for minutes), and S (for seconds).

To distinguish a unit of minutes(M) from a unit of months(M), you must specify PT before indicating the value in minutes. For example, a **lifetime** parameter of 5M indicates an interval of five months, while a **lifetime** parameter of PT5M indicates an interval of five minutes.

The **lifetime** parameter is checked at intervals defined by the **cluster-recheck-interval** cluster property. By default this value is 15 minutes. If your configuration requires that you check this parameter more frequently, you can reset this value with the following command.

```
pcs property set cluster-recheck-interval=value
```

You can optionally configure a **--wait[=n]** parameter for the **pcs resource move** command to indicate the number of seconds to wait for the resource to start on the destination node before returning 0 if the resource is started or 1 if the resource has not yet started. If you do not specify n, the default resource timeout will be used.

The following command moves the resource **resource1** to node **example-node2** and prevents it from moving back to the node on which it was originally running for one hour and thirty minutes.

```
pcs resource move resource1 example-node2 lifetime=PT1H30M
```

The following command moves the resource **resource1** to node **example-node2** and prevents it from moving back to the node on which it was originally running for thirty minutes.

```
pcs resource move resource1 example-node2 lifetime=PT30M
```

For information on resource constraints, see [Chapter 7, Resource Constraints](#).

8.1.2. Moving a Resource to its Preferred Node

After a resource has moved, either due to a failover or to an administrator manually moving the node, it will not necessarily move back to its original node even after the circumstances that caused the failover have been corrected. To relocate resources to their preferred node, use the following command. A preferred node is determined by the current cluster status, constraints, resource location, and other settings and may change over time.

```
pcs resource relocate run [resource1] [resource2] ...
```

If you do not specify any resources, all resource are relocated to their preferred nodes.

This command calculates the preferred node for each resource while ignoring resource stickiness. After calculating the preferred node, it creates location constraints which will cause the resources to move to their preferred nodes. Once the resources have been moved, the constraints are deleted automatically. To remove all constraints created by the **pcs resource relocate run** command, you can enter the **pcs resource relocate clear** command. To display the current status of resources and their optimal node ignoring resource stickiness, enter the **pcs resource relocate show** command.

8.2. MOVING RESOURCES DUE TO FAILURE

When you create a resource, you can configure the resource so that it will move to a new node after a defined number of failures by setting the **migration-threshold** option for that resource. Once the threshold has been reached, this node will no longer be allowed to run the failed resource until:

- The administrator manually resets the resource's **failcount** using the **pcs resource failcount** command.
- The resource's **failure-timeout** value is reached.

The value of **migration-threshold** is set to **INFINITY** by default. **INFINITY** is defined internally as a very large but finite number. A value of 0 disables the **migration-threshold** feature.



NOTE

Setting a **migration-threshold** for a resource is not the same as configuring a resource for migration, in which the resource moves to another location without loss of state.

The following example adds a migration threshold of 10 to the resource named **dummy_resource**, which indicates that the resource will move to a new node after 10 failures.

```
# pcs resource meta dummy_resource migration-threshold=10
```

You can add a migration threshold to the defaults for the whole cluster with the following command.

```
# pcs resource defaults migration-threshold=10
```

To determine the resource's current failure status and limits, use the **pcs resource failcount** command.

There are two exceptions to the migration threshold concept; they occur when a resource either fails to start or fails to stop. If the cluster property **start-failure-is-fatal** is set to **true** (which is the default), start failures cause the **failcount** to be set to **INFINITY** and thus always cause the resource to move immediately. For information on the **start-failure-is-fatal** option, see [Table 12.1, "Cluster Properties"](#).

Stop failures are slightly different and crucial. If a resource fails to stop and STONITH is enabled, then the cluster will fence the node in order to be able to start the resource elsewhere. If STONITH is not enabled, then the cluster has no way to continue and will not try to start the resource elsewhere, but will try to stop it again after the failure timeout.

8.3. MOVING RESOURCES DUE TO CONNECTIVITY CHANGES

Setting up the cluster to move resources when external connectivity is lost is a two step process.

1. Add a **ping** resource to the cluster. The **ping** resource uses the system utility of the same name to test if a list of machines (specified by DNS host name or IPv4/IPv6 address) are reachable and uses the results to maintain a node attribute called **pingd**.
2. Configure a location constraint for the resource that will move the resource to a different node when connectivity is lost.

[Table 6.1, "Resource Properties"](#) describes the properties you can set for a **ping** resource.

Table 8.1. Properties of a ping resources

Field	Description
dampen	The time to wait (dampening) for further changes to occur. This prevents a resource from bouncing around the cluster when cluster nodes notice the loss of connectivity at slightly different times.
multiplier	The number of connected ping nodes gets multiplied by this value to get a score. Useful when there are multiple ping nodes configured.
host_list	The machines to contact in order to determine the current connectivity status. Allowed values include resolvable DNS host names, IPv4 and IPv6 addresses. The entries in the host list are space separated.

The following example command creates a **ping** resource that verifies connectivity to **gateway.example.com**. In practice, you would verify connectivity to your network gateway/router. You configure the **ping** resource as a clone so that the resource will run on all cluster nodes.

```
# pcs resource create ping ocf:pacemaker:ping dampen=5s multiplier=1000
host_list=gateway.example.com clone
```

The following example configures a location constraint rule for the existing resource named **Webserver**. This will cause the **Webserver** resource to move to a host that is able to ping **gateway.example.com** if the host that it is currently running on cannot ping **gateway.example.com**.

```
# pcs constraint location Webserver rule score=-INFINITY pingd lt 1 or not_defined pingd
```

8.4. ENABLING, DISABLING, AND BANNING CLUSTER RESOURCES

In addition to the **pcs resource move** and **pcs resource relocate** commands described in [Section 8.1, “Manually Moving Resources Around the Cluster”](#), there are a variety of other commands you can use to control the behavior of cluster resources.

You can manually stop a running resource and prevent the cluster from starting it again with the following command. Depending on the rest of the configuration (constraints, options, failures, and so on), the resource may remain started. If you specify the **--wait** option, **pcs** will wait up to 'n' seconds for the resource to stop and then return 0 if the resource is stopped or 1 if the resource has not stopped. If 'n' is not specified it defaults to 60 minutes.

```
pcs resource disable resource_id [--wait[=n]]
```

You can use the following command to allow the cluster to start a resource. Depending on the rest of the configuration, the resource may remain stopped. If you specify the **--wait** option, **pcs** will wait up to 'n' seconds for the resource to start and then return 0 if the resource is started or 1 if the resource has not started. If 'n' is not specified it defaults to 60 minutes.

```
pcs resource enable resource_id [--wait[=n]]
```

Use the following command to prevent a resource from running on a specified node, or on the current node if no node is specified.

```
pcs resource ban resource_id [node] [--master] [lifetime=lifetime] [--wait[=n]]
```

Note that when you execute the **pcs resource ban** command, this adds a **-INFINITY** location constraint to the resource to prevent it from running on the indicated node. You can execute the **pcs resource clear** or the **pcs constraint delete** command to remove the constraint. This does not necessarily move the resources back to the indicated node; where the resources can run at that point depends on how you have configured your resources initially. For information on resource constraints, see [Chapter 7, Resource Constraints](#).

If you specify the **--master** parameter of the **pcs resource ban** command, the scope of the constraint is limited to the master role and you must specify *master_id* rather than *resource_id*.

You can optionally configure a **lifetime** parameter for the **pcs resource ban** command to indicate a period of time the constraint should remain. For information on specifying units for the **lifetime** parameter and on specifying the intervals at which the **lifetime** parameter should be checked, see [Section 8.1, “Manually Moving Resources Around the Cluster”](#).

You can optionally configure a **--wait[=n]** parameter for the **pcs resource ban** command to indicate the number of seconds to wait for the resource to start on the destination node before returning 0 if the resource is started or 1 if the resource has not yet started. If you do not specify n, the default resource timeout will be used.

You can use the **debug-start** parameter of the **pcs resource** command to force a specified resource to start on the current node, ignoring the cluster recommendations and printing the output from starting the resource. This is mainly used for debugging resources; starting resources on a cluster is (almost) always done by Pacemaker and not directly with a **pcs** command. If your resource is not starting, it is usually due to either a misconfiguration of the resource (which you debug in the system log), constraints that prevent the resource from starting, or the resource being disabled. You can use this command to test resource configuration, but it should not normally be used to start resources in a cluster.

The format of the **debug-start** command is as follows.

```
pcs resource debug-start resource_id
```

8.5. DISABLING A MONITOR OPERATION

The easiest way to stop a recurring monitor is to delete it. However, there can be times when you only want to disable it temporarily. In such cases, add **enabled="false"** to the operation's definition with the **pcs resource update** command. When you want to reinstate the monitoring operation, set **enabled="true"** to the operation's definition.

When you update a resource's operation with the **pcs resource update** command, any options you do not specifically call out are reset to their default values. For example, if you have configured a monitoring operation with a custom timeout value of 600, running the following commands will reset the timeout value to the default value of 20 (or whatever you have set the default value to with the **pcs resource ops default** command).

```
# pcs resource update resourceXZY op monitor enabled=false
# pcs resource update resourceXZY op monitor enabled=true
```

In order to maintain the original value of 600 for this option, when you reinstate the monitoring operation you must specify that value, as in the following example.

```
# pcs resource update resourceXZY op monitor timeout=600 enabled=true
```

8.6. MANAGED RESOURCES

You can set a resource to **unmanaged** mode, which indicates that the resource is still in the configuration but Pacemaker does not manage the resource.

The following command sets the indicated resources to **unmanaged** mode.

```
pcs resource unmanage resource1 [resource2] ...
```

The following command sets resources to **managed** mode, which is the default state.

```
pcs resource manage resource1 [resource2] ...
```


You can specify the name of a resource group with the **pcs resource manage** or **pcs resource unmanage** command. The command will act on all of the resources in the group, so that you can set all of the resources in a group to **managed** or **unmanaged** mode with a single command and then manage the contained resources individually.

CHAPTER 9. ADVANCED CONFIGURATION

This chapter describes advanced resource types and advanced configuration features that Pacemaker supports.

9.1. RESOURCE CLONES

You can clone a resource so that the resource can be active on multiple nodes. For example, you can use cloned resources to configure multiple instances of an IP resource to distribute throughout a cluster for node balancing. You can clone any resource provided the resource agent supports it. A clone consists of one resource or one resource group.



NOTE

Only resources that can be active on multiple nodes at the same time are suitable for cloning. For example, a **Filesystem** resource mounting a non-clustered file system such as **ext4** from a shared memory device should not be cloned. Since the **ext4** partition is not cluster aware, this file system is not suitable for read/write operations occurring from multiple nodes at the same time.

9.1.1. Creating and Removing a Cloned Resource

You can create a resource and a clone of that resource at the same time with the following command.

```
pcs resource create resource_id standard:provider:type|type [resource options] \
clone [meta clone options]
```

The name of the clone will be ***resource_id*-clone**.

You cannot create a resource group and a clone of that resource group in a single command.

Alternately, you can create a clone of a previously-created resource or resource group with the following command.


```
pcs resource clone resource_id | group_name [clone options]...
```

The name of the clone will be ***resource_id*-clone** or ***group_name*-clone**.



NOTE

You need to configure resource configuration changes on one node only.



NOTE

When configuring constraints, always use the name of the group or clone.

When you create a clone of a resource, the clone takes on the name of the resource with **-clone** appended to the name. The following commands creates a resource of type **apache** named **webfarm** and a clone of that resource named **webfarm-clone**.

```
# pcs resource create webfarm apache clone
```



NOTE

When you create a resource or resource group clone that will be ordered after another clone, you should almost always set the **interleave=true** option. This ensures that copies of the dependent clone can stop or start when the clone it depends on has stopped or started on the same node. If you do not set this option, if a cloned resource B depends on a cloned resource A and a node leaves the cluster, when the node returns to the cluster and resource A starts on that node, then all of the copies of resource B on all of the nodes will restart. This is because when a dependent cloned resource does not have the **Interleave** option set, all instances of that resource depend on any running instance of the resource it depends on.

Use the following command to remove a clone of a resource or a resource group. This does not remove the resource or resource group itself.

```
pcs resource unclone resource_id | group_name
```

For information on resource options, see [Section 6.1, “Resource Creation”](#).

[Table 9.1, “Resource Clone Options”](#) describes the options you can specify for a cloned resource.

Table 9.1. Resource Clone Options

Field	Description
priority, target-role, is-managed	Options inherited from resource that is being cloned, as described in Table 6.3, “Resource Meta Options” .
clone-max	How many copies of the resource to start. Defaults to the number of nodes in the cluster.
clone-node-max	How many copies of the resource can be started on a single node; the default value is 1 .
notify	When stopping or starting a copy of the clone, tell all the other copies beforehand and when the action was successful. Allowed values: false, true . The default value is false .
globally-unique	<p>Does each copy of the clone perform a different function? Allowed values: false, true</p> <p>If the value of this option is false, these resources behave identically everywhere they are running and thus there can be only one copy of the clone active per machine.</p> <p>If the value of this option is true, a copy of the clone running on one machine is not equivalent to another instance, whether that instance is running on another node or on the same node. The default value is true if the value of clone-node-max is greater than one; otherwise the default value is false.</p>
ordered	Should the copies be started in series (instead of in parallel). Allowed values: false, true . The default value is false .

Field	Description
interleave	Changes the behavior of ordering constraints (between clones/masters) so that copies of the first clone can start or stop as soon as the copy on the same node of the second clone has started or stopped (rather than waiting until every instance of the second clone has started or stopped). Allowed values: false , true . The default value is false .
clone-min	If a value is specified, any clones which are ordered after this clone will not be able to start until the specified number of instances of the original clone are running, even if the interleave option is set to true .

9.1.2. Clone Constraints

In most cases, a clone will have a single copy on each active cluster node. You can, however, set **clone-max** for the resource clone to a value that is less than the total number of nodes in the cluster. If this is the case, you can indicate which nodes the cluster should preferentially assign copies to with resource location constraints. These constraints are written no differently to those for regular resources except that the clone's id must be used.

The following command creates a location constraint for the cluster to preferentially assign resource clone **webfarm-clone** to **node1**.

```
# pcs constraint location webfarm-clone prefers node1
```

Ordering constraints behave slightly differently for clones. In the example below, because the **interleave** clone option is left to default as **false**, no instance of **webfarm-stats** will start until all instances of **webfarm-clone** that need to be started have done so. Only if no copies of **webfarm-clone** can be started then **webfarm-stats** will be prevented from being active. Additionally, **webfarm-clone** will wait for **webfarm-stats** to be stopped before stopping itself.

```
# pcs constraint order start webfarm-clone then webfarm-stats
```

Colocation of a regular (or group) resource with a clone means that the resource can run on any machine with an active copy of the clone. The cluster will choose a copy based on where the clone is running and the resource's own location preferences.

Colocation between clones is also possible. In such cases, the set of allowed locations for the clone is limited to nodes on which the clone is (or will be) active. Allocation is then performed as normally.

The following command creates a colocation constraint to ensure that the resource **webfarm-stats** runs on the same node as an active copy of **webfarm-clone**.

```
# pcs constraint colocation add webfarm-stats with webfarm-clone
```

9.1.3. Clone Stickiness

To achieve a stable allocation pattern, clones are slightly sticky by default. If no value for **resource-stickiness** is provided, the clone will use a value of 1. Being a small value, it causes minimal disturbance to the score calculations of other resources but is enough to prevent Pacemaker from needlessly moving copies around the cluster.

9.2. MULTISTATE RESOURCES: RESOURCES THAT HAVE MULTIPLE MODES

Multistate resources are a specialization of Clone resources. They allow the instances to be in one of two operating modes; these are called **Master** and **Slave**. The names of the modes do not have specific meanings, except for the limitation that when an instance is started, it must come up in the **Slave** state.

You can create a resource as a master/slave clone with the following single command.

```
pcs resource create resource_id standard:provider:type|type [resource options] master  
[master_options]
```

The name of the master/slave clone will be ***resource_id*-master**.



NOTE

For Red Hat Enterprise Linux release 7.3 and earlier, use the following format to create a master/slave clone.

```
pcs resource create resource_id standard:provider:type|type [resource options] --  
master [meta master_options]
```

Alternately, you can create a master/slave resource from a previously-created resource or resource group with the following command: When you use this command, you can specify a name for the master/slave clone. If you do not specify a name, the name of the master/slave clone will be ***resource_id*-master** or ***group_name*-master**.

```
pcs resource master master/slave_name resource_id/group_name [master_options]
```

For information on resource options, see [Section 6.1, “Resource Creation”](#).

[Table 9.2, “Properties of a Multistate Resource”](#) describes the options you can specify for a multistate resource.

Table 9.2. Properties of a Multistate Resource

Field	Description
id	Your name for the multistate resource
priority, target-role, is-managed	See Table 6.3, “Resource Meta Options” .
clone-max, clone-node-max, notify, globally-unique, ordered, interleave	See Table 9.1, “Resource Clone Options” .
master-max	How many copies of the resource can be promoted to master status; default 1.
master-node-max	How many copies of the resource can be promoted to master status on a single node; default 1.

9.2.1. Monitoring Multi-State Resources

To add a monitoring operation for the master resource only, you can add an additional monitor operation to the resource. Note, however, that every monitor operation on a resource must have a different interval.

The following example configures a monitor operation with an interval of 11 seconds on the master resource for **ms_resource**. This monitor operation is in addition to the default monitor operation with the default monitor interval of 10 seconds.

```
# pcs resource op add ms_resource interval=11s role=Master
```

9.2.2. Multistate Constraints

In most cases, a multistate resources will have a single copy on each active cluster node. If this is not the case, you can indicate which nodes the cluster should preferentially assign copies to with resource location constraints. These constraints are written no differently than those for regular resources.

For information on resource location constraints, see [Section 7.1, “Location Constraints”](#).

You can create a colocation constraint which specifies whether the resources are master or slave resources. The following command creates a resource colocation constraint.

```
pcs constraint colocation add [master|slave] source_resource with [master|slave] target_resource [score] [options]
```

For information on colocation constraints, see [Section 7.3, “Colocation of Resources”](#).

When configuring an ordering constraint that includes multistate resources, one of the actions that you can specify for the resources is **promote**, indicating that the resource be promoted from slave to master. Additionally, you can specify an action of **demote**, indicated that the resource be demoted from master to slave.

The command for configuring an order constraint is as follows.

```
pcs constraint order [action] resource_id then [action] resource_id [options]
```

For information on resource order constraints, see [Section 7.2, “Order Constraints”](#).

9.2.3. Multistate Stickiness

To achieve a stable allocation pattern, multistate resources are slightly sticky by default. If no value for **resource-stickiness** is provided, the multistate resource will use a value of 1. Being a small value, it causes minimal disturbance to the score calculations of other resources but is enough to prevent Pacemaker from needlessly moving copies around the cluster.

9.3. CONFIGURING A VIRTUAL DOMAIN AS A RESOURCE

You can configure a virtual domain that is managed by the **libvirt** virtualization framework as a cluster resource with the **pcs resource create** command, specifying **VirtualDomain** as the resource type.

When configuring a virtual domain as a resource, take the following considerations into account:

- A virtual domain should be stopped before you configure it as a cluster resource.

- Once a virtual domain is a cluster resource, it should not be started, stopped, or migrated except through the cluster tools.
- Do not configure a virtual domain that you have configured as a cluster resource to start when its host boots.
- All nodes must have access to the necessary configuration files and storage devices for each managed virtual domain.

If you want the cluster to manage services within the virtual domain itself, you can configure the virtual domain as a guest node. For information on configuring guest nodes, see [Section 9.4, “The pacemaker_remote Service”](#)

For information on configuring virtual domains, see the [Virtualization Deployment and Administration Guide](#).

[Table 9.3, “Resource Options for Virtual Domain Resources”](#) describes the resource options you can configure for a **VirtualDomain** resource.

Table 9.3. Resource Options for Virtual Domain Resources

Field	Default	Description
config		(required) Absolute path to the libvirt configuration file for this virtual domain.
hypervisor	System dependent	Hypervisor URI to connect to. You can determine the system's default URL by running the virsh --quiet uri command.
force_stop	0	Always forcefully shut down ("destroy") the domain on stop. The default behavior is to resort to a forceful shutdown only after a graceful shutdown attempt has failed. You should set this to true only if your virtual domain (or your virtualization back end) does not support graceful shutdown.
migration_transport	System dependent	Transport used to connect to the remote hypervisor while migrating. If this parameter is omitted, the resource will use libvirt 's default transport to connect to the remote hypervisor.
migration_network_suffix		Use a dedicated migration network. The migration URI is composed by adding this parameter's value to the end of the node name. If the node name is a fully qualified domain name (FQDN), insert the suffix immediately prior to the first period (.) in the FQDN. Ensure that this composed host name is locally resolvable and the associated IP address is reachable through the favored network.
monitor_scripts		To additionally monitor services within the virtual domain, add this parameter with a list of scripts to monitor. Note: When monitor scripts are used, the start and migrate_from operations will complete only when all monitor scripts have completed successfully. Be sure to set the timeout of these operations to accommodate this delay

Field	Default	Description
autoSet_utilization_cp_u	true	If set to true , the agent will detect the number of domainU 's vCPUs from virsh , and put it into the CPU utilization of the resource when the monitor is executed.
autoSet_utilization_hv_memory	true	If set to true , the agent will detect the number of Max memory from virsh , and put it into the hv_memory utilization of the source when the monitor is executed.
migrateport	random highport	This port will be used in the qemu migrate URI. If unset, the port will be a random highport.
snapshot		Path to the snapshot directory where the virtual machine image will be stored. When this parameter is set, the virtual machine's RAM state will be saved to a file in the snapshot directory when stopped. If on start a state file is present for the domain, the domain will be restored to the same state it was in right before it stopped last. This option is incompatible with the force_stop option.

In addition to the **VirtualDomain** resource options, you can configure the **allow-migrate** metadata option to allow live migration of the resource to another node. When this option is set to **true**, the resource can be migrated without loss of state. When this option is set to **false**, which is the default state, the virtual domain will be shut down on the first node and then restarted on the second node when it is moved from one node to the other.

Use the following procedure to create a **VirtualDomain** resource:

1. To create the **VirtualDomain** resource agent for the management of the virtual machine, Pacemaker requires the virtual machine's xml config file to be dumped to a file on disk. For example, if you created a virtual machine named **guest1**, dump the xml to a file somewhere on the host. You can use a file name of your choosing; this example uses **/etc/pacemaker/guest1.xml**.

```
# virsh dumpxml guest1 > /etc/pacemaker/guest1.xml
```

2. If it is running, shut down the guest node. Pacemaker will start the node when it is configured in the cluster.
3. Configure the **VirtualDomain** resource with the **pcs resource create** command. For example, The following command configures a **VirtualDomain** resource named **VM**. Since the **allow-migrate** option is set to **true** a **pcs move VM nodeX** command would be done as a live migration.

```
# pcs resource create VM VirtualDomain config=.../vm.xml \
    migration_transport=ssh meta allow-migrate=true
```

9.4. THE PACEMAKER_REMOTE SERVICE

The **pacemaker_remote** service allows nodes not running **corosync** to integrate into the cluster and have the cluster manage their resources just as if they were real cluster nodes.

Among the capabilities that the **pacemaker_remote** service provides are the following:

- The **pacemaker_remote** service allows you to scale beyond the Red Hat support limit of 32 nodes for RHEL 7.7.
- The **pacemaker_remote** service allows you to manage a virtual environment as a cluster resource and also to manage individual services within the virtual environment as cluster resources.

The following terms are used to describe the **pacemaker_remote** service.

- *cluster node* – A node running the High Availability services (**pacemaker** and **corosync**).
- *remote node* – A node running **pacemaker_remote** to remotely integrate into the cluster without requiring **corosync** cluster membership. A remote node is configured as a cluster resource that uses the **ocf:pacemaker:remote** resource agent.
- *guest node* – A virtual guest node running the **pacemaker_remote** service. The virtual guest resource is managed by the cluster; it is both started by the cluster and integrated into the cluster as a remote node.
- *pacemaker_remote* – A service daemon capable of performing remote application management within remote nodes and guest nodes (KVM and LXC) in a Pacemaker cluster environment. This service is an enhanced version of Pacemaker's local resource management daemon (LRMD) that is capable of managing resources remotely on a node not running corosync.
- *LXC* – A Linux Container defined by the **libvirt-lxc** Linux container driver.

A Pacemaker cluster running the **pacemaker_remote** service has the following characteristics.

- Remote nodes and guest nodes run the **pacemaker_remote** service (with very little configuration required on the virtual machine side).
- The cluster stack (**pacemaker** and **corosync**), running on the cluster nodes, connects to the **pacemaker_remote** service on the remote nodes, allowing them to integrate into the cluster.
- The cluster stack (**pacemaker** and **corosync**), running on the cluster nodes, launches the guest nodes and immediately connects to the **pacemaker_remote** service on the guest nodes, allowing them to integrate into the cluster.

The key difference between the cluster nodes and the remote and guest nodes that the cluster nodes manage is that the remote and guest nodes are not running the cluster stack. This means the remote and guest nodes have the following limitations:

- they do not take place in quorum
- they do not execute fencing device actions
- they are not eligible to be the cluster's Designated Controller (DC)
- they do not themselves run the full range of **pcs** commands

On the other hand, remote nodes and guest nodes are not bound to the scalability limits associated with the cluster stack.

Other than these noted limitations, the remote and guest nodes behave just like cluster nodes in respect to resource management, and the remote and guest nodes can themselves be fenced. The cluster is fully capable of managing and monitoring resources on each remote and guest node: You can

build constraints against them, put them in standby, or perform any other action you perform on cluster nodes with the **pcs** commands. Remote and guest nodes appear in cluster status output just as cluster nodes do.

9.4.1. Host and Guest Authentication

The connection between cluster nodes and pacemaker_remote is secured using Transport Layer Security (TLS) with pre-shared key (PSK) encryption and authentication over TCP (using port 3121 by default). This means both the cluster node and the node running **pacemaker_remote** must share the same private key. By default this key must be placed at **/etc/pacemaker/authkey** on both cluster nodes and remote nodes.

As of Red Hat Enterprise Linux 7.4, the **pcs cluster node add-guest** command sets up the **authkey** for guest nodes and the **pcs cluster node add-remote** command sets up the **authkey** for remote nodes.

9.4.2. Guest Node Resource Options

When configuring a virtual machine or LXC resource to act as a guest node, you create a **VirtualDomain** resource, which manages the virtual machine. For descriptions of the options you can set for a **VirtualDomain** resource, see [Table 9.3, “Resource Options for Virtual Domain Resources”](#).

In addition to the **VirtualDomain** resource options, metadata options define the resource as a guest node and define the connection parameters. As of Red Hat Enterprise Linux 7.4, you should set these resource options with the **pcs cluster node add-guest** command. In releases earlier than 7.4, you can set these options when creating the resource. [Table 9.4, “Metadata Options for Configuring KVM/LXC Resources as Remote Nodes”](#) describes these metadata options.

Table 9.4. Metadata Options for Configuring KVM/LXC Resources as Remote Nodes

Field	Default	Description
remote-node	<none>	The name of the guest node this resource defines. This both enables the resource as a guest node and defines the unique name used to identify the guest node. WARNING: This value cannot overlap with any resource or node IDs.
remote-port	3121	Configures a custom port to use for the guest connection to pacemaker_remote
remote-addr	remote-node value used as host name	The IP address or host name to connect to if remote node’s name is not the host name of the guest
remote-connect-timeout	60s	Amount of time before a pending guest connection will time out

9.4.3. Remote Node Resource Options

A remote node is defined as a cluster resource with **ocf:pacemaker:remote** as the resource agent. In Red Hat Enterprise Linux 7.4, you should create this resource with the **pcs cluster node add-remote** command. In releases earlier than 7.4, you can create this resource with the **pcs resource create** command. [Table 9.5, “Resource Options for Remote Nodes”](#) describes the resource options you can configure for a **remote** resource.

Table 9.5. Resource Options for Remote Nodes

Field	Default	Description
reconnect_interval	0	Time in seconds to wait before attempting to reconnect to a remote node after an active connection to the remote node has been severed. This wait is recurring. If reconnect fails after the wait period, a new reconnect attempt will be made after observing the wait time. When this option is in use, Pacemaker will keep attempting to reach out and connect to the remote node indefinitely after each wait interval.
server		Server location to connect to. This can be an IP address or host name.
port		TCP port to connect to.

9.4.4. Changing Default Port Location

If you need to change the default port location for either Pacemaker or **pacemaker_remote**, you can set the **PCMK_remote_port** environment variable that affects both of these daemons. This environment variable can be enabled by placing it in the **/etc/sysconfig/pacemaker** file as follows.

```
#==#==# Pacemaker Remote
...
#
# Specify a custom port for Pacemaker Remote connections
PCMK_remote_port=3121
```

When changing the default port used by a particular guest node or remote node, the **PCMK_remote_port** variable must be set in that node's **/etc/sysconfig/pacemaker** file, and the cluster resource creating the guest node or remote node connection must also be configured with the same port number (using the **remote-port** metadata option for guest nodes, or the **port** option for remote nodes).

9.4.5. Configuration Overview: KVM Guest Node

This section provides a high-level summary overview of the steps to perform to have Pacemaker launch a virtual machine and to integrate that machine as a guest node, using **libvirt** and KVM virtual guests.

1. Configure the **VirtualDomain** resources, as described in [Section 9.3, “Configuring a Virtual Domain as a Resource”](#).
2. On systems running Red Hat Enterprise Linux 7.3 and earlier, put the same encryption key with the path **/etc/pacemaker/authkey** on every cluster node and virtual machine with the following procedure. This secures remote communication and authentication.
 1. Enter the following set of commands on every node to create the **authkey** directory with secure permissions.

```
# mkdir -p --mode=0750 /etc/pacemaker
# chgrp haclient /etc/pacemaker
```

2. The following command shows one method to create an encryption key. You should create the key only once and then copy it to all of the nodes.

```
# dd if=/dev/urandom of=/etc/pacemaker/authkey bs=4096 count=1
```

3. For Red Hat Enterprise Linux 7.4, enter the following commands on every virtual machine to install **pacemaker_remote** packages, start the **pcsd** service and enable it to run on startup, and allow TCP port 3121 through the firewall.

```
# yum install pacemaker-remote resource-agents pcs
# systemctl start pcسد.service
# systemctl enable pcسد.service
# firewall-cmd --add-port 3121/tcp --permanent
# firewall-cmd --add-port 2224/tcp --permanent
# firewall-cmd --reload
```

For Red Hat Enterprise Linux 7.3 and earlier, run the following commands on every virtual machine to install **pacemaker_remote** packages, start the **pacemaker_remote** service and enable it to run on startup, and allow TCP port 3121 through the firewall.

```
# yum install pacemaker-remote resource-agents pcs
# systemctl start pacemaker_remote.service
# systemctl enable pacemaker_remote.service
# firewall-cmd --add-port 3121/tcp --permanent
# firewall-cmd --add-port 2224/tcp --permanent
# firewall-cmd --reload
```

4. Give each virtual machine a static network address and unique host name, which should be known to all nodes. For information on setting a static IP address for the guest virtual machine, see the *Virtualization Deployment and Administration Guide*.
5. For Red Hat Enterprise Linux 7.4 and later, use the following command to convert an existing **VirtualDomain** resource into a guest node. This command must be run on a cluster node and not on the guest node which is being added. In addition to converting the resource, this command copies the **/etc/pacemaker/authkey** to the guest node and starts and enables the **pacemaker_remote** daemon on the guest node.

```
pcs cluster node add-guest hostname resource_id [options]
```

For Red Hat Enterprise Linux 7.3 and earlier, use the following command to convert an existing **VirtualDomain** resource into a guest node. This command must be run on a cluster node and not on the guest node which is being added.

```
pcs cluster remote-node add hostname resource_id [options]
```

6. After creating the **VirtualDomain** resource, you can treat the guest node just as you would treat any other node in the cluster. For example, you can create a resource and place a resource constraint on the resource to run on the guest node as in the following commands, which are run from a cluster node. As of Red Hat Enterprise Linux 7.3, you can include guest nodes in groups, which allows you to group a storage device, file system, and VM.

```
# pcs resource create webserver apache configfile=/etc/httpd/conf/httpd.conf op monitor
interval=30s
# pcs constraint location webserver prefers guest1
```

9.4.6. Configuration Overview: Remote Node (Red Hat Enterprise Linux 7.4)

This section provides a high-level summary overview of the steps to perform to configure a Pacemaker Remote node and to integrate that node into an existing Pacemaker cluster environment for Red Hat Enterprise Linux 7.4.

1. On the node that you will be configuring as a remote node, allow cluster-related services through the local firewall.

```
# firewall-cmd --permanent --add-service=high-availability  
success  
# firewall-cmd --reload  
success
```



NOTE

If you are using **iptables** directly, or some other firewall solution besides **firewalld**, simply open the following ports: TCP ports 2224 and 3121.

2. Install the **pacemaker_remote** daemon on the remote node.

```
# yum install -y pacemaker-remote resource-agents pcs
```

3. Start and enable **pcsd** on the remote node.

```
# systemctl start pcsd.service  
# systemctl enable pcsd.service
```

4. If you have not already done so, authenticate **pcs** to the node you will be adding as a remote node.

```
# pcs cluster auth remote1
```

5. Add the remote node resource to the cluster with the following command. This command also syncs all relevant configuration files to the new node, starts the node, and configures it to start **pacemaker_remote** on boot. This command must be run on a cluster node and not on the remote node which is being added.

```
# pcs cluster node add-remote remote1
```

6. After adding the **remote** resource to the cluster, you can treat the remote node just as you would treat any other node in the cluster. For example, you can create a resource and place a resource constraint on the resource to run on the remote node as in the following commands, which are run from a cluster node.

```
# pcs resource create webserver apache configfile=/etc/httpd/conf/httpd.conf op monitor  
interval=30s  
# pcs constraint location webserver prefers remote1
```

**WARNING**

Never involve a remote node connection resource in a resource group, colocation constraint, or order constraint.

- Configure fencing resources for the remote node. Remote nodes are fenced the same way as cluster nodes. Configure fencing resources for use with remote nodes the same as you would with cluster nodes. Note, however, that remote nodes can never initiate a fencing action. Only cluster nodes are capable of actually executing a fencing operation against another node.

9.4.7. Configuration Overview: Remote Node (Red Hat Enterprise Linux 7.3 and earlier)

This section provides a high-level summary overview of the steps to perform to configure a Pacemaker Remote node and to integrate that node into an existing Pacemaker cluster environment in a Red Hat Enterprise Linux 7.3 (and earlier) system.

- On the node that you will be configuring as a remote node, allow cluster-related services through the local firewall.

```
# firewall-cmd --permanent --add-service=high-availability
success
# firewall-cmd --reload
success
```

**NOTE**

If you are using **iptables** directly, or some other firewall solution besides **firewalld**, simply open the following ports: TCP ports 2224 and 3121.

- Install the **pacemaker_remote** daemon on the remote node.

```
# yum install -y pacemaker-remote resource-agents pcs
```

- All nodes (both cluster nodes and remote nodes) must have the same authentication key installed for the communication to work correctly. If you already have a key on an existing node, use that key and copy it to the remote node. Otherwise, create a new key on the remote node.

Enter the following set of commands on the remote node to create a directory for the authentication key with secure permissions.

```
# mkdir -p --mode=0750 /etc/pacemaker
# chgrp haclient /etc/pacemaker
```

The following command shows one method to create an encryption key on the remote node.

```
# dd if=/dev/urandom of=/etc/pacemaker/authkey bs=4096 count=1
```

- Start and enable the **pacemaker_remote** daemon on the remote node.

```
# systemctl enable pacemaker_remote.service
# systemctl start pacemaker_remote.service
```

- On the cluster node, create a location for the shared authentication key with the same path as the authentication key on the remote node and copy the key into that directory. In this example, the key is copied from the remote node where the key was created.

```
# mkdir -p --mode=0750 /etc/pacemaker
# chgrp haclient /etc/pacemaker
# scp remote1:/etc/pacemaker/authkey /etc/pacemaker/authkey
```

- Enter the following command from a cluster node to create a **remote** resource. In this case the remote node is **remote1**.

```
# pcs resource create remote1 ocf:pacemaker:remote
```

- After creating the **remote** resource, you can treat the remote node just as you would treat any other node in the cluster. For example, you can create a resource and place a resource constraint on the resource to run on the remote node as in the following commands, which are run from a cluster node.

```
# pcs resource create webserver apache configfile=/etc/httpd/conf/httpd.conf op monitor
interval=30s
# pcs constraint location webserver prefers remote1
```



WARNING

Never involve a remote node connection resource in a resource group, colocation constraint, or order constraint.

- Configure fencing resources for the remote node. Remote nodes are fenced the same way as cluster nodes. Configure fencing resources for use with remote nodes the same as you would with cluster nodes. Note, however, that remote nodes can never initiate a fencing action. Only cluster nodes are capable of actually executing a fencing operation against another node.

9.4.8. System Upgrades and pacemaker_remote

As of Red Hat Enterprise Linux 7.3, if the **pacemaker_remote** service is stopped on an active Pacemaker Remote node, the cluster will gracefully migrate resources off the node before stopping the node. This allows you to perform software upgrades and other routine maintenance procedures without removing the node from the cluster. Once **pacemaker_remote** is shut down, however, the cluster will immediately try to reconnect. If **pacemaker_remote** is not restarted within the resource's monitor timeout, the cluster will consider the monitor operation as failed.

If you wish to avoid monitor failures when the **pacemaker_remote** service is stopped on an active Pacemaker Remote node, you can use the following procedure to take the node out of the cluster before performing any system administration that might stop **pacemaker_remote**



WARNING

For Red Hat Enterprise Linux release 7.2 and earlier, if **pacemaker_remote** stops on a node that is currently integrated into a cluster, the cluster will fence that node. If the stop happens automatically as part of a **yum update** process, the system could be left in an unusable state (particularly if the kernel is also being upgraded at the same time as **pacemaker_remote**). For Red Hat Enterprise Linux release 7.2 and earlier you must use the following procedure to take the node out of the cluster before performing any system administration that might stop **pacemaker_remote**.

1. Stop the node's connection resource with the **pcs resource disable resourcename**, which will move all services off the node. For guest nodes, this will also stop the VM, so the VM must be started outside the cluster (for example, using **virsh**) to perform any maintenance.
2. Perform the required maintenance.
3. When ready to return the node to the cluster, re-enable the resource with the **pcs resource enable**.

9.5. PACEMAKER SUPPORT FOR DOCKER CONTAINERS (TECHNOLOGY PREVIEW)



IMPORTANT

Pacemaker support for Docker containers is provided for technology preview only. For details on what "technology preview" means, see [Technology Preview Features Support Scope](#).

There is one exception to this feature being Technology Preview: As of Red Hat Enterprise Linux 7.4, Red Hat fully supports the usage of Pacemaker bundles for Red Hat Openstack Platform (RHOSP) deployments.

Pacemaker supports a special syntax for launching a Docker container with any infrastructure it requires: the *bundle*. After you have created a Pacemaker bundle, you can create a Pacemaker resource that the bundle encapsulates.

- [Section 9.5.1, "Configuring a Pacemaker Bundle Resource"](#) describes the syntax for the command to create a Pacemaker bundle and provides tables summarizing the parameters you can define for each bundle parameter.
- [Section 9.5.2, "Configuring a Pacemaker Resource in a Bundle"](#) provides information on configuring a resource contained in a Pacemaker bundle.
- [Section 9.5.3, "Limitations of Pacemaker Bundles"](#) notes the limitations of Pacemaker bundles.

- [Section 9.5.4, “Pacemaker Bundle Configuration Example”](#) provides a Pacemaker bundle configuration example.

9.5.1. Configuring a Pacemaker Bundle Resource

The syntax for the command to create a Pacemaker bundle for a Docker container is as follows. This command creates a bundle that encapsulates no other resources. For information on creating a cluster resource in a bundle see [Section 9.5.2, “Configuring a Pacemaker Resource in a Bundle”](#).

```
pcs resource bundle create bundle_id container docker [container_options] [network network_options] [port-map port_options]... [storage-map storage_options]... [meta meta_options] [--disabled] [--wait[=n]]
```

The required *bundle_id* parameter must be a unique name for the bundle. If the **--disabled** option is specified, the bundle is not started automatically. If the **--wait** option is specified, Pacemaker will wait up to **n** seconds for the bundle to start and then return 0 on success or 1 on error. If **n** is not specified it defaults to 60 minutes.

The following sections describe the parameters you can configure for each element of a Pacemaker bundle.

9.5.1.1. Docker Parameters

[Table 9.6, “Docker Container Parameters”](#) describes the **docker** container options you can set for a bundle.



NOTE

Before configuring a **docker** bundle in Pacemaker, you must install Docker and supply a fully configured Docker image on every node allowed to run the bundle.

Table 9.6. Docker Container Parameters

Field	Default	Description
image		Docker image tag (required)
replicas	Value of promoted-max if that is positive, otherwise 1.	A positive integer specifying the number of container instances to launch
replicas-per-host	1	A positive integer specifying the number of container instances allowed to run on a single node
promoted-max	0	A non-negative integer that, if positive, indicates that the containerized service should be treated as a multistate service, with this many replicas allowed to run the service in the master role
network		If specified, this will be passed to the docker run command as the network setting for the Docker container.

Field	Default	Description
run-command	<code>/usr/sbin/pacemaker_remoted</code> if the bundle contains a resource, otherwise none	This command will be run inside the container when launching it ("PID 1"). If the bundle contains a resource, this command must start the pacemaker_remoted daemon (but it could, for example, be a script that performs others tasks as well).
options		Extra command-line options to pass to the docker run command

9.5.1.2. Bundle Network Parameters

Table 9.7, “Bundle Resource Network Parameters” describes the **network** options you can set for a bundle.

Table 9.7. Bundle Resource Network Parameters

Field	Default	Description
add-host	TRUE	If TRUE, and ip-range-start is used, Pacemaker will automatically ensure that the /etc/hosts file inside the containers has entries for each replica name and its assigned IP.
ip-range-start		If specified, Pacemaker will create an implicit ocf:heartbeat:IPaddr2 resource for each container instance, starting with this IP address, using as many sequential addresses as were specified as the replicas parameter for the Docker element. These addresses can be used from the host's network to reach the service inside the container, although it is not visible within the container itself. Only IPv4 addresses are currently supported.
host-netmask	32	If ip-range-start is specified, the IP addresses are created with this CIDR netmask (as a number of bits).
host-interface		If ip-range-start is specified, the IP addresses are created on this host interface (by default, it will be determined from the IP address).

Field	Default	Description
control-port	3121	<p>If the bundle contains a Pacemaker resource, the cluster will use this integer TCP port for communication with Pacemaker Remote inside the container. Changing this is useful when the container is unable to listen on the default port, which could happen when the container uses the host's network rather than ip-range-start (in which case replicas-per-host must be 1), or when the bundle may run on a Pacemaker Remote node that is already listening on the default port. Any PCMK_remote_port environment variable set on the host or in the container is ignored for bundle connections.</p> <p>When a Pacemaker bundle configuration uses the control-port parameter, then if the bundle has its own IP address the port needs to be open on that IP address on and from all full cluster nodes running corosync. If, instead, the bundle has set the network="host" container parameter, the port needs to be open on each cluster node's IP address from all cluster nodes.</p>



NOTE

Replicas are named by the bundle ID plus a dash and an integer counter starting with zero. For example, if a bundle named **httpd-bundle** has configured **replicas=2**, its containers will be named **httpd-bundle-0** and **httpd-bundle-1**.

In addition to the network parameters, you can optionally specify **port-map** parameters for a bundle. [Table 9.8, "Bundle Resource port-map Parameters"](#) describes these **port-map** parameters.

Table 9.8. Bundle Resource port-map Parameters

Field	Default	Description
id		A unique name for the port mapping (required)
port		If this is specified, connections to this TCP port number on the host network (on the container's assigned IP address, if ip-range-start is specified) will be forwarded to the container network. Exactly one of port or range must be specified in a port-mapping.
internal-port	Value of port	If port and internal-port are specified, connections to port on the host's network will be forwarded to this port on the container network.
range		If range is specified, connections to these TCP port numbers (expressed as <i>first_port-last_port</i>) on the host network (on the container's assigned IP address, if ip-range-start is specified) will be forwarded to the same ports in the container network. Exactly one of port or range must be specified in a port mapping.

**NOTE**

If the bundle contains a resource, Pacemaker will automatically map the **control-port**, so it is not necessary to specify that port in a port mapping.


9.5.1.3. Bundle Storage Parameters

You can optionally configure **storage-map** parameters for a bundle. Table 9.9, “Bundle Resource Storage Mapping Parameters” describes these parameters.

Table 9.9. Bundle Resource Storage Mapping Parameters

Field	Default	Description
id		A unique name for the storage mapping (required)
source-dir		The absolute path on the host’s filesystem that will be mapped into the container. Exactly one of source-dir and source-dir-root parameter must be specified when configuring a storage-map parameter.
source-dir-root		The start of a path on the host’s filesystem that will be mapped into the container, using a different subdirectory on the host for each container instance. The subdirectory will be named with the same name as the bundle name, plus a dash and an integer counter starting with 0. Exactly one source-dir and source-dir-root parameter must be specified when configuring a storage-map parameter.
target-dir		The path name within the container where the host storage will be mapped (required)
options		File system mount options to use when mapping the storage

As an example of how subdirectories on a host are named using the **source-dir-root** parameter, if **source-dir-root=/path/to/my/directory**, **target-dir=/srv/appdata**, and the bundle is named **mybundle** with **replicas=2**, then the cluster will create two container instances with host names **mybundle-0** and **mybundle-1** and create two directories on the host running the containers: **/path/to/my/directory/mybundle-0** and **/path/to/my/directory/mybundle-1**. Each container will be given one of those directories, and any application running inside the container will see the directory as **/srv/appdata**.

**NOTE**

Pacemaker does not define the behavior if the source directory does not already exist on the host. However, it is expected that the container technology or its resource agent will create the source directory in that case.



NOTE

If the bundle contains a Pacemaker resource, Pacemaker will automatically map the equivalent of **source-dir=/etc/pacemaker/authkey**, **target-dir=/etc/pacemaker/authkey** and **source-dir-root=/var/log/pacemaker/bundle**, **target-dir=/var/log** into the container, so it is not necessary to specify those paths in when configuring **storage-map** parameters.



IMPORTANT

The **PCMK_authkey_location** environment variable must not be set to anything other than the default of **/etc/pacemaker/authkey** on any node in the cluster.

9.5.2. Configuring a Pacemaker Resource in a Bundle

A bundle may optionally contain one Pacemaker cluster resource. As with a resource that is not contained in a bundle, the cluster resource may have operations, instance attributes, and metadata attributes defined. If a bundle contains a resource, the container image must include the Pacemaker Remote daemon, and **ip-range-start** or **control-port** must be configured in the bundle. Pacemaker will create an implicit **ocf:pacemaker:remote** resource for the connection, launch Pacemaker Remote within the container, and monitor and manage the resource by means of Pacemaker Remote. If the bundle has more than one container instance (replica), the Pacemaker resource will function as an implicit clone, which will be a multistate clone if the bundle has configured the **promoted-max** option as greater than zero.

You create a resource in a Pacemaker bundle with the **pcs resource create** command by specifying the **bundle** parameter for the command and the bundle ID in which to include the resource. For an example of creating a Pacemaker bundle that contains a resource, see [Section 9.5.4, “Pacemaker Bundle Configuration Example”](#).



IMPORTANT

Containers in bundles that contain a resource must have an accessible networking environment, so that Pacemaker on the cluster nodes can contact Pacemaker Remote inside the container. For example, the **docker** option **--net=none** should not be used with a resource. The default (using a distinct network space inside the container) works in combination with the **ip-range-start** parameter. If the **docker** option **--net=host** is used (making the container share the host’s network space), a unique **control-port** parameter should be specified for each bundle. Any firewall must allow access to the **control-port**.

9.5.2.1. Node Attributes and Bundle Resources

If the bundle contains a cluster resource, the resource agent may want to set node attributes such as master scores. However, with containers, it is not apparent which node should get the attribute.

If the container uses shared storage that is the same no matter which node the container is hosted on, then it is appropriate to use the master score on the bundle node itself. On the other hand, if the container uses storage exported from the underlying host, then it may be more appropriate to use the master score on the underlying host. Since this depends on the particular situation, the **container-attribute-target** resource metadata attribute allows the user to specify which approach to use. If it is set to **host**, then user-defined node attributes will be checked on the underlying host. If it is anything else, the local node (in this case the bundle node) is used. This behavior applies only to user-defined attributes; the cluster will always check the local node for cluster-defined attributes such as **#uname**.

If **container-attribute-target** is set to **host**, the cluster will pass additional environment variables to the resource agent that allow it to set node attributes appropriately.

9.5.2.2. Metadata Attributes and Bundle Resources

Any metadata attribute set on a bundle will be inherited by the resource contained in a bundle and any resources implicitly created by Pacemaker for the bundle. This includes options such as **priority**, **target-role**, and **is-managed**.

9.5.3. Limitations of Pacemaker Bundles

Pacemaker bundles operate with the following limitations:

- Bundles may not be included in groups or explicitly cloned with a **pcs** command. This includes a resource that the bundle contains, and any resources implicitly created by Pacemaker for the bundle. Note, however, that if a bundle is configured with a value of **replicas** greater than one, the bundle behaves as if it were a clone.
- Restarting Pacemaker while a bundle is unmanaged or the cluster is in maintenance mode may cause the bundle to fail.
- Bundles do not have instance attributes, utilization attributes, or operations, although a resource contained in a bundle may have them.
- A bundle that contains a resource can run on a Pacemaker Remote node only if the bundle uses a distinct **control-port**.

9.5.4. Pacemaker Bundle Configuration Example

The following example creates a Pacemaker **bundle** resource with a bundle ID of **httpd-bundle** that contains an **ocf:heartbeat:apache** resource with a resource ID of **httpd**.

This procedure requires the following prerequisite configuration:

- Docker has been installed and enabled on every node in the cluster.
- There is an existing Docker image, named **pcmktest:http**
- The container image includes the Pacemaker Remote daemon.
- The container image includes a configured Apache web server.
- Every node in the cluster has directories **/var/local/containers/httpd-bundle-0**, **/var/local/containers/httpd-bundle-1**, and **/var/local/containers/httpd-bundle-2**, containing an **index.html** file for the web server root. In production, a single, shared document root would be more likely, but for the example this configuration allows you to make the **index.html** file on each host different so that you can connect to the web server and verify which **index.html** file is being served.

This procedure configures the following parameters for the Pacemaker bundle:

- The bundle ID is **httpd-bundle**.
- The previously-configured Docker container image is **pcmktest:http**.
- This example will launch three container instances.

- This example will pass the command-line option **--log-driver=journald** to the **docker run** command. This parameter is not required, but is included to show how to pass an extra option to the **docker** command. A value of **--log-driver=journald** means that the system logs inside the container will be logged in the underlying host's **systemd** journal.
- Pacemaker will create three sequential implicit **ocf:heartbeat:IPaddr2** resources, one for each container image, starting with the IP address 192.168.122.131.
- The IP addresses are created on the host interface eth0.
- The IP addresses are created with a CIDR netmask of 24.
- This example creates a port map ID of **http-port**; connections to port 80 on the container's assigned IP address will be forwarded to the container network.
- This example creates a storage map ID of **httpd-root**. For this storage mapping:
 - The value of **source-dir-root** is **/var/local/containers**, which specifies the start of the path on the host's file system that will be mapped into the container, using a different subdirectory on the host for each container instance.
 - The value of **target-dir** is **/var/www/html**, which specifies the path name within the container where the host storage will be mapped.
 - The file system **rw** mount option will be used when mapping the storage.
 - Since this example container includes a resource, Pacemaker will automatically map the equivalent of **source-dir=/etc/pacemaker/authkey** in the container, so you do not need to specify that path in the storage mapping.

In this example, the existing cluster configuration is put into a temporary file named **temp-cib.xml**, which is then copied to a file named **temp-cib.xml.deltasrc**. All modifications to the cluster configuration are made to the **tmp-cib.xml** file. When the updates are complete, this procedure uses the **diff-against** option of the **pcs cluster cib-push** command so that only the updates to the configuration file are pushed to the active configuration file.

```
# pcs cluster cib tmp-cib.xml
# cp tmp-cib.xml tmp-cib.xml.deltasrc
# pcs -f tmp-cib.xml resource bundle create httpd-bundle \
  container docker image=pcmktest:http replicas=3 \
  options="--log-driver=journald \
  network ip-range-start=192.168.122.131 host-interface=eth0 \
  host-netmask=24 port-map id=httpd-port port=80 \
  storage-map id=httpd-root source-dir-root=/var/local/containers \
  target-dir=/var/www/html options=rw \
# pcs -f tmp-cib.xml resource create httpd ocf:heartbeat:apache \
  statusurl=http://localhost/server-status bundle httpd-bundle
# pcs cluster cib-push tmp-cib.xml diff-against=tmp-cib.xml.deltasrc
```

9.6. UTILIZATION AND PLACEMENT STRATEGY

Pacemaker decides where to place a resource according to the resource allocation scores on every node. The resource will be allocated to the node where the resource has the highest score. This allocation score is derived from a combination of factors, including resource constraints, **resource-stickiness** settings, prior failure history of a resource on each node, and utilization of each node.

If the resource allocation scores on all the nodes are equal, by the default placement strategy Pacemaker will choose a node with the least number of allocated resources for balancing the load. If the number of resources on each node is equal, the first eligible node listed in the CIB will be chosen to run the resource.

Often, however, different resources use significantly different proportions of a node's capacities (such as memory or I/O). You cannot always balance the load ideally by taking into account only the number of resources allocated to a node. In addition, if resources are placed such that their combined requirements exceed the provided capacity, they may fail to start completely or they may run with degraded performance. To take these factors into account, Pacemaker allows you to configure the following components:

- the capacity a particular node provides
- the capacity a particular resource requires
- an overall strategy for placement of resources

The following sections describe how to configure these components.

9.6.1. Utilization Attributes

To configure the capacity that a node provides or a resource requires, you can use *utilization attributes* for nodes and resources. You do this by setting a utilization variable for a resource and assigning a value to that variable to indicate what the resource requires, and then setting that same utilization variable for a node and assigning a value to that variable to indicate what that node provides.

You can name utilization attributes according to your preferences and define as many name and value pairs as your configuration needs. The values of utilization attributes must be integers.

As of Red Hat Enterprise Linux 7.3, you can set utilization attributes with the **pcs** command.

The following example configures a utilization attribute of CPU capacity for two nodes, naming the attribute **cpu**. It also configures a utilization attribute of RAM capacity, naming the attribute **memory**. In this example:

- Node 1 is defined as providing a CPU capacity of two and a RAM capacity of 2048
- Node 2 is defined as providing a CPU capacity of four and a RAM capacity of 2048

```
# pcs node utilization node1 cpu=2 memory=2048
# pcs node utilization node2 cpu=4 memory=2048
```

The following example specifies the same utilization attributes that three different resources require. In this example:

- resource **dummy-small** requires a CPU capacity of 1 and a RAM capacity of 1024
- resource **dummy-medium** requires a CPU capacity of 2 and a RAM capacity of 2048
- resource **dummy-large** requires a CPU capacity of 1 and a RAM capacity of 3072

```
# pcs resource utilization dummy-small cpu=1 memory=1024
# pcs resource utilization dummy-medium cpu=2 memory=2048
# pcs resource utilization dummy-large cpu=3 memory=3072
```

A node is considered eligible for a resource if it has sufficient free capacity to satisfy the resource's requirements, as defined by the utilization attributes.

9.6.2. Placement Strategy

After you have configured the capacities your nodes provide and the capacities your resources require, you need to set the **placement-strategy** cluster property, otherwise the capacity configurations have no effect. For information on setting cluster properties, see [Chapter 12, Pacemaker Cluster Properties](#).

Four values are available for the **placement-strategy** cluster property:

- **default** – Utilization values are not taken into account at all. Resources are allocated according to allocation scores. If scores are equal, resources are evenly distributed across nodes.
- **utilization** – Utilization values are taken into account only when deciding whether a node is considered eligible (that is, whether it has sufficient free capacity to satisfy the resource's requirements). Load-balancing is still done based on the number of resources allocated to a node.
- **balanced** – Utilization values are taken into account when deciding whether a node is eligible to serve a resource and when load-balancing, so an attempt is made to spread the resources in a way that optimizes resource performance.
- **minimal** – Utilization values are taken into account only when deciding whether a node is eligible to serve a resource. For load-balancing, an attempt is made to concentrate the resources on as few nodes as possible, thereby enabling possible power savings on the remaining nodes.

The following example command sets the value of **placement-strategy** to **balanced**. After running this command, Pacemaker will ensure the load from your resources will be distributed evenly throughout the cluster, without the need for complicated sets of colocation constraints.

```
# pcs property set placement-strategy=balanced
```

9.6.3. Resource Allocation

The following subsections summarize how Pacemaker allocates resources.

9.6.3.1. Node Preference

Pacemaker determines which node is preferred when allocating resources according to the following strategy.

- The node with the highest node weight gets consumed first. Node weight is a score maintained by the cluster to represent node health.
- If multiple nodes have the same node weight:
 - If the **placement-strategy** cluster property is **default** or **utilization**:
 - The node that has the least number of allocated resources gets consumed first.
 - If the numbers of allocated resources are equal, the first eligible node listed in the CIB gets consumed first.

- If the **placement-strategy** cluster property is **balanced**:
 - The node that has the most free capacity gets consumed first.
 - If the free capacities of the nodes are equal, the node that has the least number of allocated resources gets consumed first.
 - If the free capacities of the nodes are equal and the number of allocated resources is equal, the first eligible node listed in the CIB gets consumed first.
- If the **placement-strategy** cluster property is **minimal**, the first eligible node listed in the CIB gets consumed first.

9.6.3.2. Node Capacity

Pacemaker determines which node has the most free capacity according to the following strategy.

- If only one type of utilization attribute has been defined, free capacity is a simple numeric comparison.
- If multiple types of utilization attributes have been defined, then the node that is numerically highest in the most attribute types has the most free capacity. For example:
 - If NodeA has more free CPUs, and NodeB has more free memory, then their free capacities are equal.
 - If NodeA has more free CPUs, while NodeB has more free memory and storage, then NodeB has more free capacity.

9.6.3.3. Resource Allocation Preference

Pacemaker determines which resource is allocated first according to the following strategy.

- The resource that has the highest priority gets allocated first. For information on setting priority for a resource, see [Table 6.3, “Resource Meta Options”](#).
- If the priorities of the resources are equal, the resource that has the highest score on the node where it is running gets allocated first, to prevent resource shuffling.
- If the resource scores on the nodes where the resources are running are equal or the resources are not running, the resource that has the highest score on the preferred node gets allocated first. If the resource scores on the preferred node are equal in this case, the first runnable resource listed in the CIB gets allocated first.

9.6.4. Resource Placement Strategy Guidelines

To ensure that Pacemaker's placement strategy for resources works most effectively, you should take the following considerations into account when configuring your system.

- Make sure that you have sufficient physical capacity.

If the physical capacity of your nodes is being used to near maximum under normal conditions, then problems could occur during failover. Even without the utilization feature, you may start to experience timeouts and secondary failures.

- Build some buffer into the capabilities you configure for the nodes.

Advertise slightly more node resources than you physically have, on the assumption that a Pacemaker resource will not use 100% of the configured amount of CPU, memory, and so forth all the time. This practice is sometimes called overcommit.

- Specify resource priorities.

If the cluster is going to sacrifice services, it should be the ones you care about least. Ensure that resource priorities are properly set so that your most important resources are scheduled first. For information on setting resource priorities, see [Table 6.3, “Resource Meta Options”](#).

9.6.5. The NodeUtilization Resource Agent (Red Hat Enterprise Linux 7.4 and later)

Red Hat Enterprise Linux 7.4 supports the **NodeUtilization** resource agent. The NodeUtilization agent can detect the system parameters of available CPU, host memory availability, and hypervisor memory availability and add these parameters into the CIB. You can run the agent as a clone resource to have it automatically populate these parameters on each node.

For information on the **NodeUtilization** resource agent and the resource options for this agent, run the **pcs resource describe NodeUtilization** command.

9.7. CONFIGURING STARTUP ORDER FOR RESOURCE DEPENDENCIES NOT MANAGED BY PACEMAKER (RED HAT ENTERPRISE LINUX 7.4 AND LATER)

It is possible for a cluster to include resources with dependencies that are not themselves managed by the cluster. In this case, you must ensure that those dependencies are started before Pacemaker is started and stopped after Pacemaker is stopped.

As of Red Hat Enterprise Linux 7.4, you can configure your startup order to account for this situation by means of the **systemd resource-agents-deps** target. You can create a **systemd** drop-in unit for this target and Pacemaker will order itself appropriately relative to this target.

For example, if a cluster includes a resource that depends on the external service **foo** that is not managed by the cluster, you can create the drop-in unit **/etc/systemd/system/resource-agents-deps.target.d/foo.conf** that contains the following:

```
[Unit]
Requires=foo.service
After=foo.service
```

After creating a drop-in unit, run the **systemctl daemon-reload** command.

A cluster dependency specified in this way can be something other than a service. For example, you may have a dependency on mounting a file system at **/srv**, in which case you would create a **systemd** file **srv.mount** for it according to the **systemd** documentation, then create a drop-in unit as described here with **srv.mount** in the **.conf** file instead of **foo.service** to make sure that Pacemaker starts after the disk is mounted.

9.8. QUERYING A PACEMAKER CLUSTER WITH SNMP (RED HAT ENTERPRISE LINUX 7.5 AND LATER)

As of Red Hat Enterprise Linux 7.5, you can use the **pcs_snmp_agent** daemon to query a Pacemaker cluster for data by means of SNMP. The **pcs_snmp_agent** daemon is an SNMP agent that connects to the master agent (**snmpd**) by means of **agentx** protocol. The **pcs_snmp_agent** agent does not work as

a standalone agent as it only provides data to the master agent.

The following procedure sets up a basic configuration for a system to use SNMP with a Pacemaker cluster. You run this procedure on each node of the cluster from which you will be using SNMP to fetch data for the cluster.

1. Install the **pcs-snmp** package on each node of the cluster. This will also install the **net-snmp** package which provides the **snmpd** daemon.

```
# yum install pcs-snmp
```

2. Add the following line to the **/etc/snmp/snmpd.conf** configuration file to set up the **snmpd** daemon as **master agentx**.

```
master agentx
```

3. Add the following line to the **/etc/snmp/snmpd.conf** configuration file to enable **pcs_snmp_agent** in the same SNMP configuration.

```
view systemview included .1.3.6.1.4.1.32723.100
```

4. Start the **pcs_snmp_agent** service.

```
# systemctl start pcs_snmp_agent.service
# systemctl enable pcs_snmp_agent.service
```

5. To check the configuration, display the status of the cluster with the **pcs status** and then try to fetch the data from SNMP to check whether it corresponds to the output. Note that when you use SNMP to fetch data, only primitive resources are provided.

The following example shows the output of a **pcs status** command on a running cluster with one failed action.

```
# pcs status
Cluster name: rhel75-cluster
Stack: corosync
Current DC: rhel75-node2 (version 1.1.18-5.el7-1a4ef7d180) - partition with quorum
Last updated: Wed Nov 15 16:07:44 2017
Last change: Wed Nov 15 16:06:40 2017 by hacluster via cibadmin on rhel75-node1

2 nodes configured
14 resources configured (1 DISABLED)

Online: [ rhel75-node1 rhel75-node2 ]

Full list of resources:

fencing      (stonith:fence_xvm): Started rhel75-node1
dummy5 (ocf::pacemaker:Dummy): Stopped (disabled)
dummy6 (ocf::pacemaker:Dummy): Stopped
dummy7 (ocf::pacemaker:Dummy): Started rhel75-node2
dummy8 (ocf::pacemaker:Dummy): Started rhel75-node1
dummy9 (ocf::pacemaker:Dummy): Started rhel75-node2
Resource Group: group1
```

```

dummy1  (ocf::pacemaker:Dummy): Started rhel75-node1
dummy10 (ocf::pacemaker:Dummy): Started rhel75-node1
Clone Set: group2-clone [group2]
    Started: [ rhel75-node1 rhel75-node2 ]
Clone Set: dummy4-clone [dummy4]
    Started: [ rhel75-node1 rhel75-node2 ]

```

Failed Actions:

```

* dummy6_start_0 on rhel75-node1 'unknown error' (1): call=87, status=complete,
exitreason="",
last-rc-change='Wed Nov 15 16:05:55 2017', queued=0ms, exec=20ms

```

```

# snmpwalk -v 2c -c public localhost PACEMAKER-PCS-V1-MIB::pcmkPcsV1Cluster
PACEMAKER-PCS-V1-MIB::pcmkPcsV1ClusterName.0 = STRING: "rhel75-cluster"
PACEMAKER-PCS-V1-MIB::pcmkPcsV1ClusterQuorate.0 = INTEGER: 1
PACEMAKER-PCS-V1-MIB::pcmkPcsV1ClusterNodesNum.0 = INTEGER: 2
PACEMAKER-PCS-V1-MIB::pcmkPcsV1ClusterNodesNames.0 = STRING: "rhel75-node1"
PACEMAKER-PCS-V1-MIB::pcmkPcsV1ClusterNodesNames.1 = STRING: "rhel75-node2"
PACEMAKER-PCS-V1-MIB::pcmkPcsV1ClusterCorosyncNodesOnlineNum.0 = INTEGER: 2
PACEMAKER-PCS-V1-MIB::pcmkPcsV1ClusterCorosyncNodesOnlineNames.0 = STRING:
"rhel75-node1"
PACEMAKER-PCS-V1-MIB::pcmkPcsV1ClusterCorosyncNodesOnlineNames.1 = STRING:
"rhel75-node2"
PACEMAKER-PCS-V1-MIB::pcmkPcsV1ClusterCorosyncNodesOfflineNum.0 = INTEGER: 0
PACEMAKER-PCS-V1-MIB::pcmkPcsV1ClusterPcmkNodesOnlineNum.0 = INTEGER: 2
PACEMAKER-PCS-V1-MIB::pcmkPcsV1ClusterPcmkNodesOnlineNames.0 = STRING:
"rhel75-node1"
PACEMAKER-PCS-V1-MIB::pcmkPcsV1ClusterPcmkNodesOnlineNames.1 = STRING:
"rhel75-node2"
PACEMAKER-PCS-V1-MIB::pcmkPcsV1ClusterPcmkNodesStandbyNum.0 = INTEGER: 0
PACEMAKER-PCS-V1-MIB::pcmkPcsV1ClusterPcmkNodesOfflineNum.0 = INTEGER: 0
PACEMAKER-PCS-V1-MIB::pcmkPcsV1ClusterAllResourcesNum.0 = INTEGER: 11
PACEMAKER-PCS-V1-MIB::pcmkPcsV1ClusterAllResourcesIds.0 = STRING: "fencing"
PACEMAKER-PCS-V1-MIB::pcmkPcsV1ClusterAllResourcesIds.1 = STRING: "dummy5"
PACEMAKER-PCS-V1-MIB::pcmkPcsV1ClusterAllResourcesIds.2 = STRING: "dummy6"
PACEMAKER-PCS-V1-MIB::pcmkPcsV1ClusterAllResourcesIds.3 = STRING: "dummy7"
PACEMAKER-PCS-V1-MIB::pcmkPcsV1ClusterAllResourcesIds.4 = STRING: "dummy8"
PACEMAKER-PCS-V1-MIB::pcmkPcsV1ClusterAllResourcesIds.5 = STRING: "dummy9"
PACEMAKER-PCS-V1-MIB::pcmkPcsV1ClusterAllResourcesIds.6 = STRING: "dummy1"
PACEMAKER-PCS-V1-MIB::pcmkPcsV1ClusterAllResourcesIds.7 = STRING: "dummy10"
PACEMAKER-PCS-V1-MIB::pcmkPcsV1ClusterAllResourcesIds.8 = STRING: "dummy2"
PACEMAKER-PCS-V1-MIB::pcmkPcsV1ClusterAllResourcesIds.9 = STRING: "dummy3"
PACEMAKER-PCS-V1-MIB::pcmkPcsV1ClusterAllResourcesIds.10 = STRING: "dummy4"
PACEMAKER-PCS-V1-MIB::pcmkPcsV1ClusterRunningResourcesNum.0 = INTEGER: 9
PACEMAKER-PCS-V1-MIB::pcmkPcsV1ClusterRunningResourcesIds.0 = STRING:
"fencing"
PACEMAKER-PCS-V1-MIB::pcmkPcsV1ClusterRunningResourcesIds.1 = STRING:
"dummy7"
PACEMAKER-PCS-V1-MIB::pcmkPcsV1ClusterRunningResourcesIds.2 = STRING:
"dummy8"
PACEMAKER-PCS-V1-MIB::pcmkPcsV1ClusterRunningResourcesIds.3 = STRING:
"dummy9"
PACEMAKER-PCS-V1-MIB::pcmkPcsV1ClusterRunningResourcesIds.4 = STRING:
"dummy1"
PACEMAKER-PCS-V1-MIB::pcmkPcsV1ClusterRunningResourcesIds.5 = STRING:
"dummy10"

```

```
PACEMAKER-PCS-V1-MIB::pcmkPcsV1ClusterRunningResourcesIds.6 = STRING:  
"dummy2"  
PACEMAKER-PCS-V1-MIB::pcmkPcsV1ClusterRunningResourcesIds.7 = STRING:  
"dummy3"  
PACEMAKER-PCS-V1-MIB::pcmkPcsV1ClusterRunningResourcesIds.8 = STRING:  
"dummy4"  
PACEMAKER-PCS-V1-MIB::pcmkPcsV1ClusterStoppedResroucesNum.0 = INTEGER: 1  
PACEMAKER-PCS-V1-MIB::pcmkPcsV1ClusterStoppedResroucesIds.0 = STRING:  
"dummy5"  
PACEMAKER-PCS-V1-MIB::pcmkPcsV1ClusterFailedResourcesNum.0 = INTEGER: 1  
PACEMAKER-PCS-V1-MIB::pcmkPcsV1ClusterFailedResourcesIds.0 = STRING: "dummy6"  
PACEMAKER-PCS-V1-MIB::pcmkPcsV1ClusterFailedResourcesIds.0 = No more variables  
left in this MIB View (It is past the end of the MIB tree)
```

9.9. CONFIGURING RESOURCES TO REMAIN STOPPED ON CLEAN NODE SHUTDOWN (RED HAT ENTERPRISE LINUX 7.8 AND LATER)

When a cluster node shuts down, Pacemaker’s default response is to stop all resources running on that node and recover them elsewhere, even if the shutdown is a clean shutdown. As of Red Hat Enterprise Linux 7.8, you can configure Pacemaker so that when a node shuts down cleanly, the resources attached to the node will be locked to the node and unable to start elsewhere until they start again when the node that has shut down rejoins the cluster. This allows you to power down nodes during maintenance windows when service outages are acceptable without causing that node’s resources to fail over to other nodes in the cluster.

9.9.1. Cluster Properties to Configure Resources to Remain Stopped on Clean Node Shutdown

The ability to prevent resources from failing over on a clean node shutdown is implemented by means of the following cluster properties.

shutdown-lock

When this cluster property is set to the default value of **false**, the cluster will recover resources that are active on nodes being cleanly shut down. When this property is set to **true**, resources that are active on the nodes being cleanly shut down are unable to start elsewhere until they start on the node again after it rejoins the cluster.

The **shutdown-lock** property will work for either cluster nodes or remote nodes, but not guest nodes.

If **shutdown-lock** is set to **true**, you can remove the lock on the cluster resources when a node is down so that the resource can start elsewhere by performing a manual refresh on the node with the following command.

```
pcs resource refresh resource --node node
```

Note that once the resources are unlocked, the cluster is free to move the resources elsewhere. You can control the likelihood of this occurring by using stickiness values or location preferences for the resource.

**NOTE**


A manual refresh will work with remote nodes only if you first run the following commands:

1. Run the **systemctl stop pacemaker_remote** command on the remote node to stop the node.
2. Run the **pcs resource disable *remote-connection-resource*** command.

You can then perform a manual refresh on the remote node.

shutdown-lock-limit

When this cluster property is set to a time other than the default value of 0, resources will be available for recovery on other nodes if the node does not rejoin within the specified time since the shutdown was initiated. Note, however, that the time interval will not be checked any more often than the value of the **cluster-recheck-interval** cluster property.

**NOTE**

The **shutdown-lock-limit** property will work with remote nodes only if you first run the following commands:

1. Run the **systemctl stop pacemaker_remote** command on the remote node to stop the node.
2. Run the **pcs resource disable *remote-connection-resource*** command.

After you run these commands, the resources that had been running on the remote node will be available for recovery on other nodes when the amount of time specified as the **shutdown-lock-limit** has passed.

9.9.2. Setting the shutdown-lock Cluster Property

The following example sets the **shutdown-lock** cluster property to **true** in an example cluster and shows the effect this has when the node is shut down and started again. This example cluster consists of three nodes: **z1.example.com**, **z2.example.com**, and **z3.example.com**.

1. Set the **shutdown-lock** property to **true** and verify its value. In this example the **shutdown-lock-limit** property maintains its default value of 0.

```
[root@z3.example.com ~]# pcs property set shutdown-lock=true
[root@z3.example.com ~]# pcs property list --all | grep shutdown-lock
shutdown-lock: true
shutdown-lock-limit: 0
```

2. Check the status of the cluster. In this example, resources **third** and **fifth** are running on **z1.example.com**.

```
[root@z3.example.com ~]# pcs status
...
Full List of Resources:
```

```
...
* first (ocf::pacemaker:Dummy): Started z3.example.com
* second (ocf::pacemaker:Dummy): Started z2.example.com
* third (ocf::pacemaker:Dummy): Started z1.example.com
* fourth (ocf::pacemaker:Dummy): Started z2.example.com
* fifth (ocf::pacemaker:Dummy): Started z1.example.com
...
```

- Shut down **z1.example.com**, which will stop the resources that are running on that node.

```
[root@z3.example.com ~]# pcs cluster stop z1.example.com
Stopping Cluster (pacemaker)...
Stopping Cluster (corosync)...
```

Running the **pcs status** command shows that node **z1.example.com** is offline and that the resources that had been running on **z1.example.com** are **LOCKED** while the node is down.

```
[root@z3.example.com ~]# pcs status
...
Node List:
* Online: [ z2.example.com z3.example.com ]
* OFFLINE: [ z1.example.com ]

Full List of Resources:
...
* first (ocf::pacemaker:Dummy): Started z3.example.com
* second (ocf::pacemaker:Dummy): Started z2.example.com
* third (ocf::pacemaker:Dummy): Stopped z1.example.com (LOCKED)
* fourth (ocf::pacemaker:Dummy): Started z3.example.com
* fifth (ocf::pacemaker:Dummy): Stopped z1.example.com (LOCKED)
...
```

- Start cluster services again on **z1.example.com** so that it rejoins the cluster. Locked resources should get started on that node, although once they start they will not necessarily remain on the same node.

```
[root@z3.example.com ~]# pcs cluster start z1.example.com
Starting Cluster...
```

In this example, resources third and fifth are recovered on node z1.example.com.

```
[root@z3.example.com ~]# pcs status
...
Node List:
* Online: [ z1.example.com z2.example.com z3.example.com ]

Full List of Resources:
..
* first (ocf::pacemaker:Dummy): Started z3.example.com
* second (ocf::pacemaker:Dummy): Started z2.example.com
* third (ocf::pacemaker:Dummy): Started z1.example.com
* fourth (ocf::pacemaker:Dummy): Started z3.example.com
```

* fifth (ocf::pacemaker:Dummy): Started z1.example.com

...

CHAPTER 10. CLUSTER QUORUM

A Red Hat Enterprise Linux High Availability Add-On cluster uses the **votequorum** service, in conjunction with fencing, to avoid split brain situations. A number of votes is assigned to each system in the cluster, and cluster operations are allowed to proceed only when a majority of votes is present. The service must be loaded into all nodes or none; if it is loaded into a subset of cluster nodes, the results will be unpredictable. For information on the configuration and operation of the **votequorum** service, see the **votequorum(5)** man page.

10.1. CONFIGURING QUORUM OPTIONS

There are some special features of quorum configuration that you can set when you create a cluster with the **pcs cluster setup** command. [Table 10.1, “Quorum Options”](#) summarizes these options.

Table 10.1. Quorum Options

Option	Description
--auto_tie_breaker	<p>When enabled, the cluster can suffer up to 50% of the nodes failing at the same time, in a deterministic fashion. The cluster partition, or the set of nodes that are still in contact with the nodeid configured in auto_tie_breaker_node (or lowest nodeid if not set), will remain quorate. The other nodes will be inquorate.</p> <p>The auto_tie_breaker option is principally used for clusters with an even number of nodes, as it allows the cluster to continue operation with an even split. For more complex failures, such as multiple, uneven splits, it is recommended that you use a quorum device, as described in Section 10.5, “Quorum Devices”. The auto_tie_breaker option is incompatible with quorum devices.</p>
--wait_for_all	<p>When enabled, the cluster will be quorate for the first time only after all nodes have been visible at least once at the same time.</p> <p>The wait_for_all option is primarily used for two-node clusters and for even-node clusters using the quorum device lms (last man standing) algorithm.</p> <p>The wait_for_all option is automatically enabled when a cluster has two nodes, does not use a quorum device, and auto_tie_breaker is disabled. You can override this by explicitly setting wait_for_all to 0.</p>
--last_man_standing	<p>When enabled, the cluster can dynamically recalculate expected_votes and quorum under specific circumstances. You must enable wait_for_all when you enable this option. The last_man_standing option is incompatible with quorum devices.</p>
--last_man_standing_window	<p>The time, in milliseconds, to wait before recalculating expected_votes and quorum after a cluster loses nodes.</p>

For further information about configuring and using these options, see the **votequorum(5)** man page.

10.2. QUORUM ADMINISTRATION COMMANDS (RED HAT ENTERPRISE LINUX 7.3 AND LATER)

Once a cluster is running, you can enter the following cluster quorum commands.

The following command shows the quorum configuration.

```
pcs quorum [config]
```

The following command shows the quorum runtime status.

```
pcs quorum status
```

If you take nodes out of a cluster for a long period of time and the loss of those nodes would cause quorum loss, you can change the value of the **expected_votes** parameter for the live cluster with the **pcs quorum expected-votes** command. This allows the cluster to continue operation when it does not have quorum.



WARNING

Changing the expected votes in a live cluster should be done with extreme caution. If less than 50% of the cluster is running because you have manually changed the expected votes, then the other nodes in the cluster could be started separately and run cluster services, causing data corruption and other unexpected results. If you change this value, you should ensure that the **wait_for_all** parameter is enabled.

The following command sets the expected votes in the live cluster to the specified value. This affects the live cluster only and does not change the configuration file; the value of **expected_votes** is reset to the value in the configuration file in the event of a reload.

```
pcs quorum expected-votes votes
```

10.3. MODIFYING QUORUM OPTIONS (RED HAT ENTERPRISE LINUX 7.3 AND LATER)

As of Red Hat Enterprise Linux 7.3, you can modify general quorum options for your cluster with the **pcs quorum update** command. Executing this command requires that the cluster be stopped. For information on the quorum options, see the **votequorum(5)** man page.

The format of the **pcs quorum update** command is as follows.

```
pcs quorum update [auto_tie_breaker=[0|1]] [last_man_standing=[0|1]] [last_man_standing_window=[time-in-ms]] [wait_for_all=[0|1]]
```

The following series of commands modifies the **wait_for_all** quorum option and displays the updated status of the option. Note that the system does not allow you to execute this command while the cluster is running.

```
[root@node1:~]# pcs quorum update wait_for_all=1
Checking corosync is not running on nodes...
Error: node1: corosync is running
```

```
Error: node2: corosync is running
```

```
[root@node1:~]# pcs cluster stop --all
node2: Stopping Cluster (pacemaker)...
node1: Stopping Cluster (pacemaker)...
node1: Stopping Cluster (corosync)...
node2: Stopping Cluster (corosync)...
```


```
[root@node1:~]# pcs quorum update wait_for_all=1
Checking corosync is not running on nodes...
node2: corosync is not running
node1: corosync is not running
Sending updated corosync.conf to nodes...
node1: Succeeded
node2: Succeeded
```

```
[root@node1:~]# pcs quorum config
```

Options:
wait_for_all: 1

10.4. THE QUORUM UNBLOCK COMMAND

In a situation in which you know that the cluster is inquorate but you want the cluster to proceed with resource management, you can use the following command to prevent the cluster from waiting for all nodes when establishing quorum.



NOTE

This command should be used with extreme caution. Before issuing this command, it is imperative that you ensure that nodes that are not currently in the cluster are switched off and have no access to shared resources.

```
# pcs cluster quorum unblock
```

10.5. QUORUM DEVICES

Red Hat Enterprise Linux 7.4 provides full support for the ability to configure a separate quorum device which acts as a third-party arbitration device for the cluster. Its primary use is to allow a cluster to sustain more node failures than standard quorum rules allow. A quorum device is recommended for clusters with an even number of nodes and highly recommended for two-node clusters.

You must take the following into account when configuring a quorum device.

- It is recommended that a quorum device be run on a different physical network at the same site as the cluster that uses the quorum device. Ideally, the quorum device host should be in a separate rack than the main cluster, or at least on a separate PSU and not on the same network segment as the corosync ring or rings.
- You cannot use more than one quorum device in a cluster at the same time.
- Although you cannot use more than one quorum device in a cluster at the same time, a single quorum device may be used by several clusters at the same time. Each cluster using that quorum device can use different algorithms and quorum options, as those are stored on the

cluster nodes themselves. For example, a single quorum device can be used by one cluster with an **ffsplit** (fifty/fifty split) algorithm and by a second cluster with an **lms** (last man standing) algorithm.

- A quorum device should not be run on an existing cluster node.

10.5.1. Installing Quorum Device Packages

Configuring a quorum device for a cluster requires that you install the following packages:

- Install **corosync-qdevice** on the nodes of an existing cluster.

```
[root@node1:~]# yum install corosync-qdevice
[root@node2:~]# yum install corosync-qdevice
```

- Install **pcs** and **corosync-qnetd** on the quorum device host.

```
[root@qdevice:~]# yum install pcs corosync-qnetd
```

- Start the **pcsd** service and enable **pcsd** at system start on the quorum device host.

```
[root@qdevice:~]# systemctl start pcasd.service
[root@qdevice:~]# systemctl enable pcasd.service
```

10.5.2. Configuring a Quorum Device

This section provides a sample procedure to configure a quorum device in a Red Hat high availability cluster. The following procedure configures a quorum device and adds it to the cluster. In this example:

- The node used for a quorum device is **qdevice**.
- The quorum device model is **net**, which is currently the only supported model. The **net** model supports the following algorithms:
 - **ffsplit**: fifty-fifty split. This provides exactly one vote to the partition with the highest number of active nodes.
 - **lms**: last-man-standing. If the node is the only one left in the cluster that can see the **qnetd** server, then it returns a vote.



WARNING

The LMS algorithm allows the cluster to remain quorate even with only one remaining node, but it also means that the voting power of the quorum device is great since it is the same as `number_of_nodes - 1`. Losing connection with the quorum device means losing `number_of_nodes - 1` votes, which means that only a cluster with all nodes active can remain quorate (by overvoting the quorum device); any other cluster becomes inquorate.

For more detailed information on the implementation of these algorithms, see the **corosync-qdevice(8)** man page.

- The cluster nodes are **node1** and **node2**.

The following procedure configures a quorum device and adds that quorum device to a cluster.

1. On the node that you will use to host your quorum device, configure the quorum device with the following command. This command configures and starts the quorum device model **net** and configures the device to start on boot.

```
[root@qdevice:~]# pcs qdevice setup model net --enable --start
Quorum device 'net' initialized
quorum device enabled
Starting quorum device...
quorum device started
```

After configuring the quorum device, you can check its status. This should show that the **corosync-qnetd** daemon is running and, at this point, there are no clients connected to it. The **--full** command option provides detailed output.

```
[root@qdevice:~]# pcs qdevice status net --full
QNetd address:          *:5403
TLS:                  Supported (client certificate required)
Connected clients:      0
Connected clusters:     0
Maximum send/receive size: 32768/32768 bytes
```

2. Enable the ports on the firewall needed by the **pcsd** daemon and the **net** quorum device by enabling the **high-availability** service on **firewalld** with following commands.

```
[root@qdevice:~]# firewall-cmd --permanent --add-service=high-availability
[root@qdevice:~]# firewall-cmd --add-service=high-availability
```

3. From one of the nodes in the existing cluster, authenticate user **hacluster** on the node that is hosting the quorum device.

```
[root@node1:~] # pcs cluster auth qdevice
Username: hacluster
Password:
qdevice: Authorized
```

4. Add the quorum device to the cluster.

Before adding the quorum device, you can check the current configuration and status for the quorum device for later comparison. The output for these commands indicates that the cluster is not yet using a quorum device.

```
[root@node1:~]# pcs quorum config
Options:
```

```
[root@node1:~]# pcs quorum status
Quorum information
```

```
-----
Date:      Wed Jun 29 13:15:36 2016
Quorum provider: corosync_votequorum
Nodes:      2
Node ID:    1
Ring ID:    1/8272
Quorate:   Yes
```

Votequorum information

```
-----
Expected votes: 2
Highest expected: 2
Total votes: 2
Quorum: 1
Flags: 2Node Quorate
```

Membership information

Nodeid	Votes	Qdevice Name
1	1	NR node1 (local)
2	1	NR node2

The following command adds the quorum device that you have previously created to the cluster. You cannot use more than one quorum device in a cluster at the same time. However, one quorum device can be used by several clusters at the same time. This example command configures the quorum device to use the **ffsplit** algorithm. For information on the configuration options for the quorum device, see the **corosync-qdevice(8)** man page.

```
[root@node1:~]# pcs quorum device add model net host=qdevice algorithm=ffsplit
Setting up qdevice certificates on nodes...
node2: Succeeded
node1: Succeeded
Enabling corosync-qdevice...
node1: corosync-qdevice enabled
node2: corosync-qdevice enabled
Sending updated corosync.conf to nodes...
node1: Succeeded
node2: Succeeded
Corosync configuration reloaded
Starting corosync-qdevice...
node1: corosync-qdevice started
node2: corosync-qdevice started
```

5. Check the configuration status of the quorum device.

From the cluster side, you can execute the following commands to see how the configuration has changed.

The **pcs quorum config** shows the quorum device that has been configured.

```
[root@node1:~]# pcs quorum config
Options:
Device:
```

```
Model: net
algorithm: ffsplit
host: qdevice
```

The **pcs quorum status** command shows the quorum runtime status, indicating that the quorum device is in use.

```
[root@node1:~]# pcs quorum status
Quorum information
-----
Date:      Wed Jun 29 13:17:02 2016
Quorum provider: corosync_votequorum
Nodes:      2
Node ID:    1
Ring ID:   1/8272
Quorate:   Yes

Votequorum information
-----
Expected votes: 3
Highest expected: 3
Total votes: 3
Quorum: 2
Flags: Quorate Qdevice

Membership information
-----
  Nodeid  Votes  Qdevice Name
    1      1    A,V,NMW node1 (local)
    2      1    A,V,NMW node2
    0      1        Qdevice
```

The **pcs quorum device status** shows the quorum device runtime status.

```
[root@node1:~]# pcs quorum device status
Qdevice information
-----
Model:      Net
Node ID:    1
Configured node list:
  0  Node ID = 1
  1  Node ID = 2
Membership node list: 1, 2

Qdevice-net information
-----
Cluster name:      mycluster
QNetd host:       qdevice:5403
Algorithm:        ffsplit
Tie-breaker:      Node with lowest node ID
State:           Connected
```

From the quorum device side, you can execute the following status command, which shows the status of the **corosync-qnetd** daemon.

```
[root@qdevice:~]# pcs qdevice status net --full
QNetd address:          *:5403
TLS:                  Supported (client certificate required)
Connected clients:      2
Connected clusters:     1
Maximum send/receive size: 32768/32768 bytes
Cluster "mycluster":
Algorithm:        ffsplit
Tie-breaker:      Node with lowest node ID
Node ID 2:
Client address:    ::ffff:192.168.122.122:50028
HB interval:       8000ms
Configured node list: 1, 2
Ring ID:           1.2050
Membership node list: 1, 2
TLS active:        Yes (client certificate verified)
Vote:              ACK (ACK)
Node ID 1:
Client address:    ::ffff:192.168.122.121:48786
HB interval:       8000ms
Configured node list: 1, 2
Ring ID:           1.2050
Membership node list: 1, 2
TLS active:        Yes (client certificate verified)
Vote:              ACK (ACK)
```

10.5.3. Managing the Quorum Device Service

PCS provides the ability to manage the quorum device service on the local host (**corosync-qnetd**), as shown in the following example commands. Note that these commands affect only the **corosync-qnetd** service.

```
[root@qdevice:~]# pcs qdevice start net
[root@qdevice:~]# pcs qdevice stop net
[root@qdevice:~]# pcs qdevice enable net
[root@qdevice:~]# pcs qdevice disable net
[root@qdevice:~]# pcs qdevice kill net
```

10.5.4. Managing the Quorum Device Settings in a Cluster

The following sections describe the PCS commands that you can use to manage the quorum device settings in a cluster, showing examples that are based on the quorum device configuration in [Section 10.5.2, “Configuring a Quorum Device”](#).

10.5.4.1. Changing Quorum Device Settings

You can change the setting of a quorum device with the **pcs quorum device update** command.

**WARNING**

To change the **host** option of quorum device model **net**, use the **pcs quorum device remove** and the **pcs quorum device add** commands to set up the configuration properly, unless the old and the new host are the same machine.

The following command changes the quorum device algorithm to **lms**.

```
[root@node1:~]# pcs quorum device update model algorithm=lms
Sending updated corosync.conf to nodes...
node1: Succeeded
node2: Succeeded
Corosync configuration reloaded
Reloading qdevice configuration on nodes...
node1: corosync-qdevice stopped
node2: corosync-qdevice stopped
node1: corosync-qdevice started
node2: corosync-qdevice started
```

10.5.4.2. Removing a Quorum Device

Use the following command to remove a quorum device configured on a cluster node.

```
[root@node1:~]# pcs quorum device remove
Sending updated corosync.conf to nodes...
node1: Succeeded
node2: Succeeded
Corosync configuration reloaded
Disabling corosync-qdevice...
node1: corosync-qdevice disabled
node2: corosync-qdevice disabled
Stopping corosync-qdevice...
node1: corosync-qdevice stopped
node2: corosync-qdevice stopped
Removing qdevice certificates from nodes...
node1: Succeeded
node2: Succeeded
```

After you have removed a quorum device, you should see the following error message when displaying the quorum device status.

```
[root@node1:~]# pcs quorum device status
Error: Unable to get quorum status: corosync-qdevice-tool: Can't connect to QDevice socket (is
QDevice running?): No such file or directory
```

10.5.4.3. Destroying a Quorum Device

To disable and stop a quorum device on the quorum device host and delete all of its configuration files, use the following command.

```
[root@qdevice:~]# pcs qdevice destroy net
Stopping quorum device...
quorum device stopped
quorum device disabled
Quorum device 'net' configuration files removed
```

CHAPTER 11. PACEMAKER RULES

Rules can be used to make your configuration more dynamic. One use of rules might be to assign machines to different processing groups (using a node attribute) based on time and to then use that attribute when creating location constraints.

Each rule can contain a number of expressions, date-expressions and even other rules. The results of the expressions are combined based on the rule's **boolean-op** field to determine if the rule ultimately evaluates to **true** or **false**. What happens next depends on the context in which the rule is being used.

Table 11.1. Properties of a Rule

Field	Description
role	Limits the rule to apply only when the resource is in that role. Allowed values: Started , Slave , and Master . NOTE: A rule with role="Master" cannot determine the initial location of a clone instance. It will only affect which of the active instances will be promoted.
score	The score to apply if the rule evaluates to true . Limited to use in rules that are part of location constraints.
score-attribute	The node attribute to look up and use as a score if the rule evaluates to true . Limited to use in rules that are part of location constraints.
boolean-op	How to combine the result of multiple expression objects. Allowed values: and and or . The default value is and .

11.1. NODE ATTRIBUTE EXPRESSIONS

Node attribute expressions are used to control a resource based on the attributes defined by a node or nodes.

Table 11.2. Properties of an Expression

Field	Description
attribute	The node attribute to test
type	Determines how the value(s) should be tested. Allowed values: string , integer , version . The default value is string

Field	Description
operation	The comparison to perform. Allowed values: <ul style="list-style-type: none">* lt - True if the node attribute's value is less than value* gt - True if the node attribute's value is greater than value* lte - True if the node attribute's value is less than or equal to value* gte - True if the node attribute's value is greater than or equal to value* eq - True if the node attribute's value is equal to value* ne - True if the node attribute's value is not equal to value* defined - True if the node has the named attribute* not_defined - True if the node does not have the named attribute
value	User supplied value for comparison (required)

In addition to any attributes added by the administrator, the cluster defines special, built-in node attributes for each node that can also be used, as described in [Table 11.3, “Built-in Node Attributes”](#).

Table 11.3. Built-in Node Attributes

Name	Description
#uname	Node name
#id	Node ID
#kind	Node type. Possible values are cluster , remote , and container . The value of kind is remote for Pacemaker Remote nodes created with the ocf:pacemaker:remote resource, and container for Pacemaker Remote guest nodes and bundle nodes.
#is_dc	true if this node is a Designated Controller (DC), false otherwise
#cluster_name	The value of the cluster-name cluster property, if set
#site_name	The value of the site-name node attribute, if set, otherwise identical to #cluster-name
#role	The role the relevant multistate resource has on this node. Valid only within a rule for a location constraint for a multistate resource.

11.2. TIME/DATE BASED EXPRESSIONS

Date expressions are used to control a resource or cluster option based on the current date/time. They can contain an optional date specification.

Table 11.4. Properties of a Date Expression

Field	Description
start	A date/time conforming to the ISO8601 specification.
end	A date/time conforming to the ISO8601 specification.
operation	<p>Compares the current date/time with the start or the end date or both the start and end date, depending on the context. Allowed values:</p> <ul style="list-style-type: none"> * gt - True if the current date/time is after start * lt - True if the current date/time is before end * in-range - True if the current date/time is after start and before end * date-spec - performs a cron-like comparison to the current date/time

11.3. DATE SPECIFICATIONS

Date specifications are used to create cron-like expressions relating to time. Each field can contain a single number or a single range. Instead of defaulting to zero, any field not supplied is ignored.

For example, **monthdays="1"** matches the first day of every month and **hours="09-17"** matches the hours between 9 am and 5 pm (inclusive). However, you cannot specify **weekdays="1,2"** or **weekdays="1-2,5-6"** since they contain multiple ranges.

Table 11.5. Properties of a Date Specification

Field	Description
id	A unique name for the date
hours	Allowed values: 0-23
monthdays	Allowed values: 0-31 (depending on month and year)
weekdays	Allowed values: 1-7 (1=Monday, 7=Sunday)
yeardays	Allowed values: 1-366 (depending on the year)
months	Allowed values: 1-12
weeks	Allowed values: 1-53 (depending on weekyear)
years	Year according the Gregorian calendar
weekyears	May differ from Gregorian years; for example, 2005-001 Ordinal is also 2005-01-01 Gregorian is also 2004-W53-6 Weekly
moon	Allowed values: 0-7 (0 is new, 4 is full moon).

11.4. DURATIONS

Durations are used to calculate a value for **end** when one is not supplied to in_range operations. They contain the same fields as **date_spec** objects but without the limitations (ie. you can have a duration of 19 months). Like **date_specs**, any field not supplied is ignored.

11.5. CONFIGURING RULES WITH PCS

To configure a rule using **pcs**, you can configure a location constraint that uses rules, as described in [Section 7.1.3, “Using Rules to Determine Resource Location”](#).

To remove a rule, use the following. If the rule that you are removing is the last rule in its constraint, the constraint will be removed.

```
pcs constraint rule remove rule_id
```

CHAPTER 12. PACEMAKER CLUSTER PROPERTIES

Cluster properties control how the cluster behaves when confronted with situations that may occur during cluster operation.

- [Table 12.1, “Cluster Properties”](#) describes the cluster properties options.
- [Section 12.2, “Setting and Removing Cluster Properties”](#) describes how to set cluster properties.
- [Section 12.3, “Querying Cluster Property Settings”](#) describes how to list the currently set cluster properties.

12.1. SUMMARY OF CLUSTER PROPERTIES AND OPTIONS

[Table 12.1, “Cluster Properties”](#) summaries the Pacemaker cluster properties, showing the default values of the properties and the possible values you can set for those properties.



NOTE

In addition to the properties described in this table, there are additional cluster properties that are exposed by the cluster software. For these properties, it is recommended that you not change their values from their defaults.

Table 12.1. Cluster Properties

Option	Default	Description
batch-limit	0	The number of resource actions that the cluster is allowed to execute in parallel. The "correct" value will depend on the speed and load of your network and cluster nodes.
migration-limit	-1 (unlimited)	The number of migration jobs that the cluster is allowed to execute in parallel on a node.
no-quorum-policy	stop	What to do when the cluster does not have quorum. Allowed values: * ignore - continue all resource management * freeze - continue resource management, but do not recover resources from nodes not in the affected partition * stop - stop all resources in the affected cluster partition * suicide - fence all nodes in the affected cluster partition
symmetric-cluster	true	Indicates whether resources can run on any node by default.

Option	Default	Description
stonith-enabled	true	<p>Indicates that failed nodes and nodes with resources that cannot be stopped should be fenced. Protecting your data requires that you set this true.</p> <p>If true, or unset, the cluster will refuse to start resources unless one or more STONITH resources have been configured also.</p>
stonith-action	reboot	Action to send to STONITH device. Allowed values: reboot , off . The value poweroff is also allowed, but is only used for legacy devices.
cluster-delay	60s	Round trip delay over the network (excluding action execution). The "correct" value will depend on the speed and load of your network and cluster nodes.
stop-orphan-resources	true	Indicates whether deleted resources should be stopped.
stop-orphan-actions	true	Indicates whether deleted actions should be canceled.
start-failure-is-fatal	true	<p>Indicates whether a failure to start a resource on a particular node prevents further start attempts on that node. When set to false, the cluster will decide whether to try starting on the same node again based on the resource's current failure count and migration threshold. For information on setting the migration-threshold option for a resource, see Section 8.2, “Moving Resources Due to Failure”.</p> <p>Setting start-failure-is-fatal to false incurs the risk that this will allow one faulty node that is unable to start a resource to hold up all dependent actions. This is why start-failure-is-fatal defaults to true. The risk of setting start-failure-is-fatal=false can be mitigated by setting a low migration threshold so that other actions can proceed after that many failures.</p>
pe-error-series-max	-1 (all)	The number of PE inputs resulting in ERRORS to save. Used when reporting problems.
pe-warn-series-max	-1 (all)	The number of PE inputs resulting in WARNINGS to save. Used when reporting problems.
pe-input-series-max	-1 (all)	The number of "normal" PE inputs to save. Used when reporting problems.
cluster-infrastructure		The messaging stack on which Pacemaker is currently running. Used for informational and diagnostic purposes; not user-configurable.
dc-version		Version of Pacemaker on the cluster's Designated Controller (DC). Used for diagnostic purposes; not user-configurable.

Option	Default	Description
last-lrm-refresh		Last refresh of the Local Resource Manager, given in units of seconds since epoch. Used for diagnostic purposes; not user-configurable.
cluster-recheck-interval	15 minutes	Polling interval for time-based changes to options, resource parameters and constraints. Allowed values: Zero disables polling, positive values are an interval in seconds (unless other SI units are specified, such as 5min). Note that this value is the maximum time between checks; if a cluster event occurs sooner than the time specified by this value, the check will be done sooner.
maintenance-mode	false	Maintenance Mode tells the cluster to go to a "hands off" mode, and not start or stop any services until told otherwise. When maintenance mode is completed, the cluster does a sanity check of the current state of any services, and then stops or starts any that need it.
shutdown-escalation	20min	The time after which to give up trying to shut down gracefully and just exit. Advanced use only.
stonith-timeout	60s	How long to wait for a STONITH action to complete.
stop-all-resources	false	Should the cluster stop all resources.
enable-acl	false	(Red Hat Enterprise Linux 7.1 and later) Indicates whether the cluster can use access control lists, as set with the pcs acl command.
placement-strategy	default	Indicates whether and how the cluster will take utilization attributes into account when determining resource placement on cluster nodes. For information on utilization attributes and placement strategies, see Section 9.6, “Utilization and Placement Strategy” .
fence-reaction	stop	(Red Hat Enterprise Linux 7.8 and later) Determines how a cluster node should react if notified of its own fencing. A cluster node may receive notification of its own fencing if fencing is misconfigured, or if fabric fencing is in use that does not cut cluster communication. Allowed values are stop to attempt to immediately stop Pacemaker and stay stopped, or panic to attempt to immediately reboot the local node, falling back to stop on failure.

12.2. SETTING AND REMOVING CLUSTER PROPERTIES

To set the value of a cluster property, use the following **pcs** command.

```
pcs property set property=value
```

For example, to set the value of **symmetric-cluster** to **false**, use the following command.

```
# pcs property set symmetric-cluster=false
```

You can remove a cluster property from the configuration with the following command.

```
pcs property unset property
```

Alternately, you can remove a cluster property from a configuration by leaving the value field of the **pcs property set** command blank. This restores that property to its default value. For example, if you have previously set the **symmetric-cluster** property to **false**, the following command removes the value you have set from the configuration and restores the value of **symmetric-cluster** to **true**, which is its default value.

```
# pcs property set symmetric-cluster=
```

12.3. QUERYING CLUSTER PROPERTY SETTINGS

In most cases, when you use the **pcs** command to display values of the various cluster components, you can use **pcs list** or **pcs show** interchangeably. In the following examples, **pcs list** is the format used to display an entire list of all settings for more than one property, while **pcs show** is the format used to display the values of a specific property.

To display the values of the property settings that have been set for the cluster, use the following **pcs** command.

```
pcs property list
```

To display all of the values of the property settings for the cluster, including the default values of the property settings that have not been explicitly set, use the following command.

```
pcs property list --all
```

To display the current value of a specific cluster property, use the following command.

```
pcs property show property
```

For example, to display the current value of the **cluster-infrastructure** property, execute the following command:

```
# pcs property show cluster-infrastructure
Cluster Properties:
  cluster-infrastructure: cman
```

For informational purposes, you can display a list of all of the default values for the properties, whether they have been set to a value other than the default or not, by using the following command.

```
pcs property [list|show] --defaults
```

CHAPTER 13. TRIGGERING SCRIPTS FOR CLUSTER EVENTS

A Pacemaker cluster is an event-driven system, where an event might be a resource or node failure, a configuration change, or a resource starting or stopping. You can configure Pacemaker cluster alerts to take some external action when a cluster event occurs. You can configure cluster alerts in one of two ways:

- As of Red Hat Enterprise Linux 7.3, you can configure Pacemaker alerts by means of alert agents, which are external programs that the cluster calls in the same manner as the cluster calls resource agents to handle resource configuration and operation. This is the preferred, simpler method of configuring cluster alerts. Pacemaker alert agents are described in [Section 13.1, “Pacemaker Alert Agents \(Red Hat Enterprise Linux 7.3 and later\)”](#).
- The **ocf:pacemaker:ClusterMon** resource can monitor the cluster status and trigger alerts on each cluster event. This resource runs the **crm_mon** command in the background at regular intervals. For information on the **ClusterMon** resource see [Section 13.2, “Event Notification with Monitoring Resources”](#).

13.1. PACEMAKER ALERT AGENTS (RED HAT ENTERPRISE LINUX 7.3 AND LATER)

You can create Pacemaker alert agents to take some external action when a cluster event occurs. The cluster passes information about the event to the agent by means of environment variables. Agents can do anything with this information, such as send an email message or log to a file or update a monitoring system.

- Pacemaker provides several sample alert agents, which are installed in **/usr/share/pacemaker/alerts** by default. These sample scripts may be copied and used as is, or they may be used as templates to be edited to suit your purposes. Refer to the source code of the sample agents for the full set of attributes they support. See [Section 13.1.1, “Using the Sample Alert Agents”](#) for an example of a basic procedure for configuring an alert that uses a sample alert agent.
- General information on configuring and administering alert agents is provided in [Section 13.1.2, “Alert Creation”](#), [Section 13.1.3, “Displaying, Modifying, and Removing Alerts”](#), [Section 13.1.4, “Alert Recipients”](#), [Section 13.1.5, “Alert Meta Options”](#), and [Section 13.1.6, “Alert Configuration Command Examples”](#).
- You can write your own alert agents for a Pacemaker alert to call. For information on writing alert agents, see [Section 13.1.7, “Writing an Alert Agent”](#).

13.1.1. Using the Sample Alert Agents

When you use one of the sample alert agents, you should review the script to ensure that it suits your needs. These sample agents are provided as a starting point for custom scripts for specific cluster environments. Note that while Red Hat supports the interfaces that the alert agents scripts use to communicate with Pacemaker, Red Hat does not provide support for the custom agents themselves.

To use one of the sample alert agents, you must install the agent on each node in the cluster. For example, the following command installs the **alert_file.sh.sample** script as **alert_file.sh**.

```
# install --mode=0755 /usr/share/pacemaker/alerts/alert_file.sh.sample
/var/lib/pacemaker/alert_file.sh
```

After you have installed the script, you can create an alert that uses the script.

The following example configures an alert that uses the installed **alert_file.sh** alert agent to log events to a file. Alert agents run as the user **hacluster**, which has a minimal set of permissions.

This example creates the log file **pcmk_alert_file.log** that will be used to record the events. It then creates the alert agent and adds the path to the log file as its recipient.

```
# touch /var/log/pcmk_alert_file.log
# chown hacluster:haclient /var/log/pcmk_alert_file.log
# chmod 600 /var/log/pcmk_alert_file.log
# pcs alert create id=alert_file description="Log events to a file." path=/var/lib/pacemaker/alert_file.sh
# pcs alert recipient add alert_file id=my-alert_logfile value=/var/log/pcmk_alert_file.log
```

The following example installs the **alert_snmp.sh.sample** script as **alert_snmp.sh** and configures an alert that uses the installed **alert_snmp.sh** alert agent to send cluster events as SNMP traps. By default, the script will send all events except successful monitor calls to the SNMP server. This example configures the timestamp format as a meta option. For information about meta options, see [Section 13.1.5, “Alert Meta Options”](#). After configuring the alert, this example configures a recipient for the alert and displays the alert configuration.

```
# install --mode=0755 /usr/share/pacemaker/alerts/alert_snmp.sh.sample
/var/lib/pacemaker/alert_snmp.sh
# pcs alert create id=snmp_alert path=/var/lib/pacemaker/alert_snmp.sh meta timestamp-
format="%Y-%m-%d,%H:%M:%S.%01N"
# pcs alert recipient add snmp_alert value=192.168.1.2
# pcs alert
Alerts:
Alert: snmp_alert (path=/var/lib/pacemaker/alert_snmp.sh)
Meta options: timestamp-format=%Y-%m-%d,%H:%M:%S.%01N.
Recipients:
Recipient: snmp_alert-recipient (value=192.168.1.2)
```

The following example installs the **alert_smtp.sh** agent and then configures an alert that uses the installed alert agent to send cluster events as email messages. After configuring the alert, this example configures a recipient and displays the alert configuration.

```
# install --mode=0755 /usr/share/pacemaker/alerts/alert_smtp.sh.sample
/var/lib/pacemaker/alert_smtp.sh
# pcs alert create id=smtp_alert path=/var/lib/pacemaker/alert_smtp.sh options
email_sender=donotreply@example.com
# pcs alert recipient add smtp_alert value=admin@example.com
# pcs alert
Alerts:
Alert: smtp_alert (path=/var/lib/pacemaker/alert_smtp.sh)
Options: email_sender=donotreply@example.com
Recipients:
Recipient: smtp_alert-recipient (value=admin@example.com)
```

For more information on the format of the **pcs alert create** and **pcs alert recipient add** commands, see [Section 13.1.2, “Alert Creation”](#) and [Section 13.1.4, “Alert Recipients”](#).

13.1.2. Alert Creation

The following command creates a cluster alert. The options that you configure are agent-specific

configuration values that are passed to the alert agent script at the path you specify as additional environment variables. If you do not specify a value for **id**, one will be generated. For information on alert meta options, [Section 13.1.5, “Alert Meta Options”](#).

```
pcs alert create path=path [id=alert-id] [description=description] [options [option=value]...] [meta [meta-option=value]...]
```

Multiple alert agents may be configured; the cluster will call all of them for each event. Alert agents will be called only on cluster nodes. They will be called for events involving Pacemaker Remote nodes, but they will never be called on those nodes.

The following example creates a simple alert that will call **myscript.sh** for each event.

```
# pcs alert create id=my_alert path=/path/to/myscript.sh
```

For an example that shows how to create a cluster alert that uses one of the sample alert agents, see [Section 13.1.1, “Using the Sample Alert Agents”](#).

13.1.3. Displaying, Modifying, and Removing Alerts

The following command shows all configured alerts along with the values of the configured options.

```
pcs alert [config|show]
```

The following command updates an existing alert with the specified *alert-id* value.

```
pcs alert update alert-id [path=path] [description=description] [options [option=value]...] [meta [meta-option=value]...]
```

The following command removes an alert with the specified *alert-id* value.

```
pcs alert remove alert-id
```

Alternately, you can run the **pcs alert delete** command, which is identical to the **pcs alert remove** command. Both the **pcs alert delete** and the **pcs alert remove** commands allow you to specify more than one alert to be deleted.

13.1.4. Alert Recipients

Usually alerts are directed towards a recipient. Thus each alert may be additionally configured with one or more recipients. The cluster will call the agent separately for each recipient.

The recipient may be anything the alert agent can recognize: an IP address, an email address, a file name, or whatever the particular agent supports.

The following command adds a new recipient to the specified alert.

```
pcs alert recipient add alert-id value=recipient-value [id=recipient-id] [description=description] [options [option=value]...] [meta [meta-option=value]...]
```

The following command updates an existing alert recipient.

```
pcs alert recipient update recipient-id [value=recipient-value] [description=description] [options  
[option=value]...] [meta [meta-option=value]...]
```

The following command removes the specified alert recipient.

```
pcs alert recipient remove recipient-id
```

Alternately, you can run the **pcs alert recipient delete** command, which is identical to the **pcs alert recipient remove** command. Both the **pcs alert recipient remove** and the **pcs alert recipient delete** commands allow you to remove more than one alert recipient.

The following example command adds the alert recipient **my-alert-recipient** with a recipient ID of **my-recipient-id** to the alert **my-alert**. This will configure the cluster to call the alert script that has been configured for **my-alert** for each event, passing the recipient **some-address** as an environment variable.

```
# pcs alert recipient add my-alert value=my-alert-recipient id=my-recipient-id options value=some-  
address
```

13.1.5. Alert Meta Options

As with resource agents, meta options can be configured for alert agents to affect how Pacemaker calls them. [Table 13.1, “Alert Meta Options”](#) describes the alert meta options. Meta options can be configured per alert agent as well as per recipient.

Table 13.1. Alert Meta Options

Meta-Attribute	Default	Description
timestamp-format	%H:%M:%S.%06N	Format the cluster will use when sending the event’s timestamp to the agent. This is a string as used with the date(1) command.
timeout	30s	If the alert agent does not complete within this amount of time, it will be terminated.

The following example configures an alert that calls the script **myscript.sh** and then adds two recipients to the alert. The first recipient has an ID of **my-alert-recipient1** and the second recipient has an ID of **my-alert-recipient2**. The script will get called twice for each event, with each call using a 15-second timeout. One call will be passed to the recipient **someuser@example.com** with a timestamp in the format %D %H:%M, while the other call will be passed to the recipient **otheruser@example.com** with a timestamp in the format %c.

```
# pcs alert create id=my-alert path=/path/to/myscript.sh meta timeout=15s  
# pcs alert recipient add my-alert value=someuser@example.com id=my-alert-recipient1 meta  
timestamp-format="%D %H:%M"  
# pcs alert recipient add my-alert value=otheruser@example.com id=my-alert-recipient2 meta  
timestamp-format=%c
```

13.1.6. Alert Configuration Command Examples

The following sequential examples show some basic alert configuration commands to show the format to use to create alerts, add recipients, and display the configured alerts. Note that while you must install the alert agents themselves on each node in a cluster, you need to run the `pcs` commands only once.

The following commands create a simple alert, add two recipients to the alert, and display the configured values.

- Since no alert ID value is specified, the system creates an alert ID value of **alert**.
- The first recipient creation command specifies a recipient of **rec_value**. Since this command does not specify a recipient ID, the value of **alert-recipient** is used as the recipient ID.
- The second recipient creation command specifies a recipient of **rec_value2**. This command specifies a recipient ID of **my-recipient** for the recipient.

```
# pcs alert create path=/my/path
# pcs alert recipient add alert value=rec_value
# pcs alert recipient add alert value=rec_value2 id=my-recipient
# pcs alert config
Alerts:
Alert: alert (path=/my/path)
Recipients:
Recipient: alert-recipient (value=rec_value)
Recipient: my-recipient (value=rec_value2)
```

This following commands add a second alert and a recipient for that alert. The alert ID for the second alert is **my-alert** and the recipient value is **my-other-recipient**. Since no recipient ID is specified, the system provides a recipient id of **my-alert-recipient**.

```
# pcs alert create id=my-alert path=/path/to/script description=alert_description options
option1=value1 opt=val meta timeout=50s timestamp-format="%H%B%S"
# pcs alert recipient add my-alert value=my-other-recipient
# pcs alert
Alerts:
Alert: alert (path=/my/path)
Recipients:
Recipient: alert-recipient (value=rec_value)
Recipient: my-recipient (value=rec_value2)
Alert: my-alert (path=/path/to/script)
Description: alert_description
Options: opt=val option1=value1
Meta options: timestamp-format=%H%B%S timeout=50s
Recipients:
Recipient: my-alert-recipient (value=my-other-recipient)
```

The following commands modify the alert values for the alert **my-alert** and for the recipient **my-alert-recipient**.

```
# pcs alert update my-alert options option1=newvalue1 meta timestamp-format="%H%M%S"
# pcs alert recipient update my-alert-recipient options option1=new meta timeout=60s
# pcs alert
Alerts:
Alert: alert (path=/my/path)
Recipients:
Recipient: alert-recipient (value=rec_value)
```

```

Recipient: my-recipient (value=rec_value2)
Alert: my-alert (path=/path/to/script)
Description: alert_description
Options: opt=val option1=newvalue1
Meta options: timestamp-format=%H%M%S timeout=50s
Recipients:
  Recipient: my-alert-recipient (value=my-other-recipient)
  Options: option1=new
  Meta options: timeout=60s

```

The following command removes the recipient **my-alert-recipient** from **alert**.

```

# pcs alert recipient remove my-recipient
# pcs alert
Alerts:
  Alert: alert (path=/my/path)
  Recipients:
    Recipient: alert-recipient (value=rec_value)
  Alert: my-alert (path=/path/to/script)
  Description: alert_description
  Meta options: timestamp-format="%M%B%S" timeout=50s
  Meta options: m=newval meta-option1=2
  Recipients:
    Recipient: my-alert-recipient (value=my-other-recipient)
    Options: option1=new
    Meta options: timeout=60s

```

The following command removes **myalert** from the configuration.

```

# pcs alert remove my-alert
# pcs alert
Alerts:
  Alert: alert (path=/my/path)
  Recipients:
    Recipient: alert-recipient (value=rec_value)

```

13.1.7. Writing an Alert Agent

There are three types of Pacemaker alerts: node alerts, fencing alerts, and resource alerts. The environment variables that are passed to the alert agents can differ, depending on the type of alert.

[Table 13.2, “Environment Variables Passed to Alert Agents”](#) describes the environment variables that are passed to alert agents and specifies when the environment variable is associated with a specific alert type.

Table 13.2. Environment Variables Passed to Alert Agents

Environment Variable	Description
CRM_alert_kind	The type of alert (node, fencing, or resource)
CRM_alert_version	The version of Pacemaker sending the alert
CRM_alert_recipient	The configured recipient

Environment Variable	Description
CRM_alert_node_sequence	A sequence number increased whenever an alert is being issued on the local node, which can be used to reference the order in which alerts have been issued by Pacemaker. An alert for an event that happened later in time reliably has a higher sequence number than alerts for earlier events. Be aware that this number has no cluster-wide meaning.
CRM_alert_timestamp	A timestamp created prior to executing the agent, in the format specified by the timestamp-format meta option. This allows the agent to have a reliable, high-precision time of when the event occurred, regardless of when the agent itself was invoked (which could potentially be delayed due to system load or other circumstances).
CRM_alert_node	Name of affected node
CRM_alert_desc	Detail about event. For node alerts, this is the node's current state (member or lost). For fencing alerts, this is a summary of the requested fencing operation, including origin, target, and fencing operation error code, if any. For resource alerts, this is a readable string equivalent of CRM_alert_status .
CRM_alert_nodeid	ID of node whose status changed (provided with node alerts only)
CRM_alert_task	The requested fencing or resource operation (provided with fencing and resource alerts only)
CRM_alert_rc	The numerical return code of the fencing or resource operation (provided with fencing and resource alerts only)
CRM_alert_rsc	The name of the affected resource (resource alerts only)
CRM_alert_interval	The interval of the resource operation (resource alerts only)
CRM_alert_target_rc	The expected numerical return code of the operation (resource alerts only)
CRM_alert_status	A numerical code used by Pacemaker to represent the operation result (resource alerts only)

When writing an alert agent, you must take the following concerns into account.

- Alert agents may be called with no recipient (if none is configured), so the agent must be able to handle this situation, even if it only exists in that case. Users may modify the configuration in stages, and add a recipient later.
- If more than one recipient is configured for an alert, the alert agent will be called once per recipient. If an agent is not able to run concurrently, it should be configured with only a single recipient. The agent is free, however, to interpret the recipient as a list.
- When a cluster event occurs, all alerts are fired off at the same time as separate processes. Depending on how many alerts and recipients are configured and on what is done within the alert agents, a significant load burst may occur. The agent could be written to take this into

consideration, for example by queueing resource-intensive actions into some other instance, instead of directly executing them.

- Alert agents are run as the **hacluster** user, which has a minimal set of permissions. If an agent requires additional privileges, it is recommended to configure **sudo** to allow the agent to run the necessary commands as another user with the appropriate privileges.
- Take care to validate and sanitize user-configured parameters, such as **CRM_alert_timestamp** (whose content is specified by the user-configured **timestamp-format**), **CRM_alert_recipient**, and all alert options. This is necessary to protect against configuration errors. In addition, if some user can modify the CIB without having **hacluster**-level access to the cluster nodes, this is a potential security concern as well, and you should avoid the possibility of code injection.
- If a cluster contains resources with operations for which the **on-fail** parameter is set to **fence**, there will be multiple fence notifications on failure, one for each resource for which this parameter is set plus one additional notification. Both the STONITH daemon and the **crmd** daemon will send notifications. Pacemaker performs only one actual fence operation in this case, however, no matter how many notifications are sent.



NOTE

The alerts interface is designed to be backward compatible with the external scripts interface used by the **ocf:pacemaker:ClusterMon** resource. To preserve this compatibility, the environment variables passed to alert agents are available prepended with **CRM_notify_** as well as **CRM_alert_**. One break in compatibility is that the **ClusterMon** resource ran external scripts as the root user, while alert agents are run as the **hacluster** user. For information on configuring scripts that are triggered by the **ClusterMon**, see [Section 13.2, “Event Notification with Monitoring Resources”](#).

13.2. EVENT NOTIFICATION WITH MONITORING RESOURCES

The **ocf:pacemaker:ClusterMon** resource can monitor the cluster status and trigger alerts on each cluster event. This resource runs the **crm_mon** command in the background at regular intervals.

By default, the **crm_mon** command listens for resource events only; to enable listing for fencing events you can provide the **--watch-fencing** option to the command when you configure the **ClusterMon** resource. The **crm_mon** command does not monitor for membership issues but will print a message when fencing is started and when monitoring is started for that node, which would imply that a member just joined the cluster.

The **ClusterMon** resource can execute an external program to determine what to do with cluster notifications by means of the **extra_options** parameter. [Table 13.3, “Environment Variables Passed to the External Monitor Program”](#) lists the environment variables that are passed to that program, which describe the type of cluster event that occurred.

Table 13.3. Environment Variables Passed to the External Monitor Program

Environment Variable	Description
CRM_notify_recipient	The static external-recipient from the resource definition
CRM_notify_node	The node on which the status change happened

Environment Variable	Description
CRM_notify_rsc	The name of the resource that changed the status
CRM_notify_task	The operation that caused the status change
CRM_notify_desc	The textual output relevant error code of the operation (if any) that caused the status change
CRM_notify_rc	The return code of the operation
CRM_target_rc	The expected return code of the operation
CRM_notify_status	The numerical representation of the status of the operation

The following example configures a **ClusterMon** resource that executes the external program **crm_logger.sh** which will log the event notifications specified in the program.

The following procedure creates the **crm_logger.sh** program that this resource will use.

1. On one node of the cluster, create the program that will log the event notifications.

```
# cat <<-END >/usr/local/bin/crm_logger.sh
#!/bin/sh
logger -t "ClusterMon-External" "${CRM_notify_node} ${CRM_notify_rsc} \
${CRM_notify_task} ${CRM_notify_desc} ${CRM_notify_rc} \
${CRM_notify_target_rc} ${CRM_notify_status} ${CRM_notify_recipient}";
exit;
END
```

2. Set the ownership and permissions for the program.

```
# chmod 700 /usr/local/bin/crm_logger.sh
# chown root.root /usr/local/bin/crm_logger.sh
```

3. Use the **scp** command to copy the **crm_logger.sh** program to the other nodes of the cluster, putting the program in the same location on those nodes and setting the same ownership and permissions for the program.

The following example configures the **ClusterMon** resource, named **ClusterMon-External**, that runs the program **/usr/local/bin/crm_logger.sh**. The **ClusterMon** resource outputs the cluster status to an **html** file, which is **/var/www/html/cluster_mon.html** in this example. The **pidfile** detects whether **ClusterMon** is already running; in this example that file is **/var/run/crm_mon-external.pid**. This resource is created as a clone so that it will run on every node in the cluster. The **watch-fencing** is specified to enable monitoring of fencing events in addition to resource events, including the start/stop/monitor, start/monitor, and stop of the fencing resource.

```
# pcs resource create ClusterMon-External ClusterMon user=root \
update=10 extra_options="-E /usr/local/bin/crm_logger.sh --watch-fencing" \
htmlfile=/var/www/html/cluster_mon.html \
pidfile=/var/run/crm_mon-external.pid clone
```



NOTE

The **crm_mon** command that this resource executes and which could be run manually is as follows:

```
# /usr/sbin/crm_mon -p /var/run/crm_mon-manual.pid -d -i 5 \
-h /var/www/html/crm_mon-manual.html -E "/usr/local/bin/crm_logger.sh" \
--watch-fencing
```

The following example shows the format of the output of the monitoring notifications that this example yields.

```
Aug 7 11:31:32 rh6node1pcmk ClusterMon-External: rh6node2pcmk.examplerh.com ClusterIP
st_notify_fence Operation st_notify_fence requested by rh6node1pcmk.examplerh.com for peer
rh6node2pcmk.examplerh.com: OK (ref=b206b618-e532-42a5-92eb-44d363ac848e) 0 0 0 #177
Aug 7 11:31:32 rh6node1pcmk ClusterMon-External: rh6node1pcmk.examplerh.com ClusterIP start
OK 0 0 0
Aug 7 11:31:32 rh6node1pcmk ClusterMon-External: rh6node1pcmk.examplerh.com ClusterIP
monitor OK 0 0 0
Aug 7 11:33:59 rh6node1pcmk ClusterMon-External: rh6node1pcmk.examplerh.com fence_xvms
monitor OK 0 0 0
Aug 7 11:33:59 rh6node1pcmk ClusterMon-External: rh6node1pcmk.examplerh.com ClusterIP
monitor OK 0 0 0
Aug 7 11:33:59 rh6node1pcmk ClusterMon-External: rh6node1pcmk.examplerh.com ClusterMon-
External start OK 0 0 0
Aug 7 11:33:59 rh6node1pcmk ClusterMon-External: rh6node1pcmk.examplerh.com fence_xvms
start OK 0 0 0
Aug 7 11:33:59 rh6node1pcmk ClusterMon-External: rh6node1pcmk.examplerh.com ClusterIP start
OK 0 0 0
Aug 7 11:33:59 rh6node1pcmk ClusterMon-External: rh6node1pcmk.examplerh.com ClusterMon-
External monitor OK 0 0 0
Aug 7 11:34:00 rh6node1pcmk crmd[2887]: notice: te_rsc_command: Initiating action 8: monitor
ClusterMon-External:1_monitor_0 on rh6node2pcmk.examplerh.com
Aug 7 11:34:00 rh6node1pcmk crmd[2887]: notice: te_rsc_command: Initiating action 16: start
ClusterMon-External:1_start_0 on rh6node2pcmk.examplerh.com
Aug 7 11:34:00 rh6node1pcmk ClusterMon-External: rh6node1pcmk.examplerh.com ClusterIP stop
OK 0 0 0
Aug 7 11:34:00 rh6node1pcmk crmd[2887]: notice: te_rsc_command: Initiating action 15: monitor
ClusterMon-External_monitor_10000 on rh6node2pcmk.examplerh.com
Aug 7 11:34:00 rh6node1pcmk ClusterMon-External: rh6node2pcmk.examplerh.com ClusterMon-
External start OK 0 0 0
Aug 7 11:34:00 rh6node1pcmk ClusterMon-External: rh6node2pcmk.examplerh.com ClusterMon-
External monitor OK 0 0 0
Aug 7 11:34:00 rh6node1pcmk ClusterMon-External: rh6node2pcmk.examplerh.com ClusterIP start
OK 0 0 0
Aug 7 11:34:00 rh6node1pcmk ClusterMon-External: rh6node2pcmk.examplerh.com ClusterIP
monitor OK 0 0 0
```

CHAPTER 14. CONFIGURING MULTI-SITE CLUSTERS WITH PACEMAKER

When a cluster spans more than one site, issues with network connectivity between the sites can lead to split-brain situations. When connectivity drops, there is no way for a node on one site to determine whether a node on another site has failed or is still functioning with a failed site interlink. In addition, it can be problematic to provide high availability services across two sites which are too far apart to keep synchronous.

To address these issues, Red Hat Enterprise Linux release 7.4 provides full support for the ability to configure high availability clusters that span multiple sites through the use of a Booth cluster ticket manager. The Booth *ticket manager* is a distributed service that is meant to be run on a different physical network than the networks that connect the cluster nodes at particular sites. It yields another, loose cluster, a *Booth formation*, that sits on top of the regular clusters at the sites. This aggregated communication layer facilitates consensus-based decision processes for individual Booth tickets.

A Booth *ticket* is a singleton in the Booth formation and represents a time-sensitive, movable unit of authorization. Resources can be configured to require a certain ticket to run. This can ensure that resources are run at only one site at a time, for which a ticket or tickets have been granted.

You can think of a Booth formation as an overlay cluster consisting of clusters running at different sites, where all the original clusters are independent of each other. It is the Booth service which communicates to the clusters whether they have been granted a ticket, and it is Pacemaker that determines whether to run resources in a cluster based on a Pacemaker ticket constraint. This means that when using the ticket manager, each of the clusters can run its own resources as well as shared resources. For example there can be resources A, B and C running only in one cluster, resources D, E, and F running only in the other cluster, and resources G and H running in either of the two clusters as determined by a ticket. It is also possible to have an additional resource J that could run in either of the two clusters as determined by a separate ticket.

The following procedure provides an outline of the steps you follow to configure a multi-site configuration that uses the Booth ticket manager.

These example commands use the following arrangement:

- Cluster 1 consists of the nodes **cluster1-node1** and **cluster1-node2**
- Cluster 1 has a floating IP address assigned to it of 192.168.11.100
- Cluster 2 consists of **cluster2-node1** and **cluster2-node2**
- Cluster 2 has a floating IP address assigned to it of 192.168.22.100
- The arbitrator node is **arbitrator-node** with an ip address of 192.168.99.100
- The name of the Booth ticket that this configuration uses is **apacheticket**

These example commands assume that the cluster resources for an Apache service have been configured as part of the resource group **apache** for each cluster. It is not required that the resources and resource groups be the same on each cluster to configure a ticket constraint for those resources, since the Pacemaker instance for each cluster is independent, but that is a common failover scenario.

For a full cluster configuration procedure that configures an Apache service in a cluster, see the example in *High Availability Add-On Administration*.

Note that at any time in the configuration procedure you can enter the **pcs booth config** command to display the booth configuration for the current node or cluster or the **pcs booth status** command to display the current status of booth on the local node.

1. Install the **booth-site** Booth ticket manager package on each node of both clusters.

```
[root@cluster1-node1 ~]# yum install -y booth-site  
[root@cluster1-node2 ~]# yum install -y booth-site  
[root@cluster2-node1 ~]# yum install -y booth-site  
[root@cluster2-node2 ~]# yum install -y booth-site
```

2. Install the **pcs**, **booth-core**, and **booth-arbitrator** packages on the arbitrator node.

```
[root@arbitrator-node ~]# yum install -y pcs booth-core booth-arbitrator
```

3. Ensure that ports 9929/tcp and 9929/udp are open on all cluster nodes and on the arbitrator node.

For example, running the following commands on all nodes in both clusters as well as on the arbitrator node allows access to ports 9929/tcp and 9929/udp on those nodes.

```
# firewall-cmd --add-port=9929/udp  
# firewall-cmd --add-port=9929/tcp  
# firewall-cmd --add-port=9929/udp --permanent  
# firewall-cmd --add-port=9929/tcp --permanent
```

Note that this procedure in itself allows any machine anywhere to access port 9929 on the nodes. You should ensure that on your site the nodes are open only to the nodes that require them.

4. Create a Booth configuration on one node of one cluster. The addresses you specify for each cluster and for the arbitrator must be IP addresses. For each cluster, you specify a floating IP address.

```
[cluster1-node1 ~] # pcs booth setup sites 192.168.11.100 192.168.22.100 arbitrators  
192.168.99.100
```

This command creates the configuration files **/etc/booth/booth.conf** and **/etc/booth/booth.key** on the node from which it is run.

5. Create a ticket for the Booth configuration. This is the ticket that you will use to define the resource constraint that will allow resources to run only when this ticket has been granted to the cluster.

This basic failover configuration procedure uses only one ticket, but you can create additional tickets for more complicated scenarios where each ticket is associated with a different resource or resources.

```
[cluster1-node1 ~] # pcs booth ticket add apacheticket
```

6. Synchronize the Booth configuration to all nodes in the current cluster.

```
[cluster1-node1 ~] # pcs booth sync
```

- From the arbitrator node, pull the Booth configuration to the arbitrator. If you have not previously done so, you must first authenticate **pcs** to the node from which you are pulling the configuration.

```
[arbitrator-node ~] # pcs cluster auth cluster1-node1
[arbitrator-node ~] # pcs booth pull cluster1-node1
```

- Pull the Booth configuration to the other cluster and synchronize to all the nodes of that cluster. As with the arbitrator node, if you have not previously done so, you must first authenticate **pcs** to the node from which you are pulling the configuration.

```
[cluster2-node1 ~] # pcs cluster auth cluster1-node1
[cluster2-node1 ~] # pcs booth pull cluster1-node1
[cluster2-node1 ~] # pcs booth sync
```

- Start and enable Booth on the arbitrator.



NOTE

You must not manually start or enable Booth on any of the nodes of the clusters since Booth runs as a Pacemaker resource in those clusters.

```
[arbitrator-node ~] # pcs booth start
[arbitrator-node ~] # pcs booth enable
```

- Configure Booth to run as a cluster resource on both cluster sites. This creates a resource group with **booth-ip** and **booth-service** as members of that group.

```
[cluster1-node1 ~] # pcs booth create ip 192.168.11.100
[cluster2-node1 ~] # pcs booth create ip 192.168.22.100
```

- Add a ticket constraint to the resource group you have defined for each cluster.

```
[cluster1-node1 ~] # pcs constraint ticket add apacheticket apachegroup
[cluster2-node1 ~] # pcs constraint ticket add apacheticket apachegroup
```

You can enter the following command to display the currently configured ticket constraints.

```
pcs constraint ticket [show]
```

- Grant the ticket you created for this setup to the first cluster.

Note that it is not necessary to have defined ticket constraints before granting a ticket. Once you have initially granted a ticket to a cluster, then Booth takes over ticket management unless you override this manually with the **pcs booth ticket revoke** command. For information on the **pcs booth** administration commands, see the PCS help screen for the **pcs booth** command.

```
[cluster1-node1 ~] # pcs booth ticket grant apacheticket
```

It is possible to add or remove tickets at any time, even after completing this procedure. After adding or removing a ticket, however, you must synchronize the configuration files to the other nodes and clusters as well as to the arbitrator and grant the ticket as is shown in this procedure.

For information on additional Booth administration commands that you can use for cleaning up and removing Booth configuration files, tickets, and resources, see the PCS help screen for the **pcs booth** command.

APPENDIX A. OCF RETURN CODES

This appendix describes the OCF return codes and how they are interpreted by Pacemaker.

The first thing the cluster does when an agent returns a code is to check the return code against the expected result. If the result does not match the expected value, then the operation is considered to have failed, and recovery action is initiated.

For any invocation, resource agents must exit with a defined return code that informs the caller of the outcome of the invoked action.

There are three types of failure recovery, as described in [Table A.1, “Types of Recovery Performed by the Cluster”](#).

Table A.1. Types of Recovery Performed by the Cluster

Type	Description	Action Taken by the Cluster
soft	A transient error occurred.	Restart the resource or move it to a new location .
hard	A non-transient error that may be specific to the current node occurred.	Move the resource elsewhere and prevent it from being retried on the current node.
fatal	A non-transient error that will be common to all cluster nodes occurred (for example, a bad configuration was specified).	Stop the resource and prevent it from being started on any cluster node.

[Table A.2, “OCF Return Codes”](#) provides The OCF return codes and the type of recovery the cluster will initiate when a failure code is received. Note that even actions that return 0 (OCF alias **OCF_SUCCESS**) can be considered to have failed, if 0 was not the expected return value.

Table A.2. OCF Return Codes

Return Code	OCF Label	Description
0	OCF_SUCCESS	<p>The action completed successfully. This is the expected return code for any successful start, stop, promote, and demote command.</p> <p>Type if unexpected: soft</p>
1	OCF_ERR_GENERIC	<p>The action returned a generic error.</p> <p>Type: soft</p> <p>The resource manager will attempt to recover the resource or move it to a new location.</p>

Return Code	OCF Label	Description
2	OCF_ERR_ARGS	<p>The resource's configuration is not valid on this machine. For example, it refers to a location not found on the node.</p>
		<p>Type: hard</p>
		<p>The resource manager will move the resource elsewhere and prevent it from being retried on the current node</p>
3	OCF_ERR_UNIMPLEMENTED	<p>The requested action is not implemented.</p>
		<p>Type: hard</p>
4	OCF_ERR_PERM	<p>The resource agent does not have sufficient privileges to complete the task. This may be due, for example, to the agent not being able to open a certain file, to listen on a specific socket, or to write to a directory.</p>
		<p>Type: hard</p>
		<p>Unless specifically configured otherwise, the resource manager will attempt to recover a resource which failed with this error by restarting the resource on a different node (where the permission problem may not exist).</p>
5	OCF_ERR_INSTALLED	<p>A required component is missing on the node where the action was executed. This may be due to a required binary not being executable, or a vital configuration file being unreadable.</p>
		<p>Type: hard</p>
		<p>Unless specifically configured otherwise, the resource manager will attempt to recover a resource which failed with this error by restarting the resource on a different node (where the required files or binaries may be present).</p>

Return Code	OCF Label	Description
6	OCF_ERR_CONFIGURED	<p>The resource's configuration on the local node is invalid.</p> <p>Type: fatal</p> <p>When this code is returned, Pacemaker will prevent the resource from running on any node in the cluster, even if the service configuration is valid on some other node.</p>
7	OCF_NOT_RUNNING	<p>The resource is safely stopped. This implies that the resource has either gracefully shut down, or has never been started.</p> <p>Type if unexpected: soft</p> <p>The cluster will not attempt to stop a resource that returns this for any action.</p>
8	OCF_RUNNING_MASTER	<p>The resource is running in master mode.</p> <p>Type if unexpected: soft</p>
9	OCF_FAILED_MASTER	<p>The resource is in master mode but has failed.</p> <p>Type: soft</p> <p>The resource will be demoted, stopped and then started (and possibly promoted) again.</p>
other	N/A	Custom error code.

APPENDIX B. CLUSTER CREATION IN RED HAT ENTERPRISE LINUX 6 AND RED HAT ENTERPRISE LINUX 7

Configuring a Red Hat High Availability Cluster in Red Hat Enterprise Linux 7 with Pacemaker requires a different set of configuration tools with a different administrative interface than configuring a cluster in Red Hat Enterprise Linux 6 with **rgmanager**. [Section B.1, “Cluster Creation with rgmanager and with Pacemaker”](#) summarizes the configuration differences between the various cluster components.

Red Hat Enterprise Linux 6.5 and later releases support cluster configuration with Pacemaker, using the **pcs** configuration tool. [Section B.2, “Pacemaker Installation in Red Hat Enterprise Linux 6 and Red Hat Enterprise Linux 7”](#) summarizes the Pacemaker installation differences between Red Hat Enterprise Linux 6 and Red Hat Enterprise Linux 7.

B.1. CLUSTER CREATION WITH RGMANAGER AND WITH PACEMAKER

[Table B.1, “Comparison of Cluster Configuration with rgmanager and with Pacemaker”](#) provides a comparative summary of how you configure the components of a cluster with **rgmanager** in Red Hat Enterprise Linux 6 and with Pacemaker in Red Hat Enterprise Linux 7.

Table B.1. Comparison of Cluster Configuration with rgmanager and with Pacemaker

Configuration Component	rgmanager	Pacemaker
Cluster configuration file	The cluster configuration file on each node is cluster.conf file, which can be edited directly. Otherwise, use the luci or ccs interface to define the cluster configuration.	The cluster and Pacemaker configuration files are corosync.conf and cib.xml . Do not edit the cib.xml file directly; use the pcs or pcsd interface instead.
Network setup	Configure IP addresses and SSH before configuring the cluster.	Configure IP addresses and SSH before configuring the cluster.
Cluster Configuration Tools	luci , ccs command, manual editing of cluster.conf file.	pcs or pcsd .
Installation	Install rgmanager (which pulls in all dependencies, including ricci , luci , and the resource and fencing agents). If needed, install lvm2-cluster and gfs2-utils .	Install pcs , and the fencing agents you require. If needed, install lvm2-cluster and gfs2-utils .

Configuration Component	rgmanager	Pacemaker
Starting cluster services	<p>Start and enable cluster services with the following procedure:</p> <ol style="list-style-type: none"> 1. Start rgmanager, cman, and, if needed, clvmd and gfs2. 2. Start ricci, and start luci if using the luci interface. 3. Run chkconfig on for the needed services so that they start at each runtime. <p>Alternately, you can enter ccs --start to start and enable the cluster services.</p>	<p>Start and enable cluster services with the following procedure:</p> <ol style="list-style-type: none"> 1. On every node, execute systemctl start pcsd.service, then systemctl enable pcsd.service to enable pcsd to start at runtime. 2. On one node in the cluster, enter pcs cluster start --all to start corosync and pacemaker.
Controlling access to configuration tools	For luci , the root user or a user with luci permissions can access luci . All access requires the ricci password for the node.	The pcsd gui requires that you authenticate as user hacluster , which is the common system user. The root user can set the password for hacluster .
Cluster creation	Name the cluster and define which nodes to include in the cluster with luci or ccs , or directly edit the cluster.conf file.	Name the cluster and include nodes with pcs cluster setup command or with the pcsd Web UI. You can add nodes to an existing cluster with the pcs cluster node add command or with the pcsd Web UI.
Propagating cluster configuration to all nodes	When configuration a cluster with luci , propagation is automatic. With ccs , use the --sync option. You can also use the cman_tool version -r command.	Propagation of the cluster and Pacemaker configuration files, corosync.conf and cib.xml , is automatic on cluster setup or when adding a node or resource.
Global cluster properties	<p>The following feature are supported with rgmanager in Red Hat Enterprise Linux 6:</p> <ul style="list-style-type: none"> * You can configure the system so that the system chooses which multicast address to use for IP multicasting in the cluster network. * If IP multicasting is not available, you can use UDP Unicast transport mechanism. * You can configure a cluster to use RRP protocol. 	<p>Pacemaker in Red Hat Enterprise Linux 7 supports the following features for a cluster:</p> <ul style="list-style-type: none"> * You can set no-quorum-policy for the cluster to specify what the system should do when the cluster does not have quorum. * For additional cluster properties you can set, see Table 12.1, “Cluster Properties”.
Logging	You can set global and daemon-specific logging configuration.	See the file /etc/sysconfig/pacemaker for information on how to configure logging manually.

Configuration Component	rgmanager	Pacemaker
Validating the cluster	Cluster validation is automatic with luci and with ccs , using the cluster schema. The cluster is automatically validated on startup.	The cluster is automatically validated on startup, or you can validate the cluster with pcs cluster verify .
Quorum in two-node clusters	With a two-node cluster, you can configure how the system determines quorum: * Configure a quorum disk * Use ccs or edit the cluster.conf file to set two_node=1 and expected_votes=1 to allow a single node to maintain quorum.	pcs automatically adds the necessary options for a two-node cluster to corosync .
Cluster status	On luci , the current status of the cluster is visible in the various components of the interface, which can be refreshed. You can use the --getconf option of the ccs command to see current the configuration file. You can use the clustat command to display cluster status.	You can display the current cluster status with the pcs status command.
Resources	You add resources of defined types and configure resource-specific properties with luci or the ccs command, or by editing the cluster.conf configuration file.	You add resources of defined types and configure resource-specific properties with the pcs resource create command or with the pcsd Web UI. For general information on configuring cluster resources with Pacemaker see Chapter 6, Configuring Cluster Resources .

Configuration Component	rgmanager	Pacemaker
Resource behavior, grouping, and start/stop order	Define cluster services to configure how resources interact.	<p>With Pacemaker, you use resource groups as a shorthand method of defining a set of resources that need to be located together and started and stopped sequentially. In addition, you define how resources behave and interact in the following ways:</p> <ul style="list-style-type: none"> * You set some aspects of resource behavior as resource options. * You use location constraints to determine which nodes a resource can run on. * You use order constraints to determine the order in which resources run. * You use colocation constraints to determine that the location of one resource depends on the location of another resource. <p>For more complete information on these topics, see Chapter 6, Configuring Cluster Resources and Chapter 7, Resource Constraints.</p>
Resource administration: Moving, starting, stopping resources	With luci , you can manage clusters, individual cluster nodes, and cluster services. With the ccs command, you can manage cluster. You can use the clusvadm to manage cluster services.	You can temporarily disable a node so that it cannot host resources with the pcs cluster standby command, which causes the resources to migrate. You can stop a resource with the pcs resource disable command.
Removing a cluster configuration completely	With luci , you can select all nodes in a cluster for deletion to delete a cluster entirely. You can also remove the cluster.conf from each node in the cluster.	You can remove a cluster configuration with the pcs cluster destroy command.
Resources active on multiple nodes, resources active on multiple nodes in multiple modes	No equivalent.	With Pacemaker, you can clone resources so that they can run in multiple nodes, and you can define cloned resources as master and slave resources so that they can run in multiple modes. For information on cloned resources and master/slave resources, see Chapter 9, Advanced Configuration .

Configuration Component	rgmanager	Pacemaker
Fencing -- single fence device per node	Create fencing devices globally or locally and add them to nodes. You can define post-fail delay and post-join delay values for the cluster as a whole.	Create a fencing device for each node with the pcs stonith create command or with the pcsd Web UI. For devices that can fence multiple nodes, you need to define them only once rather than separately for each node. You can also define pcmk_host_map to configure fencing devices for all nodes with a single command; for information on pcmk_host_map see Table 5.1, "General Properties of Fencing Devices" . You can define the stonith-timeout value for the cluster as a whole.
Multiple (backup) fencing devices per node	Define backup devices with luci or the ccs command, or by editing the cluster.conf file directly.	Configure fencing levels.

B.2. PACEMAKER INSTALLATION IN RED HAT ENTERPRISE LINUX 6 AND RED HAT ENTERPRISE LINUX 7

Red Hat Enterprise Linux 6.5 and later releases support cluster configuration with Pacemaker, using the **pcs** configuration tool. There are, however, some differences in cluster installation between Red Hat Enterprise Linux 6 and Red Hat Enterprise Linux 7 when using Pacemaker.

The following commands install the Red Hat High Availability Add-On software packages that Pacemaker requires in Red Hat Enterprise Linux 6 and prevent **corosync** from starting without **cman**. You must enter these commands on each node in the cluster.

```
[root@rhel6]# yum install pacemaker cman pcs
[root@rhel6]# chkconfig corosync off
[root@rhel6]# chkconfig cman off
```

On each node in the cluster, you set up a password for the **pcs** administration account named **hacluster**, and you start and enable the **pcsd** service.

```
[root@rhel6]# passwd hacluster
[root@rhel6]# service pcasd start
[root@rhel6]# chkconfig pcasd on
```

On one node in the cluster, you then authenticate the administration account for the nodes of the cluster.

```
[root@rhel6]# pcs cluster auth [node] [...] [-u username] [-p password]
```

In Red Hat Enterprise Linux 7, you run the following commands on each node in the cluster to install the Red Hat High Availability Add-On software packages that Pacemaker requires, set up a password for the **pcs** administration account named **hacluster**, and start and enable the **pcsd** service,

```
[root@rhel7]# yum install pcs pacemaker fence-agents-all  
[root@rhel7]# passwd hacluster  
[root@rhel7]# systemctl start pcsd.service  
[root@rhel7]# systemctl enable pcsd.service
```

In Red Hat Enterprise Linux 7, as in Red Hat Enterprise Linux 6, you authenticate the administration account for the nodes of the cluster by running the following command on one node in the cluster.

```
[root@rhel7]# pcs cluster auth [node] [...] [-u username] [-p password]
```

For further information on installation in Red Hat Enterprise Linux 7, see [Chapter 1, Red Hat High Availability Add-On Configuration and Management Reference Overview](#) and [Chapter 4, Cluster Creation and Administration](#).

APPENDIX C. REVISION HISTORY

Revision 8.1-1	Fri Feb 28 2020	Steven Levine
Document version for 7.8 Beta publication.		
Revision 7.1-1	Wed Aug 7 2019	Steven Levine
Document version for 7.7 GA publication.		
Revision 6.1-1	Thu Oct 4 2018	Steven Levine
Document version for 7.6 GA publication.		
Revision 5.1-2	Thu Mar 15 2018	Steven Levine
Document version for 7.5 GA publication.		
Revision 5.1-0	Thu Dec 14 2017	Steven Levine
Document version for 7.5 Beta publication.		
Revision 4.1-9	Tue Oct 17 2017	Steven Levine
Updated version for 7.4.		
Revision 4.1-5	Wed Jul 19 2017	Steven Levine
Document version for 7.4 GA publication.		
Revision 4.1-2	Wed May 10 2017	Steven Levine
Preparing document for 7.4 Beta publication.		
Revision 3.1-10	Tue May 2 2017	Steven Levine
Update to version for 7.3 GA publication.		
Revision 3.1-4	Mon Oct 17 2016	Steven Levine
Version for 7.3 GA publication.		
Revision 3.1-3	Wed Aug 17 2016	Steven Levine
Preparing document for 7.3 Beta publication.		
Revision 2.1-8	Mon Nov 9 2015	Steven Levine
Preparing document for 7.2 GA publication		
Revision 2.1-5	Mon Aug 24 2015	Steven Levine
Preparing document for 7.2 Beta publication.		
Revision 1.1-9	Mon Feb 23 2015	Steven Levine
Version for 7.1 GA release		
Revision 1.1-7	Thu Dec 11 2014	Steven Levine
Version for 7.1 Beta release		
Revision 0.1-41	Mon Jun 2 2014	Steven Levine
Version for 7.0 GA release		
Revision 0.1-2	Thu May 16 2013	Steven Levine
First printing of initial draft		

INDEX

, Cluster Creation

A

ACPI

configuring, Configuring ACPI For Use with Integrated Fence Devices

Action

Property

enabled, Resource Operations

id, Resource Operations

interval, Resource Operations

name, Resource Operations

on-fail, Resource Operations

timeout, Resource Operations

Action Property, Resource Operations

attribute, Node Attribute Expressions

Constraint Expression, Node Attribute Expressions

Attribute Expression, Node Attribute Expressions

attribute, Node Attribute Expressions

operation, Node Attribute Expressions

type, Node Attribute Expressions

value, Node Attribute Expressions

B

batch-limit, Summary of Cluster Properties and Options

Cluster Option, Summary of Cluster Properties and Options

boolean-op, Pacemaker Rules

Constraint Rule, Pacemaker Rules

C

Clone

Option

clone-max, Creating and Removing a Cloned Resource

clone-node-max, Creating and Removing a Cloned Resource

globally-unique, Creating and Removing a Cloned Resource

interleave, Creating and Removing a Cloned Resource

notify, Creating and Removing a Cloned Resource

ordered, [Creating and Removing a Cloned Resource](#)

Clone Option, Creating and Removing a Cloned Resource

Clone Resources, Resource Clones

clone-max, Creating and Removing a Cloned Resource

Clone Option, Creating and Removing a Cloned Resource

clone-node-max, Creating and Removing a Cloned Resource

Clone Option, Creating and Removing a Cloned Resource

Clones, Resource Clones

Cluster

Option

batch-limit, Summary of Cluster Properties and Options

cluster-delay, Summary of Cluster Properties and Options

cluster-infrastructure, Summary of Cluster Properties and Options

cluster-recheck-interval, Summary of Cluster Properties and Options

dc-version, Summary of Cluster Properties and Options

enable-acl, Summary of Cluster Properties and Options

fence-reaction, Summary of Cluster Properties and Options

last-lrm-refresh, Summary of Cluster Properties and Options

maintenance-mode, Summary of Cluster Properties and Options

migration-limit, Summary of Cluster Properties and Options

no-quorum-policy, Summary of Cluster Properties and Options

pe-error-series-max, Summary of Cluster Properties and Options

pe-input-series-max, Summary of Cluster Properties and Options

pe-warn-series-max, Summary of Cluster Properties and Options

placement-strategy, Summary of Cluster Properties and Options

shutdown-escalation, Summary of Cluster Properties and Options

start-failure-is-fatal, Summary of Cluster Properties and Options

stonith-action, Summary of Cluster Properties and Options

stonith-enabled, Summary of Cluster Properties and Options

stonith-timeout, Summary of Cluster Properties and Options

stop-all-resources, Summary of Cluster Properties and Options

stop-orphan-actions, Summary of Cluster Properties and Options

stop-orphan-resources, Summary of Cluster Properties and Options

symmetric-cluster, Summary of Cluster Properties and Options

Querying Properties, Querying Cluster Property Settings

Removing Properties, Setting and Removing Cluster Properties

Setting Properties, Setting and Removing Cluster Properties

cluster administration

configuring ACPI, [Configuring ACPI For Use with Integrated Fence Devices](#)

Cluster Option, [Summary of Cluster Properties and Options](#)

Cluster Properties, [Setting and Removing Cluster Properties](#), [Querying Cluster Property Settings](#)
cluster status

display, [Displaying Cluster Status](#)

cluster-delay, [Summary of Cluster Properties and Options](#)

Cluster Option, [Summary of Cluster Properties and Options](#)

cluster-infrastructure, [Summary of Cluster Properties and Options](#)

Cluster Option, [Summary of Cluster Properties and Options](#)

cluster-recheck-interval, [Summary of Cluster Properties and Options](#)

Cluster Option, [Summary of Cluster Properties and Options](#)

Colocation, [Colocation of Resources](#)

Constraint

Attribute Expression, [Node Attribute Expressions](#)

attribute, [Node Attribute Expressions](#)

operation, [Node Attribute Expressions](#)

type, [Node Attribute Expressions](#)

value, [Node Attribute Expressions](#)

Date Specification, [Date Specifications](#)

hours, [Date Specifications](#)

id, [Date Specifications](#)

monthdays, [Date Specifications](#)

months, [Date Specifications](#)

moon, [Date Specifications](#)

weekdays, [Date Specifications](#)

weeks, [Date Specifications](#)

weekyears, [Date Specifications](#)

yeardays, [Date Specifications](#)

years, [Date Specifications](#)

Date/Time Expression, [Time/Date Based Expressions](#)

end, [Time/Date Based Expressions](#)

operation, [Time/Date Based Expressions](#)

start, [Time/Date Based Expressions](#)

Duration, [Durations](#)

Rule, [Pacemaker Rules](#)

boolean-op, [Pacemaker Rules](#)

role, [Pacemaker Rules](#)

score, [Pacemaker Rules](#)
score-attribute, [Pacemaker Rules](#)

[Constraint Expression](#), [Node Attribute Expressions](#), [Time/Date Based Expressions](#)

[Constraint Rule](#), [Pacemaker Rules](#)

Constraints

[Colocation](#), [Colocation of Resources](#)

Location

id, [Basic Location Constraints](#)
score, [Basic Location Constraints](#)

[Order](#), [Order Constraints](#)

kind, [Order Constraints](#)

D

dampen, [Moving Resources Due to Connectivity Changes](#)

Ping Resource Option, [Moving Resources Due to Connectivity Changes](#)

[Date Specification](#), [Date Specifications](#)

hours, [Date Specifications](#)
id, [Date Specifications](#)
monthdays, [Date Specifications](#)
months, [Date Specifications](#)
moon, [Date Specifications](#)
weekdays, [Date Specifications](#)
weeks, [Date Specifications](#)
weekyears, [Date Specifications](#)
yeardays, [Date Specifications](#)
years, [Date Specifications](#)

[Date/Time Expression](#), [Time/Date Based Expressions](#)

end, [Time/Date Based Expressions](#)
operation, [Time/Date Based Expressions](#)
start, [Time/Date Based Expressions](#)

dc-version, [Summary of Cluster Properties and Options](#)

Cluster Option, [Summary of Cluster Properties and Options](#)

Determine by Rules, [Using Rules to Determine Resource Location](#)

Determine Resource Location, [Using Rules to Determine Resource Location](#)

disabling

resources, [Enabling and Disabling Cluster Resources](#)

Duration, [Durations](#)

E

enable-acl, [Summary of Cluster Properties and Options](#)

Cluster Option, [Summary of Cluster Properties and Options](#)

enabled, [Resource Operations](#)

Action Property, [Resource Operations](#)

enabling

resources, [Enabling and Disabling Cluster Resources](#)

end, [Time/Date Based Expressions](#)

Constraint Expression, [Time/Date Based Expressions](#)

F

failure-timeout, [Resource Meta Options](#)

Resource Option, [Resource Meta Options](#)

features, new and changed, [New and Changed Features](#)

fence-reaction, [Summary of Cluster Properties and Options](#)

Cluster Option, [Summary of Cluster Properties and Options](#)

G

globally-unique, [Creating and Removing a Cloned Resource](#)

Clone Option, [Creating and Removing a Cloned Resource](#)

Group Resources, [Resource Groups](#)

Groups, [Resource Groups](#), [Group Stickiness](#)

H

host_list, [Moving Resources Due to Connectivity Changes](#)

Ping Resource Option, [Moving Resources Due to Connectivity Changes](#)

hours, [Date Specifications](#)

Date Specification, [Date Specifications](#)

I

id, [Resource Properties](#), [Resource Operations](#), [Date Specifications](#)

Action Property, [Resource Operations](#)

Date Specification, [Date Specifications](#)

Location Constraints, [Basic Location Constraints](#)

Multi-State Property, [Multistate Resources: Resources That Have Multiple Modes](#)

[Resource](#), [Resource Properties](#)

integrated fence devices

configuring ACPI, [Configuring ACPI For Use with Integrated Fence Devices](#)

interleave, [Creating and Removing a Cloned Resource](#)

Clone Option, [Creating and Removing a Cloned Resource](#)

interval, [Resource Operations](#)

Action Property, [Resource Operations](#)

is-managed, [Resource Meta Options](#)

Resource Option, [Resource Meta Options](#)

K

kind, [Order Constraints](#)

Order Constraints, [Order Constraints](#)

L

last-lrm-refresh, [Summary of Cluster Properties and Options](#)

Cluster Option, [Summary of Cluster Properties and Options](#)

Location

Determine by Rules, [Using Rules to Determine Resource Location](#)

score, [Basic Location Constraints](#)

Location Constraints, [Basic Location Constraints](#)

Location Relative to other Resources, [Colocation of Resources](#)

M

maintenance-mode, [Summary of Cluster Properties and Options](#)

Cluster Option, [Summary of Cluster Properties and Options](#)

master-max, [Multistate Resources: Resources That Have Multiple Modes](#)

Multi-State Option, [Multistate Resources: Resources That Have Multiple Modes](#)

master-node-max, [Multistate Resources: Resources That Have Multiple Modes](#)

Multi-State Option, [Multistate Resources: Resources That Have Multiple Modes](#)

migration-limit, [Summary of Cluster Properties and Options](#)

Cluster Option, [Summary of Cluster Properties and Options](#)

migration-threshold, [Resource Meta Options](#)

Resource Option, [Resource Meta Options](#)

monthdays, [Date Specifications](#)

 Date Specification, [Date Specifications](#)

months, [Date Specifications](#)

 Date Specification, [Date Specifications](#)

moon, [Date Specifications](#)

 Date Specification, [Date Specifications](#)

Moving, [Manually Moving Resources Around the Cluster](#)

 Resources, [Manually Moving Resources Around the Cluster](#)

Multi-State

 Option

 master-max, [Multistate Resources: Resources That Have Multiple Modes](#)

 master-node-max, [Multistate Resources: Resources That Have Multiple Modes](#)

 Property

 id, [Multistate Resources: Resources That Have Multiple Modes](#)

Multi-State Option, [Multistate Resources: Resources That Have Multiple Modes](#)

Multi-State Property, [Multistate Resources: Resources That Have Multiple Modes](#)

multiple-active, [Resource Meta Options](#)

 Resource Option, [Resource Meta Options](#)

multiplier, [Moving Resources Due to Connectivity Changes](#)

 Ping Resource Option, [Moving Resources Due to Connectivity Changes](#)

Multistate, [Multistate Resources: Resources That Have Multiple Modes](#), [Multistate Stickiness](#)

N

name, [Resource Operations](#)

 Action Property, [Resource Operations](#)

no-quorum-policy, [Summary of Cluster Properties and Options](#)

 Cluster Option, [Summary of Cluster Properties and Options](#)

notify, [Creating and Removing a Cloned Resource](#)

 Clone Option, [Creating and Removing a Cloned Resource](#)

O

OCF

 return codes, [OCF Return Codes](#)

on-fail, [Resource Operations](#)

Action Property, Resource Operations

operation, [Node Attribute Expressions](#), [Time/Date Based Expressions](#)

Constraint Expression, [Node Attribute Expressions](#), [Time/Date Based Expressions](#)

Option

batch-limit, [Summary of Cluster Properties and Options](#)

clone-max, [Creating and Removing a Cloned Resource](#)

clone-node-max, [Creating and Removing a Cloned Resource](#)

cluster-delay, [Summary of Cluster Properties and Options](#)

cluster-infrastructure, [Summary of Cluster Properties and Options](#)

cluster-recheck-interval, [Summary of Cluster Properties and Options](#)

dampen, [Moving Resources Due to Connectivity Changes](#)

dc-version, [Summary of Cluster Properties and Options](#)

enable-acl, [Summary of Cluster Properties and Options](#)

failure-timeout, [Resource Meta Options](#)

fence-reaction, [Summary of Cluster Properties and Options](#)

globally-unique, [Creating and Removing a Cloned Resource](#)

host_list, [Moving Resources Due to Connectivity Changes](#)

interleave, [Creating and Removing a Cloned Resource](#)

is-managed, [Resource Meta Options](#)

last-lrm-refresh, [Summary of Cluster Properties and Options](#)

maintenance-mode, [Summary of Cluster Properties and Options](#)

master-max, [Multistate Resources: Resources That Have Multiple Modes](#)

master-node-max, [Multistate Resources: Resources That Have Multiple Modes](#)

migration-limit, [Summary of Cluster Properties and Options](#)

migration-threshold, [Resource Meta Options](#)

multiple-active, [Resource Meta Options](#)

multiplier, [Moving Resources Due to Connectivity Changes](#)

no-quorum-policy, [Summary of Cluster Properties and Options](#)

notify, [Creating and Removing a Cloned Resource](#)

ordered, [Creating and Removing a Cloned Resource](#)

pe-error-series-max, [Summary of Cluster Properties and Options](#)

pe-input-series-max, [Summary of Cluster Properties and Options](#)

pe-warn-series-max, [Summary of Cluster Properties and Options](#)

placement-strategy, [Summary of Cluster Properties and Options](#)

priority, [Resource Meta Options](#)

requires, [Resource Meta Options](#)

resource-stickiness, [Resource Meta Options](#)

shutdown-escalation, [Summary of Cluster Properties and Options](#)

start-failure-is-fatal, [Summary of Cluster Properties and Options](#)

stonith-action, [Summary of Cluster Properties and Options](#)

stonith-enabled, [Summary of Cluster Properties and Options](#)
 stonith-timeout, [Summary of Cluster Properties and Options](#)
 stop-all-resources, [Summary of Cluster Properties and Options](#)
 stop-orphan-actions, [Summary of Cluster Properties and Options](#)
 stop-orphan-resources, [Summary of Cluster Properties and Options](#)
 symmetric-cluster, [Summary of Cluster Properties and Options](#)
 target-role, [Resource Meta Options](#)

Order

kind, [Order Constraints](#)

Order Constraints, [Order Constraints](#)

symmetrical, [Order Constraints](#)

ordered, [Creating and Removing a Cloned Resource](#)

Clone Option, [Creating and Removing a Cloned Resource](#)

Ordering, [Order Constraints](#)

overview

features, new and changed, [New and Changed Features](#)

P

pe-error-series-max, [Summary of Cluster Properties and Options](#)
 Cluster Option, [Summary of Cluster Properties and Options](#)

pe-input-series-max, [Summary of Cluster Properties and Options](#)
 Cluster Option, [Summary of Cluster Properties and Options](#)

pe-warn-series-max, [Summary of Cluster Properties and Options](#)
 Cluster Option, [Summary of Cluster Properties and Options](#)

Ping Resource

Option

dampen, [Moving Resources Due to Connectivity Changes](#)
 host_list, [Moving Resources Due to Connectivity Changes](#)
 multiplier, [Moving Resources Due to Connectivity Changes](#)

Ping Resource Option, [Moving Resources Due to Connectivity Changes](#)

placement strategy, [Utilization and Placement Strategy](#)

placement-strategy, [Summary of Cluster Properties and Options](#)

Cluster Option, [Summary of Cluster Properties and Options](#)

priority, [Resource Meta Options](#)

Resource Option, [Resource Meta Options](#)

Property

enabled, [Resource Operations](#)

id, [Resource Properties](#), [Resource Operations](#), [Multistate Resources: Resources That Have Multiple Modes](#)

interval, [Resource Operations](#)

name, [Resource Operations](#)

on-fail, [Resource Operations](#)

provider, [Resource Properties](#)

standard, [Resource Properties](#)

timeout, [Resource Operations](#)

type, [Resource Properties](#)

provider, [Resource Properties](#)

Resource, [Resource Properties](#)

Q

Querying

[Cluster Properties](#), [Querying Cluster Property Settings](#)

[Querying Options](#), [Querying Cluster Property Settings](#)

R

Removing

[Cluster Properties](#), [Setting and Removing Cluster Properties](#)

[Removing Properties](#), [Setting and Removing Cluster Properties](#)

requires, [Resource Meta Options](#)

Resource, [Resource Properties](#)

Constraint

[Attribute Expression](#), [Node Attribute Expressions](#)

[Date Specification](#), [Date Specifications](#)

[Date/Time Expression](#), [Time/Date Based Expressions](#)

[Duration](#), [Durations](#)

[Rule](#), [Pacemaker Rules](#)

Constraints

[Colocation](#), [Colocation of Resources](#)

[Order](#), [Order Constraints](#)

Location

[Determine by Rules](#), [Using Rules to Determine Resource Location](#)

[Location Relative to other Resources](#), [Colocation of Resources](#)

[Moving](#), [Manually Moving Resources Around the Cluster](#)

Option

failure-timeout, [Resource Meta Options](#)
is-managed, [Resource Meta Options](#)
migration-threshold, [Resource Meta Options](#)
multiple-active, [Resource Meta Options](#)
priority, [Resource Meta Options](#)
requires, [Resource Meta Options](#)
resource-stickiness, [Resource Meta Options](#)
target-role, [Resource Meta Options](#)

Property

id, [Resource Properties](#)
provider, [Resource Properties](#)
standard, [Resource Properties](#)
type, [Resource Properties](#)

Start Order, Order Constraints

Resource Option, [Resource Meta Options](#)
resource-stickiness, [Resource Meta Options](#)
Groups, [Group Stickiness](#)
Multi-State, [Multistate Stickiness](#)
Resource Option, [Resource Meta Options](#)

Resources, Manually Moving Resources Around the Cluster

Clones, [Resource Clones](#)
Groups, [Resource Groups](#)
Multistate, [Multistate Resources: Resources That Have Multiple Modes](#)

resources

cleanup, [Cluster Resources Cleanup](#)
disabling, [Enabling and Disabling Cluster Resources](#)
enabling, [Enabling and Disabling Cluster Resources](#)

role, Pacemaker Rules

Constraint Rule, [Pacemaker Rules](#)

Rule, Pacemaker Rules

boolean-op, [Pacemaker Rules](#)
Determine Resource Location, [Using Rules to Determine Resource Location](#)
role, [Pacemaker Rules](#)
score, [Pacemaker Rules](#)
score-attribute, [Pacemaker Rules](#)

S

score, [Basic Location Constraints](#), [Pacemaker Rules](#)

Constraint Rule, [Pacemaker Rules](#)

Location Constraints, [Basic Location Constraints](#)

score-attribute, [Pacemaker Rules](#)

Constraint Rule, [Pacemaker Rules](#)

Setting

Cluster Properties, [Setting and Removing Cluster Properties](#)

Setting Properties, [Setting and Removing Cluster Properties](#)

shutdown-escalation, [Summary of Cluster Properties and Options](#)

Cluster Option, [Summary of Cluster Properties and Options](#)

standard, [Resource Properties](#)

Resource, [Resource Properties](#)

start, [Time/Date Based Expressions](#)

Constraint Expression, [Time/Date Based Expressions](#)

Start Order, [Order Constraints](#)

start-failure-is-fatal, [Summary of Cluster Properties and Options](#)

Cluster Option, [Summary of Cluster Properties and Options](#)

status

display, [Displaying Cluster Status](#)

stonith-action, [Summary of Cluster Properties and Options](#)

Cluster Option, [Summary of Cluster Properties and Options](#)

stonith-enabled, [Summary of Cluster Properties and Options](#)

Cluster Option, [Summary of Cluster Properties and Options](#)

stonith-timeout, [Summary of Cluster Properties and Options](#)

Cluster Option, [Summary of Cluster Properties and Options](#)

stop-all-resources, [Summary of Cluster Properties and Options](#)

Cluster Option, [Summary of Cluster Properties and Options](#)

stop-orphan-actions, [Summary of Cluster Properties and Options](#)

Cluster Option, [Summary of Cluster Properties and Options](#)

stop-orphan-resources, [Summary of Cluster Properties and Options](#)

Cluster Option, [Summary of Cluster Properties and Options](#)

symmetric-cluster, [Summary of Cluster Properties and Options](#)

Cluster Option, [Summary of Cluster Properties and Options](#)

symmetrical, [Order Constraints](#)

[Order Constraints](#), [Order Constraints](#)

T

target-role, [Resource Meta Options](#)

[Resource Option](#), [Resource Meta Options](#)

Time Based Expressions, [Time/Date Based Expressions](#)

timeout, [Resource Operations](#)

Action Property, [Resource Operations](#)

type, [Resource Properties](#), [Node Attribute Expressions](#)

Constraint Expression, [Node Attribute Expressions](#)

Resource, [Resource Properties](#)

U

utilization attributes, [Utilization and Placement Strategy](#)

V

value, [Node Attribute Expressions](#)

Constraint Expression, [Node Attribute Expressions](#)

W

weekdays, [Date Specifications](#)

Date Specification, [Date Specifications](#)

weeks, [Date Specifications](#)

Date Specification, [Date Specifications](#)

weekyears, [Date Specifications](#)

Date Specification, [Date Specifications](#)

Y

yeardays, [Date Specifications](#)

Date Specification, [Date Specifications](#)

years, [Date Specifications](#)

Date Specification, [Date Specifications](#)

